

Effective C++ Programming

NIE-EPC (v. 2021):

INTRODUCTION, OBSERVABLE BEHAVIOR, AS-IF RULE,
OPTIMIZATIONS, TRANSLATION UNITS, ABI

© 2021 DANIEL LANGR, ČVUT (CTU) FIT

Programming languages

- **Programming** — writing **some functionality** with a special syntax (*source code*).
- “Normal” programs — execution of the functionality **by a CPU**.
- CPUs understand only a single “language” — **machine code**.
- *Implication*: source code must be **translated** into machine code.
 - The role of **programming languages implementations** (*compilers, interpreters, libraries, linkers,...*).
- *Example*:

```
// C++ function source code
int absmax(int a, int b)
{
    int abs_a = std::abs(a);
    int abs_b = std::abs(b);
    return std::max(abs_a, abs_b);
}
```

```
// function machine code in hex...
// ... generated by GCC on x86_64:
89 fa 89 f0 c1
fa 1f 31 d7 29
d7 99 31 d0 29
d0 39 c7 0f 4d
c7 c3
```

Abstraction

- How do programming languages **differ**?
- *Among others (different syntax, target application domain,...), they provide **different sets of abstraction mechanisms**.*
- **Abstraction:**
 - A **programming language concept** that makes expressing the required functionality easier.
 - *Examples:* variables, function parameters, function return values, loop constructs, structures, classes, inheritance, polymorphism, templates, reflection,...
 - **None of these concepts exist at a machine code level.**

```
// C++ function source code
int absmax(int a, int b)
{
    int abs_a = std::abs(a);
    int abs_b = std::abs(b);
    return std::max(abs_a, abs_b);
}
```

```
89 fa 89 f0 c1 fa
1f 31 d7 29 d7 99
31 d0 29 d0 39 c7
0f 4d c7 c3
// machine instruction ONLY for...
// ... register operations...
// ... plus final ret
```

Abstraction levels

- **Classification** of programming languages **according to abstraction mechanisms** they provide:
- **Low-level programming languages:**
 - **Low-level** = **simple** abstraction mechanisms (*variables, function parameters,...*).
- **High-level programming languages:**
 - **High-level** = **complex and more powerful** abstraction mechanisms (*classes, polymorphism, templates,...*).
 - Allow us to **express some functionality more easily...**
 - ...at a price of **loosing a control over a system**.
- **Low-level programming languages (cont.):**
 - **More effort** to express some functionality...
 - ...but provide **better system control**.
- **Generally:** the higher level of abstraction, the less system control, and vice versa.

Exemplary languages

- **Assembler**

- the lowest-level programming language,
- **total system control** — 1-to-1 correspondence between assembler instructions and machine code instructions,
- **extreme effort to write some functionality**,
- **not portable**.



- **C**

- portable,
- **large to moderate effort to write some functionality**,
- relatively **high level of system control** due to low-level abstraction mechanisms (*pointers, atomics, memory model control*).



- **Python**

- **easy to learn, easy to write some functionality with**,
- **almost no control of what happens at a machine code level** (no pointers, no atomics,...).






- **C++ unique feature:**
 - widely adopted language that...
 - ... has high-level abstraction mechanisms...
 - ... and provide high-level of system control at the same time.
- **High level of abstraction** — a relatively easy way how to put together some functionality, **if all necessary building blocks are available**.
 - *Example:* Writing a program that works with strings is easy by **using the `std::string` class** (no need to understand what happens at the machine code level).
- **High level of system control** — **development of building blocks with maximum performance and efficiency**.
 - *Example (cont.):* **implementing `std::string`** efficiently requires high level of system control (*pointers, stack/heap allocations, alignment, cache system, memory representation bits manipulation,...*).

C++ (cont.)

- **Control = opposite of safety:**
 - More control \Rightarrow more we can do with the system \Rightarrow **more bad things we can do with the system.**
 - *"Writing in C or C++ is like running a chain saw with all the safety guards removed." [Bob Gray]*
- **Control over system = control over performance \neq performance:**
 - *"C++ doesn't give you performance, it gives you control over performance." [Chandler Carruth (?)]*
 - **Performance is not provided automatically** by C++; it's easy to write an inefficient and slow program.
 - Doing the opposite requires a **high level of knowledge and experience.**
- **Mastering C++ is very hard!**
 - The latest C++ Standard has over **1850 pages.**



C++ standards

- Definition of C++ — C++ standards.
- Evolution of C++ — multiple standards:
 - “Old C++”: C++98, C++03
 - “Modern C++”: C++11, C++14, C++17, C++20 (latest), C++23 (upcoming)
- New standards are mostly backwards-compatible.
- “Release” versions are paid ISO/IEC documents.
- Working drafts are available, e.g., here: [[hyperlink](#)]
 - Prior to publication, drafts were/are referred to as C++1x, ..., C++2b.
- Standards are (try to be) comprehensive and exact.
 - *Implication:* they are hard to read (like a law code).
- Better for learning: books, tutorials, lectures, language references, Stack Overflow,...
- These sources may contain imprecise, misleading, or even incorrect information.
- When you are in trouble or have doubts, always check with the C++ standard.



C++ standards (cont.)

C++ Standards have 2 main parts:

1. C++ language

- Defines **core language**.
- Exemplary entities: *types, variables, functions, classes, templates,...*

2. C++ standard library

- Defines some useful **building blocks** for different functionality domains.
- Exemplary entities: *algorithms, containers, iterators, smart pointers, threading, I/O streams,...*
- **No implementation details** — C++ standards define only **API** (*application programming interface*) for library entities and optionally some requirements for their implementation.
- *Example: `std::sort` — **API + the requirement for $O(n \log n)$ time complexity on average** \Rightarrow can be implemented with *quicksort, heapsort,...* but not with *insertion sort, bubblesort* etc.*

How do C++ standards define C++?

- C++ Standards define how source code should **behave when it is translated and executed on an abstract machine**.
- **Abstract machine** — an abstract (*imaginary, virtual*) computer system with a **relatively small set of capabilities**:
 - **abstract CPU** — arithmetic operations, bitwise operations,...;
 - **memory** — static allocations, dynamic allocations, addressing,...;
 - **files** — reading/writing data;
 - **file system** (*since C++17*);
 - **standard input/output/error streams**;
 - **threading and memory model** (*since C++11*),
 - **some other abstract operation system mechanisms** (program arguments, program exit code, signals,...).
- **No screen, no keyboard, no mouse, no camera, no GPU, no networking, no heap, no stack, no instruction set, no caches, no cores,...**

Observable behavior

- Prescribed **behavior** for some code = effects of this code **observable from outside of this code, that is, from the perspective of its exterior/surroundings.**
- The term “observable behavior” is typically used.
- *Example:*

```
void f()  
{  
    std::cout << "some string";  
}
```

- Observable behavior of function f:
 1. writing the string “some string” into the program output stream,
 2. implicit return = returning the control flow to the function caller.
- This behavior **may be observed from outside of the function:**
 1. by reading the standard output stream on the abstract machine,
 2. by the function caller.

Observable behavior (cont.)

- *Another example:*

```
void g() {  
    int a = 1, b = 2, c; // declaration (+initialization) of local variables  
    c = a + b;           // addition and assignment  
}
```

- Observable behavior of function g:
 - implicit return = returning the control flow to the function caller;
 - nothing of the explicitly written code inside the function body contributes to its observable behavior (function-local non-static variables cannot be accessed from outside of this function).
- *Yet another example:*

```
int h() {  
    int a = 1, b = 2, c;  
    c = a + b;  
    return c;  
}
```

- Observable behavior of function h:
 - Returning the value 3 to the function caller.
 - Nothing else inside its definition contributes to its observable behavior.

Observable behavior (*cont.*)

- The boundary between observable and non-observable behavior **is not always clear**.
- **Observable behavior:**
 - reading/writing standard input/output/error streams;
 - writing data to files;
 - passing return values from functions to their callers;
 - modifying data in memory, if this memory may be accessed from outside;
 - updates of objects of *volatile* and *atomic* data types;
 - program termination with some exit code;
 - etc.
- **Hard case example:**
 - dynamic memory allocations.



Observable behavior (*cont.*)

- *Example:*

```
int main() {  
    int* ptr = new int(1);  
    delete ptr;  
}
```

- **Observable behavior of function main:**
 - Implicit “return 0” at the end (*since C++11*).
- **Nothing else?** — The allocated memory cannot be accessed from outside of the function.
- *Controversy:* memory allocation **may be observed from outside of the code** (e.g., by measuring program memory consumption, by memory debuggers like *Valgrind* or *Heaptrack*, etc.).
- **In the past**, dynamic memory allocations **were mostly considered** as a part of the observable behavior.
- **Now, “mostly” not:**
 - *Experiments:* GCC since version 10, Clang since version 3.2.

Observable behavior (*cont.*)

- *...example from the previous slide:*

```
int main() {  
    int* ptr = new int(1);  
    delete ptr;  
}
```

- *Another example:*

```
int* f() {  
    int* ptr = new int(1);  
    return ptr;  
}
```

- Observable behavior of function `f`:
 - Returning a pointer to dynamically allocated object of type `int` with value 1.
- The **effect** of both:
 - dynamic memory allocation and
 - initialization of the allocated object to 1
- **is observable from outside of the function.**

Observable behavior (*cont.*)

- *Yet another example:*

```
int* f() {  
    int* ptr = new int(1);  
    return ptr;  
}  
  
int main() {  
    int* ptr = f();  
    delete ptr;  
}
```

- Observable behavior of function `f` **is the very same** as on the previous slide:
 - Returning a pointer to dynamically allocated object of type `int` with value 1.
- **Observable behavior of function `main`:**
 - **Implicit “return 0” only**, provided that allocations themselves are not considered to contribute to the observable behavior (*as with latest GCC and Clang*).
 - Calling `f` itself **cannot be observed** from outside of `main`.

C++ implementation

- **C++ standard** — defines **observable behavior** for some code executed on the **abstract machine**.
- How does C++ work on a **real computer system**?
- There, we need a **C++ implementation**:
 - A tool that for some source code **yields on a particular computer system the same observable behavior** as is prescribed to it by the C++ standard.
- Typically, a **collection of tools** rather than a single tool only.
 - CPUs do not understand C++ source code \Rightarrow **it must be translated into machine code**.
 - Such translation is the main purpose of a **C++ compiler**.
 - Mainstream C++ compiler vendors:
GNU/GCC, LLVM/Clang, Microsoft/MSVC, Intel, IBM, PGI, Cray,...



C++ implementation (*cont.*)

- C++ implementation = C++ compiler only? Not at all...
- **Preprocessor:**
 - Processes **preprocessor directives and symbols** (`#include`, `#define`, *macros, conditional compilation,...*).
- **Linker:**
 - Links **different parts of machine code together**.
 - Generates the **final binary executable file** (with a particular structure given by the system, such as ELF on Linux).
- **Implementation of the C++ standard library:**
 - **Implements the C++ Standard library API** defined by C++ standards with respect to their requirements.
 - May or may not be **bound to a compiler**.
 - Mainstream implementations: *libstdc++ (GNU)*, *libc++ (LLVM)*, *Microsoft STL* (*note: their source code is available on GitHub*).
- **And possibly others...**

As-if rule

- One of the **most fundamental and essential C++ concept**.
- C++ standard prescribes **observable behavior** for some source code when it is executed on the **abstract machine**.
- For the same source code, a C++ implementation generates **machine code** executable on a particular **computer system**.
- **Conforming implementation**: the observable behavior of the generated machine code **must be the same** as the observable behavior prescribed to the source code by the C++ standard.
- The implementation must guarantee that the machine code will **observably behave as if** the source code was executed on the abstract machine \Rightarrow known as the "**as-if rule**".



```
// abstract machine
int absmax(int a, int b)
{
    int abs_a = std::abs(a);
    int abs_b = std::abs(b);
    return std::max(abs_a, abs_b);
}
```



```
// GCC/x86_64/Linux
89 fa 89 f0 c1 fa
1f 31 d7 29 d7 99
31 d0 29 d0 39 c7
0f 4d c7 c3
```

Optimizations

- The **as-if rule**:
 - The standard tells the implementation **what behavior it should provide** for some source code.
 - It does not tell anything about **how this behavior should be provided**.
- How implementations work in this regard? Generally, they are two options:

1. Disabled optimizations

- *Semantics*: a user **does not care about performance and efficiency**.
- *Typical situation*: **debugging** (debug builds).
- *Example*: g++ or clang++ with -O0 command line option.

2. Enabled optimizations

- *Semantics*: a user **needs performance and efficiency**.
- *Typical situations*: **release (production) builds**.
- *Example*: g++ or clang++ with -O2 command line option.

Optimizations (cont.)

- Effects:

1. Disabled optimizations:

- a compiler generates machine code for **everything that is in the source code**.

2. Enabled optimizations:

- a compiler generates optimal machine code **for observable behavior only**.

- *Example (Clang/Linux/x86_64):*

```
void g()
{
    int a = 1;
    int b = 2;
    int c;
    c = a + b;
    // implicit return
}
```

```
// machine code
// compiled with
// clang++ -O0:
55 48 89 e5 c7 45
fc 01 00 00 00 c7
45 f8 02 00 00 00
8b 45 fc 03 45 f8
89 45 f4 5d c3
```



```
// machine code
// compiled with
// clang++ -O2:
c3
```



- *Recall:* observable behavior of function g = **returning the control flow to the function caller**.

Disassembling

- Reading machine code is **hard for humans**.
- Each **machine code instruction** (*opcode*) can be translated into a **corresponding assembler instruction**, which are much easier to read.
- This process is called **disassembling**.
- Special tools = **disassemblers**:
 - objdump, gdb, Compiler Explorer (*online*),...
- **Compiler explorer** (<https://godbolt.org/>):
 - excellent easy-to-use online disassembler;
 - many programming languages including **C++**;
 - **many C++ implementations** ("compilers") and their versions (*GCC, Clang, Intel, MSVC,...*);
 - **multiple target architectures** (*x86, x86_64, ARM64, MIPS64, power64, RISC-V,...*);
 - created and maintained by *Matt Godbolt*.

Compiler Explorer

- *Example:*

The screenshot displays the Compiler Explorer interface. On the left, the C++ source code is shown in a text editor, with a red circle highlighting the code block. A red arrow points from this circle to the label "source code". Above the code editor, a dropdown menu shows "C++" selected, with a red circle around it and a red arrow pointing to the label "language selection". To the right of the source code, the compiler selection dropdown shows "x86-64 clang 11.0.1", circled in red, with a red arrow pointing to the label "implementation selection". Further right, the optimization level dropdown shows "-O0", also circled in red, with a red arrow pointing to the label "compiler arguments". On the right side of the interface, the disassembled machine code is displayed, with a red circle around the assembly block and a red arrow pointing to the label "disassembled generated machine code".

source code

language selection

implementation selection

compiler arguments

disassembled generated machine code

Optimizations (cont.)

- *Example (cont.)* — **disassembled** generated machine code with **disabled optimizations**:

```
void g()
{
    int a = 1;
    int b = 2;
    int c;
    c = a + b;
    // implicit return
}
```

```
// clang++ -O0:
g():
    // initialize stack frame:
    push rbp
    mov rbp, rsp
    // initialize a mapped to address (rbp-4):
    mov dword ptr [rbp - 4], 1
    // initialize b mapped to address (rbp-8):
    mov dword ptr [rbp - 8], 2
    mov eax, dword ptr [rbp - 4] // :load a from memory
    // load b from memory and add with a:
    add eax, dword ptr [rbp - 8]
    // store result to c mapped to address (rbp-12):
    mov dword ptr [rbp - 12], eax
    pop rbp // :restore stack frame
    ret    // :return control flow to caller
```

- *Observation*: machine code reflects **all the source code**, independently of whether it **does or does not contribute to the observable behavior**.
- *Live demo*: <https://godbolt.org/z/6fGGEd>.

Optimizations (cont.)

- *Example (cont.)* — **disassembled** generated machine code with **enabled optimizations**:

```
void g() {  
    int a = 1;  
    int b = 2;  
    int c;  
    c = a + b;  
    // implicit return  
}
```

```
// clang++ -O2:  
g():  
    ret    // :return control flow to caller
```

- *Observation*: machine code reflects **only observable function behavior**.
- *Live demo*: <https://godbolt.org/z/aWx7c5>.
- *Common C++ (or C) myth*: **function-local (non-static) variables are stored on the stack** → **false**.
- *Note*: CPUs don't know the concept of a variable ⇒ better considering **variable binary representation being stored** somewhere than variables themselves.
- *Truth*: Variable binary representation may be stored **on the stack**, **somewhere else** (e.g., in registers), **nowhere at all**.

**Myth or
Fact?**

Use case — benchmarking

- The as-if rule is highly relevant to **writing benchmarks**.
- **Benchmarking** — measuring **runtime or performance** of some operations.
- *Example:* measuring single-core **floating-point performance** on small (*fit-into-cache*) matrix multiplication:

```
void benchmark() {  
    double A[32][32] = {};  
    double B[32][32] = {};  
    double C[32][32] = {};  
  
    for (int n = 0; n < 1'000'000'000; n++) // billion times...  
        for (int i = 0; i < 32; i++) // ... multiply C ← C + A * B  
            for (int j = 0; j < 32; j++)  
                for (int k = 0; k < 32; k++)  
                    C[i][j] += A[i][k] * B[k][j];  
}
```

$$C_{i,j} \leftarrow C_{i,j} + \sum_k A_{i,k} B_{k,j}$$

- **Total number of FP operations:** $N = 65.536$ trillions.
- **FP performance** = measured function runtime / N .
- **Measurement** on Microsoft Surface (i5, 2017): **40.96 EFlop/s**.

Use case — benchmarking (cont.)

- *Microsoft Surface*, single core: 40.96 EFlop/s.
- *Fugaku supercomputer*, 7.6 million cores: 0.442 EFlop/s.
- *Apparent conclusion*: my Surface is 93× faster than the fastest supercomputer in the world (in Q1/2021).



- *Explanation*: observable behavior of benchmark function is only **implicit return** \Rightarrow measured number of FP operations = 0.



```
void benchmark() {  
    double A[32][32] = {};  
    double B[32][32] = {};  
    double C[32][32] = {};  
  
    for (int n = 0; n < 1'000'000'000; n++)  
        for (int i = 0; i < 32; i++)  
            for (int j = 0; j < 32; j++)  
                for (int k = 0; k < 32; k++)  
                    C[i][j] += A[i][k] * B[k][j];  
}
```

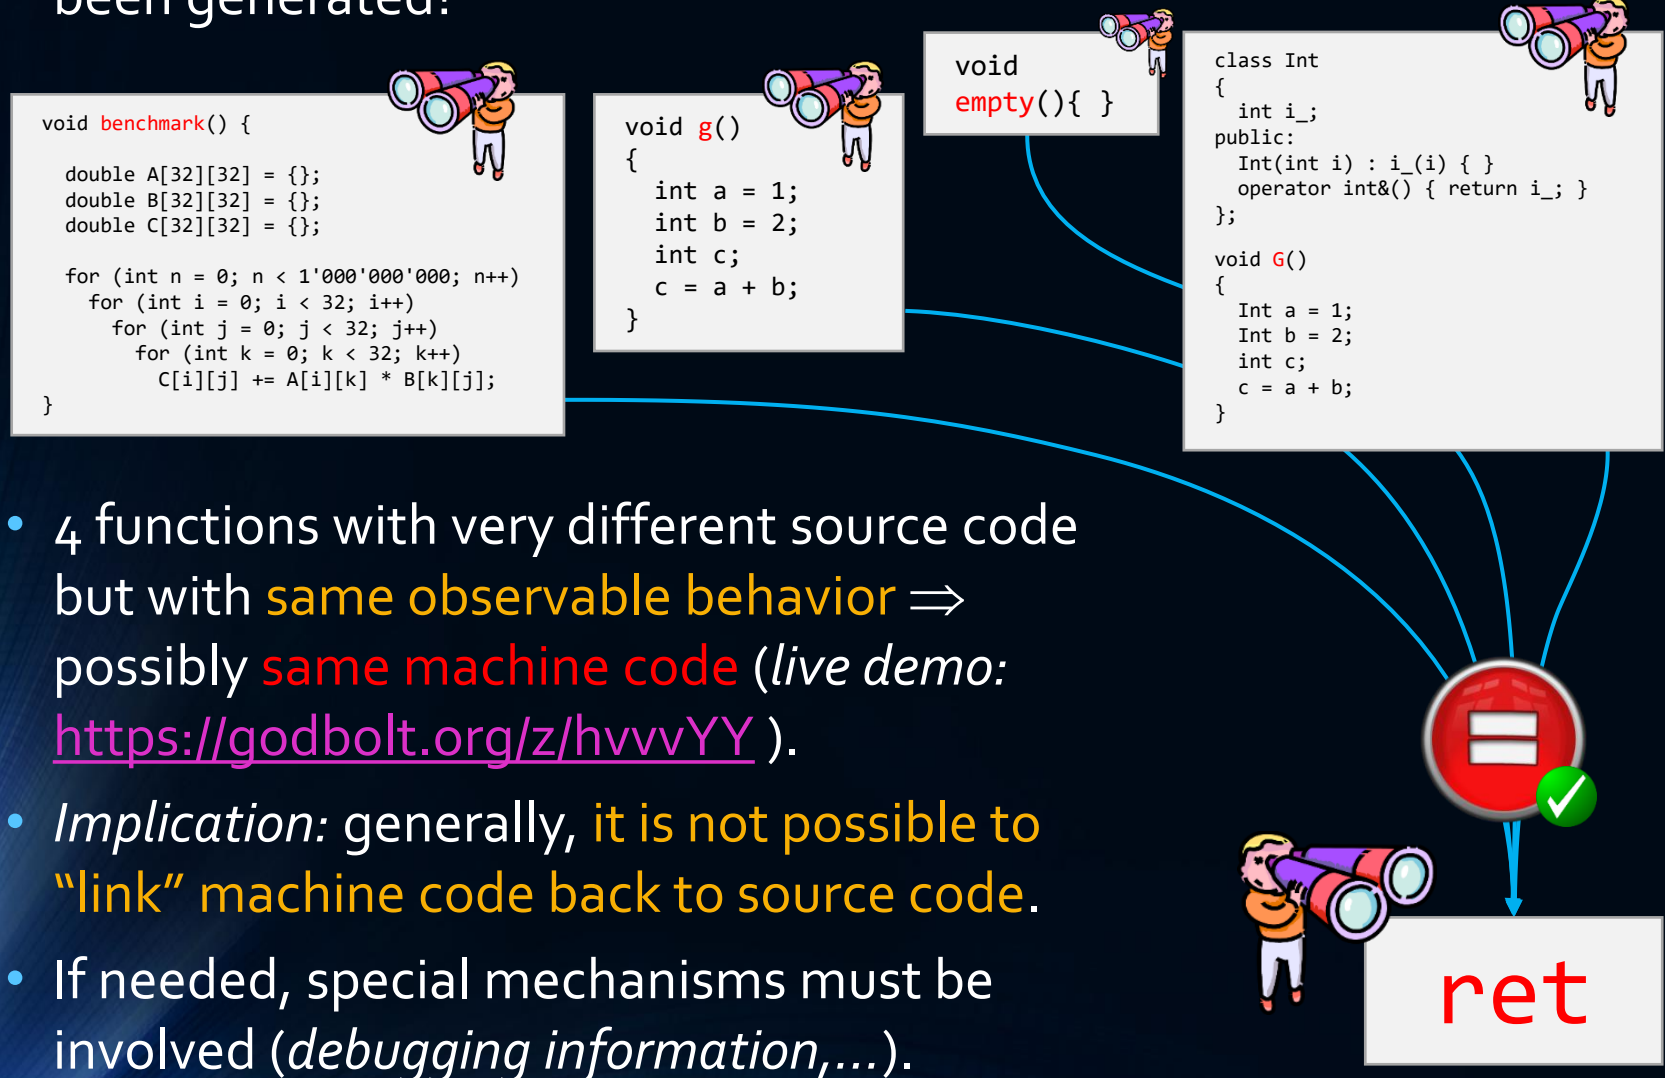


```
// g++ -O2:  
benchmark():  
    ret
```

- *Conclusion*: benchmarked operations **must contribute to the observable behavior** to avoid them being optimized away.

Machine code → source code?

- Can we obtain source code from which some machine code has been generated?



- 4 functions with very different source code but with **same observable behavior** ⇒ possibly **same machine code** (live demo: <https://godbolt.org/z/hvYvYY>).
- Implication:* generally, it is not possible to “link” machine code back to source code.
- If needed, special mechanisms must be involved (debugging information,...).

Translation units

- **Observable behavior** is related to translation units.
- **Translation unit** — a piece of source code that **is translated into machine code (compiled) at once**.
- Basically: a translation unit = **source file**:
 - **without comments**
 - and **processed by a C++ preprocessor**.
- *Example: **source file** (left), **its translation unit** with undefined INT (right above) and INT defined as long (right below):*

```
// add.cpp source file
#ifndef INT
    #define INT int
#endif

// returns sum of its arguments:
INT add(INT a, INT b) {
    return a + b;
}
```

```
int add(int a, int b) {
    return a + b;
}
```

```
long add(long a, long b) {
    return a + b;
}
```

- *Note: translation unit may be obtained by some implementation tools, such as „**g++ -DINT=long -E -P add.cpp**“.*

Translation units (*cont.*)

- Compilers translate **each translation unit separately**.
- **Machine code for a single translation unit** is typically stored in so-called **object file** (.o or .obj file extension).
- To generate an object file with the machine code translated from a translation unit related to some source file, GCC or Clang uses:
 - **-c** command line option,
 - **source file name** as a command line argument.
- *Example:* `g++ -O2 -c add.cpp`

// add.cpp source file

```
#ifndef INT
#define INT int
#endif
```

// returns sum of its arguments:

```
INT add(INT a, INT b) {
    return a + b;
}
```

**source file
= source code**

```
int add(int a, int b) {
    return a + b;
}
```

translation unit

add(int, int):
8d 04 37 c3

**object file
= machine code**

add(int, int):
lea eax, [rdi+rsi]
ret

**disassembled
object file**

GCC/Clang workflow

- How do `g++` and `clang++` commands work?

a) *Input* = `source file(s)` and `-c` option \Rightarrow *translation/compilation*:

1. transforms the `source file` into a `translation unit`,
2. translates the source code of the `translation unit` into `machine code`,
3. stores that `machine code` in an `object file`.

- *Example*: `g++ -O2 -c main.cpp`.

b) *Input* = `object file(s)` \Rightarrow *linking*:

1. merges/links `machine code` from object file(s),
 2. stores it into an `executable binary file`.
- *Note*: one of the object files must contain `machine code of the main function`.

- *Example*: `g++ main.o` (executable file is called `a.out` by default).

c) *Input* = `source file(s)`, no `-c` option \Rightarrow *translation + linking*:

- *Ad a)* and *ad b)* at once.
- *Note*: an object file **is not explicitly stored** (it is stored only temporarily, for example in `/tmp` directory with some random file name).

GCC/Clang workflow (cont.)

- *Example* — source code:

```
// main.cpp
int main() {
    return 12;
}
```

- **Separate** compilation and linking:

```
$ g++ -O2 -c main.cpp
$ g++ main.o
$ ./a.out
$ echo $?
12
```

- **Single-command** alternative:

```
$ g++ -O2 main.cpp
$ ./a.out
$ echo $?
12
```



- *Observable behavior of main*: returns 12.
- *C++ standard*: value returned from `main` = program exit status.
- *Bash*: exit status of terminated program — `$?` shell parameter.

Translation units (*cont.*)

- *Previous example:* only a single source file \Rightarrow a single translation unit.
- Source code of programs usually consists of multiple translation units.
- *Problems:*
 - Multiple translation units must cooperate at the source code level.
 - Machine code translated from multiple translation units must cooperate at the machine code level.
 - How is observable behavior resolved in the context of multiple translation units?
- *Example:*

```
// add.cpp source file
int add(int a, int b) {
    return a + b;
}
```

```
// main.cpp source file ver.1
int main() {
    return add(1, 2) + 3;
}
```

- In one source file, we want to use a function defined in another source file.

Translation units (*cont.*)


- Translation units for `add.cpp` and `main.cpp`, respectively:

```
int add(int a, int b) {  
    return a + b;  
}
```

```
int main() {  
    return add(1, 2) + 3;  
}
```


- Compilation of `add.cpp` is **ok**:

```
$ g++ -O2 -c add.cpp
```



- *Problem*: compilation of `main.cpp` is **not ok**:

```
$ g++ -O2 -c main.cpp  
error: 'add' was not declared in this scope
```



- **Observable behavior of `main`:**
 - **calling function** named `add` with 1 and 2 integer arguments,
 - **adding** returned value with integer 3,
 - **returning** the result out of `main`.
- A compiler needs to generate machine code for this observable behavior. **Why did it fail?**

Translation units (*cont.*)

- During **translation**, a compiler always works only with a **single translation unit**.

```
int main() { return add(1, 2) + 3; }
```

- When a compiler translates `main.cpp`, **what does it “see”**?
 1. **add function call**,
 2. **passed arguments** (integer literals having values 1 and 2),
 3. **what is done with the returned value** (addition with 3).
- This is **not enough** for generating machine code for the observable behavior. **Why?**
 - For example, are the arguments **passed to add by value or by reference**?
 - Or, is there any **implicit conversion** required?
 - Such things can make **difference at the machine code level**.
- *Implication*: to be able to translate a function call, a compiler **must see the type of its parameters as well as its return type**.

Function declaration

- Type of function parameters and its return type makes part of its **interface (API)**.
- In C++, function API = **function declaration** (a.k.a. *prototype*).
- **Using (= calling) a function** in some translation unit requires its **declaration to be present in that translation unit** as well.
- *Note:* function definition is also its declaration.
- Minimal **declaration of add function**:

```
int add(int, int);
```

- Addition into **main.cpp**:

```
// main.cpp source file ver.2  
int add(int, int);  
int main() { return add(1, 2) + 3; }
```

- Compilation is **now ok**:
 - types of parameters **are known**, return type **is known**.

```
$ g++ -O2 -c main.cpp
```



Translation units (*cont.*)

- **API** (such as declarations for functions) resolves **cooperation of translation units at a source code level**.
- What about the **machine code level**?
- Translation unit of `main.cpp`:

```
int add(int, int);  
int main() { return add(1, 2) + 3; }
```

- Its compilation with...

```
$ g++ -O2 -c main.cpp
```

- ...results in an **object file `main.o`** that contains **machine code of function `main`** that reflects its **observable behavior**.
- If we try to **build an executable program** file from it, **it fails**:

```
$ g++ main.o  
undefined reference to `add(int, int)'
```

- *Why?*



Translation units (cont.)

- The part of the **observable behavior of main** is **calling add**.
- This observable behavior is **reflected in its machine code**:

```
main:
  sub    rsp, 8
  mov    esi, 2
  mov    edi, 1
  call   add(int, int)
  add    rsp, 8
  add    eax, 3
  ret
```

- To **build the executable binary**, the linker needs to **connect (link)** this **call** instruction with the machine code of the **add** function.
- However, the linker **cannot find this machine code**.
- **Machine code of add** cannot be generated from the translation unit of **main.cpp** \Rightarrow it **is not included in main.o** object file.
- This problem triggers the “well-known” **undefined reference error**:

```
undefined reference to `add(int, int)'
```

Translation units (*cont.*)

- The only way the **machine code of add** can be generated is from its definition.
- **Function definition** basically = function declaration + its body.
- In our case, this definition is **in the translation unit of add.cpp (left)** \Rightarrow its compilation therefore generates **machine code of add function in add.o object file (right)**:

```
int add(int a, int b) {  
    return a + b;  
}
```

```
add(int, int):  
    lea eax, [rdi+rsi]  
    ret
```

- To generate the **final executable file**, we need to provide **both main.o and add.o** object files to the linker:

```
$ g++ main.o add.o
```



ABI

```
$ g++ main.o add.o
```

- The final executable file `a.out` contains **machine code of both functions `main` and `add`**:

```
int add(int, int);  
int main()  
{  
    return add(1, 2) + 3;  
}
```

```
int add(int a, int b)  
{  
    return a + b;  
}
```

- Both pieces of machine code **need to cooperate at a machine code level**.
- Why? Because of **passing arguments from `main` to `add` and the return value from `add` back to `main`**.

```
main:  
    sub    rsp, 8  
    mov    esi, 2  
    mov    edi, 1  
    call   add(int, int)  
    add    rsp, 8  
    add    eax, 3  
    ret
```

```
add(int, int):  
    lea    eax, [rdi+rsi]  
    ret
```

- The **`call`** instruction **doesn't provide such a functionality** (it only stores the current value of the *instruction pointer* on the stack and sets its new value to the address of the machine code of **`add`**).
- The arguments and the return value **must be passed by some explicit mechanism**.

ABI (cont.)

- In order to **allow different functions to cooperate at a machine code level**, there must be some **agreement about how to pass arguments and return values to/from functions**.
- This agreement (*calling conventions*) is a part of so-called **application binary interface (ABI)**.
- **API:**
 - enables cooperation of translation units at a source code (*programming*) level;
 - is **defined by C++ standards**;
 - *example*: function declaration allows to compile function calls.
- **ABI:**
 - enables cooperation of machine code generated from different translation units at a machine code (*binary*) level;
 - is **not defined by C++ standards** (typically, it is defined by the architecture + operating system + C++ implementation);
 - *example*: specification of how to pass arguments to functions.

ABI (cont.)

- In my case — **x86_64 + Linux + GCC** — the **applied ABI** specifies that:
 - **int** is a **32-bit signed integer**;
 - **argument for the first parameter of type int** is passed into a function through the **edi register**;
 - **argument for the second parameter of type int** is passed into a function through the **esi register**;
 - **return value of type int** is passed from a function through the **eax register**.

```
int main()
{
    return add(1, 2) + 3;
}
```

```
main:
    sub    rsp, 8
    mov    esi, 2
    mov    edi, 1
    call   add(int, int)
    add    rsp, 8
    add    eax, 3
    ret
```

```
add(int, int):
    lea    eax, [rdi+rsi]
    ret
```

- *Note:* the latter holds for **add as well as for main**.
- *Note:* the ABI requires **stack frames to be 16-byte aligned** \Rightarrow **rsp register manipulations**.

ABI (cont.)

- **Importance of ABI** — it allows linking the machine code generated:
 - by **different C++ implementations**,
 - at **different places** (different computers),
 - at **different times**.

```
$ clang++ -O2 -c add.cpp
$ g++ -O2 -c main.cpp
$ g++ main.o add.o
$ ./a.out
$ echo $?
6
```

- ABI is a must for using **libraries**.
- **Library:**
 - A **collection of object files** that implements some functionality bundled in a **file with specific format** (.a, .so, .lib, .dll,...).
 - Plus **API** (usually *header files*) that enables to work with that functionality at a source code level (for example, library function declarations).