

# Effective C++ Programming

NIE-EPC (v. 2021):  
INTRODUCTION, OBSERVABLE BEHAVIOR, AS-IF RULE,  
OPTIMIZATIONS, TRANSLATION UNITS, ABI  
© 2021 DANIEL LANGR, ČVUT (CTU) FIT

1

## Programming languages

- Programming — writing some functionality with a special syntax (*source code*).
- “Normal” programs — execution of the functionality by a CPU.
- CPUs understand only a single “language” — machine code.
- *Implication*: source code must be translated into machine code.
  - The role of programming languages implementations (*compilers, interpreters, libraries, linkers,...*).
- *Example*:

```
// C++ function source code
int absmax(int a, int b)
{
    int abs_a = std::abs(a);
    int abs_b = std::abs(b);
    return std::max(abs_a, abs_b);
}
```

```
// function machine code in hex...
// ... generated by GCC on x86_64:
89 fa 89 f0 c1
fa 1f 31 d7 29
d7 99 31 d0 29
d0 39 c7 0f 4d
c7 c3
```

2

## Abstraction

- How do programming languages differ?
- *Among others (different syntax, target application domain,...), they provide different sets of abstraction mechanisms.*
- **Abstraction**:
  - A programming language concept that makes expressing the required functionality easier.
  - *Examples*: variables, function parameters, function return values, loop constructs, structures, classes, inheritance, polymorphism, templates, reflection,...
  - None of these concepts exist at a machine code level.

```
// C++ function source code
int absmax(int a, int b)
{
    int abs_a = std::abs(a);
    int abs_b = std::abs(b);
    return std::max(abs_a, abs_b);
}
```

```
89 fa 89 f0 c1 fa
1f 31 d7 29 d7 99
31 d0 29 d0 39 c7
0f 4d c7 c3
// machine instruction ONLY for...
// ... register operations...
// ... plus final ret
```

3

## Abstraction levels

- Classification of programming languages according to abstraction mechanisms they provide:
- **Low-level programming languages**:
  - Low-level = simple abstraction mechanisms (*variables, function parameters,...*).
- **High-level programming languages**:
  - High-level = complex and more powerful abstraction mechanisms (*classes, polymorphism, templates,...*).
  - Allow us to express some functionality more easily...
  - ...at a price of losing a control over a system.
- **Low-level programming languages (cont.)**:
  - More effort to express some functionality...
  - ...but provide better system control.
- *Generally*: the higher level of abstraction, the less system control, and vice versa.

4

## Exemplary languages

- **Assembler**
  - the lowest-level programming language,
  - total system control — 1-to-1 correspondence between assembler instructions and machine code instructions,
  - extreme effort to write some functionality,
  - not portable.
- **C**
  - portable,
  - large to moderate effort to write some functionality,
  - relatively high level of system control due to low-level abstraction mechanisms (*pointers, atomics, memory model control*).
- **Python**
  - easy to learn, easy to write some functionality with,
  - almost no control of what happens at a machine code level (no pointers, no atomics,...).



5

## C++

- **C++ unique feature**:
  - widely adopted language that...
  - ... has high-level abstraction mechanisms...
  - ... and provide high-level of system control at the same time.
- **High level of abstraction** — a relatively easy way how to put together some functionality, if all necessary building blocks are available.
  - *Example*: Writing a program that works with strings is easy by using the `std::string` class (no need to understand what happens at the machine code level).
- **High level of system control** — development of building blocks with maximum performance and efficiency.
  - *Example (cont.)*: implementing `std::string` efficiently requires high level of system control (*pointers, stack/heap allocations, alignment, cache system, memory representation bits manipulation,...*).



6

## C++ (cont.)

- **Control = opposite of safety:**
  - More control  $\Rightarrow$  more we can do with the system  $\Rightarrow$  more bad things we can do with the system.
  - "Writing in C or C++ is like running a chain saw with all the safety guards removed." [Bob Gray]
- **Control over system = control over performance  $\neq$  performance:**
  - "C++ doesn't give you performance, it gives you control over performance." [Chandler Carruth (?)]
  - Performance is not provided automatically by C++; it's easy to write an inefficient and slow program.
  - Doing the opposite requires a high level of knowledge and experience.
- **Mastering C++ is very hard!**
  - The latest C++ Standard has over 1850 pages.



7

## C++ standards

- Definition of C++ — C++ standards.
- Evolution of C++ — multiple standards:
  - "Old C++": C++98, C++03
  - "Modern C++": C++11, C++14, C++17, C++20 (latest), C++23 (upcoming)
- New standards are mostly backwards-compatible.
- "Release" versions are paid ISO/IEC documents.
- Working drafts are available, e.g., here: [\[hyperlink\]](#)
  - Prior to publication, drafts were referred to as C++1X, ..., C++2b.
- Standards are (try to be) comprehensive and exact.
  - *Implication:* they are hard to read (*like a law code*).
- Better for learning: books, tutorials, lectures, language references, Stack Overflow, ...
  - These sources may contain imprecise, misleading, or even incorrect information.
  - When you are in trouble or have doubts, always check with the C++ standard.



8

## C++ standards (cont.)

C++ Standards have 2 main parts:

1. **C++ language**
  - Defines core language.
  - Exemplary entities: *types, variables, functions, classes, templates,...*
2. **C++ standard library**
  - Defines some useful building blocks for different functionality domains.
  - Exemplary entities: *algorithms, containers, iterators, smart pointers, threading, I/O streams,...*
  - No implementation details — C++ standards define only API (*application programming interface*) for library entities and optionally some requirements for their implementation.
  - *Example:* `std::sort` — API + the requirement for  $O(n \log n)$  time complexity on average  $\Rightarrow$  can be implemented with *quicksort*, *heapsort*, ... but not with *insertion sort*, *bubblesort* etc.

9

## How do C++ standards define C++?

- C++ Standards define how source code should behave when it is translated and executed on an abstract machine.
- **Abstract machine** — an abstract (*imaginary, virtual*) computer system with a relatively small set of capabilities:
  - abstract CPU — arithmetic operations, bitwise operations, ...;
  - memory — static allocations, dynamic allocations, addressing, ...;
  - files — reading/writing data;
  - file system (*since C++17*);
  - standard input/output/error streams;
  - threading and memory model (*since C++11*);
  - some other abstract operation system mechanisms (program arguments, program exit code, signals, ...).
- No screen, no keyboard, no mouse, no camera, no GPU, no networking, no heap, no stack, no instruction set, no caches, no cores, ...

10

## Observable behavior

- Prescribed behavior for some code = effects of this code observable from outside of this code, that is, from the perspective of its exterior/surroundings.
- The term "observable behavior" is typically used.
- *Example:*

```
void f()
{
    std::cout << "some string";
}
```

- Observable behavior of function f:
  1. writing the string "some string" into the program output stream,
  2. implicit return = returning the control flow to the function caller.
- This behavior may be observed from outside of the function:
  1. by reading the standard output stream on the abstract machine,
  2. by the function caller.

11

## Observable behavior (cont.)

- *Another example:*

```
void g() {
    int a = 1, b = 2, c; // declaration (+initialization) of local variables
    c = a + b;           // addition and assignment
}
```

- Observable behavior of function g:
  - implicit return = returning the control flow to the function caller;
  - nothing of the explicitly written code inside the function body contributes to its observable behavior (function-local non-static variables cannot be accessed from outside of this function).

- *Yet another example:*

```
int h() {
    int a = 1, b = 2, c;
    c = a + b;
    return c;
}
```

- Observable behavior of function h:
  - Returning the value 3 to the function caller.
  - Nothing else inside its definition contributes to its observable behavior.

12

### Observable behavior (cont.)

- The boundary between observable and non-observable behavior is not always clear.
- Observable behavior:**
  - reading/writing standard input/output/error streams;
  - writing data to files;
  - passing return values from functions to their callers;
  - modifying data in memory, if this memory may be accessed from outside;
  - updates of objects of *volatile* and *atomic* data types;
  - program termination with some exit code;
  - etc.
- Hard case example:**
  - dynamic memory allocations.



13

### Observable behavior (cont.)

• **Example:**

```
int main() {
    int* ptr = new int(1);
    delete ptr;
}
```

- Observable behavior of function `main`:
  - Implicit "return 0" at the end (since C++11).
- Nothing else?** — The allocated memory cannot be accessed from outside of the function.
- Controversy:** memory allocation may be observed from outside of the code (e.g., by measuring program memory consumption, by memory debuggers like *Valgrind* or *Heaptrack*, etc.).
- In the past, dynamic memory allocations were mostly considered as a part of the observable behavior.
- Now, "mostly" not:
  - Experiments:* GCC since version 10, Clang since version 3.2.

14

### Observable behavior (cont.)

• ...example from the previous slide:

```
int main() {
    int* ptr = new int(1);
    delete ptr;
}
```

• **Another example:**

```
int* f() {
    int* ptr = new int(1);
    return ptr;
}
```

- Observable behavior of function `f`:
  - Returning a pointer to dynamically allocated object of type `int` with value 1.
- The effect of both:
  - dynamic memory allocation and
  - initialization of the allocated object to 1
- is observable from outside of the function.

15

### Observable behavior (cont.)

• **Yet another example:**

```
int* f() {
    int* ptr = new int(1);
    return ptr;
}

int main() {
    int* ptr = f();
    delete ptr;
}
```

- Observable behavior of function `f` is the very same as on the previous slide:
  - Returning a pointer to dynamically allocated object of type `int` with value 1.
- Observable behavior of function `main`:
  - Implicit "return 0" only, provided that allocations themselves are not considered to contribute to the observable behavior (as with latest GCC and Clang).
  - Calling `f` itself cannot be observed from outside of `main`.

16

### C++ implementation

- C++ standard — defines observable behavior for some code executed on the abstract machine.
- How does C++ work on a real computer system?
- There, we need a **C++ implementation**:
  - A tool that for some source code yields on a particular computer system the same observable behavior as is prescribed to it by the C++ standard.
- Typically, a collection of tools rather than a single tool only.
  - CPUs do not understand C++ source code  $\Rightarrow$  it must be translated into machine code.
- Such translation is the main purpose of a **C++ compiler**.
- Mainstream C++ compiler vendors:
  - GNU/GCC, LLVM/Clang, Microsoft/MSVC, Intel, IBM, PGI, Cray,...



17

### C++ implementation (cont.)

- C++ implementation = C++ compiler only? Not at all...
- Preprocessor:**
  - Processes preprocessor directives and symbols (`#include`, `#define`, *macros*, *conditional compilation*, ...).
- Linker:**
  - Links different parts of machine code together.
  - Generates the final binary executable file (with a particular structure given by the system, such as ELF on Linux).
- Implementation of the C++ standard library:**
  - Implements the C++ Standard library API defined by C++ standards with respect to their requirements.
  - May or may not be bound to a compiler.
  - Mainstream implementations: *libstdc++* (GNU), *libc++* (LLVM), *Microsoft STL* (note: their source code is available on *GitHub*).
- And possibly others...**

18

### As-if rule

- One of the most fundamental and essential C++ concept.
- C++ standard prescribes observable behavior for some source code when it is executed on the abstract machine.
- For the same source code, a C++ implementation generates machine code executable on a particular computer system.
- Conforming implementation: the observable behavior of the generated machine code **must be the same** as the observable behavior prescribed to the source code by the C++ standard.
- The implementation must guarantee that the machine code will **observably behave as if** the source code was executed on the abstract machine  $\Rightarrow$  know as the "**as-if rule**".



19


### Optimizations

- The as-if rule:
  - The standard tells the implementation what behavior it should provide for some source code.
  - It does not tell anything about how this behavior should be provided.
- How implementations work in this regard? Generally, they are two options:
  1. **Disabled optimizations**
    - *Semantics*: a user does not care about performance and efficiency.
    - *Typical situation*: debugging (debug builds).
    - *Example*: g++ or clang++ with -O0 command line option.
  2. **Enabled optimizations**
    - *Semantics*: a user needs performance and efficiency.
    - *Typical situations*: release (production) builds.
    - *Example*: g++ or clang++ with -O2 command line option.

20

### Optimizations (cont.)

- Effects:
  1. **Disabled optimizations:**
    - a compiler generates machine code for everything that is in the source code.
  2. **Enabled optimizations:**
    - a compiler generates optimal machine code for observable behavior only.
- *Example (Clang/Linux/x86\_64):*



- *Recall*: observable behavior of function g = returning the control flow to the function caller.

21

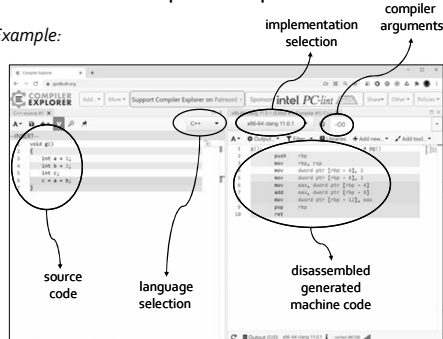
### Disassembling

- Reading machine code is hard for humans.
- Each machine code instruction (*opcode*) can be translated into a corresponding assembler instruction, which are much easier to read.
- This process is called **disassembling**.
- Special tools = disassemblers:
  - objdump, gdb, Compiler Explorer (online),...
- Compiler explorer (<https://godbolt.org/>):
  - excellent easy-to-use online disassembler;
  - many programming languages including C++;
  - many C++ implementations ("compilers") and their versions (GCC, Clang, Intel, MSVC,...);
  - multiple target architectures (x86, x86\_64, ARM64, MIPS64, power64, RISC-V,...);
  - created and maintained by Matt Godbolt.

22

### Compiler Explorer

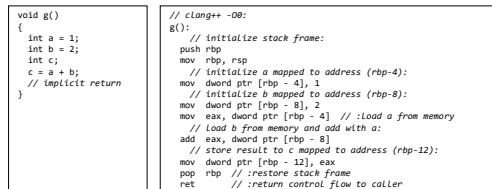
- *Example:*



23

### Optimizations (cont.)

- *Example (cont.)* — disassembled generated machine code with disabled optimizations:



- *Observation*: machine code reflects all the source code, independently of whether it does or does not contribute to the observable behavior.
- *Live demo*: <https://godbolt.org/z/6fGGEd>.

24

## Optimizations (cont.)

- *Example (cont.)* — disassembled generated machine code with enabled optimizations:

```
void g() {
    int a = 1;
    int b = 2;
    int c;
    c = a + b;
    // implicit return
}
```

```
// clang++ -O2:
ret // :return control flow to caller
```

- **Observation:** machine code reflects only observable function behavior.
- **Live demo:** <https://godbolt.org/z/aWx7c5>.
- **Common C++ (or C) myth:** function-local (non-static) variables are stored on the stack → *false*.
- **Note:** CPUs don't know the concept of a variable ⇒ better considering variable *binary representation* being stored somewhere than variables themselves.
- **Truth:** Variable binary representation may be stored on the stack, somewhere else (e.g., in registers), nowhere at all.

25

## Use case — benchmarking

- The as-if rule is highly relevant to writing benchmarks.
- **Benchmarking** — measuring runtime or performance of some operations.
- **Example:** measuring single-core floating-point performance on small (*fit-into-cache*) matrix multiplication:

```
void benchmark() {
    double A[32][32] = {};
    double B[32][32] = {};
    double C[32][32] = {};

    for (int n = 0; n < 1'000'000'000; n++) // Billion times...
        for (int i = 0; i < 32; i++)
            for (int j = 0; j < 32; j++)
                for (int k = 0; k < 32; k++)
                    C[i][j] += A[i][k] * B[k][j];
}
```

$$C_{i,j} \leftarrow C_{i,j} + \sum_k A_{i,k} B_{k,j}$$

- Total number of FP operations:  $N = 65.536$  trillions.
- **FP performance** = *measured function runtime* /  $N$ .
- Measurement on Microsoft Surface (i5, 2017): 40.96 EFlop/s.

26

## Use case — benchmarking (cont.)

- **Microsoft Surface**, single core: 40.96 EFlop/s.
- **Fugaku supercomputer**, 7.6 million cores: 0.442 EFlop/s.
- **Apparent conclusion:** my Surface is 93x faster than the fastest supercomputer in the world (in Q1/2021).



- **Explanation:** observable behavior of benchmark function is only implicit return ⇒ measured number of FP operations = 0.

```
// add.cpp source file
// returns sum of its arguments:
INT add(INT a, INT b) {
    return a + b;
}
```

```
// g++ -O2:
benchmark():
ret
```

- **Conclusion:** benchmarked operations must contribute to the observable behavior to avoid them being optimized away.

27

## Machine code → source code?

- Can we obtain source code from which some machine code has been generated?

```
void benchmark() {
    double A[32][32] = {};
    double B[32][32] = {};
    double C[32][32] = {};

    for (int n = 0; n < 1'000'000'000; n++)
        for (int i = 0; i < 32; i++)
            for (int j = 0; j < 32; j++)
                for (int k = 0; k < 32; k++)
                    C[i][j] += A[i][k] * B[k][j];
}
```

```
void g() {
    int a = 1;
    int b = 2;
    int c;
    c = a + b;
}
```

```
class Int {
    int i;
public:
    Int(int i) : i(i) {}
    operator int&() { return i; }
};

void h() {
    Int a = 1;
    Int b = 2;
    int c;
    c = a + b;
}
```

- 4 functions with very different source code but with same observable behavior ⇒ possibly same machine code (*live demo*: <https://godbolt.org/z/hvvvYY>).
- **Implication:** generally, it is not possible to "link" machine code back to source code.
- If needed, special mechanisms must be involved (*debugging information*, ...).

28

## Translation units

- Observable behavior is related to translation units.
- **Translation unit** — a piece of source code that is translated into machine code (compiled) at once.
- Basically: a translation unit = source file:
  - without comments
  - and processed by a C++ preprocessor.
- **Example:** source file (*left*), its translation unit with undefined `INT` (*right above*) and `INT` defined as `long` (*right below*):

```
// add.cpp source file
#ifndef INT
#define INT int
#endif

// returns sum of its arguments:
INT add(INT a, INT b) {
    return a + b;
}
```

```
int add(int a, int b) {
    return a + b;
}
```

```
long add(long a, long b) {
    return a + b;
}
```

- **Note:** translation unit may be obtained by some implementation tools, such as `g++ -DINT=long -E -P add.cpp`.

29

## Translation units (cont.)

- Compilers translate each translation unit separately.
- Machine code for a single translation unit is typically stored in so-called object file (`.o` or `.obj` file extension).
- To generate an object file with the machine code translated from a translation unit related to some source file, GCC or Clang uses:
  - `-c` command line option,
  - source file name as a command line argument.
- **Example:** `g++ -O2 -c add.cpp`

```
// add.cpp source file
#ifndef INT
#define INT int
#endif

// returns sum of its arguments:
INT add(INT a, INT b) {
    return a + b;
}
```

```
int add(int a, int b) {
    return a + b;
}
```

translation unit

```
add(int, int):
8d 04 37 c3
```

object file = machine code

```
add(int, int):
lea eax, [rdi+rsi]
ret
```

disassembled object file

30

## GCC/Clang workflow

- How do g++ and clang++ commands work?
- a) *Input* = source file(s) and -c option  $\Rightarrow$  *translation/compilation*:
  - transforms the source file into a translation unit,
  - translates the source code of the translation unit into machine code,
  - stores that machine code in an object file.
  - Example*: g++ -O2 -c main.cpp.
- b) *Input* = object file(s)  $\Rightarrow$  *linking*:
  - merges/links machine code from object file(s),
  - stores it into an executable binary file.
  - Note*: one of the object files must contain machine code of the main function.
  - Example*: g++ main.o (executable file is called a.out by default).
- c) *Input* = source file(s), no -c option  $\Rightarrow$  *translation + linking*:
  - Ad a) and ad b) at once.
  - Note*: an object file is not explicitly stored (it is stored only temporarily, for example in /tmp directory with some random file name).

31

## GCC/Clang workflow (cont.)

- Example* — source code:

```
// main.cpp
int main() {
    return 12;
}
```

- Separate compilation and linking:

```
$ g++ -O2 -c main.cpp
$ g++ main.o
$ ./a.out
$ echo $?
12
```

- Single-command alternative:

```
$ g++ -O2 main.cpp
$ ./a.out
$ echo $?
12
```



- Observable behavior of main*: returns 12.
- C++ standard*: value returned from main = program exit status.
- Bash*: exit status of terminated program — \$? shell parameter.

32

## Translation units (cont.)

- Previous example*: only a single source file  $\Rightarrow$  a single translation unit.
- Source code of programs usually consists of multiple translation units.
- Problems*:
  - Multiple translation units must cooperate at the source code level.
  - Machine code translated from multiple translation units must cooperate at the machine code level.
  - How is observable behavior resolved in the context of multiple translation units?
- Example*:

```
// add.cpp source file
int add(int a, int b) {
    return a + b;
}
```

```
// main.cpp source file ver.1
int main() {
    return add(1, 2) + 3;
}
```

- In one source file, we want to use a function defined in another source file.

33

## Translation units (cont.)

- Translation units for add.cpp and main.cpp, respectively:

```
int add(int a, int b) {
    return a + b;
}
```

```
int main() {
    return add(1, 2) + 3;
}
```

- Compilation of add.cpp is ok:

```
$ g++ -O2 -c add.cpp
```

- Problem*: compilation of main.cpp is not ok:

```
$ g++ -O2 -c main.cpp
error: 'add' was not declared in this scope
```

- Observable behavior of main:

- calling function named add with 1 and 2 integer arguments,
- adding returned value with integer 3,
- returning the result out of main.
- A compiler needs to generate machine code for this observable behavior. Why did it fail?

34

## Translation units (cont.)

- During translation, a compiler always works only with a single translation unit.
- ```
int main() { return add(1, 2) + 3; }
```
- When a compiler translates main.cpp, what does it “see”?
    - add function call,
    - passed arguments (integer literals having values 1 and 2),
    - what is done with the returned value (addition with 3).
  - This is not enough for generating machine code for the observable behavior. Why?
    - For example, are the arguments passed to add by value or by reference?
    - Or, is there any implicit conversion required?
    - Such things can make difference at the machine code level.
  - Implication*: to be able to translate a function call, a compiler must see the type of its parameters as well as its return type.

35

## Function declaration

- Type of function parameters and its return type makes part of its interface (API).
- In C++, function API = **function declaration** (a.k.a. *prototype*).
- Using (= calling) a function in some translation unit requires its declaration to be present in that translation unit as well.
- Note*: function definition is also its declaration.
- Minimal declaration of add function:

```
int add(int, int);
```

- Addition into main.cpp:

```
// main.cpp source file ver.2
int add(int, int);
int main() { return add(1, 2) + 3; }
```

- Compilation is now ok:

- types of parameters are known, return type is known.

```
$ g++ -O2 -c main.cpp
```

36

### Translation units (cont.)

- API (such as declarations for functions) resolves cooperation of translation units at a source code level.
- What about the machine code level?
- Translation unit of `main.cpp`:

```
int add(int, int);
int main() { return add(1, 2) + 3; }
```

- Its compilation with...

```
$ g++ -O2 -c main.cpp
```

- ...results in an object file `main.o` that contains machine code of function `main` that reflects its observable behavior.
- If we try to build an executable program file from it, it fails:

```
$ g++ main.o
undefined reference to `add(int, int)'
```

- Why?

37

### Translation units (cont.)

- The part of the observable behavior of `main` is calling `add`.
- This observable behavior is reflected in its machine code:

```
main:
  sub    rsp, 8
  mov    esi, 2
  mov    edi, 1
  call   add(int, int)
  add    rsp, 8
  add    eax, 3
  ret
```

- To build the executable binary, the linker needs to connect (link) this `call` instruction with the machine code of the `add` function.
- However, the linker cannot find this machine code.
- Machine code of `add` cannot be generated from the translation unit of `main.cpp`  $\Rightarrow$  it is not included in `main.o` object file.
- This problem triggers the "well-known" *undefined reference error*:

```
undefined reference to `add(int, int)'
```

38

### Translation units (cont.)

- The only way the machine code of `add` can be generated is from its definition.
- Function definition** basically = function declaration + its body.
- In our case, this definition is in the translation unit of `add.cpp` (left)  $\Rightarrow$  its compilation therefore generates machine code of `add` function in `add.o` object file (right):

```
int add(int a, int b) {
  return a + b;
}
```

```
add(int, int):
  lea    eax, [rdi+rsi]
  ret
```

- To generate the final executable file, we need to provide both `main.o` and `add.o` object files to the linker:

```
$ g++ main.o add.o
```

39

### ABI

```
$ g++ main.o add.o
```

- The final executable file `a.out` contains machine code of both functions `main` and `add`:

```
int add(int, int);
int main()
{
  return add(1, 2) + 3;
}
```

```
int add(int a, int b)
{
  return a + b;
}
```

- Both pieces of machine code need to cooperate at a machine code level.
- Why? Because of passing arguments from `main` to `add` and the return value from `add` back to `main`.

```
main:
  sub    rsp, 8
  mov    esi, 2
  mov    edi, 1
  call   add(int, int)
  add    rsp, 8
  add    eax, 3
  ret

add(int, int):
  lea    eax, [rdi+rsi]
  ret
```

- The `call` instruction doesn't provide such a functionality (it only stores the current value of the *instruction pointer* on the stack and sets its new value to the address of the machine code of `add`).
- The arguments and the return value must be passed by some explicit mechanism.

40

### ABI (cont.)

- In order to allow different functions to cooperate at a machine code level, there must be some agreement about how to pass arguments and return values to/from functions.
- This agreement (*calling conventions*) is a part of so-called **application binary interface (ABI)**.
- API:
  - enables cooperation of translation units at a source code (*programming*) level;
  - is defined by C++ standards;
  - example*: function declaration allows to compile function calls.
- ABI:
  - enables cooperation of machine code generated from different translation units at a machine code (*binary*) level;
  - is not defined by C++ standards (typically, it is defined by the architecture + operating system + C++ implementation);
  - example*: specification of how to pass arguments to functions.

41

### ABI (cont.)

- In my case — `x86_64` + Linux + GCC — the applied ABI specifies that:
  - `int` is a 32-bit signed integer;
  - argument for the first parameter of type `int` is passed into a function through the `edi` register;
  - argument for the second parameter of type `int` is passed into a function through the `esi` register;
  - return value of type `int` is passed from a function through the `eax` register.

```
int main()
{
  return add(1, 2) + 3;
}
```

```
main:
  sub    rsp, 8
  mov    esi, 2
  mov    edi, 1
  call   add(int, int)
  add    rsp, 8
  add    eax, 3
  ret
```

```
add(int, int):
  lea    eax, [rdi+rsi]
  ret
```

- Note*: the latter holds for `add` as well as for `main`.
- Note*: the ABI requires stack frames to be 16-byte aligned  $\Rightarrow$  `rsp` register manipulations.

42

### ABI (cont.)

- Importance of ABI — it allows linking the machine code generated:
  - by different C++ implementations,
  - at different places (different computers),
  - at different times.

```
$ clang++ -O2 -c add.cpp
$ g++ -O2 -c main.cpp
$ g++ main.o add.o
$ ./a.out
$ echo $?
6
```

- ABI is a must for using libraries.
- **Library:**
  - A collection of object files that implements some functionality bundled in a file with specific format (.a, .so, .lib, .dll,...).
  - Plus API (usually *header files*) that enables to work with that functionality at a source code level (for example, library function declarations).

43