

Effective C++ Programming

NIE-EPC (v. 2021):

TRANSLATION UNITS AND OBSERVABLE BEHAVIOR,
ODR, INTERNAL LINKAGE, INLINE, STATIC AND
DYNAMIC LINKING, PIMPL

© 2021 DANIEL LANGR, ČVUT (CTU) FIT

Translation units + observable behavior

- Is **observable behavior** related to **translation units**?
- **Yes** — observable behavior of some code = its **observable behavior** from the perspective of the translation unit where the code is.
- *Recapitulation:* **translation units** for `add.cpp` and `main.cpp`:
- **Observable behavior** of function `add` from the perspective of its translation unit:
 - returns the sum of its parameters.
- **Observable behavior** of function `main` from the perspective of its translation unit:
 - calls `add` with arguments 1 and 2, adds the returned value with 3, and returns the result.
- The **generated machine code** in object files `add.o` and `main.o` matches to this behavior.
- The **linker** only **merges this machine code** into a **single executable binary file**.

```
int add(int a, int b) {  
    return a + b;  
}
```

```
int add(int, int);  
  
int main() {  
    return add(1, 2) + 3;  
}
```

```
main:  
    sub    rsp, 8  
    mov    esi, 2  
    mov    edi, 1  
    call   add(int, int)  
    add    rsp, 8  
    add    eax, 3  
    ret
```

```
add(int, int):  
    lea    eax, [rdi+rsi]  
    ret
```

Translation units + observable behavior (*cont.*)

- *Alternative option:* what if we put **all the code into a single translation unit**? Source file (*left*), translation unit (*right*):

```
// main.cpp source file ver.3
int add(int a, int b) {
    return a + b;
}
int main() {
    return add(1, 2) + 3;
}
```

```
int add(int a, int b) {
    return a + b;
}
int main() {
    return add(1, 2) + 3;
}
```

- **Does anything change?**
- Observable behavior of `main` **with respect to this translation unit** — **returning the value 6.**
- **Generated machine code** with `g++ -O2`:

```
main:
    mov eax, 6
    ret
```

Translation units + observable behavior (cont.)

- In both cases, the **very same source code**:

```
// main.cpp source file ver.2
int add(int, int);
int main() { return add(1, 2) + 3; }
```

```
// add.cpp
int add(int a, int b)
{ return a + b; }
```



```
// main.cpp source file ver.3
int add(int a, int b) {
    return a + b;
}
int main() {
    return add(1, 2) + 3;
}
```

- In both cases, compiled **with optimizations**:

```
$ g++ -O2 -c add.cpp
$ g++ -O2 -c main.cpp
$ g++ main.o add.o
```

```
$ g++ -O2 -c main.cpp
$ g++ main.o
```

- In both cases, the **same program runtime behavior**:

```
$ ./a.out
$ echo $?
6
```

```
$ ./a.out
$ echo $?
6
```

- In both cases, the **very different runtime (and memory) efficiency**:

```
main:
    sub    rsp, 8
    mov    esi, 2
    mov    edi, 1
    call   add(int, int)
    add    rsp, 8
    add    eax, 3
    ret
add(int, int):
    lea    eax, [rdi+rsi]
    ret
```



```
main:
    mov    eax, 6
    ret
```



Translation units and efficiency



- Should we put definitions of all the functions called in some translation unit into that translation unit? **NO!**
- However, it may be **profitable** for “short” functions (w.r.t. their runtime) that are **called frequently** (at runtime).
- *Example: fused vector multiply-add (FMADD) operation defined in terms of scalar FMADD:*

```
void scalar_fmadd( float& x, float y, float alpha ) {  
    x += y * alpha;  
}
```

```
void vector_fmadd( std::vector<float>& x, const std::vector<float>& y, float alpha ) {  
    for (size_t i = 0; i < x.size(); i++)  
        scalar_fmadd(x[i], y[i], alpha);  
}
```

- *Experiment:* scalar_fmadd defined in
(1) the same / (2) a different translation unit as/than vector_fmadd.
- *Observation:* in (1) the operation was **2.2x faster** than in (2).
- *Setup:* GCC, **enabled optimizations**, Intel Core i9.

Matrix benchmark revisited

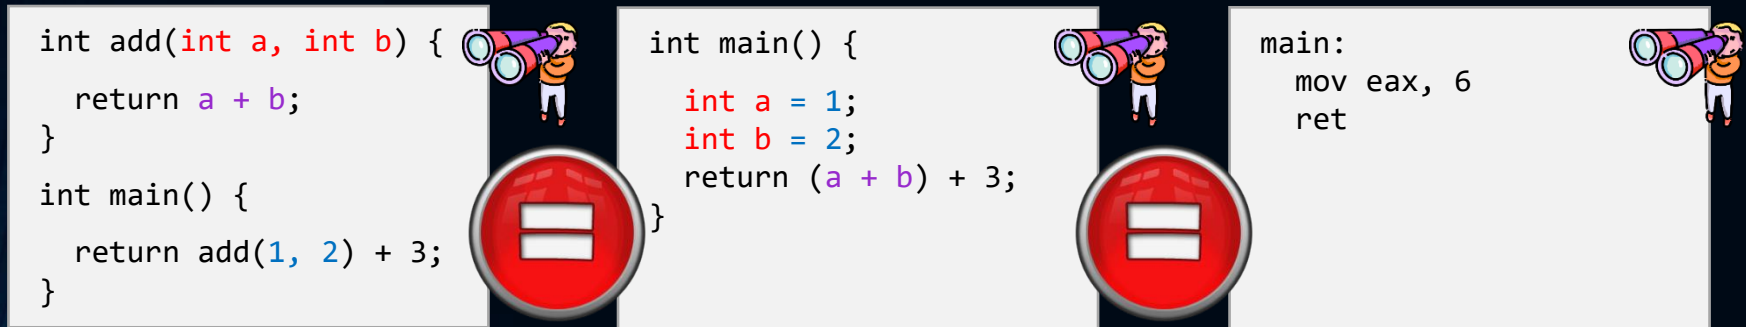
- *Original experiment:* the benchmark function was **defined in the same translation unit where it was called**.
- Putting benchmark function into a **separate translation unit**...

```
// benchmark.cpp  
  
void benchmark(const double A[][32], const double B[][32], double C[][32]) {  
    for (int n = 0; n < 1'000'000; n++)  
        for (int i = 0; i < 32; i++)  
            for (int j = 0; j < 32; j++)  
                for (int k = 0; k < 32; k++)  
                    C[i][j] += A[i][k] * B[k][j];  
}
```

- ...and passing arrays (matrices) as arguments **assure that the floating-point operations will be performed at runtime**.
- Their effect is **observable from outside of this translation unit** (by reading elements of the array passed as an argument for C).
- *Measured performance on Surface:* **0.79 GFlop/s**.
- *Fugaku:* **0.442 EFlop/s \Rightarrow 559 million times more capable**.

Function *inlining*

- Resolution of observable behavior **across function boundaries** is an optimization technique called *inlining*.
- Inlining** treats a function as a “macro” that is **expanded at the place where the function is called**.



- Some implementations allow to **selectively control inlining**.
- Disabling inlining** restricts the resolution of observable behavior across function boundaries \Rightarrow **each function call statement** generates a corresponding **call** instruction.
- Example:* compilation with **g++ -O2 -fno-inline** \rightarrow
- Note:* **no relation** with the **inline** keyword.
- Experiment:* `std::sort` **2.7x slower** with **disabled inlining**.

```
main:  
    sub    rsp, 8  
    mov    esi, 2  
    mov    edi, 1  
    call   add(int, int)  
    add    rsp, 8  
    add    eax, 3  
    ret  
  
add(int, int):  
    lea    eax, [rdi+rsi]  
    ret
```

Translation units vs source code files

- If a function **needs to be called in some translation unit**, its **declaration must be in the same translation unit** as well (*or, its definition, which is also a declaration*).
- *Implication?* Does it mean that we need to **put this declaration into all source code files**, where the function is called? **No**.
- *Recall:* **translation unit = preprocessed source code file**.
- One of **preprocessor abilities**: **inclusion of some source code file into another source code file**.
- *Preprocessor directive: #include:*
 - This directive **replaces itself with the content of the included file**.
 - It does that **recursively** (included file may also contain #include directives).
- *Example:* 2 source code files + **translation unit for main.cpp (right)**:

```
// add.cpp
```

```
int add(int a, int b) {  
    return a + b;  
}
```

```
// main.cpp
```

```
#include "add.cpp"  
  
int main() {  
    return add(1, 2) + 3;  
}
```

```
int add(int a, int b) {  
    return a + b;  
}  
  
int main() {  
    return add(1, 2) + 3;  
}
```


Translation units vs source code files (cont.)

- *Convention:*
 - Files to be included into other files are typically called „**header files**“ (shortly **headers**) and have file extensions as `.h`, `.hpp`, `.hxx`, `.hh`, *no one*.
 - Files to be transformed into translation units are typically called „**source files**“ and have file extensions as `.cpp`, `.cxx`, `.cc`, etc.
- This is just a **convention**; compilers and preprocessors **do not care** about how we call files or their extensions.
- Typical **intent of headers**: **API**.

```
// add.h
int add(int a, int b);
```

```
// main.cpp
#include "add.h"

int main() {
    return add(1, 2) + 3;
}
```

```
int add(int a, int b);
int main() {
    return add(1, 2) + 3;
}
```

```
// add.cpp
int add(int a, int b) {
    return a + b;
}
```

```
int add(int a, int b) {
    return a + b;
}
```

Translation units vs source code files (*cont.*)

- Some “piece of functionality” — **interface** (API) + **implementation**.
- *Ideal world:*
 - **API** in a **header file** (header files),
 - **implementation** in a **source file** (source files).
- Advantage of such **separation**:
 - With **dynamically-linked libraries** (.so, .dll), **changing implementation details does not require recompilation of programs' source code**.
- *Reality:*
 - **Implementation details** are frequently in **header files**.
- *Examples:*
 - **Inlining** — function bodies (implementation) must be visible to enable their inlining.
 - **Class non-public members** — must be at least **declared** at the place where the class is defined.
 - **Templates** — definitions must be visible to enable their instantiation.
 - **Others** — functions with return type deduction, etc.

ODR

- *Example:* **function definition in a header file** to enable its inlining:

```
// add.h  
int add(int a, int b) { return a + b; }
```

- *Example (cont.):* **multiple source files that call this function:**

```
// sub.cpp  
#include "add.h"  
int sub(int a, int b) {  
    return add(a, -b);  
}
```

```
// main.cpp  
#include "add.h"  
int main() {  
    return add(1, 2) + 3;  
}
```

- *Example (cont.):* **compilation and linkage...**

```
$ g++ -O2 -c sub.cpp  
$ g++ -O2 -c main.cpp  
$ g++ main.o sub.o
```

- ...resulted in the **following linker error:**

```
multiple definition of `add(int, int)'
```



ODR (cont.)

- **Translation units** for `sub.cpp` (left) and `main.cpp` (right):

```
int add(int a, int b) {  
    return a + b;  
}
```

```
int sub(int a, int b) {  
    return add(a, -b);  
}
```



```
int add(int a, int b) {  
    return a + b;  
}
```

```
int main() {  
    return add(1, 2) + 3;  
}
```

- *Problem:*
 - C++ rule: an **entity** (function, variable,...) **generally cannot be defined in multiple translation units**.
 - It can be defined **only once** — “one definition rule” (ODR).
 - *Exceptions/workarounds:* static, anonymous namespace, inline, templates.
- *Simplified explanation:*
 - Function `add` is **defined in some translation unit** and may be **used in other translation units** as well → it has so-called “**external linkage**”.
 - Linker “sees” **multiple machine code of `add`** — in `sub.o` and `main.o`.
 - It cannot decide **to which one it should link the `add` calls**.

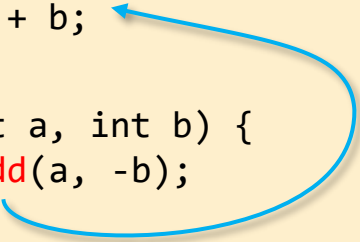
Specifier `static`

- How to enable to **define a function inside a header file** (for example, for inlining purposes) **without breaking ODR?**
- *First option: making this function static.*

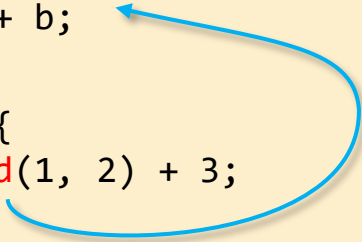
```
// add.h  
static int add(int a, int b) { return a + b; }
```

- **Translation units:**

```
static int add(int a, int b) {  
    return a + b;  
}  
  
int sub(int a, int b) {  
    return add(a, -b);  
}
```



```
static int add(int a, int b) {  
    return a + b;  
}  
  
int main() {  
    return add(1, 2) + 3;  
}
```



- Functions `add` now have “**internal linkage**” — the **definition** in one translation unit is “**private**” to this unit \Rightarrow no ambiguity for a linker.
- These **two definitions** define **two different functions** (separate entities) **regardless of their equivalent forms** \Rightarrow **no violation of ODR.**
- **Without `static`** they would define a **single function** (entity).

Anonymous namespaces

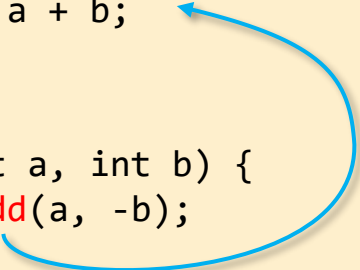
- *Alternative option:* putting add into **anonymous namespace**:

```
// add.h
namespace {
    int add(int a, int b) { return a + b; }
}
```

- **Translation units:**

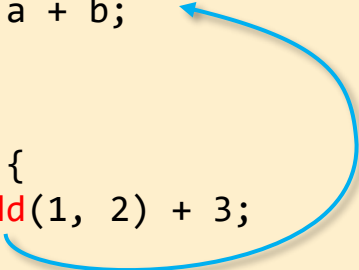
```
namespace {
    int add(int a, int b) {
        return a + b;
    }
}

int sub(int a, int b) {
    return add(a, -b);
}
```



```
namespace {
    int add(int a, int b) {
        return a + b;
    }
}

int main() {
    return add(1, 2) + 3;
}
```



- *Our case:* **the same effect** — add now have **internal linkage**.
- **Anonymous namespace** vs **static**:
 - Generally, **anonymous namespaces** are more “powerful”.
 - *For example*, we can put **type definitions** into anonymous namespace but cannot make types static.
 - Also, static has **multiple meanings in different contexts**, which might be **confusing**.

Confusion with static: example

- Example of **confusedness of static**:

```
// X.h
class X {
    static void f();
};
void X::f() { }
```


```
// main.cpp
#include "X.h"
int main() { }
```

```
// other.cpp
#include "X.h"
```

```
$ g++ -O2 -c other.cpp
$ g++ -O2 -c main.cpp
$ g++ main.o other.o
```

multiple definition of `X::f()'

- **Static member function** `f` has **external linkage**
⇒ **violation of ODR**.
- Here, `static` indicates that `f` is related to class `X`, not to its instances.



Specifier `inline`

- Another option: `inline` functions.

```
// add.h  
inline int add(int a, int b) { return a + b; }
```

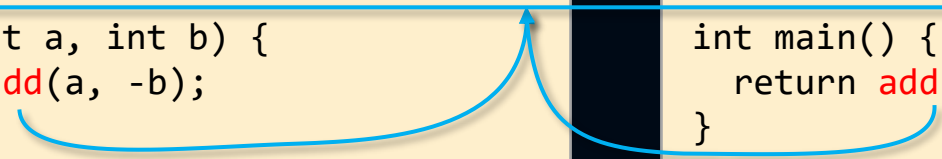
- Translation units:

```
inline int add(int a, int b) {  
    return a + b;  
}
```

```
int sub(int a, int b) {  
    return add(a, -b);  
}
```

```
inline int add(int a, int b) {  
    return a + b;  
}
```

```
int main() {  
    return add(1, 2) + 3;  
}
```



- Both definitions of `add` define a single function (entity).
- With `inline`, they do not break ODR — exception.
- Generally, definitions of the same `inline` definition in multiple translation units define a single entity \Rightarrow all these definitions must be identical.
- An inline function must be defined in each translation unit where it is used (called).

Internal linkage vs inline

1. **Internal linkage** (static, anonymous namespace):
 - Entities “private” for a translation unit.
 - May have different definitions in different translation units.
 - Usually defined in **source files**.
 - *Typical use case:* auxiliary entities private for some source file for which we don't want to care about breaking ODR.
2. **inline**
 - Reusable entities \Rightarrow usually defined in **header files**.
 - *Typical use case:* functions for which we want to enable their inlining.
 - *Note:* inline itself does not enable or even force inlining; it just helps avoiding ODR violation.

Member functions and inline

- Class **non-static member functions** are:
 - **implicitly inline** if they are defined **within the class definition** (body),
 - **implicitly not-inline** if they are defined **outside of the class definition**.
- *Example:*

```
// Y.h ver. 1
class Y {
    void f() { } // implicit inline
};
```

```
// main.cpp
#include "Y.h"
int main() { }
```

```
$ g++ -O2 -c other.cpp
$ g++ -O2 -c main.cpp
$ g++ main.o other.o
```



```
// Y.h ver. 2
class Y {
    void f();
};
void X::f() { } // not inline
```

```
// other.cpp
#include "Y.h"
```

multiple definition of `Y::f()'



Libraries

- **Library** — reusable coded functionality.
- *Example:* Library that contains a **single function** for adding two unsigned integer numbers:

```
unsigned add(unsigned a, unsigned b) {  
    return a + b;  
}
```

- **In which form** can this library be implemented?
- **First option** — putting the function(ality) **into a library header file** \Rightarrow “**header-only library**”:

```
// unsigned_add.h  
#ifndef UNSIGNED_ADD_H  
#define UNSIGNED_ADD_H  
inline unsigned add(unsigned a, unsigned b) { return a + b; }  
#endif
```

Header-only libraries



- **Advantages:**
 - Functions may be inlined \Rightarrow potentially **higher performance**.
 - Using of this library **involves only inclusion of the header file(s)**.
 - **No linking** is involved.
 - **Library users** (developers of programs built upon this library) need only the header file(s) \Rightarrow **no machine code (object files)** are required.
 - **Programs users** **do not need anything regarding the library** (do not need to know about its existence at all).
- **Disadvantages:**
 - Library **source code** is **"copy-pasted"** into the programs source code.
 - \Rightarrow **Machine code** related to the library functionality **is in the program binary file**.
 - \Rightarrow **When the implementation of the library functionality is updated** and we want to get this updated functionality into programs, these **need to be rebuilt** (recompiled and relinked).

Header-only libraries (*cont.*)

- Real-world scenario:
 1. Someone **develops a library**; role — *library developer*.
 2. Someone else **develops a program** by using this library; role — *program developer*.
 3. The program is **installed on computer systems and run**; role — *program user*.
 4. *Library developer* **updates the library code** (*bug fixes, improved performance,...*).
 5. *Program users* want this **updated library functionality in their programs**.
- Library form — **header only**:
 - \Rightarrow Program **needs to be rebuilt**.
 - If the **program source code is available**, this can be done by *program users*.
 - Otherwise, it **needs to be done by program developer**.

Binary libraries

- Another option is to **split the interface and implementation details between header and source files**:

```
// unsigned_add.h  
#ifndef UNSIGNED_ADD_H  
#define UNSIGNED_ADD_H  
unsigned add(unsigned, unsigned); // declaration only  
#endif
```

```
// unsigned_add.cpp  
unsigned add(unsigned a, unsigned b) { return a + b; } // definition
```

- **Developers of programs** using this library now **do not need the library source code file** which implements its functionality.
- Instead, they need only:
 - **the header file** (to be able to compile the program code that calls the library function),
 - the **object file with the machine code** of the function (to be able to link the final executable program together).

Binary libraries (*cont.*)

- The **machine code** of the program contains **call instructions** for the library function.
- These calls **needs to be linked with the machine code of the function** — the library machine code.
- **When this linking happens?**
- There are **two options**:
 1. The linking happens **when the program binary executable file is built**.
 2. The linking happens **not until the program is executed**.
- **Ad 1.** This approach is called **“static linking”** and the libraries used this way are called **“statically linked libraries”**.
- **Ad 2.** This approach is called **“dynamic linking”** and the libraries used this way are called **“dynamically linked libraries”**.
- In both cases, **library machine code** (object files) is **“packed/archived”** into **special file formats**:
 - Ad 1. **.a or .lib** extensions (Windows / Linux),
 - Ad 2. **.dll / .so** extensions (ditto).

Statically linked libraries

- Linking happens **when the program is built** such that both:
 - the **machine code of the program**,
 - the **machine code of the library**,
- is together **“copied” into the final program executable file.**
- \Rightarrow The **library machine code is “hard-wired” into the program.**
- **Advantages:**
 - **Program users do not need to care about the library at all**; they do not need to know about its existence.
- **Disadvantages:**
 - Reflection of **library functionality updates** into programs still **requires their rebuilding** (namely, only relinking is required without program source code recompilation).
 - **Program users** still need to **update program binary files** to reflect the library updates.

Statically linked libraries (*cont.*)

- *Example:* building a static version of the libunsigned library:

```
// unsigned_add.cpp
```

```
unsigned add(unsigned a, unsigned b) { return a + b; }
```

```
$ g++ -O2 -c unsigned_add.cpp
```

```
$ ar rcs libunsigned.a unsigned_add.o
```

- The **ar tool** “archives” object files **into a single library file**.
- **Program developers** now need the **library header file for program compilation...**

```
// main.cpp
```

```
#include <iostream>
```

```
#include <unsigned_add.h>
```

```
int main() { std::cout << add(1u, 2u) << std::endl; }
```

- ...and the **library binary file for program linking:**

```
$ g++ -O2 -c -I"/path/to/unsigned_add.h" main.cpp
```

```
$ g++ -O2 main.o -L"path/to/libunsigned.a" -lunsigned
```

- **Program users do not need that library file** to run the program.

Dynamically linked libraries

- Linking happens when the program is executed.
- The system where the program is run needs to find the library machine code and link it with the library function call instructions in the program.
- Disadvantages:
 - Program users need to install binary library files (.dll, .so).
- Advantages:
 - When the library source code is updated and the new machine code is redistributed to systems (for example, under system packages updates), then programs immediately reflect the library updates.
 - No program change is required; the program binary executable files remain the very same (no recompilation, no relinking).

Dynamically linked libraries (*cont.*)

- *Example:*

```
// unsigned_add.cpp
unsigned add(unsigned a, unsigned b) { return a + b; }
```

```
$ g++ -O2 -c -fPIC unsigned_add.cpp
$ g++ -shared -o libunsigned.so unsigned_add.o
```

- The **library binary file** is “archived” by GCC.
- **Program developers** need the **library header file** for program compilation...

```
// main.cpp
#include <iostream>
#include <unsigned_add.h>
int main() { std::cout << add(0xFFFFFFFFu, 1u) << std::endl; }
```

- ...and the **library binary file** as well to “register” linking:

```
$ g++ -O2 -c -I"/path/to/unsigned_add.h" main.cpp
$ g++ -O2 main.o -L"path/to/libunsigned.so" -lunsigned
```

Dynamically linked libraries (cont.)

- The **library** binary file **libunsigned.so** is now needed by **users** of the **program** when this gets executed:

```
$ ./a.out
./a.out: error while loading shared libraries: libunsigned.so: cannot open shared object
file: No such file or directory
$ export LD_LIBRARY_PATH="path/to/libunsigned.so:$LD_LIBRARY_PATH"
$ ./a.out
0
```

- Now, **changing the implementation of the library...**

```
// unsigned_add.cpp
#include <iostream>
#include <limits>

unsigned add(unsigned a, unsigned b) {
    if (std::numeric_limits<unsigned>::max() - a < b)
        std::cerr << "WARNING: unsigned overflow occurred" << std::endl;
    return a + b;
}
```

- ...and **rebuilding and redistribution of new libunsigned.so** is **automatically reflected** in the program:

```
$ ./a.out
WARNING: unsigned overflow occurred
0
```


Library updates

- Does any update of library code need rebuilding of programs with dynamically-linked libraries?
- 1. **Updates of library interface:**
 - Machine code of programs that interacts with the library machine code is generated based on the code in library headers.
 - \Rightarrow Changes in these headers usually require program machine code regeneration.
- *Examples:*
 - *Compatible changes:* adding new entities (functions, global variables,...).
 - *Incompatible changes:* changing numbers and types of function parameters, types of returned values, names of entities, removing entities,...
- \Rightarrow It is important to design a library interface in such a way that it would require minimum changes in the future.
- *Note:* Sometimes, it's very hard, since developers do not know which way their libraries will evolve.

Library updates (*cont.*)

2. Updates of library implementation:

- *Ideal world:*

- **Header files** — functionality **interface** (what does library do).
- **Source files/machine code** — **implementation** (how does library do that).
- \Rightarrow Changes in the implementation do not require rebuilding of programs.

- *Reality:*

- **Implementation details** sometimes **need to be in header files**.
- Their **updates may break binary compatibility** between program and library machine code.
- \Rightarrow **Program rebuilds are required**.

Library updates (*cont.*)

- *Example: library code:*

```
// X.h
class X {
    long i_;
public:
    X();
};
```

```
// X.cpp
#include "X.h"
X::X() : i_(0x12L) { }
```

- *Program code (left) and its translation unit (right):*

```
// main.cpp
#include <X.h>
int main() {
    X x;
}
```

```
class X {
    long i_;
public:
    X();
};
int main() {
    X x;
}
```

- *Observable behavior of main:*
 - calls **default constructor of X**,
 - calls **destructor of X**,
 - **returns 0** to the caller (implicit return for main).

C++ object model

- **Object** in C++ = **instance of some type** (not necessarily class).
- **Constructor** call is a part of **class object initialization**.
- To **initialize an object** of type T, a **storage** for its **binary representation** must be provided first.
- This storage must satisfy **two requirements**:
 - It must be **large enough** — `sizeof(T)` bytes.
 - It must be **properly aligned** — address divisible by `alignof(T)`.
- *Example:*

```
std::cout<< sizeof(long) << alignof(long); // printed '88' on x86_64/Linux
```

- *Observation:*
 - A storage for **any object of type long** must **have 8 bytes** and be **aligned to an address divisible by 8** on this system.
 - *Note:* It's **prescribed by its ABI** \Rightarrow it holds for any compiler.

C++ object model (*cont.*)

- Binary representation of class objects consists of binary representation of its subobjects.
- Class subobjects:
 - base class objects,
 - non-static member variables.
- Binary representation of class objects is generally implementation-defined.
- Class X has only one subobject — member variable of type long.
- \Rightarrow In our case, the binary representation of X objects was the same as the binary representation of long objects.

```
class X {  
    long i_; // member variable  
public:  
    X();  
};
```

```
std::cout << sizeof(X) << alignof(X); // printed '88' on x86_64/Linux/Clang
```

C++ object model (*cont.*)

- **Non-static member function:**
 - A function that is called “on an object of some class”.
 - This function needs a “reference” to the object on which it is called.
 - *Note:* This reference is available inside the function body in the form of the pointer **this**.
 - This reference represents a hidden parameter of all non-static member functions.
 - Its form is implementation-dependent.
 - Typical implementations pass an address of the object as argument for this parameter.
- **Constructor** = (special) member function:
 - Object does not yet exist; the purpose of constructors is to create it.
 - \Rightarrow The caller passes the address of the storage where the object should be initialized \Rightarrow this storage must be prepared by the caller.
 - Constructor then initializes a binary representation of the object in this storage.

C++ object model (*cont.*)


- Translation unit (*left*), machine code x86_64/Linux/Clang (*right*):

```
class X {  
    long i_;  
public:  
    X();  
};  
  
int main() {  
    X x;  
}
```

```
main:  
    push    rax  
    mov     rdi, rsp  
    call    X::X()  
    xor     eax, eax  
    pop     rax  
    ret
```

- In main, x is a **non-static local variable**.
 - \Rightarrow x is a **new object** created in each main call.
- main needs to:
 - **prepare a storage for x**,
 - **call a default constructor of X**
 - **while passing the address of the storage to the constructor.**
- *Observation:*
 - Storage for x is **allocated on the stack** (at address rsp decreased by 8).
 - Its address is to constructor **passed in rdi register** (as **prescribed by ABI**).

C++ object model (*cont.*)

- *Other considerations:*
- **Definition of the constructor** is not in the program translation unit.
 - \Rightarrow Its **call cannot be "inlined"**; it must be called explicitly at the machine code level.
- Variable `x` is an object with so-called **"automatic storage duration"**:
 - Its lifetime ends **when the function is leaved**.
 - \Rightarrow Its **destructor must be called**.
- **Destructor of `X`:**
 - There is **no custom declaration/definition**.
 - \Rightarrow The destructor is **automatically defined**.
 - \Rightarrow As if it was defined (*simplification*) **this way** 
- \Rightarrow Destruction of `x` **may be inlined** and has **no observable behavior**.
 - \Rightarrow It **does not generate any machine code** under optimizations.

```
class X {  
    long i_;  
  
public:  
    X();  
    ~X() { }  
};
```

C++ object model (*cont.*)

- Translation unit for the library source file / machine code:

```
class X {  
    long i_;  
public:  
    X();  
};  
X::X() : i_(0x12L) { }
```

```
X::X()  
    mov    qword ptr [rdi], 0x12  
    ret
```

- *Recall:* the address of the storage for the constructed object is passed in `rdi`.
- Constructor initializes the binary representation of the constructed object in this storage.
- Binary representation of an object of `X` consists of the binary representation of its member variable `i_`.
- This member variable is initialized with hex value `12`.
- \Rightarrow This value needs to be stored to the address pointed to by `rdi`.

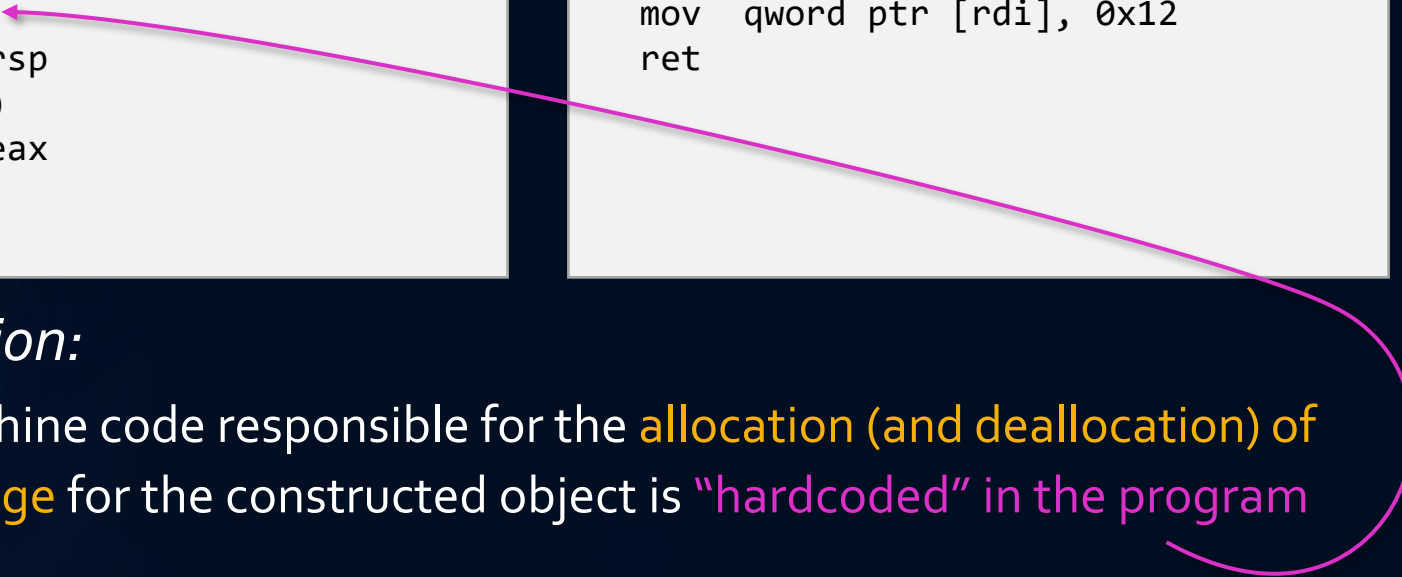
Library updates (cont.)

- **Summary:**

- *Left:* machine code stored in the **program binary file**.
- *Right:* machine code stored in the **library binary file (.so)**.

```
main:
  push  rax
  mov    rdi, rsp
  call   X::X()
  xor     eax, eax
  pop     rcx
  ret
```

```
X::X()
  mov    qword ptr [rdi], 0x12
  ret
```



- **Observation:**

- The machine code responsible for the **allocation (and deallocation) of the storage** for the constructed object is **"hardcoded"** in the **program file**.
- \Rightarrow Any **change of the library** that requires a **different way of storage allocation** will **break the binary compatibility**.

Library updates (*cont.*)

- Example of **binary incompatible library update** — adding another member variable:

```
// X.h
class X {
    long i_, j_;
public:
    X();
};
```

```
// X.cpp
#include "X.h"
X::X() : i_(0x12L), j_(0x34L) { }
```

- Machine code of program (*left*) and **new library** (*right*):

```
main:
    push    rax
    mov     rdi, rsp
    call    X::X()
    xor     eax, eax
    pop     rcx
    ret
```

```
X::X()
    mov     qword ptr [rdi], 0x12
    mov     qword ptr [rdi+8], 0x34
    ret
```

- Problem:**

- Program **allocates 8 bytes on the stack**, but constructor internally **overwrites 16 bytes!**

Library updates (*cont.*)

```
main:  
    push    rax  
    mov     rdi, rsp  
    call    X::X()  
    xor     eax, eax  
    pop     rcx  
    ret
```

```
X::X()  
    mov     qword ptr [rdi], 0x12  
    mov     qword ptr [rdi+8], 0x34  
    ret
```

- *Consequence:*
 - A constructor overwrites some **part of the stack that does not belong to the constructed object**.
 - In our case, it **overwrites the return address** where the ret instruction “jumps” when the constructor is finished.
 - In *best case*, the program just **crashes**.
 - In *worst case*, it starts **executing some (un)predicted machine code**.
 - ⇒ Such a **vulnerability** might be generally **exploited**.

Library updates (*cont.*)

- *Recapitulation:*
 - **Original** library code (*black*), **updates** (with *red*):

```
// X.h
class X {
    long i_, j_;
public:
    X();
};
```

```
// X.cpp
#include "X.h"
X::X() : i_(0x12L), j_(0x34L) { }
```

- Updates **logically** involved only **implementation library details**:
 - Whatever is **private to a class** is **logically related to implementation** of its functionality.
- However, even private members **must be declared within class definition** (body).
 - \Rightarrow These declarations **must appear in header files**.
- C++ generally **does not allow a complete logical separation of interface and implementation details** between header and source files, respectively.
- *Workaround*: **PIMPL idiom**.

PRIVATE

PIMPL idiom

- *PIMPL* = *Pointer-to-IMPLementation*.
- Technique for “true” hiding of all implementation details out of header files.
- *Functionality*:
 - Original class is “split” into two classes.
 - *First* — “implementation” class — contains the original class implementation details including all its member variables.
 - *Second* — “interface” class — contains:
 - 1) a pointer to the implementation class object,
 - 2) interface of the original class = its public member functions.
 - Public member functions of the interface class need to invoke member functions of the implementation class.
 - This is done internally inside library source files.
- *Outcome*:
 - Interface class does not need to be changed when implementation class member variables and member functions are updated in any way.
 - ⇒ If the “true interface” — public member functions — is not changed, there is no need for programs rebuilding with dynamically linked libraries.

PIMPL idiom (cont.)

- *Example* — original class:

```
// X.h
class X {
    long i_;
public:
    X();
    void do_something();
};
```

```
// X.cpp
#include <iostream>
#include "X.h"

X::X() : i_(0x12L) { }

X::do_something() { std::cout << i_; }
```

- *"PIMPLed" solution:*

```
// X.h
class X { // interface class
    class Impl;
    // -> implementation...
    // ...class declaration

    Impl* pimpl_;
    // -> pointer...
    // ...to implementation

public:
    X();
    ~X();
    void do_something();
};
```

```
// X.cpp
#include <iostream>
#include "X.h"

class X::Impl { // implementation class
    long i_;
public:
    Impl::Impl() : i_(0x12L) { }
    void do_something() { std::cout << i_; }
};

// interface-implementation binding:
X::X() : pimpl_( new Impl{} ) { }
X::~~X() { delete pimpl_; }
X::do_something() { pimpl_->do_something(); }
```

PIMPL idiom (cont.)

- *Library update:*

```
// X.h
class X {
    class Impl;
    Impl* pimpl_;
public:
    X();
    ~X();
    void do_something();
};
```

```
// X.cpp
#include <iostream>
#include "X.h"

class X::Impl {
    long i_, j_;
public:
    Impl::Impl() : i_(0x12L), j_(0x34L) { }
    void do_something() {
        std::cout << i_ << " " << j_;
    }
};

X::X() : pimpl_( new Impl{} ) { }
X::~X() { delete pimpl_; }
X::do_something() { pimpl_->do_something(); }
```

- *Observation:*

- **No change in the header file**
⇒ no need for program rebuild
(with dynamic library linking).

- *Notes:*

- In X.h, class X::impl is only declared ⇒ it is so-called “*incomplete type*”.
- It is **legal** to declare a **pointer-to-incomplete type** (pimpl_ in our case).
- All the operations that **require this type to be complete** (such as new and delete expressions) **are in the source file**.
 - The class gets complete **after its definition**.