

Effective C++ Programming

NIE-EPC (v. 2021):
TRANSLATION UNITS AND OBSERVABLE BEHAVIOR,
ODR, INTERNAL LINKAGE, INLINE, STATIC AND
DYNAMIC LINKING, PIMPL
© 2021 DANIEL LANGR, ČVUT (CTU) FIT

1

Translation units + observable behavior

- Is observable behavior related to translation units?
- Yes** — observable behavior of some code = its observable behavior from the perspective of the translation unit where the code is.
- Recapitulation:** translation units for `add.cpp` and `main.cpp`:
- Observable behavior of function `add` from the perspective of its translation unit:
 - returns the sum of its parameters.
- Observable behavior of function `main` from the perspective of its translation unit:
 - calls `add` with arguments 1 and 2, adds the returned value with 3, and returns the result.
- The generated machine code in object files `add.o` and `main.o` matches to this behavior.
- The linker only merges this machine code into a single executable binary file.

```
int add(int a, int b) {
    return a + b;
}

int add(int, int);
int main() {
    return add(1, 2) + 3;
}

main:
    sub    rsp, 8
    mov    esi, 2
    mov    edi, 1
    call   add(int, int)
    add    rsp, 8
    add    eax, 3
    ret

add(int, int):
    lea    eax, [rdi+rsi]
    ret
```

2

Translation units + observable behavior (cont.)

- Alternative option:** what if we put all the code into a single translation unit? Source file (left), translation unit (right):

```
// main.cpp source file ver.3
int add(int a, int b) {
    return a + b;
}
int main() {
    return add(1, 2) + 3;
}
```

```
int add(int a, int b) {
    return a + b;
}
int main() {
    return add(1, 2) + 3;
}
```

- Does anything change?
- Observable behavior of `main` with respect to this translation unit — returning the value 6.
- Generated machine code with `g++ -O2`:

```
main:
    mov    eax, 6
    ret
```

3

Translation units + observable behavior (cont.)

- In both cases, the very same source code:

```
// main.cpp source file ver.2
int add(int, int);
int main() { return add(1, 2) + 3; }
```

```
// add.cpp
int add(int a, int b)
{ return a + b; }
```

```
// main.cpp source file ver.3
int add(int a, int b) {
    return a + b;
}
int main() {
    return add(1, 2) + 3;
}
```

- In both cases, compiled with optimizations:

```
$ g++ -O2 -c add.cpp
$ g++ -O2 -c main.cpp
$ g++ main.o add.o
```

```
$ g++ -O2 -c main.cpp
$ g++ main.o
```

- In both cases, the same program runtime behavior:

```
$ ./a.out
$ echo $?
6
```

```
$ ./a.out
$ echo $?
6
```

- In both cases, the very different runtime (and memory) efficiency:

```
main:
    sub    rsp, 8
    mov    edi, 2
    mov    esi, 1
    call   add(int, int)
    add    rsp, 8
    add    eax, 3
    ret
add(int, int):
    lea    eax, [rdi+rsi]
    ret
```

```
main:
    mov    eax, 6
    ret
```

4

Translation units and efficiency

- Should we put definitions of all the functions called in some translation unit into that translation unit? **NO!**
- However, it may be profitable for "short" functions (w.r.t. their runtime) that are called frequently (at runtime).
- Example: fused vector multiply-add (FMADD) operation defined in terms of scalar FMADD:**

```
void scalar_fmadd( float& x, float y, float alpha ) {
    x += y * alpha;
}
```

```
void vector_fmadd( std::vector<float>& x, const std::vector<float>& y, float alpha ) {
    for (size_t i = 0; i < x.size(); i++)
        scalar_fmadd(x[i], y[i], alpha);
}
```

- Experiment:** `scalar_fmadd` defined in (1) the same / (2) a different translation unit as/than `vector_fmadd`.
- Observation:** in (1) the operation was **2.2x faster** than in (2).
- Setup:** GCC, enabled optimizations, Intel Core i9.

5

Matrix benchmark revisited

- Original experiment:** the benchmark function was defined in the same translation unit where it was called.
- Putting benchmark function into a separate translation unit...


```
// benchmark.cpp
void benchmark(const double A[1][32], const double B[1][32], double C[1][32]) {
    for (int n = 0; n < 1'000'000; n++)
        for (int i = 0; i < 32; i++)
            for (int j = 0; j < 32; j++)
                for (int k = 0; k < 32; k++)
                    C[i][j] += A[i][k] * B[k][j];
}
```

- ...and passing arrays (matrices) as arguments assure that the floating-point operations will be performed at runtime.
- Their effect is observable from outside of this translation unit (by reading elements of the array passed as an argument for C).
- Measured performance on Surface:** 0.79 GFlop/s.
- Fugaku:** 0.442 EFlop/s \Rightarrow 559 million times more capable.

6

Function inlining

- Resolution of observable behavior across function boundaries is an optimization technique called *inlining*.
- Inlining** treats a function as a "macro" that is expanded at the place where the function is called.



- Some implementations allow to selectively control inlining.
- Disabling inlining restricts the resolution of observable behavior across function boundaries \Rightarrow each function call statement generates a corresponding call instruction.
- Example:** compilation with `g++ -O2 -fno-inline` \rightarrow
- Note:** no relation with the `inline` keyword.
- Experiment:** `std::sort` 2.7x slower with disabled inlining.

```

main:
    mov     rax, 6
    call    add@PLT
    ret
  
```

7

Translation units vs source code files

- If a function needs to be called in some translation unit, its declaration must be in the **same** translation unit as well (or, its *definition, which is also a declaration*).
- Implication?** Does it mean that we need to put this declaration into all **source code files**, where the function is called? **No**.
- Recall:** translation unit = preprocessed source code file.
- One of preprocessor abilities: **inclusion of some source code file into another source code file**.
- Preprocessor directive: #include:**
 - This directive replaces itself with the content of the included file.
 - It does that recursively (included file may also contain #include directives).
- Example:** 2 source code files + translation unit for `main.cpp` (right):

```

// add.cpp
int add(int a, int b) {
    return a + b;
}

// main.cpp
#include "add.cpp"
int main() {
    return add(1, 2) + 3;
}


int add(int a, int b) {
    return a + b;
}

int main() {
    return add(1, 2) + 3;
}
  
```

8

Translation units vs source code files (cont.)

- Convention:**
 - Files to be included into other files are typically called **"header files"** (shortly headers) and have file extensions as `.h`, `.hpp`, `.hxx`, `.hh`, *no one*.
 - Files to be transformed into translation units are typically called **"source files"** and have file extensions as `.cpp`, `.cxx`, `.cc`, etc.
- This is just a convention; compilers and preprocessors do not care about how we call files or their extensions.
- Typical intent of headers: API.



9

Translation units vs source code files (cont.)

- Some "piece of functionality" — interface (API) + implementation.
- Ideal world:**
 - API in a header file (header files),
 - implementation in a source file (source files).
- Advantage of such separation:**
 - With dynamically-linked libraries (`.so`, `.dll`), changing implementation details does not require recompilation of programs' source code.
- Reality:**
 - Implementation details are frequently in header files.
- Examples:**
 - Inlining** — function bodies (implementation) must be visible to enable their inlining.
 - Class non-public members** — must be at least declared at the place where the class is defined.
 - Templates** — definitions must be visible to enable their instantiation.
 - Others** — functions with return type deduction, etc.

10

ODR

- Example:** function definition in a header file to enable its inlining:

```

// add.h
int add(int a, int b) { return a + b; }
  
```

- Example (cont.):** multiple source files that call this function:

```

// sub.cpp
#include "add.h"
int sub(int a, int b) {
    return add(a, -b);
}

// main.cpp
#include "add.h"
int main() {
    return add(1, 2) + 3;
}
  
```

- Example (cont.):** compilation and linkage...

```

$ g++ -O2 -c sub.cpp
$ g++ -O2 -c main.cpp
$ g++ main.o sub.o
  
```

- ...resulted in the following **linker error**:


```

multiple definition of `add(int, int)'
  
```

11

ODR (cont.)

- Translation units for `sub.cpp` (left) and `main.cpp` (right):



- Problem:**
 - C++ rule:** an entity (function, variable,...) generally cannot be defined in multiple translation units.
 - It can be defined only once — "one definition rule" (ODR).
 - Exceptions/workarounds:** static, anonymous namespace, inline, templates.
- Simplified explanation:**
 - Function `add` is defined in some translation unit and may be used in other translation units as well \rightarrow it has so-called **"external linkage"**.
 - Linker "sees" multiple machine code of `add` — in `sub.o` and `main.o`.
 - It cannot decide to which one it should link the `add` calls.

12

Specifier static

- How to enable to define a function inside a header file (for example, for inlining purposes) without breaking ODR?
- First option:** making this function static.

```
// add.h
static int add(int a, int b) { return a + b; }
```

- Translation units:

```
static int add(int a, int b) {
    return a + b;
}
int sub(int a, int b) {
    return add(a, -b);
}
```

```
static int add(int a, int b) {
    return a + b;
}
int main() {
    return add(1, 2) + 3;
}
```

- Functions add now have "internal linkage" — the definition in one translation unit is "private" to this unit \Rightarrow no ambiguity for a linker.
- These two definitions define two different functions (separate entities) regardless of their equivalent forms \Rightarrow no violation of ODR.
- Without static they would define a single function (entity).

13

Anonymous namespaces

- Alternative option:** putting add into anonymous namespace:

```
// add.h
namespace {
    int add(int a, int b) { return a + b; }
}
```

- Translation units:

```
namespace {
    int add(int a, int b) {
        return a + b;
    }
}
int sub(int a, int b) {
    return add(a, -b);
}
```

```
namespace {
    int add(int a, int b) {
        return a + b;
    }
}
int main() {
    return add(1, 2) + 3;
}
```

- Our case:** the same effect — add now have internal linkage.
- Anonymous namespace vs static:**
 - Generally, anonymous namespaces are more "powerful".
 - For example, we can put type definitions into anonymous namespace but cannot make types static.
 - Also, static has multiple meanings in different contexts, which might be confusing.

14

Confusion with static: example

- Example of confusedness of static:**

```
// X.h
class X {
    static void f();
};
void X::f() { }
```

```
// main.cpp
#include "X.h"
int main() { }
```

```
// other.cpp
#include "X.h"
```

```
$ g++ -O2 -c other.cpp
$ g++ -O2 -c main.cpp
$ g++ main.o other.o
```

multiple definition of 'X::f()'

- Static member function f has external linkage \Rightarrow violation of ODR.
- Here, static indicates that f is related to class X, not to its instances.

15

Specifier inline

- Another option:** inline functions.

```
// add.h
inline int add(int a, int b) { return a + b; }
```

- Translation units:

```
inline int add(int a, int b) {
    return a + b;
}
int sub(int a, int b) {
    return add(a, -b);
}
```

```
inline int add(int a, int b) {
    return a + b;
}
int main() {
    return add(1, 2) + 3;
}
```

- Both definitions of add define a single function (entity).
- With inline, they do not break ODR — exception.
- Generally, definitions of the same inline definition in multiple translation units define a single entity \Rightarrow all these definitions must be identical.
- An inline function must be defined in each translation unit where it is used (called).

16

Internal linkage vs inline

- Internal linkage** (static, anonymous namespace):

- Entities "private" for a translation unit.
- May have different definitions in different translation units.
- Usually defined in **source files**.
- Typical use case:** auxiliary entities private for some source file for which we don't want to care about breaking ODR.

- inline**

- Reusable entities \Rightarrow usually defined in **header files**.
- Typical use case:** functions for which we want to enable their inlining.
- Note:** inline itself does not enable or even force inlining; it just helps avoiding ODR violation.

17

Member functions and inline

- Class non-static member functions are:
 - implicitly inline** if they are defined within the class definition (body),
 - implicitly not-inline** if they are defined outside of the class definition.

- Example:**

```
// Y.h ver. 1
class Y {
    void f() { } // implicit inline
};
```

```
// Y.h ver. 2
class Y {
    void f();
};
void Y::f() { } // not inline
```

```
// main.cpp
#include "Y.h"
int main() { }
```

```
// other.cpp
#include "Y.h"
```

```
$ g++ -O2 -c other.cpp
$ g++ -O2 -c main.cpp
$ g++ main.o other.o
```

multiple definition of 'Y::f()'

18

Libraries

- Library — reusable coded functionality.
- Example:* Library that contains a single function for adding two unsigned integer numbers:

```
unsigned add(unsigned a, unsigned b) {
    return a + b;
}
```

- In which form can this library be implemented?
- First option — putting the function(ality) into a library header file ⇒ “header-only library”:

```
// unsigned_add.h
#ifndef UNSIGNED_ADD_H
#define UNSIGNED_ADD_H
inline unsigned add(unsigned a, unsigned b) { return a + b; }
#endif
```

19

Header-only libraries



- Advantages:
 - Functions may be inlined ⇒ potentially higher performance.
 - Using of this library involves only inclusion of the header file(s).
 - No linking is involved.
 - Library users (developers of programs built upon this library) need only the header file(s) ⇒ no machine code (object files) are required.
 - Programs users do not need anything regarding the library (do not need to know about its existence at all).
- Disadvantages:
 - Library source code is “copy-pasted” into the programs source code.
 - ⇒ Machine code related to the library functionality is in the program binary file.
 - ⇒ When the implementation of the library functionality is updated and we want to get this updated functionality into programs, these need to be rebuilt (recompiled and relinked).

20

Header-only libraries (cont.)

- Real-world scenario:
 - Someone develops a library; role — *library developer*.
 - Someone else develops a program by using this library; role — *program developer*.
 - The program is installed on computer systems and run; role — *program user*.
 - Library developer updates the library code (*bug fixes, improved performance,...*).
 - Program users want this updated library functionality in their programs.
- Library form — header only:
 - ⇒ Program needs to be rebuilt.
 - If the program source code is available, this can be done by program users.
 - Otherwise, it **needs to be done by program developer**.

21

Binary libraries

- Another option is to split the interface and implementation details between header and source files:

```
// unsigned_add.h
#ifndef UNSIGNED_ADD_H
#define UNSIGNED_ADD_H
unsigned add(unsigned, unsigned); // declaration only
#endif
```

```
// unsigned_add.cpp
unsigned add(unsigned a, unsigned b) { return a + b; } // definition
```

- Developers of programs using this library now do not need the library source code file which implements its functionality.
- Instead, they need only:
 - the header file (to be able to compile the program code that calls the library function),
 - the object file with the machine code of the function (to be able to link the final executable program together).

22

Binary libraries (cont.)

- The machine code of the program contains call instructions for the library function.
- These calls need to be linked with the machine code of the function — the library machine code.
- When** this linking happens?
- There are two options:
 - The linking happens when the program binary executable file is built.
 - The linking happens not until the program is executed.
- Ad 1. This approach is called “static linking” and the libraries used this way are called “statically linked libraries”.
- Ad 2. This approach is called “dynamic linking” and the libraries used this way are called “dynamically linked libraries”.
- In both cases, library machine code (object files) is “packed/archived” into special file formats:
 - Ad 1. .a or .lib extensions (Windows / Linux),
 - Ad 2. .dll / .so extensions (ditto).

23

Statically linked libraries

- Linking happens when the program is built such that both:
 - the machine code of the program,
 - the machine code of the library,
- is together “copied” into the final program executable file.
- ⇒ The library machine code is “hard-wired” into the program.
- Advantages:
 - Program users do not need to care about the library at all; they do not need to know about its existence.
- Disadvantages:
 - Reflection of library functionality updates into programs still **requires their rebuilding** (namely, only relinking is required without program source code recompilation).
 - Program users still need to update program binary files to reflect the library updates.

24

Statically linked libraries (cont.)

- Example: building a static version of the libunsigned library:

```
// unsigned_add.cpp
unsigned add(unsigned a, unsigned b) { return a + b; }

$ g++ -O2 -c unsigned_add.cpp
$ ar rcs libunsigned.a unsigned_add.o
```

- The ar tool "archives" object files into a single library file.
- Program developers now need the library header file for program compilation...

```
// main.cpp
#include <iostream>
#include <unsigned_add.h>
int main() { std::cout << add(1u, 2u) << std::endl; }
```

- ...and the library binary file for program linking:

```
$ g++ -O2 -c -I"/path/to/unsigned_add.h" main.cpp
$ g++ -O2 main.o -L"/path/to/libunsigned.a" -lunsigned
```

- Program users do not need that library file to run the program.

25

Dynamically linked libraries

- Linking happens when the program is executed.
- The system where the program is run needs to find the library machine code and link it with the library function call instructions in the program.
- Disadvantages:
 - Program users need to install binary library files (.dll, .so).
- Advantages:
 - When the library source code is updated and the new machine code is redistributed to systems (for example, under system packages updates), then programs immediately reflect the library updates.
 - No program change is required; the program binary executable files remain the very same (no recompilation, no relinking).

26

Dynamically linked libraries (cont.)

- Example:

```
// unsigned_add.cpp
unsigned add(unsigned a, unsigned b) { return a + b; }

$ g++ -O2 -c -fPIC unsigned_add.cpp
$ g++ -shared -o libunsigned.so unsigned_add.o
```

- The library binary file is "archived" by GCC.
- Program developers need the library header file for program compilation...

```
// main.cpp
#include <iostream>
#include <unsigned_add.h>
int main() { std::cout << add(0xFFFFFFFFu, 1u) << std::endl; }
```

- ...and the library binary file as well to "register" linking:

```
$ g++ -O2 -c -I"/path/to/unsigned_add.h" main.cpp
$ g++ -O2 main.o -L"/path/to/libunsigned.so" -lunsigned
```

27

Dynamically linked libraries (cont.)

- The library binary file libunsigned.so is now needed by users of the program when this gets executed:

```
$ ./a.out
./a.out: error while loading shared libraries: libunsigned.so: cannot open shared object
file: No such file or directory
$ export LD_LIBRARY_PATH="/path/to/libunsigned.so:$LD_LIBRARY_PATH"
$ ./a.out
0
```

- Now, changing the implementation of the library...

```
// unsigned_add.cpp
#include <iostream>
#include <limits>
unsigned add(unsigned a, unsigned b) {
    if (std::numeric_limits<unsigned>::max() - a < b)
        std::cerr << "WARNING: unsigned overflow occurred" << std::endl;
    return a + b;
}
```

- ...and rebuilding and redistribution of new libunsigned.so is automatically reflected in the program:

```
$ ./a.out
WARNING: unsigned overflow occurred
0
```

28

Library updates

- Does any update of library code need rebuilding of programs with dynamically-linked libraries?

1. Updates of library interface:

- Machine code of programs that interacts with the library machine code is generated based on the code in library headers.
- ⇒ Changes in these headers usually require program machine code regeneration.

Examples:

- *Compatible changes*: adding new entities (functions, global variables,...).
- *Incompatible changes*: changing numbers and types of function parameters, types of returned values, names of entities, removing entities,...

- ⇒ It is important to design a library interface in such a way that it would require minimum changes in the future.

- *Note*: Sometimes, it's very hard, since developers do not know which way their libraries will evolve.

29

Library updates (cont.)

2. Updates of library implementation:

Ideal world:

- Header files — functionality interface (what does library do).
- Source files/machine code — implementation (how does library do that).
- ⇒ Changes in the implementation do not require rebuilding of programs.

Reality:

- Implementation details sometimes need to be in header files.
- Their updates may break binary compatibility between program and library machine code.
- ⇒ Program rebuilds are required.

30

Library updates (cont.)

- Example: library code:

```
// X.h
class X {
    long i_;
public:
    X();
};
```

```
// X.cpp
#include "X.h"
X::X() : i_(0x12L) { }
```

- Program code (left) and its translation unit (right):

```
// main.cpp
#include <X.h>
int main() {
    X x;
}
```

```
class X {
    long i_;
public:
    X();
};
int main() {
    X x;
}
```

- Observable behavior of main:
 - calls default constructor of X,
 - calls destructor of X,
 - returns 0 to the caller (implicit return for main).

32

31

C++ object model

- **Object** in C++ = instance of some type (not necessarily class).
- Constructor call is a part of class object initialization.
- To **initialize an object** of type T, a **storage** for its binary representation must be provided first.
- This storage must satisfy two requirements:
 - It must be **large enough** — `sizeof(T)` bytes.
 - It must be **properly aligned** — address divisible by `alignof(T)`.
- Example:

```
std::cout << sizeof(long) << alignof(long); // printed '88' on x86_64/Linux
```

- **Observation:**
 - A storage for any object of type long must have 8 bytes and be aligned to an address divisible by 8 on this system.
 - Note: It's prescribed by its ABI \Rightarrow it holds for any compiler.

32

32

C++ object model (cont.)

- Binary representation of class objects consists of binary representation of its subobjects.
- Class subobjects:
 - base class objects,
 - non-static member variables.
- Binary representation of class objects is generally implementation-defined.
- Class X has only one subobject — member variable of type long.
- \Rightarrow In our case, the binary representation of X objects was the same as the binary representation of long objects.

```
class X {
    long i_; // member variable
public:
    X();
};
```

```
std::cout << sizeof(X) << alignof(X); // printed '88' on x86_64/Linux/Clang
```

33

33

C++ object model (cont.)

- **Non-static member function:**
 - A function that is called "on an object of some class".
 - This function needs a "reference" to the object on which it is called.
 - Note: This reference is available inside the function body in the form of the pointer `this`.
 - This reference represents a hidden parameter of all non-static member functions.
 - Its form is implementation-dependent.
 - Typical implementations pass an address of the object as argument for this parameter.
- **Constructor** = (special) member function:
 - Object does not yet exist; the purpose of constructors is to create it.
 - \Rightarrow The caller passes the address of the storage where the object should be initialized \Rightarrow this storage must be prepared by the caller.
 - Constructor then initializes a binary representation of the object in this storage.

34

34

C++ object model (cont.)

- Translation unit (left), machine code x86_64/Linux/Clang (right):

```
class X {
    long i_;
public:
    X();
};
int main() {
    X x;
}
```

```
main:
    push    rax
    mov     rdi, rsp
    call    X::X()
    xor     eax, eax
    pop     rax
    ret
```

- In main, x is a non-static local variable.
 - \Rightarrow x is a new object created in each main call.
- main needs to:
 - prepare a storage for x,
 - call a default constructor of X
 - while passing the address of the storage to the constructor.
- **Observation:**
 - Storage for x is allocated on the stack (at address `rsp` decreased by 8).
 - Its address is to constructor passed in `rdi` register (as prescribed by ABI).

35

35

C++ object model (cont.)

- **Other considerations:**
- Definition of the constructor is not in the program translation unit.
 - \Rightarrow Its call cannot be "inlined"; it must be called explicitly at the machine code level.
- Variable x is an object with so-called "automatic storage duration":
 - Its lifetime ends when the function is leaved.
 - \Rightarrow Its destructor must be called.
- **Destructor of X:**
 - There is no custom declaration/definition.
 - \Rightarrow The destructor is automatically defined.
 - \Rightarrow As if it was defined (simplification) this way
- \Rightarrow Destruction of x may be inlined and has no observable behavior.
 - \Rightarrow It does not generate any machine code under optimizations.

```
class X {
    long i_;
public:
    X();
    ~X() { }
```

36

36

C++ object model (cont.)

- Translation unit for the library source file / machine code:

```
class X {
    long i_;
public:
    X();
};
X::X() : i_(0x12L) { }
```

```
X::X()
    mov     qword ptr [rdi], 0x12
    ret
```

- Recall:** the address of the storage for the constructed object is passed in `rdi`.
- Constructor initializes the binary representation of the constructed object in this storage.
- Binary representation of an object of `X` consists of the binary representation of its member variable `i_`.
- This member variable is initialized with hex value 12.
- ⇒ This value needs to be stored to the address pointed to by `rdi`.

37

Library updates (cont.)

- Summary:

- Left:** machine code stored in the program binary file.
- Right:** machine code stored in the library binary file (.so).

```
main:
    push    rax
    mov     rdi, rsp
    call    X::X()
    xor     eax, eax
    pop     rcx
    ret
```

```
X::X()
    mov     qword ptr [rdi], 0x12
    ret
```

- Observation:**

- The machine code responsible for the allocation (and deallocation) of the storage for the constructed object is "hardcoded" in the program file.
- ⇒ Any change of the library that requires a different way of storage allocation will break the binary compatibility.

38

Library updates (cont.)

- Example of binary incompatible library update — adding another member variable:

```
// X.h
class X {
    long i_, j_;
public:
    X();
};
```

```
// X.cpp
#include "X.h"
X::X() : i_(0x12L), j_(0x34L) { }
```

- Machine code of program (left) and new library (right):

```
main:
    push    rax
    mov     rdi, rsp
    call    X::X()
    xor     eax, eax
    pop     rcx
    ret
```

```
X::X()
    mov     qword ptr [rdi], 0x12
    mov     qword ptr [rdi+8], 0x34
    ret
```

- Problem:**

- Program allocates 8 bytes on the stack, but constructor internally overwrites 16 bytes!

39

Library updates (cont.)

```
main:
    push    rax
    mov     rdi, rsp
    call    X::X()
    xor     eax, eax
    pop     rcx
    ret
```

```
X::X()
    mov     qword ptr [rdi], 0x12
    mov     qword ptr [rdi+8], 0x34
    ret
```

- Consequence:**

- A constructor overwrites some **part of the stack that does not belong to the constructed object**.
- In our case, it overwrites the return address where the `ret` instruction "jumps" when the constructor is finished.
 - In *best case*, the program just crashes.
 - In *worst case*, it starts executing some (un)predicted machine code.
- ⇒ Such a vulnerability might be generally exploited.

40

Library updates (cont.)

- Recapitulation:**

- Original library code (black), updates (with red):

```
// X.h
class X {
    long i_, j_;
public:
    X();
};
```

```
// X.cpp
#include "X.h"
X::X() : i_(0x12L), j_(0x34L) { }
```

- Updates logically involved only implementation library details:
 - Whatever is private to a class is logically related to implementation of its functionality.
- However, even private members **must be declared within class definition** (body).
 - ⇒ These declarations must appear in header files.
- C++ generally does not allow a complete logical separation of interface and implementation details between header and source files, respectively.
- Workaround:** PIMPL idiom.

PRIVATE

41

PIMPL idiom

- PIMPL = Pointer-to-IMPLementation.**
- Technique for "true" hiding of all implementation details out of header files.
- Functionality:**
 - Original class is "split" into two classes.
 - First** — "implementation" class — contains the original class implementation details including all its member variables.
 - Second** — "interface" class — contains:
 - a pointer to the implementation class object,
 - interface of the original class = its public member functions.
 - Public member functions of the interface class need to invoke member functions of the implementation class.
 - This is done internally inside library source files.
- Outcome:**
 - Interface class does not need to be changed when implementation class member variables and member functions are updated in any way.
 - ⇒ If the "true interface" — public member functions — is not changed, there is no need for programs rebuilding with dynamically linked libraries.

42

PIMPL idiom (cont.)

• Example — original class:

```
// X.h
class X {
    long i_;
public:
    X();
    void do_something();
};

// X.cpp
#include <iostream>
#include "X.h"
X::X() : i_(0x12L) { }
X::do_something() { std::cout << i_; }
```

• "PIMPLed" solution:

```
// X.h
class X { // interface class
    class Impl;
    // -> implementation...
    // ...class declaration
    Impl* pimpl_;
    // -> pointer...
    // ...to implementation
public:
    X();
    ~X();
    void do_something();
};

// X.cpp
#include <iostream>
#include "X.h"
class X::Impl { // implementation class
    long i_;
public:
    Impl(): i_(0x12L) { }
    void do_something() { std::cout << i_; }
};
// interface-implementation binding:
X::X() : pimpl_( new Impl() ) { }
X::~X() { delete pimpl_; }
X::do_something() { pimpl_->do_something(); }
```

43

PIMPL idiom (cont.)

• Library update:

```
// X.h
class X {
    class Impl;
    Impl* pimpl_;
public:
    X();
    ~X();
    void do_something();
};

// X.cpp
#include <iostream>
#include "X.h"
class X::Impl {
    long i_, j_;
public:
    Impl(): i_(0x12L), j_(0x34L) { }
    void do_something() {
        std::cout << i_ << " " << j_;
    }
};
X::~X() : pimpl_( new Impl() ) { }
X::~X() { delete pimpl_; }
X::do_something() { pimpl_->do_something(); }
```

• Observation:

- No change in the header file
⇒ no need for program rebuild
(with dynamic library linking).

• Notes:

- In X.h, class X::Impl is only declared ⇒ it is so-called "incomplete type".
- It is legal to declare a pointer-to-incomplete type (pimpl_ in our case).
- All the operations that require this type to be complete (such as new and delete expressions) are in the source file.
- The class gets complete after its definition.

44