

Effective C++ Programming

NIE-EPC (v. 2021):

UNDEFINED BEHAVIOR, OBJECT LIFETIME, STATIC
AND DYNAMIC STORAGE ALLOCATION,
INITIALIZATION AND DESTRUCTION

© 2021 DANIEL LANGR, ČVUT (CTU) FIT

Undefined behavior

- *Recall:*
 - C++ standards define observable behavior of a C++ program on the abstract machine.
 - C++ implementations provide the same observable behavior on a particular computer system.
- Is the observable behavior **always defined**?
 - **No**, it is defined only if C++ standard rules/requirements are satisfied.
- What happens **if they are violated**?
 - Then, C++ standards **do not define program/code behavior**.
- This case is referred to as “**undefined behavior**” (**UB**).
- *Examples:*
 - signed integer **overflow**, dereferencing **invalid pointer** (pointer that does not point to an existing object of a given type), accessing **out-of-bound array element**, creating an object in **not properly aligned** storage, **data race**, ...

Undefined behavior (*cont.*)

- Once some standard rule is violated, the C++ standard no longer define behavior on the abstract machine.
- In practice, this violation mostly happens at runtime, but it can also happen at compile time.
 - In the first case, the behavior is undefined for the executed program.
 - In the second case, the behavior is undefined even for compile-time C++ implementation tools (compiler, linker,...).
- *Example program source file:*

```
#include <iostream>
#include <limits>

int main() {
    int i = std::numeric_limits<int>::max();
    i++; // causes signed integer overflow => undefined behavior
    std::cout << i;
}
```

- *Question:* When does the overflow happen?

Undefined behavior (*cont.*)

```
int i = std::numeric_limits<int>::max();  
i++; // causes signed integer overflow => undefined behavior
```

- When does the overflow happen?
 1. With **disabled optimizations**, the **calculation will likely be performed at runtime** \Rightarrow the **overflow will happen at runtime** \Rightarrow the **program behavior is undefined at runtime** once the overflow takes place.
 2. With **enabled optimizations**, the **calculation will likely be performed at compile time** \Rightarrow the **overflow will happen at compile time** \Rightarrow the **behavior is undefined already at compile time**.
- **Ad 1.** Once the overflow occurs, the **behavior of the program** is no longer defined by the C++ standard.
 - It may crash, it may behave as expected, it may behave unexpectedly,...
- **Ad 2.** Once the overflow occurs, the **behavior of the implementation tools** is no longer defined by the C++ standard.
 - It may throw a compilation error or warning, it may crash, it may generate some machine code with some expected or unexpected behavior,...

Undefined behavior (*cont.*)

- *Example: GCC/x86_64/Linux implementation:*
 - In both cases (enabled/disabled optimization) a compiler **generated program machine code** without any error/warning message.
 - Relevant part of generated machine code for **enabled optimizations**:

```
mov    esi, -2147483648  
mov    edi, OFFSET FLAT:_ZSt4cout  
call   std::basic_ostream<char, std::char_traits<char> >::operator<<(int)
```

- ⇒ The **overflow occurred at compile time.**
- Relevant part of generated machine code for **disabled optimizations**:

```
mov    DWORD PTR [rbp-4], 2147483647  
add    DWORD PTR [rbp-4], 1  
mov    eax, DWORD PTR [rbp-4]  
mov    esi, eax  
mov    edi, OFFSET FLAT:_ZSt4cout  
call   std::basic_ostream<char, std::char_traits<char> >::operator<<(int)
```

- ⇒ The **overflow occurs at runtime** when add instruction is executed.

Undefined behavior (*cont.*)

```
int i = std::numeric_limits<int>::max();  
i++; // causes signed integer overflow => undefined behavior
```

- With this code, **overflow happens always** (**unconditionally**) and may happen **both at compile time or at runtime**.
- What if we **change this code to?**

```
std::cin >> i;  
i++; // causes overflow only if i equals max value of int
```

- Then, the **overflow**:
 - **cannot happen at compile time** (the calculation cannot be performed at compile time when the input value is not known);
 - happens **only if 2147483647 is read from the standard input**.
- **Conclusion:** The **behavior of the program** is now:
 - **undefined** only if 2147483647 is read from the standard input,
 - **(well) defined** otherwise (if user inputs 1, i will be incremented to 2).
- **Note:** This implies that, generally, undefined behavior **cannot be detected by static code analysis** (*more details later*).

Undefined behavior (*cont.*)

- The C++ standard does not define behavior for this program once the overflow occurs:

```
#include <iostream>
#include <limits>

int main() {
    int i = std::numeric_limits<int>::max();
    i++; // causes signed integer overflow => undefined behavior
    std::cout << i;
}
```

- Yet, a C++ implementation generated a program binary executable file with some machine code:

```
mov     esi, -2147483648
mov     edi, OFFSET FLAT:_ZSt4cout
call    std::basic_ostream<char...
```

```
mov     DWORD PTR [rbp-4], 2147483647
add     DWORD PTR [rbp-4], 1
mov     eax, DWORD PTR [rbp-4]
mov     esi, eax
mov     edi, OFFSET FLAT:_ZSt4cout
call    std::basic_ostream<char...
```

- In both cases, execution renders same behavior:

```
$ ./a.out
-2147483648
```

Undefined behavior (*cont.*)

- Once some **condition** for undefined behavior is met, from the perspective of the **C++ standard**, **anything can happen** (*undefined = anything is allowed*).
- This “anything” includes also the **possibility of the observed behavior...**

ANARCHAOS

```
$ ./a.out  
-2147483648
```

- ...just it is **not defined by the C++ standard**.
- Instead, it is **defined by a given implementation** (in our case, by *GCC/x86_64/Linux*).
- *Implication:*
 - When the behavior is **implementation-defined**, there is **no portability**.
 - With **another C++ implementation / on other system / with other built setup / ...**, the **observed behavior may be completely different**.

Undefined behavior (*cont.*)

- (Partial) **summary**:

1. When the **rules of the C++ standard are not violated**:

- *C++ standard* defines behavior observable on the abstract machine.
- *C++ implementation* must provide the same behavior observable on a particular computer system.
- This must hold for all C++ implementations on all systems (*portability*).

2. When the **rules of the C++ standard are violated**:

- *C++ standard* does not define any behavior on the abstract machine.
- *C++ implementation* can do whatever it wants; typically, it provides some behavior.
- This behavior is implementation-specific; generally, it is different for different implementation/systems (*no portability*).

- *Conclusion*:

- To have correctly-running portable programs, one should avoid undefined behavior as much as possible.

Perils of undefined behavior

- When a condition for undefined behavior is met, the **actually-observed behavior** (defined by a given C++ implementation) may **correspond with some “expected behavior”**.

```
int i = std::numeric_limits<int>::max();  
i++;
```

```
$ ./a.out  
-2147483648
```

- Is this result expected?
 - Mathematically, no** — adding two positive numbers **cannot give negative result**.
 - However, this outcome **might be expected** under the **assumption of two’s complement** binary representation of integer numbers.
 - With **other representations** (such as *one’s complement*, where negative zero exists), the result **might be different**.
- “Expected behavior” is **subjective** — **different people** may have **different expectations**.
- Frequent **argumentation** that “*the code is correct since it behaves as I expect*” is **very dangerous and wrong once UB is involved**.

Perils of undefined behavior (*cont.*)

- In our example, the implementation-specific behavior **was always the same** (hardcoded in the program machine code):

```
mov     esi, -2147483648
mov     edi, OFFSET FLAT:_ZSt4cout
call    std::basic_ostream<char...
```

```
mov     DWORD PTR [rbp-4], 2147483647
add     DWORD PTR [rbp-4], 1
mov     eax, DWORD PTR [rbp-4]
mov     esi, eax
mov     edi, OFFSET FLAT:_ZSt4cout
call    std::basic_ostream<char...
```

```
$ ./a.out
-2147483648
```

- Once this machine code is generated, each execution of the program **is guaranteed to have the same behavior**.
- However, **generally, this does not hold**.
- *Example:*
 - **Undefined behavior caused by data race** (unsynchronized read-write access to the same memory location from multiple threads).
 - **Each time** data race occurs, it may have **different unpredictable impact on observable program behavior**.

Perils of undefined behavior (*cont.*)

- *Example:*

```
int i = 0; // global variable shared by threads
void inc_i() { i++; }
```

- When function `inc_i` is executed N times by each of T threads, one might **expect** that the value of `i` would be incremented by $N \times T$ in total.
- This “expected” **behavior is undefined**:
 - According to the C++ standard, this is **data race**, which causes **undefined behavior**.
- *Note:* To **get rid of data race** / undefined behavior, `i` would need to have **atomic data type**, such as `std::atomic<int>`, instead of **int**.
- However, in practice, this **expected behavior may be observed in most cases** and, only **occasionally**, it may be **different**.

Perils of undefined behavior (*cont.*)

- Undefined behavior when the program **mostly behaves as expected**:
 - Extremely dangerous for **production environments and mission-critical applications** (*server applications, databases, banking, embedded systems in planes, cars, space crafts,...*).
- Program can **pass all tests** and **run without problems for long time**, until once...:
 - In **best cases**, occasional unexpected behavior **causes program to crash**.
 - In **worst cases**, it **continues to run with incorrect state**.
- Once a **problem is identified**, it might be **very hard to find its cause**:
 - We would need to replicate this **unexpected behavior during testing / debugging**.
 - Since it happens only **occasionally**, it may be **almost impossible**.

Detection undefined behavior

- Can undefined behavior be **detected**?
 - We have seen, that this is **generally impossible at compile time** (with static code analysis).
- Is it possible **at runtime**?
 - In theory, maybe.
 - **In practice, no.**
- Why?
 - Even if we could **detect violation of C++ standard rules at runtime**, it would **impose a large runtime and memory overhead into C++ programs**.
 - The **purpose of programming in C++** is mostly **performance and efficiency**.
- \Rightarrow Generally, in programming, we can **either have safety or performance**. It's **not possible to have both at once**.
- **C++ goes for performance.**

Detection undefined behavior (*cont.*)

- *Example:* (library) **source file** (*left*), **translation unit** (*right*):

```
// deref_int.cpp:  
int deref_int(int* ptr) {  
    return *ptr;  
}
```

```
int deref_int(int* ptr) {  
    return *ptr;  
}
```

- According to the C++ standard, **dereferencing of a pointer** (*ptr here) **is valid only if this pointer actually points to an existing object of a given type** (int here).
- If **this requirement is not satisfied**, the **behavior is undefined** once the dereferencing takes place.
- \Rightarrow **To detect undefined behavior** with this code, a C++ implementation would need to **test the “validity” of a pointer**.
- *How?*

Detection undefined behavior (*cont.*)

- The **most efficient representation of a pointer** is just the **value of the address**.
- This representation allows generating the following **optimal machine code** (*GCC/x86_64/Linux*):

```
int deref_int(int* ptr) {  
    return *ptr;  
}
```

```
deref_int(int*):  
    mov eax, DWORD PTR [rdi]  
    ret
```

- A compiler has **no idea what will be passed as a function argument**.
- \Rightarrow **Detection of UB** would require some **mechanism of how to find out whether a pointer actually points to an existing object**.
- Is such a mechanism even **possible**? **Likely not**.
- Even if it was, it would have **tremendous negative impact on performance**:
 - When some **object would be destroyed**, all pointers that pointed to it **would need to be informed about this destruction**.
 - In the function, there **would need to be a test for whether the pointer is valid or not**.

Non-standard C++ extensions

- What happens in case of undefined behavior is mostly implementation-specific.
- Implementations can do anything they want and, typically, they do not care much about what happens:

```
int deref_int(int* ptr) {  
    return *ptr;  
}
```

```
deref_int(int*):  
    mov eax, DWORD PTR [rdi]  
    ret
```

- *Observation:*
 - GCC does not care whether the passed pointer actually points to an existing object of type `int`.
 - This is a function user (caller) responsibility to guarantee that.
 - Machine code simply reads a 32-bit value from the passed address.
 - In case of undefined behavior, if this address is accessible, the function will just return some value.
 - Otherwise, it will likely cause the running program (process) to crash (with something like *segfault*).

Non-standard C++ extensions (*cont.*)

- In **some (rare) cases**, C++ implementations **explicitly specify what happens** even when **C++ standard rules are violated**.
- These cases form **non-standard C++ extensions**.
- *Example:*

```
union U {  
    int i;  
    float f;  
};
```

- *Union* is a **class-like type**.
- In contrast to *classes* and *structs*, **union's non-static member variables share the same storage**:

```
struct S { // same with class  
    int i;  
    float f;  
};
```

```
std::cout << sizeof(S) << sizeof(U); // prints '84' (x86_64/Linux)
```


Non-standard C++ extensions (*cont.*)

- C++ standard rules:
 - Only one union's member variable is so-called "*active*", namely the one most recently written into.
 - Only the active member variable may be read.
 - \Rightarrow Reading the inactive member variable = undefined behavior.

```
union U {  
    int i;  
    float f;  
};  
  
int main() {  
    U u;  
    u.f = 1.0f;  
    return u.i;    // undefined behavior according to C++ standards  
}
```

- However, GCC explicitly allows this code and specifies a particular behavior for it:
 - "The relevant bytes of the representation of the object *are treated as an object of the type used for the access.*"
 - \Rightarrow u.i has the value of an object of type int that has the same binary representation as the float object with value 1.0f.

Non-standard C++ extensions (*cont.*)

- *GCC/x86_64/Linux*:
 - Floating-point values are represented in IEEE 754 format.
 - `float` is a single-precision IEEE 754 floating-point data type.
 - Value 1 has binary representation `0x3f800000` in hex with this data type.

```
int main() {  
    U u;  
    u.f = 1.0f;  
    return u.i;  
}
```

```
main:  
    mov eax, 0x3F800000  
    ret
```

- The behavior of the resulting program is:
 - undefined by the C++ standard,
 - but guaranteed by the used GCC implementation.
- \Rightarrow This source code is generally not portable.
- *Similar case:*



```
int main() {  
    float f = 1.0f;  
    return *reinterpret_cast<int*>(&f);  
}
```

```
main:  
    mov eax, 0x3F800000  
    ret
```

Object lifetime

- **Object lifetime** — time interval during which the **object exist at runtime** (when the program is run).
- *Simplified rules:*
- **Lifetime begins with object initialization**, which **constructs** the object.
- In case of **class types**:
 - object initialization involves **calling a class constructor**,
 - object lifetime ends when its **destructor is called**.
- Before an object can be constructed/initialized, a **storage must be allocated for it**.
 - **Storage** = place where a **binary representation** of the object will be stored.
 - *Recall* — this storage must be **properly aligned and sized**.



Storage allocations

- **How** is object storage allocated?
- **Two separate cases:**
 1. Object existence is **known when the program is compiled**:
 - For storage allocation for this object, a **compiler** is responsible.
 - This case is referred to as a **"static allocation"**.
 - Meaning of *static* here = **resolvable-at-compile-time**.
 - *Effect* — allocation is **hard-code in the program** (it's machine code).
 2. Object existence is **not known until runtime**:
 - Storage allocation must be **explicitly** initiated by a programmer.
 - This case is referred to as a **"dynamic allocation"**.
 - Meaning of *dynamic* = **resolvable-not-until-runtime**.
 - *Effect* — allocation needs to be resolved at runtime **by some external mechanisms** (memory allocation functions).

Ambiguity of “static”



- “*Static*” vs “*static*” **ambiguity**:
 - Static allocations are **not related to the static specifier/keyword**:

```
void f() {  
    int i = 1;  
    static int j = 2;  
    ...  
}
```

- **Non-static function-local variable i**:
 - It's lifetime is **restricted to the function body**.
 - ⇒ **Each time f is called** the lifetime of i starts (with its initialization) and ends (when the function is left).
- **Static function-local variable j**:
 - It's lifetime basically **covers the whole program run**.
 - Typically, it starts when **f is called first time** (at the latest) and ends when the **program is terminated**.
- Storage for both i and j variables are **statically allocated** — allocation is “hard-coded” in the program at compile time.

Ambiguity of “static” (cont.)

- The word “*static*” can have many different meaning in C++.
- We have seen:
 - *Static entities* (functions, variables) with the meaning of being local to translation units with respect to linking (having *internal linkage*).
 - *Static class members* (functions, variables) with the meaning of belonging to the class itself instead of to its instances.
 - *Static allocations* with the meaning of being resolvable at compile time.
- The last meaning of the word “static” is used more generally.
- *Example* — static vs dynamic polymorphism:
 - *Polymorphism* = writing a unified code that can operate on objects of different types.
 - *Dynamic polymorphism* is provided by virtual member functions.
 - Function calls are resolved at runtime according to the actual (*dynamic*) object type.
 - *Static polymorphism* is provided by templates.
 - Templates are resolved at compile time (they do not exist at runtime).

Static allocations — *example*

- **Translation unit** (*left*) and generated **machine code** (*right*):


```
void g(int*, int*);  
void f() {  
    int i = 1;  
    static int j = 2;  
    g(&i, &j);  
}
```

```
f():  
    push    rax  
    mov     dword ptr [rsp + 4], 1  
    lea     rdi, [rsp + 4]  
    mov     esi, offset f()::j  
    call    g(int*, int*)  
    pop     rax  
    ret  
  
f()::j:  
    .long   2
```

- Storage for *i* and *j* variables is **statically allocated**:
 - Allocations in both cases are **hard-coded in the program file / generated machine code** (at compile time).
- *Note*: observable behavior of *f* involves **pointers to *i* and *j***.
 - Function *f* passes these pointers as **arguments to function *g***.
 - \Rightarrow A compiler is “forced” to **allocate storage for *i* and *j* in memory**.
 - According to **ABI**, pointers are **passed through *rdi* and *rsi* registers**.
 - On the contrary, **pointers *&i* and *&j* are not stored in memory**.

Static allocations — *example (cont.)*

```
void g(int*, int*);  
void f() {  
    int i = 1;  
    static int j = 2;  
    g(&i, &j);  
}
```

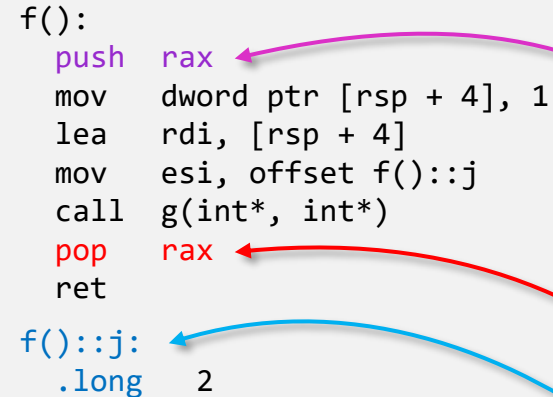


- Where in memory are `i` and `j` variables stored?
- C++ standards do not specify this; they “do not care”.
- C++ standards only prescribe so-called “*storage duration*”.
 - `i` has “*automatic storage duration*” — storage is allocated when **the code block** is entered and deallocated when it is exited.
 - `j` has “*static storage duration*” — storage is allocated when **the program** is started and deallocated when it is exited.
- Storage in memory (address space) is specified by a given C++ implementation.
- Implementations need to conform to a system ABI.

Static allocations — *example (cont.)*

```
void g(int*, int*);  
void f() {  
    int i = 1;  
    static int j = 2;  
    g(&i, &j);  
}
```

```
f():  
    push    rax  
    mov     dword ptr [rsp + 4], 1  
    lea     rdi, [rsp + 4]  
    mov     esi, offset f()::j  
    call    g(int*, int*)  
    pop     rax  
    ret  
  
f()::j:  
    .long   2
```



- Used implementation: *Clang/x86_64/Linux*:
 - Storage for **non-static function-local variable** `i` is allocated **on the stack**.
 - **Allocation** is effectively performed by “reserving” stack space by lowering the **stack pointer register** `rsp` (this is what the `push` instruction does).
 - **Deallocation** is performed by **raising the stack pointer back** to the original value (by using the `pop` instruction).
 - Storage for **static function-local variable** `j` is allocated **on the so-called program data segment**.
 - Memory for data segment is **allocated** when the program is executed and **deallocated** when it is exited.

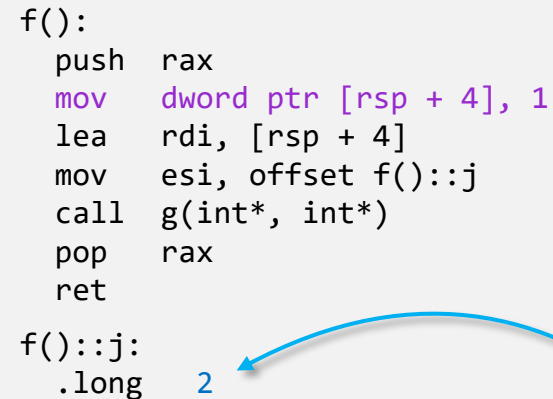
Allocation and initialization

- Once there is a storage allocated for some object, **its lifetime begins with its initialization**.
- **The object lifetime ends with its destruction**, which takes place before its storage is deallocated.
- There are **many forms of initialization**, which, generally, have **different effect for objects of different types**.
- “Rough” distinction of **two different cases**.
 1. *Class types*:
 - **Initialization** involves a **constructor call**.
 - **Destruction** involves a **destructor call**.
 2. *Non-class types* (such as “*fundamental*” language-intrinsic types):
 - **Initialization** requires **setting a binary representation to a desired value** (determined by the initialization expression).
 - **Destruction** **does not require any action**.

Allocation and initialization (*cont.*)

```
void g(int*, int*);  
void f() {  
    int i = 1;  
    static int j = 2;  
    g(&i, &j);  
}
```

```
f():  
    push    rax  
    mov     dword ptr [rsp + 4], 1  
    lea     rdi, [rsp + 4]  
    mov     esi, offset f()::j  
    call    g(int*, int*)  
    pop     rax  
    ret  
  
f()::j:  
    .long   2
```



- In our case, **int** is a fundamental **non-class type**.
- **i** is initialized by **= 1** expression:
 - It is initialized each time **f** is executed **after its storage has been allocated**.
 - **Initialization** = **setting its binary representation to "have" value 1**.
 - Its lifetime ends when the function is left (*no action required*).
- **j** is initialized by **= 2** expression:
 - **Initialization** = **setting its binary representation to "have" value 2**.
 - This value is **hard-code in the program file** and loaded to the data segment when the program is executed.

Allocation and initialization (cont.)

- For a *class type*:
 - *allocation/deallocation* works the very same way,
 - *initialization involves constructor, destruction involves destructor.*

```
struct X {  
    X(int);  
    ~X();  
};  
void g(X*) noexcept;  
void f() {  
    X x(1);  
    g(&x);  
}
```

```
f():  
    push    rbx  
    sub     rsp, 16  
    lea     rbx, [rsp + 8] ; address of the storage  
    mov     rdi, rbx      ; passed by rdi to constructor  
    mov     esi, 1        ; initialization argument passed by rsi  
    call    X::X(int)  
    mov     rdi, rbx      ; address of x passed to g by rdi  
    call    g(X*)  
    mov     rdi, rbx      ; to destructor as well  
    call    X::~~X()  
    add     rsp, 16  
    pop     rbx  
    ret
```

- *Notes:*
 - A *static* function-local variable would be initialized *only the first time f is called* and *destroyed when the program exits*.
 - Destructor for x would need to be called *even if g threw an exception*.
 - \Rightarrow In case of *not-noexcept* g, machine code would be *more complicated*.

Static vs dynamic allocations

- **Static allocations** — **always take place** when the program is run
⇒ can be hard-coded in the program machine code.
- In some cases, allocations **cannot be resolved until runtime**.
- *Examples:*
 - **Conditional allocations** dependent on some runtime-evaluable condition.
 - **Allocations of storage for multiple objects**, where their number is not known until runtime.
 - *Etc...*
- Such allocations **need to be resolved at runtime**, i.e., dynamically.
- *Example:*

```
void g(int*, int*);  
void f(bool b, long n) {  
    int* pi = b ? new int(1) : nullptr;  
    int* pa = new int[n]{};  
    g(pi, pa);  
}
```

Dynamic allocations

- “Dynamically-allocated objects” — the very same concept:
 - 1) First, a storage for an object needs to be allocated.
 - 2) Second, that object is in this storage initialized.
- Ad 1) How is a storage allocated in such a case?
 - Again, C++ standard do not specify this.
 - They only define so-called *allocation functions*.
 - These functions are called `operator new` and return a block of “uninitialized” memory of a desired size of bytes.
 - *Uninitialized* = there is no object in this memory yet.
 - How do these functions work is implementation-defined.
- Typical *implementations*:
 - Internally `allocate memory on the heap` (a system-provided mechanism for dynamic memory allocations).
 - \Rightarrow `operator new` internally calls some heap allocation function, such as `malloc`.

operator new vs new expression

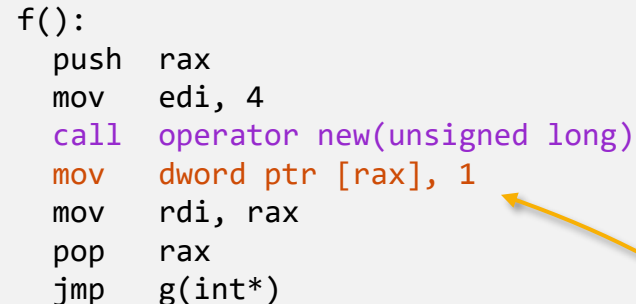
- Allocation function operator new only provides storage for dynamically-allocated objects.
- Then, objects need to be in that storage constructed / initialized.
- Usually, both these steps are required (by programmers) to be performed together.
- \Rightarrow They are “bound” in a mechanism called new expression.
- new expression is the most common form of new.
- It is responsible for both:
 - 1) allocation storage for an object,
 - 2) its initialization in this storage.
- For allocation, it internally uses operator new.

operator new vs new expression (cont.)

- Example with a **fundamental type**...

```
void g(int*);  
void f() {  
    int* pi = new int(1); // new expression  
    g(pi);  
}
```

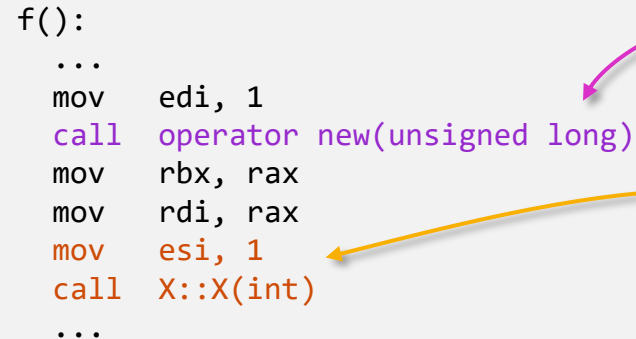
```
f():  
    push    rax  
    mov     edi, 4  
    call    operator new(unsigned long)  
    mov     dword ptr [rax], 1  
    mov     rdi, rax  
    pop     rax  
    jmp     g(int*)
```



- and with a **class type**:

```
void g(X*);  
void f() {  
    X* px = new X(1); // new expression  
    g(px);  
}
```

```
f():  
    ...  
    mov     edi, 1  
    call    operator new(unsigned long)  
    mov     rbx, rax  
    mov     rdi, rax  
    mov     esi, 1  
    call    X::X(int)  
    ...
```



- *Observation:*
 - new expression **internally allocates storage with operator new**,
 - and then **initializes the object in this storage**.

Dynamic allocations (*cont.*)

- With *statically-allocated objects*:

- objects are *automatically destructed*,
- storage is *automatically deallocated*.

- The same *does not hold for dynamically-allocated objects*:

- object is *not destructed* (no destructor call),
- storage is *not deallocated* (no deallocation function call).

- Objects “manually” *allocated and initialized with expression new* need to be “manually” *destructed and deallocated with its counterpart — expression delete*.

```
struct X {  
    X(int);  
    ~X();  
};  
void g(X*)  
    noexcept;  
void f() {  
    X x(1);  
    g(&x);  
}
```

```
f():  
    push    rbx  
    sub     rsp, 16  
    lea     rbx, [rsp + 8]  
    mov     rdi, rbx  
    mov     esi, 1  
    call    X::X(int)  
    mov     rdi, rbx  
    call    g(X*)  
    mov     rdi, rbx  
    call    X::~~X()  
    add     rsp, 16  
    pop     rbx  
    ret
```

```
void g(X*);  
void f() {  
    X* px  
        = new X(1);  
    g(px);  
}
```

```
f():  
    .  
    call    X::X(int)  
    mov     rdi, rbx  
    add     rsp, 8  
    pop     rbx  
    pop     r14  
    jmp     g(X*)
```

operator delete vs delete expression

- **new expression:**
 - 1) **allocates** storage for an object,
 - 2) **initializes** object in this storage.
- **delete expression:**
 - a) **destructs** object,
 - b) **deallocates** its storage.
- **operator new:**
 - **C++ allocation function** internally called by new expression in ad 1).
 - Typically implemented by calling some **heap allocation function** (such as malloc).
- **operator delete:**
 - **C++ deallocation function** internally called by delete expression in ad b).
 - Typically implemented by calling a **heap deallocation function** (free).

operator delete vs delete expr. (cont.)

```
void g(X*);  
void f() {  
    X* px = new X(1);  
    g(px);  
    delete px; // delete expression  
}
```

```
f():  
    ...  
    mov     edi, 1  
    call    operator new(unsigned long)  
    mov     rbx, rax  
    mov     rdi, rax  
    mov     esi, 1  
    call    X::X(int)  
    mov     rdi, rbx  
    call    g(X*)  
    mov     rdi, rbx  
    call    X::~X()  
    mov     rdi, rbx  
    add     rsp, 8  
    pop     rbx  
    pop     r14  
    jmp     operator delete(void*)
```

- *Observation* — **delete expression**:

- **destructs object** by **calling its destructor**,
- **deallocates its storage** by **calling C++ deallocation function operator delete**.

- *Notes*:

- (De)allocation functions **operator new** and **operator delete** are provided by the C++ standard library.
- \Rightarrow In practice, they are provided by **library implementations**.

- *Example* — operator new in **libc++** invokes **malloc**:

<https://github.com/llvm/llvm-project/blob/main/libcxx/src/new.cpp>.