

Effective C++ Programming

NIE-EPC (v. 2021):
UNDEFINED BEHAVIOR, OBJECT LIFETIME, STATIC
AND DYNAMIC STORAGE ALLOCATION,
INITIALIZATION AND DESTRUCTION
© 2021 DANIEL LANGR, ČVUT (CTU) FIT

1

Undefined behavior

- **Recall:**
 - C++ standards define observable behavior of a C++ program on the abstract machine.
 - C++ implementations provide the same observable behavior on a particular computer system.
- Is the observable behavior always defined?
 - **No**, it is defined only if C++ standard rules/requirements are satisfied.
- What happens if they are violated?
 - Then, C++ standards **do not define program/code behavior**.
- This case is referred to as "*undefined behavior*" (UB).
- **Examples:**
 - signed integer overflow, dereferencing invalid pointer (pointer that does not point to an existing object of a given type), accessing out-of-bound array element, creating an object in not properly aligned storage, data race, ...

2

Undefined behavior (cont.)

- Once some standard rule is violated, the C++ standard no longer define behavior on the abstract machine.
- In practice, this violation mostly happens at runtime, but it can also happen at compile time.
 - In the first case, the behavior is undefined for the executed program.
 - In the second case, the behavior is undefined even for compile-time C++ implementation tools (compiler, linker,...).
- **Example program source file:**

```
#include <iostream>
#include <limits>

int main() {
    int i = std::numeric_limits<int>::max();
    i++; // causes signed integer overflow => undefined behavior
    std::cout << i;
}
```

- **Question:** When does the overflow happen?

3

Undefined behavior (cont.)

```
int i = std::numeric_limits<int>::max();
i++; // causes signed integer overflow => undefined behavior
```

- When does the overflow happen?
 1. With disabled optimizations, the calculation will likely be performed at runtime \Rightarrow the overflow will happen at runtime \Rightarrow the program behavior is undefined at runtime once the overflow takes place.
 2. With enabled optimizations, the calculation will likely be performed at compile time \Rightarrow the overflow will happen at compile time \Rightarrow the behavior is undefined already at compile time.
- Ad 1. Once the overflow occurs, the behavior of the program is no longer defined by the C++ standard.
 - It may crash, it may behave as expected, it may behave unexpectedly,...
- Ad 2. Once the overflow occurs, the behavior of the implementation tools is no longer defined by the C++ standard.
 - It may throw a compilation error or warning, it may crash, it may generate some machine code with some expected or unexpected behavior,...

4

Undefined behavior (cont.)

- **Example:** GCC/x86_64/Linux implementation:
 - In both cases (enabled/disabled optimization) a compiler generated program machine code without any error/warning message.
 - Relevant part of generated machine code for enabled optimizations:

```
mov     esi, -2147483648
mov     edi, OFFSET FLAT:_ZSt4cout
call    std::basic_ostream<char, std::char_traits<char> >::operator<<(int)
```

- \Rightarrow The overflow occurred at compile time.
- Relevant part of generated machine code for disabled optimizations:

```
mov     DWORD PTR [rbp-4], 2147483647
add     DWORD PTR [rbp-4], 1
mov     eax, DWORD PTR [rbp-4]
mov     esi, eax
mov     edi, OFFSET FLAT:_ZSt4cout
call    std::basic_ostream<char, std::char_traits<char> >::operator<<(int)
```

- \Rightarrow The overflow occurs at runtime when add instruction is executed.

5

Undefined behavior (cont.)

```
int i = std::numeric_limits<int>::max();
i++; // causes signed integer overflow => undefined behavior
```

- With this code, overflow happens always (unconditionally) and may happen both at compile time or at runtime.
- What if we change this code to?

```
std::cin >> i;
i++; // causes overflow only if i equals max value of int
```

- Then, the overflow:
 - cannot happen at compile time (the calculation cannot be performed at compile time when the input value is not known);
 - happens only if 2147483647 is read from the standard input.
- **Conclusion:** The behavior of the program is now:
 - undefined only if 2147483647 is read from the standard input,
 - (well) defined otherwise (if user inputs 1, i will be incremented to 2).
- **Note:** This implies that, generally, undefined behavior cannot be detected by static code analysis (*more details later*).

6

Undefined behavior (cont.)

- The C++ standard does not define behavior for this program once the overflow occurs:

```
#include <iostream>
#include <limits>
int main() {
    int i = std::numeric_limits<int>::max();
    i++; // causes signed integer overflow => undefined behavior
    std::cout << i;
}
```

- Yet, a C++ implementation generated a program binary executable file with some machine code:

```
mov esi, -2147483648
mov edi, OFFSET FLAT:_ZSt4cout
call std::basic_ostream<char...
```

```
mov DWORD PTR [rbp-4], 2147483647
add DWORD PTR [rbp-4], 1
mov eax, DWORD PTR [rbp-4]
mov esi, eax
mov edi, OFFSET FLAT:_ZSt4cout
call std::basic_ostream<char...
```

- In both cases, execution renders same behavior:

```
$ ./a.out
-2147483648
```

7

Undefined behavior (cont.)

- Once some condition for undefined behavior is met, from the perspective of the C++ standard, **anything can happen** (*undefined = anything is allowed*).
- This "anything" includes also the possibility of the observed behavior...

```
$ ./a.out
-2147483648
```

- ...just it is not defined by the C++ standard.
- Instead, it is defined by a given implementation (in our case, by *GCC/x86_64/Linux*).
- Implication:**
 - When the behavior is implementation-defined, there is **no portability**.
 - With another C++ implementation / on other system / with other built setup / ..., the **observed behavior may be completely different**.

8

Undefined behavior (cont.)

- (Partial) summary:
- 1. When the **rules of the C++ standard are not violated**:
 - C++ *standard* defines behavior observable on the abstract machine.
 - C++ *implementation* must provide the same behavior observable on a particular computer system.
 - This must hold for all C++ implementations on all systems (*portability*).
- 2. When the **rules of the C++ standard are violated**:
 - C++ *standard* does not define any behavior on the abstract machine.
 - C++ *implementation* can do whatever it wants; typically, it provides some behavior.
 - This behavior is implementation-specific; generally, it is different for different implementation/systems (*no portability*).
- Conclusion:**
 - To have correctly-running portable programs, one should avoid undefined behavior as much as possible.

9

Perils of undefined behavior

- When a condition for undefined behavior is met, the actually-observed behavior (defined by a given C++ implementation) may correspond with some "**expected behavior**".

```
int i = std::numeric_limits<int>::max();
i++;
```

```
$ ./a.out
-2147483648
```

- Is this result expected?
 - Mathematically, no — adding two positive numbers cannot give negative result.
 - However, this outcome might be expected under the assumption of two's complement binary representation of integer numbers.
 - With other representations (such as *one's complement*, where negative zero exists), the result might be different.
- "Expected behavior" is **subjective** — different people may have different expectations.
- Frequent argumentation that "*the code is correct since it behaves as I expect*" is very dangerous and wrong once UB is involved.

10

Perils of undefined behavior (cont.)

- In our example, the implementation-specific behavior was always the same (hardcoded in the program machine code):

```
mov esi, -2147483648
mov edi, OFFSET FLAT:_ZSt4cout
call std::basic_ostream<char...
```

```
mov DWORD PTR [rbp-4], 2147483647
add DWORD PTR [rbp-4], 1
mov eax, DWORD PTR [rbp-4]
mov esi, eax
mov edi, OFFSET FLAT:_ZSt4cout
call std::basic_ostream<char...
```

```
$ ./a.out
-2147483648
```

- Once this machine code is generated, each execution of the program is guaranteed to have the same behavior.
- However, **generally, this does not hold**.
- Example:**
 - Undefined behavior caused by *data race* (unsynchronized read-write access to the same memory location from multiple threads).
 - Each time data race occurs, it may have different unpredictable impact on observable program behavior.

11

Perils of undefined behavior (cont.)

- Example:**

```
int i = 0; // global variable shared by threads
void inc_i() { i++; }
```

- When function `inc_i` is executed N times by each of T threads, one might **expect** that the value of `i` would be incremented by $N \times T$ in total.
- This "expected" behavior is undefined:
 - According to the C++ standard, this is *data race*, which causes undefined behavior.
- Note:** To get rid of data race / undefined behavior, `i` would need to have atomic data type, such as `std::atomic<int>`, instead of `int`.
- However, in practice, this expected behavior may be observed in most cases and, only occasionally, it may be different.

12

Perils of undefined behavior (*cont.*)

- Undefined behavior when the program mostly behaves as expected:
 - Extremely dangerous for production environments and mission-critical applications (*server applications, databases, banking, embedded systems in planes, cars, space crafts, ...*).
- Program can pass all tests and run without problems for long time, until once...:
 - In best cases, occasional unexpected behavior causes program to crash.
 - In worst cases, it continues to run with incorrect state.
- Once a problem is identified, it might be very hard to find its cause:
 - We would need to replicate this unexpected behavior during testing / debugging.
 - Since it happens only occasionally, it may be almost impossible.

33

13

Detection undefined behavior

- Can undefined behavior be detected?
 - We have seen, that this is generally impossible at compile time (with static code analysis).
- Is it possible at runtime?
 - In theory, maybe.
 - In practice, no.**
- Why?
 - Even if we could detect violation of C++ standard rules at runtime, it would **impose a large runtime and memory overhead into C++ programs**.
 - The purpose of programming in C++ is mostly **performance and efficiency**.
- ⇒ Generally, in programming, we can either have safety or performance. It's not possible to have both at once.
- C++ goes for performance.**

34

14

Detection undefined behavior (*cont.*)

- Example:** (library) source file (*left*), translation unit (*right*):

```
// deref_int.cpp:
int deref_int(int* ptr) {
    return *ptr;
}
```

```
int deref_int(int* ptr) {
    return *ptr;
}
```

- According to the C++ standard, dereferencing of a pointer (*ptr here) is valid only if this pointer actually points to an existing object of a given type (int here).
- If this requirement is not satisfied, the behavior is undefined once the dereferencing takes place.
- ⇒ To detect undefined behavior with this code, a C++ implementation would need to **test the "validity" of a pointer**.
- How?*

35

15

Detection undefined behavior (*cont.*)

- The most efficient representation of a pointer is just the **value of the address**.
- This representation allows generating the following optimal machine code (*GCC/x86_64/Linux*):

```
int deref_int(int* ptr) {
    return *ptr;
}
```

```
deref_int(int*):
    mov eax, DWORD PTR [rdi]
    ret
```

- A compiler has no idea what will be passed as a function argument.
- ⇒ Detection of UB would require some mechanism of how to find out whether a pointer actually points to an existing object.
- Is such a mechanism even possible? Likely not.
- Even if it was, it would have tremendous negative impact on performance:
 - When some object would be destroyed, all pointers that pointed to it would need to be informed about this destruction.
 - In the function, there would need to be a test for whether the pointer is valid or not.

36

16

Non-standard C++ extensions

- What happens in case of undefined behavior is mostly implementation-specific.
- Implementations can do anything they want and, typically, they do not care much about what happens:

```
int deref_int(int* ptr) {
    return *ptr;
}
```

```
deref_int(int*):
    mov eax, DWORD PTR [rdi]
    ret
```

- Observation:**
 - GCC does not care whether the passed pointer actually points to an existing object of type int.
 - This is a function user (caller) responsibility to guarantee that.
 - Machine code simply reads a 32-bit value from the passed address.
 - In case of undefined behavior, if this address is accessible, the function will just return some value.
 - Otherwise, it will likely cause the running program (process) to crash (with something like *segfault*).

37

17

Non-standard C++ extensions (*cont.*)

- In some (rare) cases, C++ implementations explicitly specify what happens even when C++ standard rules are violated.
- These cases form non-standard C++ extensions.
- Example:**

```
union U {
    int i;
    float f;
};
```

- Union is a class-like type.**
- In contrast to *classes* and *structs*, union's non-static member variables **share the same storage**:

```
struct S { // same with class
    int i;
    float f;
};
```

```
std::cout << sizeof(S) << sizeof(U); // prints '84' (x86_64/Linux)
```

38

18

Non-standard C++ extensions (cont.)

- C++ standard rules:
 - Only one union's member variable is so-called "active", namely the one most recently written into.
 - Only the active member variable may be read.
 - ⇒ Reading the inactive member variable = undefined behavior.

```
union U {
    int i;
    float f;
};

int main() {
    U u;
    u.f = 1.0f;
    return u.i; // undefined behavior according to C++ standards
}
```

- However, GCC explicitly allows this code and specifies a particular behavior for it:
 - "The relevant bytes of the representation of the object are treated as an object of the type used for the access."
 - ⇒ u.i has the value of an object of type `int` that has the same binary representation as the `float` object with value `1.0f`.

19

Non-standard C++ extensions (cont.)

- *GCC/x86_64/Linux*:
 - Floating-point values are represented in IEEE 754 format.
 - `float` is a single-precision IEEE 754 floating-point data type.
 - Value `1` has binary representation `0x3f800000` in hex with this data type.

```
int main() {
    U u;
    u.f = 1.0f;
    return u.i;
}
```

```
main:
    mov eax, 0x3f800000
    ret
```

- The behavior of the resulting program is:
 - undefined by the C++ standard,
 - but guaranteed by the used GCC implementation.
 - ⇒ This source code is generally not portable.
- *Similar case:*

```
int main() {
    float f = 1.0f;
    return *reinterpret_cast<int*>(&f);
}
```

```
main:
    mov eax, 0x3f800000
    ret
```



20

Object lifetime

- **Object lifetime** — time interval during which the object exist at runtime (when the program is run).
- *Simplified rules:*
- Lifetime begins with object initialization, which constructs the object.
- In case of class types:
 - object initialization involves calling a class constructor,
 - object lifetime ends when its destructor is called.
- Before an object can be constructed/initialized, a storage must be allocated for it.
 - Storage = place where a binary representation of the object will be stored.
 - *Recall* — this storage must be properly aligned and sized.



21

Storage allocations

- How is object storage allocated?
- Two separate cases:
 1. Object existence is known when the program is compiled:
 - For storage allocation for this object, a compiler is responsible.
 - This case is referred to as a "**static allocation**".
 - Meaning of *static* here = *resolvable-at-compile-time*.
 - *Effect* — allocation is hard-coded in the program (it's machine code).
 2. Object existence is not known until runtime:
 - Storage allocation must be *explicitly* initiated by a programmer.
 - This case is referred to as a "**dynamic allocation**".
 - Meaning of *dynamic* = *resolvable-not-until-runtime*.
 - *Effect* — allocation needs to be resolved at runtime by some external mechanisms (memory allocation functions).

22

Ambiguity of "static"

- "Static" vs "static" ambiguity:
 - Static allocations are not related to the `static` specifier/keyword:

```
void f() {
    int i = 1;
    static int j = 2;
    ...
}
```

- **Non-static** function-local variable `i`:
 - It's lifetime is restricted to the function body.
 - ⇒ Each time `f` is called the lifetime of `i` starts (with its initialization) and ends (when the function is left).
- **Static** function-local variable `j`:
 - It's lifetime basically covers the whole program run.
 - Typically, it starts when `f` is called first time (at the latest) and ends when the program is terminated.
- Storage for both `i` and `j` variables are statically allocated — allocation is "hard-coded" in the program at compile time.



23

Ambiguity of "static" (cont.)

- The word "static" can have many different meaning in C++.
- We have seen:
 - *Static entities* (functions, variables) with the meaning of being local to translation units with respect to linking (having *internal linkage*).
 - *Static class members* (functions, variables) with the meaning of belonging to the class itself instead of to its instances.
 - *Static allocations* with the meaning of being resolvable at compile time.
- The last meaning of the word "static" is used more generally.
- **Example** — static vs dynamic polymorphism:
 - *Polymorphism* = writing a unified code that can operate on objects of different types.
 - *Dynamic polymorphism* is provided by virtual member functions.
 - Function calls are resolved at runtime according to the actual (*dynamic*) object type.
 - *Static polymorphism* is provided by templates.
 - Templates are resolved at compile time (they do not exist at runtime).

24

Static allocations — example

- Translation unit (*left*) and generated machine code (*right*):

```
void g(int*, int*);
void f() {
    int i = 1;
    static int j = 2;
    g(&i, &j);
}
```

```
f():
    push    rax
    mov     dword ptr [rsp + 4], 1
    lea     rdi, [rsp + 4]
    mov     esi, offset f()::j
    call    g(int*, int*)
    pop     rax
    ret
f()::j:
    .long   2
```

- Storage for *i* and *j* variables is statically allocated:
 - Allocations in both cases are hard-coded in the program file / generated machine code (at compile time).
- Note:** observable behavior of *f* involves pointers to *i* and *j*.
 - Function *f* passes these pointers as arguments to function *g*.
 - ⇒ A compiler is "forced" to allocate storage for *i* and *j* in memory.
 - According to ABI, pointers are passed through *rdi* and *rsi* registers.
 - On the contrary, pointers *&i* and *&j* are not stored in memory.

25

Static allocations — example (cont.)

```
void g(int*, int*);
void f() {
    int i = 1;
    static int j = 2;
    g(&i, &j);
}
```

- Where in memory are *i* and *j* variables stored?
- C++ standards do not specify this; they "do not care".
- C++ standards only prescribe so-called "storage duration".
 - i* has "automatic storage duration" — storage is allocated when the code block is entered and deallocated when it is exited.
 - j* has "static storage duration" — storage is allocated when the program is started and deallocated when it is exited.
- Storage in memory (address space) is specified by a given C++ implementation.
- Implementations need to conform to a system ABI.

26

Static allocations — example (cont.)

```
void g(int*, int*);
void f() {
    int i = 1;
    static int j = 2;
    g(&i, &j);
}
```

```
f():
    push    rax
    mov     dword ptr [rsp + 4], 1
    lea     rdi, [rsp + 4]
    mov     esi, offset f()::j
    call    g(int*, int*)
    pop     rax
    ret
f()::j:
    .long   2
```

- Used implementation: *Clang/x86_64/Linux*:
 - Storage for non-static function-local variable *i* is allocated on the stack.
 - Allocation is effectively performed by "reserving" stack space by lowering the stack pointer register *rsp* (this is what the push instruction does).
 - Deallocation is performed by raising the stack pointer back to the original value (by using the pop instruction).
 - Storage for static function-local variable *j* is allocated on the so-called program data segment.
 - Memory for data segment is allocated when the program is executed and deallocated when it is exited.

27

Allocation and initialization

- Once there is a storage allocated for some object, its lifetime begins with its initialization.
- The object lifetime ends with its destruction, which takes place before its storage is deallocated.
- There are many forms of initialization, which, generally, have different effect for objects of different types.
- "Rough" distinction of two different cases.

- Class types:**
 - Initialization involves a constructor call.
 - Destruction involves a destructor call.
- Non-class types** (such as "fundamental" language-intrinsic types):
 - Initialization requires setting a binary representation to a desired value (determined by the initialization expression).
 - Destruction does not require any action.

28

Allocation and initialization (cont.)

```
void g(int*, int*);
void f() {
    int i = 1;
    static int j = 2;
    g(&i, &j);
}
```

```
f():
    push    rax
    mov     dword ptr [rsp + 4], 1
    lea     rdi, [rsp + 4]
    mov     esi, offset f()::j
    call    g(int*, int*)
    pop     rax
    ret
f()::j:
    .long   2
```

- In our case, *int* is a fundamental non-class type.
- i* is initialized by `= 1` expression:
 - It is initialized each time *f* is executed after its storage has been allocated.
 - Initialization = setting its binary representation to "have" value 1.
 - Its lifetime ends when the function is left (no action required).
- j* is initialized by `= 2` expression:
 - Initialization = setting its binary representation to "have" value 2.
 - This value is hard-code in the program file and loaded to the data segment when the program is executed.

29

Allocation and initialization (cont.)

- For a **class type**:
 - allocation/deallocation works the very same way,
 - initialization involves constructor, destruction involves destructor.

```
struct X {
    X(int);
    ~X();
};
void g(X*) noexcept;
void f() {
    X x(1);
    g(&x);
}
```

```
f():
    push    rbx, 16
    sub     rsp, [rsp + 8] ; address of the storage
    mov     rdi, rbx ; passed by rdi to constructor
    mov     esi, 1 ; initialization argument passed by rsi
    call    X::X(int) ; address of x passed to g by rdi
    mov     rdi, rbx
    call    g(X*) ; to destructor as well
    mov     rdi, rbx
    call    X::~X()
    add     rsp, 16
    pop     rbx
    ret
```

- Notes:**
 - A static function-local variable would be initialized only the first time *f* is called and destructed when the program exits.
 - Destructor for *x* would need to be called even if *g* threw an exception.
 - ⇒ In case of *not-noexcept* *g*, machine code would be more complicated.

30

Static vs dynamic allocations

- Static allocations — always take place when the program is run
⇒ can be hard-coded in the program machine code.
- In some cases, allocations cannot be resolved until runtime.
- *Examples:*
 - Conditional allocations dependent on some runtime-evaluable condition.
 - Allocations of storage for multiple objects, where their number is not known until runtime.
 - *Etc...*
- Such allocations need to be resolved at runtime, i.e., dynamically.
- *Example:*

```
void g(int*, int*);
void f(bool b, long n) {
    int* pi = b ? new int(1) : nullptr;
    int* pa = new int[n]();
    g(pi, pa);
}
```

31

Dynamic allocations

- “Dynamically-allocated objects” — the very same concept:
 - 1) First, a storage for an object needs to be allocated.
 - 2) Second, that object is in this storage initialized.
- Ad 1) How is a storage allocated in such a case?
 - Again, C++ standard do not specify this.
 - They only define so-called *allocation functions*.
 - These functions are called operator new and return a block of “uninitialized” memory of a desired size of bytes.
 - *Uninitialized* = there is no object in this memory yet.
 - How do these functions work is implementation-defined.
- *Typical implementations:*
 - Internally allocate memory on the heap (a system-provided mechanism for dynamic memory allocations).
 - ⇒ operator new internally calls some heap allocation function, such as malloc.

32

operator new vs new expression

- Allocation function operator new only provides storage for dynamically-allocated objects.
- Then, objects need to be in that storage constructed / initialized.
- Usually, both these steps are required (by programmers) to be performed together.
- ⇒ They are “bound” in a mechanism called *new expression*.
- new expression is the most common form of new.
- It is responsible for both:
 - 1) allocation storage for an object,
 - 2) its initialization in this storage.
- For allocation, it internally uses operator new.

33

operator new vs new expression (cont.)

- Example with a fundamental type...

```
void g(int*);
void f() {
    int* pi = new int(1); // new expression
    g(pi);
}
```

```
f():
    push    rax, 4
    call    operator new(unsigned long)
    mov     dword ptr [rax], 1
    mov     rdi, rax
    pop     rax
    jmp     g(int*)
```

- and with a class type:

```
void g(X*);
void f() {
    X* px = new X(1); // new expression
    g(px);
}
```

```
f():
    ...
    mov     edi, 1
    call    operator new(unsigned long)
    mov     rdx, rax
    mov     esi, 1
    call    X::X(int)
    ...
```

- *Observation:*

- new expression internally allocates storage with operator new,
- and then initializes the object in this storage.

34

Dynamic allocations (cont.)

- With *statically-allocated objects*:

- objects are automatically destructed,
- storage is automatically deallocated.

- The same does not hold for *dynamically-allocated objects*:

- object is not destructed (no destructor call),

- storage is not deallocated (no deallocation function call).

- Objects “manually” allocated and initialized with expression new need to be “manually” destructed and deallocated with its counterpart — expression delete.

```
struct X {
    X(int);
    ~X();
};
void g(X*);
noexcept;
void f() {
    X x(1);
    g(&x);
}
void g(X*);
void f() {
    X* px = new X(1);
    g(px);
}
f():
    push    rax, 16
    sub     rsp, 16
    lea     rdx, [rsp + 8]
    mov     rdi, rdx
    mov     esi, 1
    call    X::X(int)
    mov     rdi, rdx
    call    g(X*)
    mov     rdi, rdx
    call    X::~X()
    add     rsp, 16
    pop     rax
    ret
```

35

operator delete vs delete expression

- new expression:

- 1) allocates storage for an object,
- 2) initializes object in this storage.

- delete expression:

- a) destructs object,
- b) deallocates its storage.

- operator new:

- C++ allocation function internally called by new expression in ad 1).
- Typically implemented by calling some heap allocation function (such as malloc).

- operator delete:

- C++ deallocation function internally called by delete expression in ad b).
- Typically implemented by calling a heap deallocation function (free).

36

operator delete vs delete expr. (cont.)

```

void g(X*);
void f() {
    X* px = new X(1);
    g(px);
    delete px; // delete expression
}

```

- **Observation**—delete expression:

- destructs object by calling its destructor,
- deallocates its storage by calling C++ deallocation function operator delete.

```

f():
...
mov     edi, 1
call    operator new(unsigned long)
mov     rbx, rax
mov     rdi, rax
mov     esi, 1
call    X::X(int)
mov     rdi, rbx
call    g(X*)
mov     rdi, rbx
call    X::~X()
mov     rdi, rbx
add     rsp, 8
pop     rbx
pop     r14
jmp     operator delete(void*)

```

- **Notes:**

- (De)allocation functions operator new and operator delete are provided by the C++ standard library.
- \Rightarrow In practice, they are provided by library implementations.
- **Example**—operator new in *libc++* invokes `malloc`:
<https://github.com/llvm/llvm-project/blob/main/libcxx/src/new.cpp>.

37