

# Effective C++ Programming

NIE-EPC (v. 2021):

COSTS OF DYNAMIC ALLOCATIONS, ALIGNMENT,  
PADDING, PLACEMENT NEW, OWNING OBJECTS

© 2021 DANIEL LANGR, ČVUT (CTU) FIT

# Runtime costs of dynamic allocations

- Statically and dynamically allocated objects **differ only in the way how their storage is (de)allocated**.
  - Initialization and destruction is the same.
- **Dynamic allocations typically involve a *heap* — a very complex mechanism** that needs to be able to:
  - satisfy request for **allocations of different sizes**,
  - **communicate with a kernel** to ask for memory blocks (such as pages),
  - **synchronize request** in multi-threaded programs (which, generally, brings some overhead into single-threaded programs as well),
  - ...
- **⇒ Allocation of the storage on the heap is a relatively complex task**, when compared, for instance to **allocation on the stack**.
  - *Recall* — allocation on the stack involves just **decrementing the stack pointer register** (which is “super-fast”).

# Runtime costs of dynamic allocs (*cont.*)

- *Experiment* — **comparison of static and dynamic allocation** of a storage for an object of type `int`:

```
// I. static allocation experiment:  
int i = 1;  
... // make sure storage for i is allocated in memory (on the stack)
```

```
// II. dynamic allocation experiment:  
int* pi = new int(1);  
... // make sure storage for *pi is allocated in memory (on the heap)
```

- *Benchmark:*
  - Used tool — **Google Benchmark**, namely its online form **Quick C++ Benchmark**.
  - Repeated execution in the loop.
  - Relative comparison of the **average time of a single iteration**.
  - *Link:* [https://quick-bench.com/q/d\\_ENZrzvO\\_jwby2Qy7bisugNVr8](https://quick-bench.com/q/d_ENZrzvO_jwby2Qy7bisugNVr8).
- *Results:*
  - **Dynamic allocation** was **20-40x slower** than **static allocation**.



# Memory costs of dynamic allocations

- Is slower performance **the only drawback of dynamic storage allocation?**
- Unfortunately, **NO** :(
- Another disadvantage — **memory overhead**.
- Two causes:
  - 1) **Alignment:**
    - Padding caused by alignment produces **wasted memory bytes**.
  - 2) **Housekeeping data:**
    - Heap needs some **auxiliary data** to be stored in memory for each allocation to keep track of them.
- *Effect* — with typical heap implementations:
  - Each allocation — even of a single byte — **consumes at least 32 bytes** of memory on 64-bit systems.

WHAT?

# Storage alignment

- Storage for an object (= each object) of type  $T$  needs to be aligned to `alignof(T)`-byte address.
  - Alignment to an  $N$ -byte address means alignment to an address  $A$ , where  $A = 0 \pmod{N}$ .
- Alignment requirements for particular types are **implementation-defined**.
- *Example — GCC/Linux/x86\_64:*

```
std::cout << alignof( char    ) ; // prints out "1"
std::cout << alignof( int     ) ; // prints out "4"
std::cout << alignof( long    ) ; // prints out "8"
std::cout << alignof( float   ) ; // prints out "4"
std::cout << alignof( double  ) ; // prints out "8"
std::cout << alignof( long double ) ; // prints out "16"
std::cout << alignof( void*   ) ; // prints out "8"
```

- *Class types* — alignment requirements typically equal the **maximum of alignment requirements for types of its subobjects**.
- *Note: Subobjects* = non-static member variables and base class objects.



# Storage alignment (*cont.*)

- **Fundamental alignment** (*simplified definition*):
  - Alignment less than or equal to the maximum alignment for “fundamental” language-intrinsic types.
  - Is equal to the `alignof(std::max_align_t)` expression.
  - In typical C++ implementations, it equals alignment requirements for the `long double` type:

```
std::cout << alignof( long double      ); // prints out "16"  
std::cout << alignof( std::max_align_t ); // prints out "16"
```

- **Over-aligned types:**
  - Types with alignment requirements that are higher than the maximal fundamental alignment.
  - *Exemplary use cases:*
    - *Cache line aligned data* (efficiency reasons, avoiding *false sharing*,...) — typically, 64-byte alignment.
    - *SIMD-processed data* (data processed by vectorization instructions), for example — AVX2 (32-byte alignment), AVX512 (64-byte alignment),...

# Statically-allocated storage alignment

- **How is alignment provided** when a storage for an object is allocated?
- 1) **Function non-static local variables** (*automatic storage duration*) + typical **stack-based implementation**:
  - *Recall* — storage allocation = **decreasing stack pointer (SP) register**.
  - **How much it needs to be decreased?**
    - For example, decreasing by 4 bytes **does not guarantee that SP will point to a 4-byte aligned address**.
    - Generally, **calculation of a properly aligned address would need additional instructions**.
    - This would impose into functions **additional (static) allocation overhead**.
  - **Common real-world solution of this problem**:
    - ABI prescribes that when the function is executed, **SP must point to an  $A$ -byte aligned address**, where  $A$  is defined in that ABI.
    - *Example* — *Linux/x86\_64* —  $A = 16$ :
      - When the `call` instruction calls some function, **RSP must be 16-byte aligned**.

# Statically-allocated storage alignment (*cont.*)

- *Example:*

```
void g(int*);  
void f() {  
    int i = 1;  
    g(&i);  
}
```

```
f():  
    push    rax  
    mov     dword ptr [rsp + 4], 1  
    lea     rdi, [rsp + 4]  
    call    g(int*)  
    pop     rax  
    ret
```

- When the function `f` is called somewhere, `rsp` is guaranteed to be **16-bytes aligned**  $\Rightarrow$  `rsp = 0 modulo 16`.
- The corresponding `call` instruction **decreases `rsp` by 8** (it “pushes” the return instruction address to the stack).
  - $\Rightarrow$  **When the function `f` is started, `rsp = 8 modulo 16`.**
- Before `g` is called, **`rsp` needs to be again 16-byte aligned**.
  - $\Rightarrow$  **`rsp` needs to be decreased by another 8 bytes.**
- Those 8 bytes are used **for allocation of storage for `i`**:
  - They are **guaranteed to be 8-byte aligned**, `int` requires 4-byte alignment.
  - $\Rightarrow$  Storage for `i` **can be allocated in upper 4 bytes or lower 4 bytes**.
  - In our case, **compiler decided to use the upper 4 bytes (`rsp+4`)**.
  - Lower 4 bytes will remain **unused** (wasted).



# Statically-allocated storage alignment (*cont.*)

- *Modified example:*

```
void g(long*);  
void f() {  
    long i = 1;  
    g(&i);  
}
```

```
f():  
    push    rax  
    mov     qword ptr [rsp], 1  
    mov     rdi, rsp  
    call    g(long*)  
    pop     rax  
    ret
```

- Storage for an object of type **long** takes 8 bytes and requires 8-byte alignment  $\Rightarrow$  there are no bytes wasted.
- *Summary:*
  - Thanks to the **ABI-required stack pointer alignment** when functions are called, (static) allocations of storage for function-local variables can be performed **extremely efficiently** — namely, by **decreasing stack pointer**.
  - This holds for all *non-over-aligned* types.
  - For *over-aligned* types, some **additional calculation at runtime is still needed** (for instance, 16-byte alignment may or may not be 32-byte aligned at the same time).

# Dynamically-allocated storage alignment

- How is alignment provided when a storage for an object is allocated (*cont.*)?
- 2) Dynamically-allocated objects (*dynamic storage duration*) + typical heap-based implementation:
  - Recall — storage allocation = calling some heap allocation function (through operator `new` C++ allocation function).
  - Typically, `malloc` is used.
- `malloc` takes a single argument — number of allocated bytes.
  - ⇒ It allows a caller to specify the storage (byte) size :).
  - ⇒ It does not allow a caller to specify alignment requirements :(.
- Common real-world solution of this problem:
  - Each dynamically-allocated block of memory is guaranteed to be aligned at least to the *maximal fundamental alignment*.

# Dynamically-allocated storage alignment (cont.)

- Our *Linux/x86\_64* system:
  - *Heap implementation* is (by default) provided by *GNU C library* (*GLIBC*).
    - This is a common case in Linux.
  - *Maximum fundamental alignment* `alignof(std::max_align_t)` is **16**.
  - Accordingly, `malloc`-allocated memory blocks are **guaranteed to be 16-byte aligned**.
    - Note — heap implementations **may guarantee even higher/stronger alignment** (for details, see `__STDCPP_DEFAULT_NEW_ALIGNMENT__` C++ macro).
- What about over-aligned types?
  - Since **C++17**, operator `new` has overloaded versions where **alignment requirements may be specified**.
    - Internally, **other function than `malloc` needs to be used** by a C++ standard library implementation (such as the `posix_memalign` allocation function).
  - **Before C++17**, there was **no portable way to dynamically-allocate storage for over-aligned types**.

# Memory costs of dynamic allocations (*cont.*)

- Consequences (*our platform*):
  - All dynamic allocations are **16-byte aligned**.
  - $\Rightarrow$  Even if a storage for an object of some type **requires only 1 byte**, its dynamic allocation will **effectively consume 16 bytes**.
  - The **remaining bytes are wasted**.
- Moreover, heap implementations **require for each allocation some housekeeping data**  $\Rightarrow$  additional memory overhead.
  - In our case, these data take **another 16 bytes per allocation**.
- *Experiment:*

```
void g(int*); // NOOP function defined in another TU

int main() {
    for (int i = 0; i < 100'000'000; i++) {
        int* pi = new int{};
        g(pi);
    } }
```

- “Effective data” = **100M objects** of type `int`  $\Rightarrow$  **400M bytes** of storage.
- **Measured memory consumption** (*maximum RSS*): **3200M bytes**.
- $\Rightarrow$  **Each allocation consumes 32 bytes**:
  - **4 effective bytes** (storage of `int` object), **12 bytes wasted**, **16 bytes housekeeping**.

# Intermezzo — padding

- *Previous observation* — alignment can cause wasted bytes in memory in case of both:
  - static storage allocation (on the *stack*),
  - dynamic storage allocation (on the *heap*).
- These bytes are *out of the allocated storage* itself.
- Alignment requirements can cause wasted bytes even inside storage of a single object:
  - In the storage of a class type object, its *subobjects* are stored (*base class objects* and *non-static member variables*).
  - Due to alignment requirements, they may not be store next to each other in memory.
  - “Gaps” between subobjects are referred to as *padding*.
  - Padding represents wasted bytes inside the storage of a single object.



# Intermezzo — padding (cont.)

- Example — Clang/x86\_64/Linux:

```
struct X { void* v; char c; int i; float f; };
```

- Storage for an object of type X starts at address  $A$ .
- $X::v$  requires 8-byte aligned storage, which is 8 bytes long.
  - $\Rightarrow A$  is 8-byte aligned.
  - $\Rightarrow X::v$  is stored from address  $A$  to  $A + 7$ .
- $X::c$  requires 1-byte aligned storage, which is 1 byte long.
  - $\Rightarrow X::c$  is stored at address  $A + 8$  only.
- $X::i$  requires 4-byte aligned storage, which is 4 byte long.
  - $\Rightarrow X::i$  cannot be stored next to  $X::c$  at address  $A + 9$ ; this address is not 4-byte aligned.
  - $\Rightarrow$  The lowest 4-byte aligned address where  $X::i$  may be stored is  $A + 12$ .
  - $\Rightarrow X::i$  is stored from address  $A + 12$  to  $A + 15$ .
  - Bytes from  $A + 9$  to  $A + 11$  represent padding (are wasted).
- $X::f$  requires 4-byte aligned storage, which is 4 byte long.
  - $\Rightarrow X::f$  is stored from address  $A + 16$  to  $A + 19$ .

# Intermezzo — padding (cont.)

```
struct X { void* v; char c; int i; float f; };
```

- *Previous analysis:* Subobjects — member variables —  $X::v$ ,  $X::c$ ,  $X::i$ , and  $X::f$  are stored at addresses  $A$  to  $A + 19$ .
- If two objects of type  $X$  are stored in an array next to each other, and the first one is stored at address  $A$ , is the next element stored at  $A + 20$ ? **NO!**
  - Its first member  $X::v$  needs to be 8-byte aligned.
  - $\Rightarrow$  The lowest available 8-byte aligned address higher than  $A + 19$  is  $A + 24$ .
- Those 4 wasted bytes from  $A + 20$  to  $A + 23$  at the end also represent padding.
  - Even this final padding is a part of a binary representation of  $X$ .
  - $\Rightarrow$  These 4 bytes are part of the storage for each object of type  $X$ .
- *Proof:*

```
std::cout << sizeof(X); // prints out 24
```

- $\Rightarrow$  Storage of  $X$  consists of 17 data bytes and 7 padding bytes.

# Intermezzo — padding (cont.)

- *C++ rule:*
  - Class **member variables with same access rights** (*private, protected, public*) **must be stored at growing addresses in order of their declaration.**
  - $\Rightarrow$  Implementation **may not reorder** their placement in the storage of the class object.
- $\Rightarrow$  Padding **may be sometimes reduced by “better” order of (non-static) member variable declarations.**
- *Example:*

```
struct Y {  
    char c;  
    int i;  
    bool b;  
    float f;  
};
```

```
std::cout << sizeof(Y); // prints out 16
```

```
struct Z {  
    char c;  
    bool b;  
    int i;  
    float f;  
};
```

```
std::cout << sizeof(Z); // prints out 12
```

# Static vs dynamic allocations

- *Summary* — in comparison with static allocations, allocation of storage for dynamically-allocated objects:
  - is much slower,
  - consumes much more memory.
- Efficiency with *collection of objects*:
  - = objects stored in *dynamic data structures*.
  - C++ standard library defines different types of such data structures called *containers*.
  - Comparison of containers with respect to the number of required allocations for insertion of  $n$  objects:
    - `std::vector` without reserving space  $\rightarrow \log(n)$ ,
    - `std::vector` with reserving space  $\rightarrow 1$ ,
    - `std::list`  $\rightarrow n$ ,
    - `std::set` or `std::map`  $\rightarrow n$ ,
    - `std::unordered_set` or `std::unordered_map`  $\rightarrow n + \log(n)$ .

# Static vs dynamic allocations (*cont.*)

- *Experimental evaluation:*
  - Comparison of time and memory consumption when the same number of objects is inserted into different C++ library containers.
  - Measurements are *normalized* to the case where a vector was used without reserving memory.

container	time	memory (RSS)
std::vector "as-is" ( <i>baseline</i> )	1	1
std::vector + reserve	2.8× faster	1.3× lower
std::list	17× slower	5.9× higher
std::set	95× slower	8.9× higher
std::unordered_set	46× slower	7.4× higher

- *Analysis:*
  - Vector is so efficient since it allocates storage for multiple objects (its elements) at once.
  - All other containers are *node-based* — each new element requires allocation of a separate storage for a single node, in which it is then stored.



# Static vs dynamic allocations (*cont.*)

- **Vector** stores its elements **contiguously in memory**.
- **Nodes** are placed in memory at **unrelated locations**.
- $\Rightarrow$  Vector provides efficient element access.
- **Vector** vs **set/map**:
  - **Set/map** = binary search tree  $\Rightarrow O(\log(n))$  lookup time.
  - Sorted vector with binary search  $\Rightarrow O(\log(n))$  lookup time as well.

container	insertion (+sort)	lookup
std::vector "as-is" ( <i>baseline</i> ) + sorting	1	1
std::vector + reserve + sorting	1.2× faster	1
std::set	16× slower	2.1× slower
std::unordered_set	7.3× slower	4.6× faster

- **Fastest lookup** provides **unordered set** = *hash table*.
  - $\Rightarrow O(1)$  lookup time.
  - At a price of significantly **slower data structure construction** and **higher memory consumption**.

# Placement new

- Up to now, all mechanisms we have seen for both static and dynamic object allocations **automatically “attached”**:
  - **storage allocations** together with **object initialization**,
- and:
  - **object destruction** together with **storage deallocation**.
- Vector container:
  - Allocates storage **for multiple objects** (its elements) **at once**.
  - ⇒ Needs to **“detach” storage allocation from object initialization** (as well as object destruction from storage deallocation).
- How can this be done in C++?
  - For storage (de)allocation **alone**, we have **(de)allocation functions** `operator new` and `operator delete`.
  - ⇒ What we need is to **initialize object in this storage** (*uninitialized memory*), and — in the end — to **destruct it**.

# Placement new (cont.)

- Vector — highly complex data structure.
- *Simpler case for the sake of explanation:*
  - Optional object existence — object that optionally may or may not exist at runtime.
- In case of dynamically-allocated object, the solution is simple:
  - A pointer itself can contain information whether the object exist or not:

```
int* pi = (some_condition == true) ? new int(1) : nullptr;
```

- Or, better:

```
std::unique_ptr<int> pi = (some_condition == true) ? new int(1) : nullptr;
```

- But what if we want to achieve the same without the overhead of dynamic storage allocation?
- $\Rightarrow$  We need optional object initialization in a statically-allocated storage.

# Placement new (cont.)

- **Class wrapper** for such an optionally existing object:

```
class optional_int {  
    ??? buffer_[ ??? ];  
    bool exist_;  
public:  
    optional_int() : exists_(false) { }  
    optional_int(int i) : exists_(true) { ??? }  
    operator int&() { return ???; }  
};
```

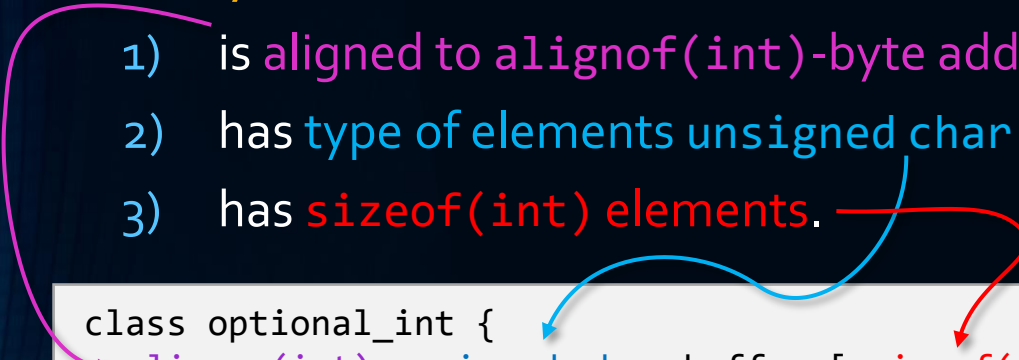
- **Storage for optional object** is a class member variable.
  - $\Rightarrow$  It is a part of the storage of the class object itself.
  - $\Rightarrow$  **No need to dynamically allocate it.**
- This storage needs to be:
  - **Large enough** to hold an object of type `int`  $\Rightarrow$  **`sizeof(int)` bytes long.**
  - **Aligned suitably** to hold an object of type `int`  $\Rightarrow$  **`alignof(int)` byte-aligned.**

# Placement new (cont.)

- Such storage can be provided as a **buffer of a form of an array**, which:

- 1) is **aligned to alignof(int)-byte address**,
- 2) has **type of elements unsigned char or std::byte** (since **C++17**),
- 3) has **sizeof(int) elements**.

```
class optional_int {  
    alignas(int) unsigned char buffer_[ sizeof(int) ];  
    bool exist_  
public:  
    optional_int() : exists_(false) { }  
    optional_int(int i) : exists_(true) { ??? }  
    operator int&() { return ???; }  
};
```



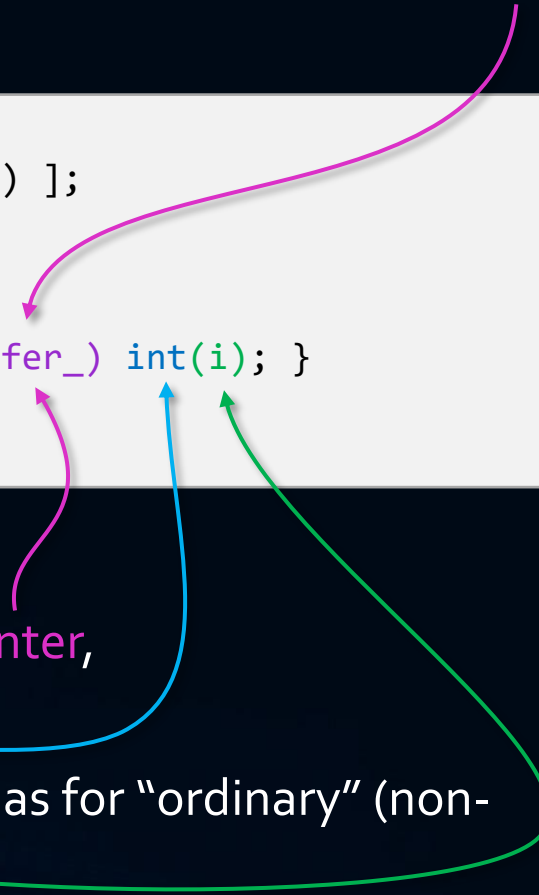
- Such a storage is now **suitable to hold an object of type int**.
- Next, we need to:
  - Construct** (=initialize) such an object in this storage (buffer).
  - Provide access** to it (to allow its update and reading).



# Placement new (cont.)

- Initialization of an object in existing storage = *placement new*.
  - Another form of new.
  - Basically, it is a new expression with an additionally provided pointer as its “argument”.

```
class optional_int {  
    alignas(int) unsigned char buffer_[ sizeof(int) ];  
    bool exist_;  
public:  
    optional_int() : exists_(false) { }  
    optional_int(int i) : exists_(true) { new (buffer_) int(i); }  
    operator int&() { return ???; }  
};
```



- Placement new initializes an object:
  - in memory at the address passed as a pointer,
  - of the specified type,
  - by initialization expression the same way as for “ordinary” (non-placement) new expression.

# Placement new (cont.)

- **Object access:**
  - There is **no direct access** (through some name/identifier).
  - But we know, where it is stored  $\Rightarrow$  we can **derive its pointer** (a *pointer-to-this-object*).
  - `buffer_` is an array  $\Rightarrow$  is **implicitly convertible to a pointer-to-unsigned char**.
  - We **need a pointer-to-int**.
  - $\Rightarrow$  *Solution* = **cast (conversion)** of these pointers:

```
class optional_int {
    alignas(int) unsigned char buffer_[ sizeof(int) ];
    bool exist_;
    int* ptr() { return reinterpret_cast<int*>(buffer_); } // helper function
public:
    optional_int() : exist_(false) { }
    optional_int(int i) : exist_(true) { new (buffer_) int(i); }
    operator int&() { return *ptr(); }
};
```

- *Note:* since the operator `int&()` **returns a reference to the stored object**, the **casted pointer is dereferenced**.

# Placement new (cont.)

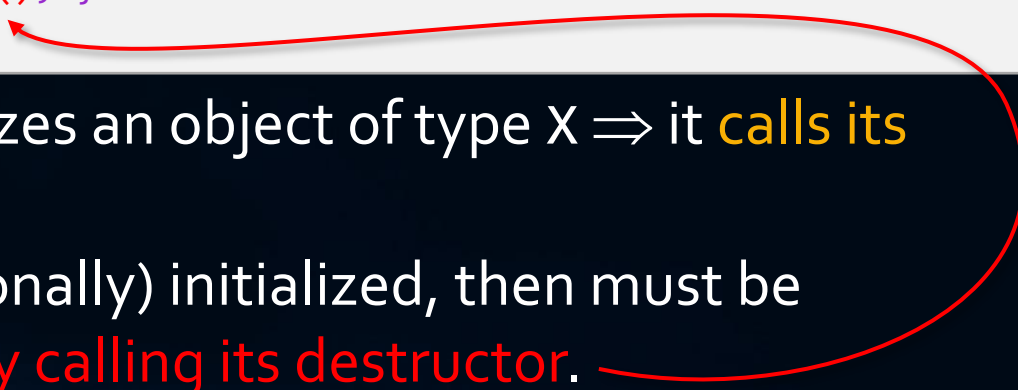
- Object destruction:
  - *Placement new* initializes an object — starts its lifetime.
  - **How to end this lifetime?**
  - Is there any “*placement delete*”? **No.**
- Two different cases:
  - 1) *Class types* — objects lifetime needs to be ended explicitly by calling their destructors.
  - 2) *Non-class types* — objects lifetime ends automatically when their storage is deallocated.
- `int` belongs to **ad 2)**
  - ⇒ **No explicit action is needed** for destruction of the object initialized by *placement new*.
  - ⇒ The class `optional_int` **does not require manual definition of a destructor** (it is *auto-generated*).
- What about **ad 1)** cases?

# Placement new (cont.)

- **Class wrapper** for such an optionally existing **class object**:

```
struct X {
    X(int);
    ~X();
};

class optional_X {
    alignas(X) unsigned char buffer_[ sizeof(X) ];
    bool exist_;
    X* ptr() { return reinterpret_cast<X*>(buffer_); }
public:
    optional_X() : exists_(false) { }
    optional_X(int i) : exists_(true) { new (buffer_) X(i) }
    operator X&() { return *ptr(); }
    ~X() { if (exist_) ptr()->~X(); }
};
```



- Placement new initializes an object of type X  $\Rightarrow$  it **calls its constructor**.
- If the object was (optionally) initialized, then must be **explicitly destructed by calling its destructor**.
  - Destructor can be called **as any other member function**.

# Placement new (cont.)

- Comparison of *non-class* (int; left) and *class* (X; right) cases:

```
optional_int::optional_int(int):  
    mov    BYTE PTR [rdi+4], 1  
    mov    DWORD PTR [rdi], esi  
    ret  
optional_int::~~optional_int():  
    ret
```

```
optional_X::optional_X(int) :  
    mov    BYTE PTR [rdi+1], 1  
    jmp    X::X(int)  
optional_X::~~optional_X() :  
    cmp    BYTE PTR [rdi+1], 0  
    jne    .L10  
    ret  
.L10:  
    jmp    X::~~X()
```

## 1. Non-class case:

- *Initialization* of an object of type int requires just setting its memory representation to a desired value (value of constructor argument passed in the esi register).
- *Destruction* does not require any action.

## 2. Class case:

- *Initialization* of an object of type X requires calling its constructor.
- *Destruction* requires calling its destructor.



# Placement new (cont.)

- Problem with explicit destructor call:

```
class optional_string {  
    alignas(std::string) unsigned char buffer_[ sizeof(std::string) ];  
    bool exist_  
    std::string* ptr() { return reinterpret_cast<std::string*>(buffer_); }  
public:  
    optional_string() : exists_(false) { }  
    optional_string(char* s) : exists_(true) { new (buffer_) std::string(s) }  
    operator std::string&() { return *ptr(); }  
    ~optional_string() { if (exist_) ptr()->~string(); } // compilation error  
};
```

```
error: expected class-name before '(' token  
    ~optional_string() { if (exist_) ptr()->~string(); }
```

- `std::string` = a type alias for an instance of `std::basic_string<char>` class template.
- $\Rightarrow$  There is no destructor named `~string()`.
- $\Rightarrow$  The destructor is called `~basic_string()`.

```
~X() { if (exist_) operator X&().~basic_string(); } // OK
```

# Placement new (*cont.*)

- Generally, this is a problem, if we do not know **whether some type is a type name or a type alias**.
- First possible solution — **templates**:

```
template <typename T>
class optional {
    alignas(T) unsigned char buffer_[ sizeof(T) ];
    bool exist_;
    T* ptr() { return reinterpret_cast<T*>(buffer_); }
public:
    optional() : exists_(false) { }
    ??? // constructor – to be explained later

    operator T&() { return *reinterpret_cast<T*>(buffer_); }
    ~optional() { if (exist_) operator ptr()->~T(); }
};
```

- Generic solution — **will work for any type** (template argument).
- For **class types**, **correct destructor name** is “encoded” in `~T()` call;
- For **non-class types**, **there is no destructor**, such as `~int()`.
  - In such case, `~T()` call is valid, but will have no effect.
  - This mechanism is called **pseudo-destructor call** and **allows us to write unified generic code** where objects need to be explicitly destructed.

# Teaser — perfect forwarding

- How to write a **constructor of optional**?
- Constructors of T may take **any number of arguments having different types and different value categories**.
- $\Rightarrow$  We need a constructor of optional class that will:
  - take **any number of argument of any types and any value categories**,
  - and “pass” them *as-they-are* to the constructor of T.
- This can be solved by combining three fundamental modern C++ techniques:
  - *variadic templates*,
  - *forwarding references*,
  - *std::forward function*,
- which is together called **perfect forwarding**.
- *More: later lectures.*

```
template <typename T>
class optional {
    ...
    template <typename Ts...>
    optional(Ts&&... args) : exists_(true)
    {
        new (buffer_) T(std::forward<Ts>(args)...)
    }
    ...
};
```

# Placement new (cont.)

- Second possible solution: `std::destroy_at()` library function:

```
~optional_string() { if (exist_) std::destroy_at(ptr()); }
```

- `std::destroy_at()` is a **function template**.
- It **derives the type name** (template parameter T) from the **type of the passed pointer** (*template argument deduction*).
  - Internally, it calls it destructor as `~T()`.
- **Advantage** — the code is **more readable and explicit**.
- **Counterpart** of `std::destroy_at` is `std::construct_at`:

```
optional_string(char* s) : exists_(true) {  
    std::construct_at<std::string>(ptr(), s);  
}
```

- Internally, it **constructs an object with placement new**:
  - its **template argument** represents a **type of the constructed object**,
  - its **first function argument** is a **pointer to the storage**,
  - **all other function arguments** are **perfectly forwarded to the initializer**.
- **Availability**: `std::destroy_at` **C++17**, `std::construct_at` **C++20**.

# Placement new (cont.)

- Using **standard library** improves readability and understandability of code; *seen examples*:
  - **placement new** vs `std::construct_at`,
  - **explicit destructor call** vs `std::destroy_at`.
- Another useful replacement from the standard library for our optimal class — **`std::aligned_storage`**.
  - Class (struct) template, which member type “type” provides a **type of storage that is aligned to the desired value and have desired length**.
  - Length and alignment are **specified as template arguments**.

```
alignas(T) unsigned char buffer_[ sizeof(T) ];
```

- can be replaced by:

```
typename std::aligned_storage<sizeof(T), alignof(T)>::type buffer_;
```

- or, **since C++14**, by:

```
std::aligned_storage_t<sizeof(T), alignof(T)>::type buffer_; // shortcut
```

# Placement new (cont.)

- optional class in C++20 with **library entities**:

```
template <typename T>
class optional {
    std::aligned_storage_t<sizeof(T), alignof(T)> buffer_;
    bool exist_;
    T* ptr() { return reinterpret_cast<T*>(buffer_); }
public:
    optional() : exist_(false) { }

    template<typename... Ts>
    optional(Ts&&... args) : exist_(true) {
        std::construct_at<T>(ptr(), std::forward<Ts>(args)...); }

    operator T&() { return *ptr(); }
    ~optional() { if (exist_) std::destroy_at(ptr()); }
};
```

- Such a class template in **a more sophisticated form** is available in the **C++ standard library** itself **since C++17** as **std::optional<T>**.
- *Common use case* — **optional returning** an object from a function:

```
std::string f1();    // (1) return object always
std::string* f2();   // (2) may return optionally, but requires heap allocation
std::optional<std::string> f3(); // (3) may return optionally and...
                                // ...does not require heap allocation
```



# Placement new (cont.)

- Drawback of (3) vs (1):
  - An optional class **needs a flag** with the **information about whether it does or does not “hold” an object.**
  - This flag, for example, **controls destruction.**

```
template <typename T>
class optional {
    bool exist_;
    ...
    optional() : exist_(false) { }
    template<typename... Ts> optional(Ts&&... args) : exist_(true) { ... }
    ~optional() { if (exist_) std::destroy_at(ptr()); }
};
```

- Boolean flag generally **requires only 1 byte.**
- However, due to alignment requirements, **there may be some padding introduced.**
- Flag + padding represent **“housekeeping” memory overhead:**

```
std::cout << sizeof( std::string );           // prints out 32
std::cout << sizeof( std::optional<std::string> ); // prints out 40 (+1B flag, +7B padding)
```

# “Owning” objects — “owners”

- An object of type `std::optional<T>` — a “wrapper” that (optionally) “owns/manages” an object of type `T`.
  - $\Rightarrow$  The “content” of `std::optional<T>` is an (optional) ownership of an object of type `T`.
- The same holds, for instance, for an object of type `std::unique_ptr<T>`.

```
std::optional< std::string > os( "..." );  
std::unique_ptr< std::string > us( new std::string( "..." ) );
```

- *Further assumption* — stack/heap-based C++ implementation.
- Difference between `optional` and `unique_ptr`:
- `unique_ptr` — the storage of the owned object is **allocated dynamically**.
  - $\Rightarrow$  The owned object itself is stored on the heap.
- `optional` — the storage of the owned object is “included” in the storage of the owning (`optional`) object itself.
  - $\Rightarrow$  For instance, if the `optional` object is stored on the stack, so is also the owned object.

# “Owners” (cont.)

- *Physical ownership* — owned object is stored in the storage of its owner (= *included storage*).
- *Logical ownership* — owner takes care about owned object(s) (manages lifetime etc.).
- `std::optional<T>` — both *physical* and *logical* ownership:
  - Object is, for example, initialized and destructed by its optional owner.
  - *Included storage* — example with GCC/x86\_64:

```
std::optional<int> oi(1);  
std::cout << &oi;           // address of owner: "0x7ffc2e7146a8"  
std::cout << &oi.value();    // address of owned: "0x7ffc2e7146a8"
```

- ⇒ The owner `oi` is stored on the *stack*, so is the owned object.
- ⇒ Storage size of the optional owner depends on the storage size of the owned object.

```
std::cout << sizeof( std::optional< int > );    // "8"  
std::cout << sizeof( std::optional< std::string > ); // "40"
```

# "Owners" (cont.)

- `std::unique_ptr<T>` — *logical ownership only*:
  - Object is, for example, *destroyed* by its `unique_ptr` owner.
  - *Dynamically allocated storage*:

```
std::unique_ptr<int> ui(new int(1));  
std::cout << &ui;    // address of owner: "0x7ffe276c8450"  
std::cout << &*ui;   // address of owned: "0x23faec0"
```

- $\Rightarrow$  The owner `ui` is stored on the *stack*, but the *owned object* is stored on the *heap*.
- $\Rightarrow$  Storage size of the `unique_ptr` owner *does not depend* on the storage size of the owned object.

```
std::cout << sizeof( std::unique_ptr< int > );    // "8"  
std::cout << sizeof( std::unique_ptr< std::string > );    // "8"
```

- *Note:* `std::unique_ptr<T>` is typically implemented as a *wrapper* of the *raw pointer member variable* of type `T*`.

# “Owners” (*cont.*)

- *C++ standard library* provides multiple generic “owner” types.
- Some *single-object owners*...
  - ...with *included storage*:
    - `std::optional<T>`, `std::variant<T1,T2,...>`;
  - ...with *dynamically-allocated storage*:
    - `std::unique_ptr<T>`, `std::shared_ptr<T>`, `std::weak_ptr<T>`.
- Some *multiple-object owners*...
  - ...with *included storage*:
    - `std::pair<T1,T2>`, `std::tuple<T1,T2,...>`, `std::array<T,N>`;
  - ...with *dynamically-allocated storage*:
    - *dynamic containers* — `std::vector<T>`, `std::list<T>`, `std::set<T>`,...
- What about `std::string`?
  - An object of type `std::string` can handle/own a *string of characters of any length*, generally unknown at compile time.
    - $\Rightarrow$ ? **Must** the *string of characters* be stored in *dynamically-allocated storage*?

# Owners — std::string case

- *Experiment — Clang/x86\_64/libc++:*

```
std::string s("...");  
std::cout << &s           // address of the string object:      "0x7ffef075e4c0"  
  << (void*)s.data();     // address of the string-of-characters: "0x7ffef075e4c1"
```

- *Observation:*
  - The **owner** `s` is stored **on the stack**.
  - The **string of characters** `"..."` managed/owned by `s` is stored also on the **stack**.
  - Namely, it is stored **in the included storage** of its owner `s`.
- *Similar experiment:*

```
std::string s("... some different string ...");  
std::cout << &s           // address of the string object:      "0x7fffd772d280"  
  << (void*)s.data();     // address of the string-of-characters: "0x18bf2a0"
```

- *Observation:*
  - Now, the **string of characters** `"... some different string ..."` is stored **on the heap** (in a **dynamically allocated storage**).
- *Explanation* — see next presentation.