

# Effective C++ Programming

NIE-EPC (v. 2021):  
COSTS OF DYNAMIC ALLOCATIONS, ALIGNMENT,  
PADDING, PLACEMENT NEW, OWNING OBJECTS  
© 2021 DANIEL LANGR, ČVUT (CTU) FIT

1

## Runtime costs of dynamic allocations

- Statically and dynamically allocated objects differ only in the way how their storage is (de)allocated.
  - Initialization and destruction is the same.
- Dynamic allocations typically involve a *heap* — a very complex mechanism that needs to be able to:
  - satisfy request for allocations of different sizes,
  - communicate with a kernel to ask for memory blocks (such as pages),
  - synchronize request in multi-threaded programs (which, generally, brings some overhead into single-threaded programs as well),
  - ...
- ⇒ Allocation of the storage on the heap is a relatively complex task, when compared, for instance to allocation on the stack.
  - Recall* — allocation on the stack involves just decrementing the stack pointer register (which is "super-fast").

2

## Runtime costs of dynamic allocs (cont.)

- Experiment* — comparison of static and dynamic allocation of a storage for an object of type `int`:

```
// I. static allocation experiment:
int i = 1;
... // make sure storage for i is allocated in memory (on the stack)
```

```
// II. dynamic allocation experiment:
int* pi = new int(1);
... // make sure storage for *pi is allocated in memory (on the heap)
```

- Benchmark:**
  - Used tool — *Google Benchmark*, namely its online form *Quick C++ Benchmark*.
  - Repeated execution in the loop.
  - Relative comparison of the average time of a single iteration.
  - Link: [https://quick-bench.com/q/d\\_ENZrzo\\_jwby2Qy7bisugNVr8](https://quick-bench.com/q/d_ENZrzo_jwby2Qy7bisugNVr8).
- Results:**
  - Dynamic allocation was **20-40x slower** than static allocation.

3

## Memory costs of dynamic allocations

- Is slower performance the only drawback of dynamic storage allocation?
- Unfortunately, NO :(
- Another disadvantage — memory overhead.
- Two causes:
  - Alignment:**
    - Padding caused by alignment produces wasted memory bytes.
  - Housekeeping data:**
    - Heap needs some auxiliary data to be stored in memory for each allocation to keep track of them.
- Effect** — with typical heap implementations:
  - Each allocation — even of a single byte — **consumes at least 32 bytes** of memory on 64-bit systems.

WHAT?

4

## Storage alignment

- Storage for an object (= each object) of type `T` needs to be aligned to `alignof(T)`-byte address.
  - Alignment to an  $N$ -byte address means alignment to an address  $A$ , where  $A \equiv 0 \pmod{N}$ .
- Alignment requirements for particular types are implementation-defined.
- Example** — GCC/Linux/x86\_64:

```
std::cout << alignof( char )   ; // prints out "1"
std::cout << alignof( int  )   ; // prints out "4"
std::cout << alignof( long  )   ; // prints out "8"
std::cout << alignof( float )   ; // prints out "4"
std::cout << alignof( double )  ; // prints out "8"
std::cout << alignof( long double ) ; // prints out "16"
std::cout << alignof( void* )   ; // prints out "8"
```

- Class types** — alignment requirements typically equal the maximum of alignment requirements for types of its subobjects.
- Note:** *Subobjects* = non-static member variables and base class objects.

5

## Storage alignment (cont.)

- Fundamental alignment (simplified definition):**
  - Alignment less than or equal to the maximum alignment for "fundamental" language-intrinsic types.
  - Is equal to the `alignof(std::max_align_t)` expression.
  - In typical C++ implementations, it equals alignment requirements for the `long double` type:

```
std::cout << alignof( long double ) ; // prints out "16"
std::cout << alignof( std::max_align_t ) ; // prints out "16"
```

- Over-aligned types:**
  - Types with alignment requirements that are higher than the maximal fundamental alignment.
- Exemplary use cases:**
  - Cache line aligned data* (efficiency reasons, avoiding *false sharing*,...) — typically, 64-byte alignment.
  - SIMD-processed data* (data processed by vectorization instructions), for example — AVX2 (32-byte alignment), AVX512 (64-byte alignment),...

6

## Statically-allocated storage alignment

- How is alignment provided when a storage for an object is allocated?
- 1) Function non-static local variables (*automatic storage duration*) + typical stack-based implementation:
  - Recall** — storage allocation = decreasing stack pointer (SP) register.
  - How much it needs to be decreased?
    - For example, decreasing by 4 bytes does not guarantee that SP will point to a 4-byte aligned address.
    - Generally, calculation of a properly aligned address would need additional instructions.
    - This would impose into functions additional (static) allocation overhead.
  - Common real-world solution of this problem:
    - ABI prescribes that when the function is executed, SP must point to an *A*-byte aligned address, where *A* is defined in that ABI.
    - Example** — *Linux/x86\_64* — *A* = 16:
      - When the `call` instruction calls some function, RSP must be 16-byte aligned.

7

## Statically-allocated storage alignment (cont.)

### • Example:

```
void g(int*);
void f() {
  int i = 1;
  g(&i);
}
```

```
f():
  push rax
  mov dword ptr [rsp + 4], 1
  lea rdi, [rsp + 4]
  call g(int*)
  pop rax
  ret
```

- When the function `f` is called somewhere, `rsp` is guaranteed to be 16-bytes aligned  $\Rightarrow \text{rsp} = 0 \text{ modulo } 16$ .
- The corresponding `call` instruction decreases `rsp` by 8 (it "pushes" the return instruction address to the stack).
  - $\Rightarrow$  When the function `f` is started, `rsp` = 8 modulo 16.
- Before `g` is called, `rsp` needs to be again 16-byte aligned.
  - $\Rightarrow$  `rsp` needs to be decreased by another 8 bytes.
- Those 8 bytes are used for allocation of storage for `i`:
  - They are guaranteed to be 8-byte aligned, `int` requires 4-byte alignment.
  - $\Rightarrow$  Storage for `i` can be allocated in upper 4 bytes or lower 4 bytes.
  - In our case, compiler decided to use the upper 4 bytes (`rsp+4`).
  - Lower 4 bytes will remain unused (wasted).

8

## Statically-allocated storage alignment (cont.)

### • Modified example:

```
void g(long*);
void f() {
  long i = 1;
  g(&i);
}
```

```
f():
  push rax
  mov qword ptr [rsp], 1
  mov rdi, rsp
  call g(long*)
  pop rax
  ret
```

- Storage for an object of type `long` takes 8 bytes and requires 8-byte alignment  $\Rightarrow$  there are no bytes wasted.
- Summary:**
  - Thanks to the ABI-required stack pointer alignment when functions are called, (static) allocations of storage for function-local variables can be performed extremely efficiently — namely, by decreasing stack pointer.
  - This holds for all *non-over-aligned* types.
  - For *over-aligned* types, some additional calculation at runtime is still needed (for instance, 16-byte alignment may or may not be 32-byte aligned at the same time).

9

## Dynamically-allocated storage alignment

- How is alignment provided when a storage for an object is allocated (cont.)?
- 2) Dynamically-allocated objects (*dynamic storage duration*) + typical heap-based implementation:
  - Recall** — storage allocation = calling some heap allocation function (through operator `new` C++ allocation function).
  - Typically, `malloc` is used.
  - `malloc` takes a single argument — number of allocated bytes.
    - $\Rightarrow$  It allows a caller to specify the storage (byte) size :).
    - $\Rightarrow$  It does not allow a caller to specify alignment requirements :(.
  - Common real-world solution of this problem:
    - Each dynamically-allocated block of memory is guaranteed to be aligned at least to the *maximal fundamental alignment*.

10

## Dynamically-allocated storage alignment (cont.)

- Our *Linux/x86\_64* system:
  - Heap implementation is (by default) provided by *GNU C library (GLIBC)*.
    - This is a common case in Linux.
  - Maximum fundamental alignment** `alignof(std::max_align_t)` is 16.
  - Accordingly, `malloc`-allocated memory blocks are guaranteed to be 16-byte aligned.
    - Note** — heap implementations may guarantee even higher/stronger alignment (for details, see `__STDCPP_DEFAULT_NEW_ALIGNMENT__` C++ macro).
- What about over-aligned types?
  - Since C++17, operator `new` has overloaded versions where alignment requirements may be specified.
  - Internally, other function than `malloc` needs to be used by a C++ standard library implementation (such as the `posix_memalign` allocation function).
  - Before C++17, there was no portable way to dynamically-allocate storage for over-aligned types.

11

## Memory costs of dynamic allocations (cont.)

- Consequences (*our platform*):
  - All dynamic allocations are 16-byte aligned.
  - $\Rightarrow$  Even if a storage for an object of some type requires only 1 byte, its dynamic allocation will effectively consume 16 bytes.
  - The remaining bytes are wasted.
- Moreover, heap implementations require for each allocation some housekeeping data  $\Rightarrow$  additional memory overhead.
  - In our case, these data take another 16 bytes per allocation.
- Experiment:**

```
void g(int*); // NOOP function defined in another TU
int main() {
  for (int i = 0; i < 100'000'000; i++) {
    int* pi = new int();
    g(pi);
  }
}
```

  - "Effective data" = 100M objects of type `int`  $\Rightarrow$  **400M bytes** of storage.
  - Measured memory consumption (*maximum RSS*): **3200M bytes**.
  - $\Rightarrow$  Each allocation consumes 32 bytes:
    - 4 effective bytes (storage of `int` object), 12 bytes wasted, 16 bytes housekeeping.

12

### Intermezzo — padding

- *Previous observation* — alignment can cause wasted bytes in memory in case of both:
  - static storage allocation (on the *stack*),
  - dynamic storage allocation (on the *heap*).
- These bytes are *out of* the allocated storage itself.
- Alignment requirements can cause wasted bytes even inside storage of a single object:
  - In the storage of a class type object, its *subobjects* are stored (*base class objects* and *non-static member variables*).
  - Due to alignment requirements, they may not be store next to each other in memory.
  - "Gaps" between subobjects are referred to as **padding**.
  - Padding represents wasted bytes inside the storage of a single object.

13

### Intermezzo — padding (cont.)

- *Example* — Clang/x86\_64/Linux:

```
struct X { void* v; char c; int i; float f; };
```

- Storage for an object of type X starts at address A.
- X::v requires 8-byte aligned storage, which is 8 bytes long.
  - $\Rightarrow$  A is 8-byte aligned.
  - $\Rightarrow$  X::v is stored from address A to A + 7.
- X::c requires 1-byte aligned storage, which is 1 byte long.
  - $\Rightarrow$  X::c is stored at address A + 8 only.
- X::i requires 4-byte aligned storage, which is 4 byte long.
  - $\Rightarrow$  X::i cannot be stored next to X::c at address A + 9; this address is not 4-byte aligned.
  - $\Rightarrow$  The lowest 4-byte aligned address where X::i may be stored is A + 12.
  - $\Rightarrow$  X::i is stored from address A + 12 to A + 15.
  - Bytes from A + 9 to A + 11 represent padding (are *wasted*).
- X::f requires 4-byte aligned storage, which is 4 byte long.
  - $\Rightarrow$  X::f is stored from address A + 16 to A + 19.

14

### Intermezzo — padding (cont.)

```
struct X { void* v; char c; int i; float f; };
```

- *Previous analysis*: Subobjects — member variables — X::v, X::c, X::i, and X::f are stored at addresses A to A + 19.
- If two objects of type X are stored in an array next to each other, and the first one is stored at address A, is the next element stored at A + 20? **NO!**
  - Its first member X::v needs to be 8-byte aligned.
  - $\Rightarrow$  The lowest available 8-byte aligned address higher than A + 19 is A + 24.
- Those 4 wasted bytes from A + 20 to A + 23 at the end also represent padding.
  - Even this final padding is a part of a binary representation of X.
  - $\Rightarrow$  These 4 bytes are part of the storage for each object of type X.
- *Proof*:


```
std::cout << sizeof(X); // prints out 24
```

- $\Rightarrow$  Storage of X consists of 17 data bytes and 7 padding bytes.

15

### Intermezzo — padding (cont.)

- *C++ rule*:
  - Class member variables with same access rights (*private*, *protected*, *public*) must be stored at growing addresses in order of their declaration.
  - $\Rightarrow$  Implementation may not reorder their placement in the storage of the class object.
- $\Rightarrow$  Padding may be sometimes reduced by "better" order of (non-static) member variable declarations.
- *Example*:



16

### Static vs dynamic allocations

- *Summary* — in comparison with static allocations, allocation of storage for dynamically-allocated objects:
  - is much slower,
  - consumes much more memory.
- Efficiency with *collection of objects*:
  - = objects stored in *dynamic data structures*.
  - C++ standard library defines different types of such data structures called *containers*.
  - Comparison of containers with respect to the number of required allocations for insertion of  $n$  objects:
    - `std::vector` without reserving space  $\rightarrow \log(n)$ ,
    - `std::vector` with reserving space  $\rightarrow 1$ ,
    - `std::list`  $\rightarrow n$ ,
    - `std::set` or `std::map`  $\rightarrow n$ ,
    - `std::unordered_set` or `std::unordered_map`  $\rightarrow n + \log(n)$ .

17

### Static vs dynamic allocations (cont.)

- *Experimental evaluation*:
  - Comparison of time and memory consumption when the same number of objects is inserted into different C++ library containers.
  - Measurements are *normalized* to the case where a vector was used without reserving memory.

| container                                       | time        | memory (RSS) |
|---|-------------|--------------|
| <code>std::vector</code> "as-is" (baseline)     | 1           | 1            |
| <code>std::vector</code> + <code>reserve</code> | 2.8× faster | 1.3× lower   |
| <code>std::list</code>                          | 17× slower  | 5.9× higher  |
| <code>std::set</code>                           | 95× slower  | 8.9× higher  |
| <code>std::unordered_set</code>                 | 46× slower  | 7.4× higher  |

- *Analysis*:
  - Vector is so efficient since it allocates storage for multiple objects (its elements) at once.
  - All other containers are *node-based* — each new element requires allocation of a separate storage for a single node, in which it is then stored.

18

### Static vs dynamic allocations (cont.)

- Vector stores its elements contiguously in memory.
- Nodes are placed in memory at unrelated locations.
- $\Rightarrow$  Vector provides efficient element access.
- **Vector vs set/map:**
  - *Set/map* = binary search tree  $\Rightarrow O(\log(n))$  lookup time.
  - Sorted vector with binary search  $\Rightarrow O(\log(n))$  lookup time as well.

| container                                | insertion (+sort) | lookup      |
|--|-------------------|-------------|
| std::vector "as-is" (baseline) + sorting | 1                 | 1           |
| std::vector + reserve + sorting          | 1.2× faster       | 1           |
| std::set                                 | 16× slower        | 2.1× slower |
| std::unordered_set                       | 7.3× slower       | 4.6× faster |

- Fastest lookup provides *unordered set* = *hash table*.
  - $\Rightarrow O(1)$  lookup time.
  - At a price of significantly slower data structure construction and higher memory consumption.

19

### Placement new

- Up to now, all mechanisms we have seen for both static and dynamic object allocations automatically "attached":
  - storage allocations together with object initialization,
- **and:**
  - object destruction together with storage deallocation.
- **Vector container:**
  - Allocates storage for multiple objects (its elements) at once.
  - $\Rightarrow$  Needs to "detach" storage allocation from object initialization (as well as object destruction from storage deallocation).
- **How can this be done in C++?**
  - For storage (de)allocation alone, we have (de)allocation functions operator new and operator delete.
  - $\Rightarrow$  What we need is to initialize object in this storage (*uninitialized memory*), and — in the end — to destruct it.

20

### Placement new (cont.)

- Vector — highly complex data structure.
- **Simpler case** for the sake of explanation:
  - Optional object existence — object that optionally may or may not exist at runtime.
- In case of dynamically-allocated object, the solution is simple:
  - A pointer itself can contain information whether the object exist or not:

```
int* pi = (some_condition == true) ? new int(1) : nullptr;
```

- Or, better:

```
std::unique_ptr<int> pi = (some_condition == true) ? new int(1) : nullptr;
```

- But what if we want to achieve the same without the overhead of dynamic storage allocation?
- $\Rightarrow$  We need optional object initialization in a statically-allocated storage.

21

### Placement new (cont.)

- Class wrapper for such an optionally existing object:

```
class optional_int {
    ??? buffer[ ??? ];
    bool exist_;
public:
    optional_int() : exists_(false) { }
    optional_int(int i) : exists_(true) { ??? }
    operator int&() { return ???; }
};
```

- Storage for optional object is a class member variable.
  - $\Rightarrow$  It is a part of the storage of the class object itself.
  - $\Rightarrow$  No need to dynamically allocate it.
- This storage needs to be:
  - Large enough to hold an object of type `int`  $\Rightarrow$  `sizeof(int)` bytes long.
  - Aligned suitably to hold an object of type `int`  $\Rightarrow$  `alignof(int)` byte-aligned.

22

### Placement new (cont.)

- Such storage can be provided as a buffer of a form of an array, which:

- 1) is aligned to `alignof(int)`-byte address,
- 2) has type of elements `unsigned char` or `std::byte` (since C++17),
- 3) has `sizeof(int)` elements.

```
class optional_int {
    alignas(int) unsigned char buffer_[ sizeof(int) ];
    bool exist_;
public:
    optional_int() : exists_(false) { }
    optional_int(int i) : exists_(true) { ??? }
    operator int&() { return ???; }
};
```

- Such a storage is now suitable to hold an object of type `int`.
- Next, we need to:
  - Construct (=initialize) such an object in this storage (buffer).
  - Provide access to it (to allow its update and reading).

23

### Placement new (cont.)

- Initialization of an object in existing storage = *placement new*.
  - Another form of new.
  - Basically, it is a new expression with an additionally provided pointer as its "argument".

```
class optional_int {
    alignas(int) unsigned char buffer_[ sizeof(int) ];
    bool exist_;
public:
    optional_int() : exists_(false) { }
    optional_int(int i) : exists_(true) { new (buffer_) int(i); }
    operator int&() { return ???; }
};
```

- **Placement new initializes an object:**
  - in memory at the address passed as a pointer,
  - of the specified type,
  - by initialization expression the same way as for "ordinary" (non-placement) new expression.

24

## Placement new (cont.)

- Object access:
  - There is no direct access (through some name/identifier).
  - But we know, where it is stored  $\Rightarrow$  we can derive its pointer (a *pointer-to-this-object*).
  - `buffer_` is an array  $\Rightarrow$  is implicitly convertible to a pointer-to-unsigned char.
  - We need a pointer-to-int.
  - $\Rightarrow$  *Solution* = cast (conversion) of these pointers:

```
class optional_int {
    alignas(int) unsigned char buffer_ [ sizeof(int) ];
    bool exist_;
    int* ptr() { return reinterpret_cast<int*>(buffer_); } // helper function
public:
    optional_int() : exist_(false) { }
    optional_int(int i) : exist_(true) { new (buffer_) int(i); }
    operator int&() { return *ptr(); }
};
```

- Note: since the operator `int&()` returns a reference to the stored object, the casted pointer is dereferenced.

25

## Placement new (cont.)

- Object destruction:
  - Placement new* initializes an object — starts its lifetime.
  - How to end this lifetime?**
  - Is there any “*placement delete*”? No.
- Two different cases:
  - Class types* — objects lifetime needs to be ended explicitly by calling their destructors.
  - Non-class types* — objects lifetime ends automatically when their storage is deallocated.
- `int` belongs to ad 2)
  - $\Rightarrow$  No explicit action is needed for destruction of the object initialized by *placement new*.
  - $\Rightarrow$  The class `optional_int` does not require manual definition of a destructor (it is *auto-generated*).
- What about ad 1) cases?

26

## Placement new (cont.)

- Class wrapper for such an optionally existing class object:

```
struct X {
    X(int);
    ~X();
};

class optional_X {
    alignas(X) unsigned char buffer_ [ sizeof(X) ];
    bool exist_;
    X* ptr() { return reinterpret_cast<X*>(buffer_); }
public:
    optional_X() : exist_(false) { }
    optional_X(int i) : exist_(true) { new (buffer_) X(i); }
    operator X&() { return *ptr(); }
    ~X() { if (exist_) ptr()->~X(); }
};
```

- Placement new initializes an object of type `X`  $\Rightarrow$  it calls its constructor.
- If the object was (optionally) initialized, then must be explicitly destructed by calling its destructor.
  - Destructor can be called as any other member function.

27

## Placement new (cont.)

- Comparison of *non-class* (`int`; left) and *class* (`X`; right) cases:

```
optional_int:optional_int(int):
    mov BYTE PTR [rdi+4], 1
    mov DWORD PTR [rdi], esi
    ret
optional_int::~optional_int():
    ret

optional_X:optional_X(int) :
    mov BYTE PTR [rdi+1], 1
    jmp X::X(int)
optional_X::~optional_X() :
    cmp BYTE PTR [rdi+1], 0
    jne .L10
    ret
.L10:
    jmp X::~~X()
```

- Non-class case:*
  - Initialization of an object of type `int` requires just setting its memory representation to a desired value (value of constructor argument passed in the `esi` register).
  - Destruction does not require any action.
- Class case:*
  - Initialization of an object of type `X` requires calling its constructor.
  - Destruction requires calling its destructor.

28

## Placement new (cont.)

- Problem with explicit destructor call:

```
class optional_string {
    alignas(std::string) unsigned char buffer_ [ sizeof(std::string) ];
    bool exist_;
    std::string* ptr() { return reinterpret_cast<std::string*>(buffer_); }
public:
    optional_string() : exist_(false) { }
    optional_string(char* s) : exist_(true) { new (buffer_) std::string(s); }
    operator std::string&() { return *ptr(); }
    ~optional_string() { if (exist_) ptr()->~string(); } // compilation error
};
```

```
error: expected class-name before '(' token
~optional_string() { if (exist_) ptr()->~string(); }
```

- `std::string` = a type alias for an instance of `std::basic_string<char>` class template.
- $\Rightarrow$  There is no destructor named `~string()`.
- $\Rightarrow$  The destructor is called `~basic_string()`.

```
~X() { if (exist_) operator X&().~basic_string(); } // OK
```

29

## Placement new (cont.)

- Generally, this is a problem, if we do not know whether some type is a type name or a type alias.
- First possible solution — templates:

```
template <typename T>
class optional {
    alignas(T) unsigned char buffer_ [ sizeof(T) ];
    bool exist_;
    T* ptr() { return reinterpret_cast<T*>(buffer_); }
public:
    optional() : exist_(false) { }
    ??? // constructor — to be explained later
    operator T&() { return *reinterpret_cast<T*>(buffer_); }
    ~optional() { if (exist_) operator ptr()->~T(); }
};
```

- Generic solution — will work for any type (template argument).
- For *class types*, correct destructor name is “encoded” in `~T()` call;
- For *non-class types*, there is no destructor, such as `~int()`.
  - In such case, `~T()` call is valid, but will have no effect.
- This mechanism is called *pseudo-destructor call* and allows us to write unified generic code where objects need to be explicitly destructed.

30

## Teaser — perfect forwarding

- How to write a constructor of `optional`?
- Constructors of `T` may take any number of arguments having different types and different value categories.
- ⇒ We need a constructor of `optional` that will:
  - take any number of argument of any types and any value categories,
  - and "pass" them *as-they-are* to the constructor of `T`.
- This can be solved by combining three fundamental modern C++ techniques:
  - variadic templates*,
  - forwarding references*,
  - `std::forward` function,
- which is together called *perfect forwarding*.
- More: later lectures.

```
template <typename T>
class optional {
...
    template <typename Ts...>
    optional(Ts&&... args) : exists_(true)
    {
        new (buffer_) T(std::forward<Ts>(args)...)
    }
    ...
};
```

31

## Placement new (cont.)

- Second possible solution: `std::destroy_at()` library function:
- ```
~optional_string() { if (exist_) std::destroy_at(ptr()); }
```
- `std::destroy_at()` is a function template.
  - It derives the type name (template parameter `T`) from the type of the passed pointer (*template argument deduction*).
    - Internally, it calls its destructor as `~T()`.
  - Advantage** — the code is more readable and explicit.
  - Counterpart of `std::destroy_at` is `std::construct_at`:

```
optional_string(char* s) : exists_(true) {
    std::construct_at(std::string>(ptr()), s);
}
```

- Internally, it constructs an object with *placement new*:
  - its template argument represents a type of the constructed object,
  - its first function argument is a pointer to the storage,
  - all other function arguments are perfectly forwarded to the initializer.
- Availability**: `std::destroy_at` C++17, `std::construct_at` C++20.

32

## Placement new (cont.)

- Using standard library improves readability and understandability of code; *see examples*:
  - placement new* vs `std::construct_at`,
  - explicit destructor call vs `std::destroy_at`.
- Another useful replacement from the standard library for our `optional` class — `std::aligned_storage`.
- Class (struct) template, which member type "type" provides a type of storage that is aligned to the desired value and have desired length.
  - Length and alignment are specified as template arguments.

```
alignas(T) unsigned char buffer_[ sizeof(T) ];
```

- can be replaced by:

```
typename std::aligned_storage<sizeof(T), alignof(T)>::type buffer_;
```

- or, since C++14, by:

```
std::aligned_storage<sizeof(T), alignof(T)>::type buffer_; // shortcut
```

33

## Placement new (cont.)

- `optional` class in C++20 with library entities:

```
template <typename T>
class optional {
    std::aligned_storage<sizeof(T), alignof(T)> buffer_;
    bool exist_;
    T* ptr() { return reinterpret_cast<T*>(buffer_); }
public:
    optional() : exists_(false) {}
    template<typename... Ts>
    optional(Ts&&... args) : exist_(true) {
        std::construct_at<T>(ptr(), std::forward<Ts>(args)...) ;
    }
    operator T&() { return *ptr(); }
    ~optional() { if (exist_) std::destroy_at(ptr()); }
};
```

- Such a class template in a more sophisticated form is available in the C++ standard library itself since C++17 as `std::optional<T>`.
- Common use case** — `optional` returning an object from a function:

```
std::string f1(); // (1) return object always
std::string* f2(); // (2) may return optionally, but requires heap allocation
std::optional<std::string> f3(); // (3) may return optionally and...
// ...does not require heap allocation
```

34

## Placement new (cont.)

- Drawback of (3) vs (1):
  - An `optional` class needs a flag with the information about whether it does or does not "hold" an object.
  - This flag, for example, controls destruction.

```
template <typename T>
class optional {
    bool exist_;
    ...
    optional() : exists_(false) {}
    template<typename... Ts> optional(Ts&&... args) : exist_(true) { ... }
    ~optional() { if (exist_) std::destroy_at(ptr()); }
};
```

- Boolean flag generally requires only 1 byte.
- However, due to alignment requirements, there may be some padding introduced.
- Flag + padding represent "housekeeping" memory overhead:

```
std::cout << sizeof( std::string ); // prints out 32
std::cout << sizeof( std::optional<std::string> ); // prints out 40 (+1B flag, +7B padding)
```

35

## "Owning" objects — "owners"

- An object of type `std::optional<T>` — a "wrapper" that (optionally) "owns/manages" an object of type `T`.
  - ⇒ The "content" of `std::optional<T>` is an (optional) ownership of an object of type `T`.
- The same holds, for instance, for an object of type `std::unique_ptr<T>`.

```
std::optional< std::string > os( "... " );
std::unique_ptr< std::string > us( new std::string( "... " ) );
```

- Further assumption** — stack/heap-based C++ implementation.
- Difference between `optional` and `unique_ptr`:
- `unique_ptr` — the storage of the owned object is **allocated dynamically**.
  - ⇒ The owned object itself is stored on the heap.
- `optional` — the storage of the owned object is "included" in the storage of the owning (`optional`) object itself.
  - ⇒ For instance, if the `optional` object is stored on the stack, so is also the owned object.

36

### “Owners” (cont.)

- *Physical ownership* — owned object is stored in the storage of its owner (= *included* storage).
- *Logical ownership* — owner takes care about owned object(s) (manages lifetime etc.).
- `std::optional<T>` — both *physical* and *logical* ownership:
  - Object is, for example, initialized and destructed by its optional owner.
  - Included storage — example with `GCC/x86_64`:

```
std::optional<int> oi(1);
std::cout << &oi;           // address of owner: "0x7ffc2e7146a8"
std::cout << &oi.value();    // address of owned: "0x7ffc2e7146a8"
```

- ⇒ The owner `oi` is stored on the *stack*, so is the owned object.
- ⇒ Storage size of the optional owner depends on the storage size of the owned object.

```
std::cout << sizeof( std::optional< int > ); // "8"
std::cout << sizeof( std::optional< std::string > ); // "40"
```

37

37

### “Owners” (cont.)

- `std::unique_ptr<T>` — *logical* ownership only:
  - Object is, for example, destructed by its `unique_ptr` owner.
  - Dynamically allocated storage:

```
std::unique_ptr<int> ui(new int(1));
std::cout << &ui;           // address of owner: "0x7ffe276c8450"
std::cout << &*ui;           // address of owned: "0x23faec0"
```

- ⇒ The owner `ui` is stored on the *stack*, but the owned object is stored on the *heap*.
- ⇒ Storage size of the `unique_ptr` owner does not depend on the storage size of the owned object.

```
std::cout << sizeof( std::unique_ptr< int > ); // "8"
std::cout << sizeof( std::unique_ptr< std::string > ); // "8"
```

- *Note*: `std::unique_ptr<T>` is typically implemented as a wrapper of the raw pointer member variable of type `T*`.

38

38

### “Owners” (cont.)

- C++ standard library provides multiple generic “owner” types.
- Some *single-object* owners...
  - ...with *included* storage:
    - `std::optional<T>`, `std::variant<T1,T2,...>`;
  - ...with *dynamically-allocated* storage:
    - `std::unique_ptr<T>`, `std::shared_ptr<T>`, `std::weak_ptr<T>`.
- Some *multiple-object* owners...
  - ...with *included* storage:
    - `std::pair<T1,T2>`, `std::tuple<T1,T2,...>`, `std::array<T,N>`;
  - ...with *dynamically-allocated* storage:
    - *dynamic containers* — `std::vector<T>`, `std::list<T>`, `std::set<T>`, ...
- What about `std::string`?
  - An object of type `std::string` can handle/own a *string of characters* of any length, generally unknown at compile time.
    - ⇒ Must the *string of characters* be stored in dynamically-allocated storage?

39

39

### Owners — `std::string` case

- *Experiment* — `Clang/x86_64/libc++`:

```
std::string s("...");
std::cout << &s;           // address of the string object: "0x7ffef075e4c0"
<< (void*)s.data();        // address of the string-of-characters: "0x7ffef075e4c1"
```

- *Observation*:
  - The owner `s` is stored on the *stack*.
  - The string of characters “...” managed/owned by `s` is stored also on the *stack*.
  - Namely, it is stored in the *included* storage of its owner `s`.
- *Similar experiment*:

```
std::string s("... some different string ...");
std::cout << &s;           // address of the string object: "0x7fffd772d280"
<< (void*)s.data();        // address of the string-of-characters: "0x18bf2a0"
```

- *Observation*:
  - Now, the string of characters “... some different string ...” is stored on the *heap* (in a *dynamically allocated storage*).
- *Explanation* — see next presentation.

40

40