

Effective C++ Programming

NIE-EPC (v. 2021):

EMPLACE SEMANTICS, PERFECT FORWARDING,
FUNCTION CALL DISPATCHING

© 2021 DANIEL LANGR, ČVUT (CTU) FIT

Implicit conversions

- Consider the following *example*:

```
struct X {  
    X(int) { std::cout << "converting constructor\n"; }  
    X(X&&) { std::cout << "move constructor\n"; }  
};
```

```
std::vector< X > v;  
v.push_back( X(1) );
```

- Here, `std::vector<X>::push_back(X&&)` is called where:
 - type of parameter* is `X&&`,
 - argument* is an *expression* that:
 - refers to an *object of type X*,
 - its *value category* is *rvalue*.
- Modified example*:

```
std::vector< X > v;  
v.push_back( 1 );
```

- What has changed?
 - Argument is an *expression* that refers to an object of **type other than X**.
 - (Namely, it refers to an object of type `int` and its category is *rvalue*.)

Implicit conversions (*cont.*)

- When:
 - initialization **requires** expression of some type **T1**, and
 - the **type of the initialization** expression is different than **T1** (say, **T2**),
- then a compiler will **try to “implicitly” convert** the initialization expression into an object of type **T1**.
- *In our case:*
 - *Parameter* of `push_back` (of type `X&&`) **expects** expression of type **X**.
 - *Initialization expression* (argument) **1** **has type** `int`.
- \Rightarrow Compiler will try to **implicitly convert 1 to an object of type X**.
- Is such conversion **possible**?
 - **Yes, it is**, due to the availability of matching **non-explicit converting constructor** (constructor with a parameter of type `int`):

```
X(int) { std::cout << "converting constructor\n"; }
```

- *Note* — **explicit specifier** does **disallow such implicit conversions**.
 - Conversions are then still possible but **only explicitly written in code**.

Emplacing-back into vector

- *Consequence* — both two options are *effectively equivalent*:

```
v.push_back( X(1) ); // option #1  
v.push_back( 1 ); // option #2
```

- In *#2*, the temporary object of type X *created automatically due to an implicit conversion*.
- What *constructors of X* are involved in...?

```
std::vector<X> v;  
v.push_back(1);
```

- 1) A *temporary of type X* is created from 1 \Rightarrow *converting constructor*.
 - 2) Then, `push_back` internally *initializes a new element* in its storage and *move content into it* from the temporary \Rightarrow *move constructor*.
- The program output agrees with that:

```
converting constructor  
move constructor
```

- *Live demo* — <https://godbolt.org/z/3Y3KaTnEx>.

Emplacing-back into vector (*cont.*)

```
std::vector<X> v;  
v.push_back(1);
```


converting constructor
move constructor

- There is something “wrong” with that from the perspective of *performance/efficiency*.
 - Why to *create a temporary* and then, immediately, *move content* from it?
 - Wouldn't be better to *create the vector element* itself with the *converting constructor* “directly” from the argument 1?
- Clearly, `push_back` itself *cannot provide that* — it accepts only *arguments of type X*.
- *Instead* — we want a *vector member function* that
 - 1) accepts *argument(s)* of *any type* and *any value category*,
 - 2) initializes a *new vector element* with *expression* that:
 - *represents the same object(s)* as the function argument(s),
 - *has the same value category(ies)* as the function argument(s).

Emplacing-back into vector (*cont.*)

- Such functionality is called “*emplacing*” and `std::vector` provides corresponding function called `emplace_back`.
- *Our first attempt* — *single-argument* implementation.
- We want to accept an *argument of any type*...
 - \Rightarrow the designed function actually *needs to be a function template*;
- ...and *any value category*, which will be *recognizable*...
 - \Rightarrow the function parameter *needs to be a forwarding reference*.

```
// template <typename T> class Vector { ...  
template <typename U>  
void emplace_back(U&& param) {  
    if (size_ == capacity_)  
        reserve(capacity_ ? 2 * capacity_ : 1);  
    new (data_ + size_) T( ??? );  
    size_++;  
}
```



- *Remains to be resolved* — how *initialization expression* should look like?

Emplacing-back into vector (cont.)

- *Example:*

```
template <typename U>
void emplace_back(U&& param) {
    ...
    new (data_ + size_) T( ??? );
```

```
std::vector<X> v;
int i = 1;
v.emplace_back( i );    // (1)
v.emplace_back( i+1 );  // (2)
```

- In **case (1)**, we need to **initialize vector element with an expression** that:
 - a) **refers** to the same object as the argument \Rightarrow **to variable *i***,
 - b) **has** the same **value category** as the argument \Rightarrow ***lvalue***.
- In **case (2)**, we need to **initialize vector element with an expression** that:
 - c) **refers** to the same object as the argument \Rightarrow **to temporary *i+1***,
 - d) **has** the same **value category** as the argument \Rightarrow ***rvalue***.
- *Possible solutions?*

```
new (data_ + size_) T( param ); // NO! param is always lvalue => breaks d)
```

```
new (data_ + size_) T( std::move(param) );
                        // NO! std::move(param) is always rvalue => breaks b)
```

Emplacing-back into vector (*cont.*)

```
new (data_ + size_) T( param );
```

```
new (data_ + size_) T( std::move(param) );
```

- *Problem:*
 - `param` is *always lvalue* (named entity case)
 - `std::move(param)` is *always rvalue*.
- Instead, we need an expression that:
 - represents the *same object* as `param`,
 - its *value category* is the *same* as the value category of the `emplace_back` argument.
- *Quick solution* — `std::forward`.

```
new (data_ + size_) T( std::forward<U>(param) );
```

- *Example:*

```
std::vector<X> v;  
v.emplace_back(1);
```

`converting` constructor

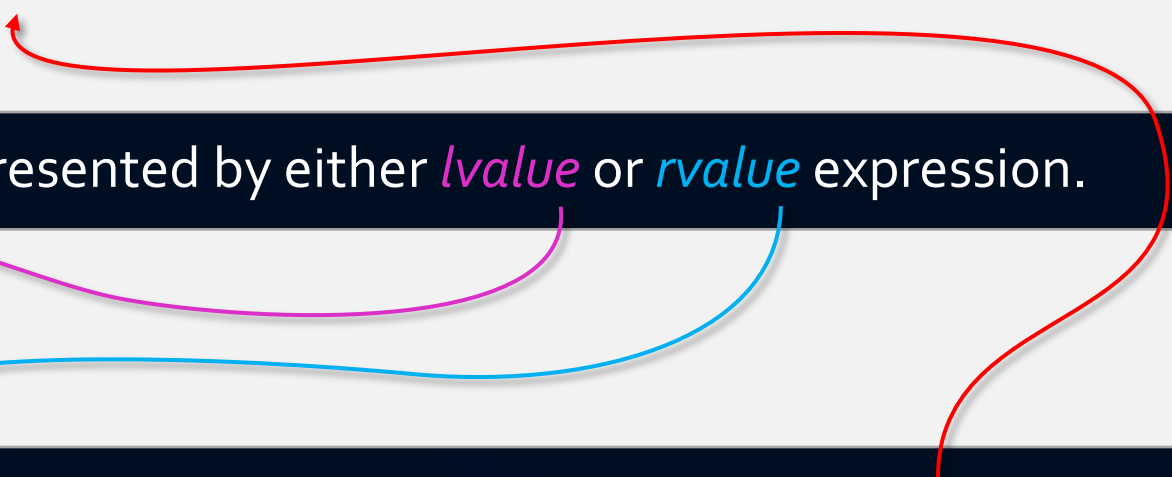
Pushing- vs emplacing-back into vector

- *Benchmark:*
 - Insertion of elements with `push_back` and `emplace_back` into a *vector of strings* (`std::string`), where *argument* is a *string literal*.
- *Results:*
 - With `GCC/libstdc++`, using `emplace_back` was **2.3× faster**.
 - Link: https://quick-bench.com/q/bVc7e_EjenbAVWqAzSK22reeqYc.
 - With `Clang/libc++`, using `emplace_back` was **3.2× faster**.
 - Link: <https://quick-bench.com/q/QtdeRHqkhARu-m9XOsJ4JeyxHr4>.
- *Alternative benchmark:*
 - The same with a *vector of integers* (`int`) and *integer literal* argument.
 - *Results* — **same** measured runtime/performance.
- *Analysis:*
 - With `std::string`, *move constructors* are **additionally involved** with `push_back`.
 - On the contrary, `int` is a *non-class type* and additional initialization is **effectively eliminated under optimizations**.

Perfect forwarding


- How does `std::forward` function work?
- *Recall* — param is a **forwarding reference** that is **bound to some object**.

```
// template <typename T> class Vector { ...  
template <typename U> void emplace_back(U&& param) {  
    if (size_ == capacity_) reserve(capacity_ ? 2 * capacity_ : 1);  
    new (data_ + size_) T( ... );  
    size_++;  
}
```



- This object is represented by either *lvalue* or *rvalue* expression.

```
Vector<X> v;  
int i = 1;  
v.emplace_back( i );  
v.emplace_back( i+1 );
```



- We want to initialize new vector element with **expression that**:
 - represents the **same object**,
 - has the **same value category**.

Perfect forwarding (cont.)

```
template <typename U>
void emplace_back(U&& param) {
    if (size_ == capacity_)
        reserve(capacity_ ? 2 * capacity_ : 1);
    new (data_ + size_) T( ... );
    size_++;
}
```

```
Vector<X> v;
int i = 1;
v.emplace_back( i );    // (1)
v.emplace_back( i+1 ); // (2)
```

- Recall:
 - In the *first (lvalue) case*, **U** is deduced as **int&** (and type of param is **int&**).
 - In the *second (rvalue) case*, **U** is deduced as **int** (and the type of param is **int&&**).
- Generally, in case of
 - *lvalue* → **U** is a (lvalue) *reference* type,
 - *rvalue* → **U** is a *non-reference* type.
- ⇒ *Initialization expression* needs to be:
 - *lvalue* expression if **U** is a *reference* type,
 - *rvalue* expression *otherwise*.
- Possible cast-based solution?

```
new (data + size_) T( static_cast<U&&>(param) );
```

Perfect forwarding — casting

```
template <typename U> void emplace_back(U&& param) {  
    if (size_ == capacity_) reserve(capacity_ ? 2 * capacity_ : 1);  
    new (data_ + size_) T( static_cast<U&&>(param) );  
    size_++;  
}
```

- *Lvalue-case example:*

```
Vector<X> v;  
int i = 1;  
v.emplace_back( i );
```

- U is deduced as `int&` \Rightarrow after *substitution* and *reference collapsing*:

```
new (data_ + size_) T( static_cast< int& >(param) ); // cast expression is lvalue
```

- \Rightarrow Initialization expression represents variable `i` and its category is *lvalue*.

- *Rvalue-case example:*

```
v.emplace_back( i+1 );
```

- U is deduced as `int` \Rightarrow cast turns into:

```
new (data_ + size_) T( static_cast< int&& >(param) ); // cast expression is rvalue
```

- \Rightarrow Initialization expression represents temporary `i+1` and its category is *rvalue*.



Perfect forwarding — `std::forward`

- Delegation of cast to a **separate function**:

```
template <typename T>
T&& forward(T& param) { return static_cast<T&&>(param); }
```

- *Application* to **`std::vector<T>::emplace_back`**:

```
template <typename U> void emplace_back(U&& param) {
    if (size_ == capacity_) reserve(capacity_ ? 2 * capacity_ : 1);
    new (data_ + size_) T( forward<U>(param) );
    size_++;
}
```

- *Lvalue-case example*:

```
Vector<X> v; int i = 1; v.emplace_back( i );
```

- \Rightarrow T in forward is **`int&`** \Rightarrow after *substitution* and *collapsing*:

```
int& forward(int& param) { return static_cast<int&>(param); } // its call is lvalue
```

- \Rightarrow forward call expression **represents variable `i`** and *is lvalue*.



- *Rvalue-case example*:

```
Vector<X> v; int i = 1; v.emplace_back( i+1 );
```

- \Rightarrow T in forward is **`int`** \Rightarrow after *substitution*:

```
int&& forward(int& param) { return static_cast<int&&>(param); } // its call is rvalue
```

- \Rightarrow forward call expression **represents temporary `i+1`** and *is rvalue*.



Perfect forwarding — `std::forward` (cont.)

```
template <typename T> T&& forward(T& param)  // accepts only lvalues
{ return static_cast<T&&>(param); }
```

- *Note* — our forward function (template) **accepts only lvalue arguments**.
- This was **ok for our `emplace_back`**.
 - Its **parameter `param`** used as an argument of forward **is always lvalue** (named-entity case).

```
new (data_ + size_) T( forward<U>(param) );  // param expression is always lvalue
```

- Generally, we need to “**cover**” cases where forward:
 - *function argument* is **either lvalue or rvalue**,
 - *template argument* is **either reference or non-reference type**.
- \Rightarrow This makes **4 different cases**, which **cannot be resolved with a single definition of forward function template**.
- An **additional overload** would need to be added for **rvalues**:

```
template <template T>
T&& forward(typename remove_reference<T>::type&& param)  // overload for rvalues
{ return static_cast<T&&>(param); }
```

Perfect forwarding — `std::forward` (*cont.*)

- Such two overloads are provided by the C++ standard library as a function template `std::forward`.
- \Rightarrow Explanation of the *original solution*:

```
new (data_ + size_) T( std::forward<U>(param) );
```

- *Relevant notes:*
 - Template argument for `(std::)forward` call must be explicitly specified.
 - Template argument deduction rules wouldn't work here with the desired functionality.
 - \Rightarrow Implementations of `std::forward` are more “robust” (than our `forward`), and enforce explicit template argument provision.
 - Note — in application to our problem, there is effectively no difference.
 - Relevant *Stack Overflow* discussions:
 - [Why does `std::forward` have two overloads?](#)
 - [The implementation of `std::forward`](#)

Perfect forwarding (cont.)

- Generalization:
 - We have a **function** (template) that has a *forwarding reference parameter*.
 - This function is called with some *function argument (= expression)*.
 - This argument **represents some object** and **has some value category**.
 - Inside the function, there is **another expression created** that:
 - **represents the same object** as the function argument,
 - **has the same value category** as the function argument.
- Such technique is generally called “**perfect forwarding**”.
- It is effectively used for **passing/“forwarding”** of arguments of “**outer**” **function call** into some **internal another function call** (for instance, constructor in case of initialization).
- “**Perfect**” = it **preserves all properties of arguments** (representation of particular object, its type, and value category).
- Note — “**forwarding**” gave rise of term “*forwarding reference*”

std::forward vs std::move

- *Perfect forwarding* is typically implemented with the help of library utility function (template) `std::forward`.
- \Rightarrow This is the reason of its name.
- `std::forward` and `std::move`:
 - 1) `std::move` — *recall*:
 - It *does not* “move” anything.
 - Instead, its call *creates an expression* that represents the same objects as its argument and has category *rvalue*.
 - 2) `std::forward` — *similarly*:
 - It *does not* “forward” anything.
 - Instead, its call *creates an expression* that represents the same object as its argument (*forwarding reference*) and has a desired value category (*same as the expression “bound” to that forwarding reference*).
- \Rightarrow Both functions are *technically similar*; they generally *differ only in the value category* they “produce”.
- \Rightarrow However, they both have *very different use cases*.

Emplacing back — multiple arguments

- Final solution:

```
// template <typename T> class Vector { ...  
template <typename U> void emplace_back(U&& param) {  
    if (size_ == capacity_) reserve(capacity_ ? 2 * capacity_ : 1);  
    new (data_ + size_) T( std::forward<U>(param) );  
    size_++;  
}
```

- It works as required, but...
- ...only for a single argument.

```
struct X {  
    X(int) { } // converting constructor  
};
```

```
Vector<X> v;  
  
v.emplace_back( 1 ); // ok
```

- \Rightarrow This version of `emplace_back` cannot perfectly-forward multiple arguments to the vector elements initialization.
 - This might be required for constructors with multiple parameters.

```
struct Y {  
    Y(int, int) { } // converting constructor  
};
```

```
Vector<Y> v;  
  
int i = 1;  
v.emplace_back( i, i+1 ); // error
```


Emplacing back — multiple args (*cont.*)

- We would like to “perfectly-forward” — in our case to the initialization expression of the added vector element — **not only a single argument**, but **any number of arguments**.
- *Quick solution* — making `emplace_back` a **variadic template**:

```
template <typename... U> void emplace_back(U&&... param) {  
    if (size_ == capacity_) reserve(capacity_ ? 2 * capacity_ : 1);  
    new (data_ + size_) T( std::forward<U>(param)... );  
    size_++;  
}
```

```
Vector<Y> v;  
int i = 1;  
v.emplace_back( i, i+1 ); // ok
```

- *Effect:*
 - Each argument is (separately) perfectly-forwarded to the initialization expression of the constructed vector element.
 - The first constructor argument represents variable `i` and its category is *lvalue*.
 - The second constructor argument represents temporary `i+1` and its category is *rvalue*.

Emplacing back — multiple args (*cont.*)

- *Resolution:*

```
template <typename... U> void emplace_back(U&&... param) {  
    if (size_ == capacity_) reserve(capacity_ ? 2 * capacity_ : 1);  
    new (data_ + size_) T( std::forward<U>(param)... );  
    size_++;  
}
```

```
Vector<Y> v;  
int i = 1;  
v.emplace_back( i, i+1 );
```

- `emplace_back` has 2 arguments \Rightarrow effectively equivalent with...

```
template <typename U, typename V> void emplace_back(U&& param1, V&& param2) {  
    if (size_ == capacity_) reserve(capacity_ ? 2 * capacity_ : 1);  
    new (data_ + size_) T( std::forward<U>(param1), std::forward<V>(param2) );  
    size_++;  
}
```

- ...which is then “resolved” (*instantiated*) as:

```
void emplace_back(int& param1, int&& param2) {  
    if (size_ == capacity_) reserve(capacity_ ? 2 * capacity_ : 1);  
    new (data_ + size_) Y( std::forward<int&&>(param1), std::forward<int>(param2) );  
    size_++;  
}
```

Emplacing back — no argument

```
template <typename... U> void emplace_back(U&&... param) {  
    if (size_ == capacity_) reserve(capacity_ ? 2 * capacity_ : 1);  
    new (data_ + size_) T( std::forward<U>(param)... );  
    size_++;  
}
```

- This variadic template-based solution even **allows to forward no argument at all**:

```
Vector< std::string > v;  
v.push_back();      // error - push_back requires an argument  
v.emplace_back();   // ok - inserts empty = default-constructed string object in the vector
```

- `emplace_back` is here *instantiated* as:

```
void emplace_back() {  
    if (size_ == capacity_) reserve(capacity_ ? 2 * capacity_ : 1);  
    new (data_ + size_) std::string( /* nothing here */ );  
    size_++;  
}
```

push_back vs emplace_back

- **push_back** requires — as its input — **an object of vector's value type T**.
 - Recall, there are two overloads for lvalues and rvalues.

```
// template <typename T> class Vector { ...  
void push_back(const T& param) { // overload for lvalues  
    if (size_ == capacity_) reserve(capacity_ ? 2 * capacity_ : 1);  
    new (data_ + size_) T( param );  
    size_++;  
}  
void push_back(T&& param) { // overload for rvalues  
    if (size_ == capacity_) reserve(capacity_ ? 2 * capacity_ : 1);  
    new (data_ + size_) T( std::move( param ) );  
    size_++;  
}
```

- On the contrary, **emplace_back** can **accept any argument**:

```
template <typename... U> void emplace_back(U&&... param) {  
    if (size_ == capacity_) reserve(capacity_ ? 2 * capacity_ : 1);  
    new (data_ + size_) T( std::forward<U>(param)... );  
    size_++;  
}
```

- \Rightarrow These arguments may — of course — **be of type T as well**.

push_back vs emplace_back (cont.)

- When **argument of emplace_back** is of type **T** and its **category is lvalue**, **emplace_back** **behaves exactly as the lvalue overload of push_back**.
 - \Rightarrow We can write this overload of **push_back** in terms of **emplace_back**:

```
void push_back(const T& param) {                                     // overload for lvalues
    emplace_back(param);
}
```

- Similarly, when **argument of emplace_back** is of type **T** and its **category is rvalue**, **emplace_back** **behaves exactly as the rvalue overload of push_back**.
 - \Rightarrow We can also write this overload of **push_back** in terms of **emplace_back**:

```
void push_back(T&& param) {                                         // overload for rvalues
    emplace_back( std::move(param) ); // need to make rvalue from param
}
```

- This way — for example — is the **push_back** **implemented in Microsoft STL** (with a bit **different internal syntax**, but the **same effect**).
 - Link: <https://github.com/microsoft/STL/blob/main/stl/inc/vector#L633>.

Perfect forwarding — use cases

- *Perfect forwarding* is one of the building blocks of modern C++.
- It was introduced in C++11 (same as, for example, *content moving semantics*).
- It is commonly used in C++ standard library.
- *Use case 1.* — `std::construct_at` (seen in lecture 4):
 - Recall — this function has the functionality of *placement new*.
 - ⇒ It explicitly initializes an object of a *give type* in the *location* specified by its first function argument...
 - ...while all other arguments are perfectly forwarded to the object initialization expression.

```
new (ptr) T(arg1, arg2, arg3);           // option #1 - placement new
std::construct_at<T>(ptr, arg1, arg2, arg3); // option #2 - equivalent
```

- *Possible implementation:*

```
template <typename T, typename... A>
T* construct_at(void* ptr, A&&... params) {
    new (ptr) T( std::forward<A>(params)... );
}
```

Perfect forwarding — use cases (cont.)

- *Use case II.* — constructor of `std::optional` (seen in lecture 4):
 - Recall — *optional class* can optionally own/store an object of its value type (template argument) into its *included storage*.
 - It has a constructor that immediately initializes the owned object and perfectly forwards any arguments to its constructor.
- Possible implementation of such a constructor:

```
// template <typename T> class optional { ...  
template <typename Ts...>  
optional(std::in_place_t, Ts&&... args) : exists_(true) {  
    new (buffer_) T(std::forward<Ts>(args)...)  
}
```

- Alternative implementation with `std::construct_at` (\Rightarrow double perfect forwarding)

```
// template <typename T> class optional { ...  
template <typename Ts...>  
optional(std::in_place_t, Ts&&... args) : exists_(true) {  
    std::construct_at<T>(buffer_, std::forward<Ts>(args)...) ;  
}
```

- What about that (unnamed) parameter of type `std::in_place_t`?

Perfect forwarding — use cases (cont.)

- Why isn't the constructor defined **simply as follows**?

```
// template <typename T> class optional { ...  
template <typename Ts...>  
optional(Ts&&... args) : exists_(true) {  
    new (buffer_) T(std::forward<Ts>(args)...)  
}
```

- Consider the following **two options** when **initializing an optional object**:
 - *First*, we want to create an *"empty" optional owner object*, which **does not own any object of its value type**.
 - For this purpose, there is a *default constructor*:

```
std::optional< std::string > o1; // default constructor => no owned object
```

- *Second*, we want to **initialize an optional owner that owns an empty = default-constructed string object**.
 - \Rightarrow We would like to **invoke the "forwarding-constructor" and forward nothing**.
 - With the **above definition**, this would require **no constructor argument as well**.
- \Rightarrow We need some mechanism to **tell the compiler that it should call the forwarding constructor in this case** (discussion: [\[link\]](#)).
- Mechanism used by `std::optional` is called *tag dispatching*.

Function call dispatching

- *Situation* — *function call* which may correspond to *multiple versions (variants/definitions)* of that function.
- *Dispatching of such function call* = resolution of *which of the versions will be called*.
- C++ has *several different dispatching mechanisms* suitable for *different situations*.
- *First* dispatching option — *function overloading*.
 - The dispatch is resolved *according to the function argument expressions* (their types and value categories).
 - *Example* — *type-based* overload-dispatching:

```
f(int);    // overload #1
f(double); // overload #2
```

```
f( 1 ); // overload #1 called
f( 1.0 ); // overload #2 called
```

- *Another example* — value *category-based* overload-dispatching:

```
struct X {
    X();           // default constructor
    X(const X&);   // copy constructor
    X(X&&);       // move constructor
};
```

```
X x1;
X x2( x1 ); // copy ctor called
X x3( std::move(x1) ); // move ctor called
```

Function call dispatching (cont.)

- *Second* dispatching option — *dynamic polymorphism*.
 - The dispatch is resolved *according to the actual (dynamic) type of the pointed-to/referenced object*.
 - *Example:*

```
struct Base {  
    virtual void f() = 0;  
    virtual ~Base() = default;  
};  
  
struct Derived1 : Base {  
    virtual void f() { std::cout << "1"; }  
};  
  
struct Derived2 : Base {  
    virtual void f() { std::cout << "2"; }  
};  
  
void f(Base& obj) {  
    obj.f(); // dispatch required  
}
```

```
std::unique_ptr<Base> p_obj;  
  
int i;  
std::cin >> i;  
  
p_obj = (i == 1) ?  
    new Derived1 : new Derived2;  
  
f(*p_obj);  
  
// => calls either  
//     Derived1::f()  
// or  
//     Derived2::f()  
// based on the type of created object
```

- Dispatch resolution is based on *virtual functions*.
 - It is sometimes referred to as "*virtual dispatch*".

Function call dispatching (cont.)

- *Third* dispatching option — *static polymorphism*.
 - The dispatch is resolved according to the *actual type of the object* that is specified by a *template parameter*.
 - *Example:*

```
template <typename T>
void f(T& obj) {
    obj.f(); // dispatch required
}

struct A { void f() { std::cout << "A"; } };
struct B { void f() { std::cout << "B"; } };
```

```
A obj;
f(obj); // calls A::f()
```

- Dispatch may be resolved
 - 1) either *at compile time* — then, it is generally called “*static dispatch*”,
 - 2) or *at runtime* — then, it is generally called “*dynamic dispatch*”.
- *Examples:*
 - *Overload-* and *static polymorphism*-based dispatches are *static*.
 - *Dynamic polymorphism (virtual)*-based dispatch is *dynamic*.

Function call dispatching (cont.)

- When class object is *initialized*, constructor is dispatched *according to the initialization expression* by the *overloading mechanism*:

```
struct X {  
    X();  
    X(const X&);  
    X(X&&);  
};
```

```
X x1;                // no argument => default constructor  
X x2(x1);            // lvalue argument of type X => copy constructor  
X x3(std::move(x1)); // rvalue argument of type X => move constructor
```

- Problem* with the *following version of the forwarding constructor* of the optional class is...

```
// template <typename T> class optional { ...  
optional();                // default constructor  
template <typename Ts...> optional(Ts&&... args); // forwarding constructor (attempt)
```

- ...that the *overloading rules cannot distinguish between the described two initialization cases*:
 - initialization of *empty optional* object,
 - initialization with *no-argument forwarding*.
- Both** syntactically correspond with *initialization with no argument*:

```
std::optional< std::string > o1; // both default and forwarding constructor match
```

Overloading

```
optional(); // default constructor
template <typename Ts...> optional(Ts&&... args); // forwarding constructor (attempt)

std::optional< std::string > o1; // both default and forwarding constructor match
```

- In this case, *both constructor match the initialization form (no/empty initialization expression)*.
 - They both can be called \Rightarrow are so-called *viable "candidates"* for overload resolution.
- Which of them will be finally called?
- Generally, when they are *multiple viable overloading candidates*, the *one with the highest priority* is selected.
 - These priorities are *specified by* (relatively complex rules of) *C++ standard*.
- In our case, the *default constructor will be selected* according to these rules (*non-templates* typically have priority over *templates*).
- Note:
 - *Multiple candidates with equal highest priority* results in *compilation error* ("ambiguous call of...").

Tag dispatching

```
optional(); // default constructor
template <typename Ts...> optional(Ts&&... args); // forwarding constructor (attempt)
```

- *Overloading rules cannot distinguish* between initialization with these two constructors *in case of no-argument forwarding*.
- \Rightarrow We *need some other dispatching mechanism* how to select forwarding constructor.
- One such mechanism is so-called “*tag dispatching*”.
 - It works such that we introduce into a function *a parameter* (typically unnamed) *of a special type* created *only for tag dispatching purposes*.
 - *Example:*

```
void f() { std::cout << "1"; }

struct tag_t { }; // special new tagging-purpose type
template <typename... Ts> void f(tag_t, Ts&&...) { std::cout << "2"; }
int main() {
    f(); // prints out "1"
    f( tag_t{} ); // prints out "2"
```

- *Note* — tag dispatching is an “*explicit overloading mechanism*”.

Tag dispatching (cont.)

- *Tag dispatching* is used for **selection of forwarding constructor of `std::optional`**.
- The special **selection "tag" type** is `std::in_place_t`:

```
// template <typename T> class optional { ...  
template <typename Ts...>  
optional(std::in_place_t, Ts&&... args);
```

- Forwarding constructor is then effectively selected by providing the **first argument of this type during initialization**:

```
std::optional< std::string > o1; // default constructor called  
std::optional< std::string > o2( std::in_place_t{} ); // forwarding constructor called
```

- *Resolution:*
 - *First initialization* creates an **empty optional object**.
 - *Second one* creates an **optional object that owns an empty string object**.
- *Note:*
 - Tag argument needs to be an **expression/object of type `std::in_place_t`**.
 - Standard library **defines such an object for us**:

```
std::optional< std::string > o2( std::in_place );
```

Perfect forwarding (cont.)

- *Use case III.* — smart pointer “makers”.
 - Smart pointers do not have forwarding constructors.
 - Instead, we can initialize them by creating owned objects explicitly:

```
std::unique_ptr<X> upy = new X(1);
```



- *Drawbacks:*
 - Type needs to be written twice in the source code (duplication).
 - Its “ugly” — using new explicitly is considered bad practice in modern C++ when programming at a high level of abstraction.
- *Alternative solution...*

```
std::unique_ptr<X> upy = std::make_unique<X>(1); // eliminates explicit new
```

- ...or, even better:

```
auto upy = std::make_unique<X>(1); // eliminates both drawbacks
```

- The same approach — `std::shared_ptr` + `std::make_shared`.
- *Caveat:*
 - Smart pointers support so-called *custom deleters*, which cannot be supplied via `make_` functions.

Perfect forwarding (cont.)

- Function (template) `std::make_unique`:
 - It creates a unique pointer that owns a constructed object of a desired type...
 - ...and perfectly forwards all function arguments to the object initialization expression.
- \Rightarrow Possible implementation:

```
template <typename T, typename... Ts>
std::unique_ptr<T> make_unique(Ts&&... params) {
    return std::unique_ptr<T>( new T( std::forward<Ts>(params)... ) );
}
```

- Does `make_unique` have any *runtime overhead*?

```
std::unique_ptr<X> upy1 = new X(1); // converting constructor of std::unique_ptr called only
auto upy2 = std::make_unique<X>(1); // which constructors called here?
```

- Since C++17, it is *guaranteed to be equivalent* \Rightarrow no overhead.
- Until C++17, it is *very likely equivalent* \Rightarrow (very) likely no overhead.
- Reason = *copy elision* optimization technique.

Perfect forwarding (cont.)

- *Use case IV.* —*emplacing functions* of **library containers**.
 - We have seen `std::vector::emplace_back`.
 - Other containers (and container adapters) **also support *emplacing semantics*** for **elements insertion operations**.
 - They differ from “traditional” insertion functions such that they **construct/initialize new container elements** while **perfectly forwarding arguments to its initialization expression**.
 - *Examples:*

```
std::set<std::string> m;  
m.insert("string");    // converting + move  
m.emplace("string");   // converting only
```

```
std::list<std::string> l;  
l.push_front("string");    // converting + move  
l.emplace_front("string"); // converting only
```

```
std::stack<std::string> s;  
s.push("string");    // converting + move  
s.emplace("string"); // converting only
```