

Effective C++ Programming

NIE-EPC (v. 2021):

EXCEPTION SAFETY, ALLOCATORS, SMALL BUFFER
OPTIMIZATION, EMPTY BASE OPTIMIZATION, TYPE
ERASURE


© 2021 DANIEL LANGR, ČVUT (CTU) FIT

Vector — reallocation — *safety concerns*

- *Not-exactly-final* solution for **reserve** member function:

```
void reserve(size_t capacity) {  
    if (capacity <= capacity_) return;  
    T* data = (T*)::operator new(capacity * sizeof(T));  
  
    for (size_t i = 0; i < size_; i++)  
        new (data + i) T( std::move( *(data_ + i) ) ); // what if throws ???  
  
    clear();  
    ::operator delete(data_);  
    data_ = data; capacity_ = capacity;  
}
```

- In case of *non-trivially-copyable types*, initialization involves **calls of copy or move constructor**, which, generally, **may throw exceptions**.
- Such situations **need to be handled**:
 - some elements **may have been already constructed** before,
 - new storage **has been already allocated**.
- *Note* — for sake of simplicity, we **ignored this issue** previously.
- *Note* — initialization of objects of *trivially copyable types* with content-copying here **cannot result in exception**.



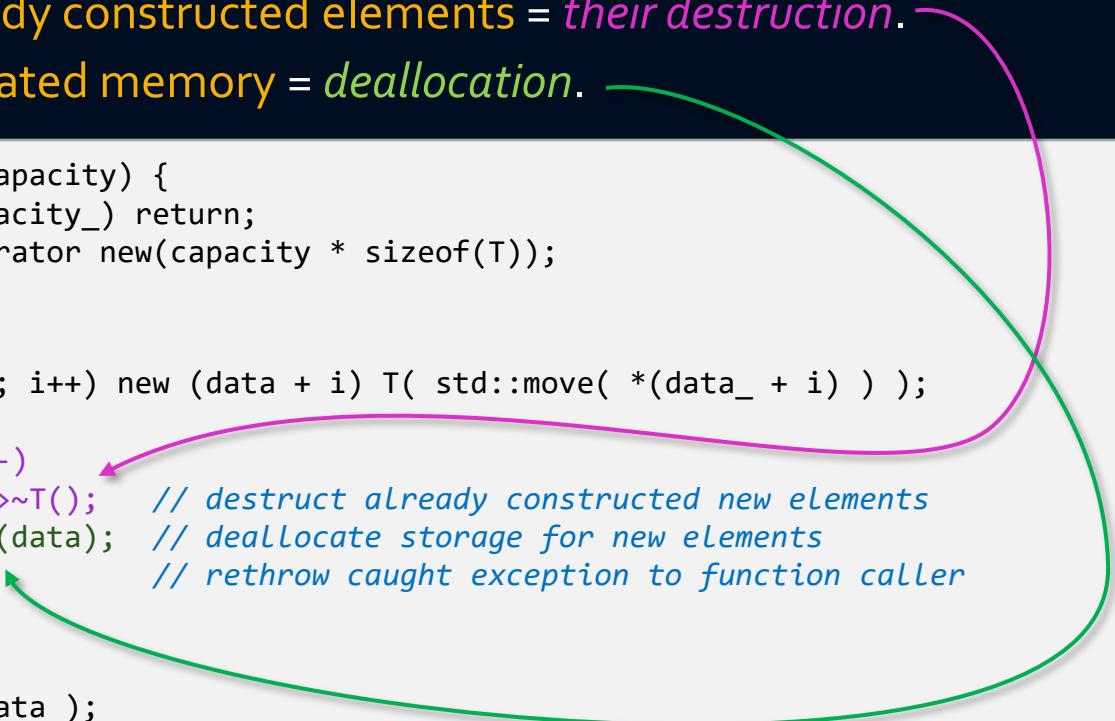
Vector — reallocation — *safety* (cont.)

- When initialization of *ith* element **throws an exception**, we need to “rollback” all the already performed operations:
 - Rollback **already constructed elements** = *their destruction*.
 - Rollback **allocated memory** = *deallocation*.

```
void reserve(size_t capacity) {
    if (capacity <= capacity_) return;
    T* data = (T*)::operator new(capacity * sizeof(T));

    size_t i = 0;
    try {
        for ( ; i < size_; i++) new (data + i) T( std::move( *(data_ + i) ) );
    } catch (...) {
        for ( ; i > 0; i--)
            (data + i - 1)->~T(); // destruct already constructed new elements
        ::operator delete(data); // deallocate storage for new elements
        throw; // rethrow caught exception to function caller
    }

    clear();
    ::operator delete(data_);
    data_ = data; capacity_ = capacity;
}
```



- *Question* — **will the state/content of the vector itself change?**
- *Answer* — **it depends.**

Vector — reallocation — *safety* (cont.)

- If new elements were initialized while *copying content* from original elements, **the answer is NO**:
 - The **original elements** — which make vector content (state) — **will remain preserved** (copying content **does not modify copied-from objects**).
- *Problem* — if new elements were initialized while *moving content* from original elements, **the answer is YES**:
 - The original elements generally **no longer have their original content** — **it was moved-from them**.
 - Instead, they are in a **"moved-from"/"empty"** state.
- Is there any solution? **No**.
 - We may **try to move content back** from new elements (before their destruction) **to the original ones**.
 - However, since *moving-content* operation may throw, it **may throw again during this "rollback"**.
 - \Rightarrow There is no guarantee that **such a rollback will succeed**.

Vector — reallocation — *safety* (cont.)

- *Conclusion:*

- *Move constructors* and *move assignment operators* should better generally not throw exceptions.
- If they do, users of such classes need to be aware of consequences.
- *Example* — when move constructor throws during Vector reallocation, (some of) its elements may end up in a modified (*empty/moved-from*) state.

- *Alternative approach* — to create vectors more “safe” with respect to exceptions, we want, during reallocation:

- 1) *move constructor* to be used for initialization of new elements if it is available and may not throw (is so-called “*non-throwing*”);
- 2) otherwise, *copy constructor* to be used if it is available;
- 3) otherwise — when there is no other possibility — “*throwing*” *move constructor* to be used (as a last unsafe instance).



Vector — reallocation — *safety* (cont.)

- Solution in our `Vector<T>::reserve` code is to **change**...

```
for ( ; i < size_; i++) new (data + i) T( std::move( *(data_ + i) ) );
```

- ...**to**:

```
for ( ; i < size_; i++) new (data + i) T( std::move_if_noexcept( *(data_ + i) ) );
```

- *Explanation:*
 - `std::move_if_noexcept` “changes a value category” of its argument to *rvalue* (which causes move constructor to be considered first) **conditionally** — namely, **only if move constructor of T is non-throwing**.
 - On the contrary, `std::move` changes a value category of its argument to rvalue *unconditionally* (in all cases).
- How to **indicate a non-throwing constructor**?

```
struct X {  
    X(X&&) noexcept { ... }  
};
```

- *Note* — if this “*non-throwing*” specification is violated (exception is thrown) program is **exited via `std::terminate` call**.

std::vector — reallocation — *safety*

- This “safer” approach is adopted by std::vector.
- *Illustrative benchmark:*
 - Insertion (“push-backing”) of elements into a vector without pre-allocation of storage (reserve).
 - *Value types* — three variants of String class without applied SSO:
 - *Variant I.* — only copy constructor (CC).
 - *Variant II.* — copy constructor + throwing move constructor (MC).
 - *Variant III.* — copy constructor + non-throwing move constructor.
 - *Results* — comparison with *Variant I.* case:
 - *Variant II.* — 1.5× faster.
 - *Variant III.* — 2.6× faster.
 - *Link* — <https://quick-bench.com/q/xv8LbXAwIhaFrR6IoQlwwHRuhEQ>.
 - *Explanation:*
 - *Variant I.* — CC is used for both *insertion* and *reallocation*.
 - *Variant II.* — MC is used for *insertion*, CC is used for *reallocation*.
 - *Variant III.* — MC is used for both *insertion* and *reallocation*.

std::vector — realloc. — safety (cont.)

- Formulation from the C++ standard regarding `std::vector::reserve` (reallocation) — [\[link\]](#):
 - *"If an exception is thrown other than by the move constructor of a non-Cpp17CopyInsertable type, there are no effects."*
- Translation to understandable form:
 - \Rightarrow *"Only if an exception is thrown ~~other than~~ by the move constructor of a non-Cpp17CopyInsertable type, there may be some effects."*
 - \Rightarrow *"Only if an exception is thrown by the move constructor of a non-copyable type, there may be some effects."*
 - \Rightarrow *"Only if an exception is thrown by the move constructor of a non-copyable type, the vector state/content may change."*
- Consequence:
 - \Rightarrow The only case when the vector state/content may change during reallocation is when its value type is non-copyable (has no copy constructor) and has throwing move constructor.
 - \Rightarrow Otherwise — if value types has either copy constructor or non-throwing move constructor, reserve that ends up with exception is guaranteed not to change the vector state.
- The same holds for inserting element at the end (`push_back`, `emplace_back`), when reallocation may happen — [\[link\]](#).

std::vector — realloc. — safety (cont.)

- Seemingly *surprising consequences*:

```
struct X {                                // non-copyable type
    X(const X&) = delete;
};

int main() {
    using T = std::map<int, X>;
    std::vector<T> v;
    v.reserve(100);
}
```

- This code:
 - **compiles well** with *libstdc++* and *libc++*,
 - **results in compilation error** with *Microsoft STL*.
 - *Live demo* — <https://godbolt.org/z/j4h14Gfae>.
- *Explanation*:
 - C++ standards do not prescribe **move constructor** of `std::map` to be *non-throwing* (noexcept).
 - However, implementations are **allowed to strengthen exception specifications** of library (member) functions.
 - *libstdc++* and *libc++* **do make** move constructor of `std::map` *non-throwing*.
 - *Microsoft STL* **does not**.

std::vector — realloc. — safety (cont.)

- *Recall* — initialization of new elements in `Vector::reserve`:

```
for ( ; i < size_; i++) new (data + i) T( std::move_if_noexcept( *(data_ + i) ) );
```

- Implementations of `std::vector::reserve` basically **use effectively the same solution** (see later).
 - ⇒ With *libstdc++* and *libc++*, the **initialization expression** inside `reserve` is **resolved as (non-throwing) move constructor call**.
 - ⇒ However, with *Microsoft STL* — due to throwing map (vector value type) move constructor — the **initialization expression** is **resolved as copy constructor call**.
- Copy constructor of `std::map` exists and **tries to copy its content**.
- This is not possible, since content of map **include objects of type X, which is non-copyable**.
- Namely, copy constructor of map **tries to call copy constructor of X, which is deleted**.
- *Note* — this is **more generic problem**:
 - **Different compilation behavior** with different implementations.
 - *Discussion*, for example — <https://stackoverflow.com/q/65140603/580083>.

std::vector — realloc. — safety (cont.)

- Error message of MSVC — 126 lines, 10909 characters.
- Important part:

```
error C2280: 'std::pair<const int,X>::pair(const std::pair<const int,X> &)': attempting to reference a deleted function
```

```
note: see declaration of 'std::pair<const int,X>::pair'
```

```
note: 'std::pair<const int,X>::pair(const std::pair<const int,X> &)': function was implicitly deleted because a data member invokes a deleted or inaccessible function 'X::X(const X &)'
```

```
note: 'X::X(const X &)': function was explicitly deleted
```

- It is very hard to find out the cause of the problem here.
 - Which is potentially throwing move constructor of std::map.
- Different behavior with different implementations kind of indicates that this is some bug in Microsoft STL.
 - Actually, it is not.
- Could concepts/constraints help here?
 - In theory yes, but even then it would not be straightforward.
 - At best, copy constructor of map could be “disabled” for non-copyable mapped type (X in our case).

Exception safety guarantees

- Important **specification regarding some operation** (such as function call) **with respect to exceptions** is:
 - whether they **may throw exceptions**,
 - **and if they do**, what **may happen with the program state** (typically, state of the involved objects).
- *Possibilities:*
 - *No-throw guarantee* — exceptions **may not be thrown**.
 - *Strong exception guarantee* — in case of exception, the **program state will not change** (operation has *no effect*).
 - *Basic exception guarantee* — in case of exception, the **program state may change but remains valid/correct**.
 - *No exception guarantee* — in case of exception, **anything may happen** (no guarantee about the program state \Rightarrow basically, it cannot continue to run safely).

Exception safety guarantees — *examples*

- *Example:*
 - With the *new terminology*, `std::vector` member functions `reserve`, `push_back`, and `emplace_back` provide strong exception guarantee if its *value type* has either copy constructor or non-throwing move constructor.
- *Example:*
 - `std::vector::clear` for *non-trivially-copyable* types involve destructor calls.
 - These destructors may, generally, throw exceptions (don't write such ones).
 - However, `clear` has `noexcept` specification \Rightarrow it provides no-throw guarantee.
 - \Rightarrow If exceptions are thrown by destructors, they must be ignored and not propagated outside of `clear` function call (which is very unsafe).
- *Example:*
 - Destructors of all library types provide no-throw guarantee:
 - “Destructor operations defined in the C++ standard library shall not throw exceptions.” [\[link\]](#)
 - *Implication* — for owner types, exceptions from destructors of owning objects are ignored.



std::vector — reallocation — *allocators*

- Final version — *really?*

```
void reserve(size_t capacity) {
    if (capacity <= capacity_) return;
    T* data = (T*)::operator new(capacity * sizeof(T));

    size_t i = 0;
    try {
        for ( ; i < size_; i++)
            new (data + i) T( std::move_if_noexcept( *(data_ + i) ) ); // can be replaced by...
                                // ... std::construct_at<T>(data + i, std::move_if_noexcept(*(data_ + i))); ...
                                // ... since C++20
    } catch (...) {
        for ( ; i > 0; i--)
            (data + i - 1)->~T(); // can be replaced by std::destroy_at(data + i - 1) since C++17
        ::operator delete(data);
        throw;
    }
    clear();
    ::operator delete(data_);
    data_ = data; capacity_ = capacity;
}
```

- Does std::vector **works this way?**
 - *By default* — **YES**.
 - *Generally* — **NO**.

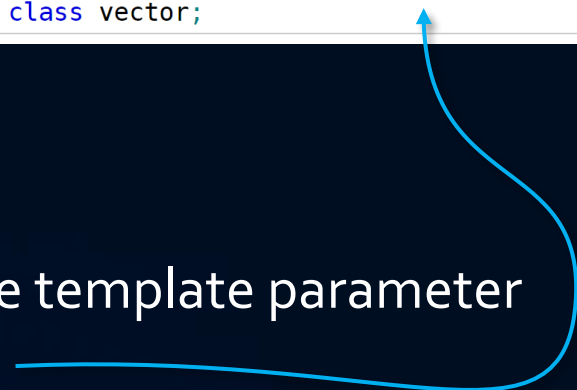
std::vector — realloc. — allocators (cont.)

- std::vector is a so-called “*allocator-aware*” container:
 - *Dynamic storage allocation/deallocation* are not performed directly by C++ allocation/deallocation functions (operator new and operator delete).
 - *Object initialization* is not performed directly by *placement-new construct*.
 - *Object destruction* is not performed directly by *(pseudo-)destructor call*.
- Instead:
 - An **allocator type** is provided as a second std::vector **template argument**.
 - This allocator type is then **used to resolve how**:
 - 1) storage will be **allocated/deallocated**,
 - 2) elements will be **initialized/destroyed**.
 - **Default template argument** for allocator type template parameter of std::vector<T> is **std::allocator<T>**.

std::vector

Defined in header <vector>

```
template<
    class T,
    class Allocator = std::allocator<T>
> class vector;
```



std::vector—realloc.—std::allocator

- With `std::allocator<T>`:
 - *Storage is allocated* by `std::allocator<T>::allocate` member function, which allocates memory by `::operator new`.
 - Link — <https://en.cppreference.com/w/cpp/memory/allocator/allocate>.
 - *Storage is deallocated* with `std::allocator<T>::deallocate` member function, which deallocates storage with `::operator delete`.
 - Link — <https://en.cppreference.com/w/cpp/memory/allocator/deallocate>.
 - *Elements are initialized* with `std::allocator_traits<std::allocator<T>>::construct`, which initializes object with *placement new* (or, `std::construct_at` since C++20).
 - Link — https://en.cppreference.com/w/cpp/memory/allocator_traits/construct.
 - *Objects are destructed* with `std::allocator_traits<std::allocator<T>>::destroy`, which destructs them with *(pseudo-)destructor call* (or, `std::destroy_at` since C++20).
 - Link — https://en.cppreference.com/w/cpp/memory/allocator_traits/destroy.
- \Rightarrow *By default* — that is with `std::allocator<T>` — `std::vector` works during reallocation *effectively the same as our Vector*.

`std::vector`—`realloc`.— *custom allocators*

- However, *customization* of all these “*special*” *allocator operations* regarding *storage* and *object-lifetime management* is possible by *providing custom allocator types*.
- *Illustrative use case 1.*:
 - *SIMD processing* — *stronger than default alignment* is required (such as 64-byte for AVX-512).
 - \Rightarrow *Custom allocate function* can provide this.
 - Instead of “*ordinary*” *operator new*, it may use:
 - *alignment-aware heap functions* (such as `posix_memalign`, `aligned_alloc`,...),
 - *special versions of operator new* that *support alignment requests* (available since C++17).
 - *Relevant discussion*: [\[Making std::vector allocate aligned memory\]](#)

std::vector — custom allocators (cont.)

- *Illustrative use case II.:*
 - For *trivially copyable types*, standard allocator causes “*zeroing out*” *binary representations* of vector elements when vector is “*resized*”.
 - For large vectors, it *may take some time* (*runtime overhead*).
 - This is *unnecessary if elements will then be rewritten anyway*.
 - A use case is, for example, *filling vector content from multiple threads* (element-insertion member functions cannot be called from multiple threads at once \Rightarrow vector needs to be “*resized*” in advance).
 - Custom allocator with *do-nothing construct* function may *resolve this issue*.
- *Relevant discussions:*
 - [\[Default-inserting into a vector isn't default initialization?\]](#)
 - [\[Should allocator construct\(\) default initialize instead of value initializing?\]](#)

Other containers — *allocators*

- *Illustrative use case III.:*
 - `std::vector` is **not the only one allocator-aware data structure** from the C++ library.
 - **Node-based allocator aware containers** (*sets, maps, their unordered variants*) may **benefit from allocators** that **use *memory pooling***.
 - *Memory pooling* employs the fact that **each node requires storage of the same size**.
 - **Memory pool-based allocators** “*merge*” heap allocations for **storage of multiple nodes into single allocation**.
 - \Rightarrow **Customization of `allocate` allocator function**.
 - Moreover, if a container is ***thread-local***, memory-pool allocators may be ***thread-local as well*** (no need for synchronization between threads).
 - *Benchmark* — **1.3× faster** insertion into ***thread-local* hash tables** (see NI-MCC for more details.)
 - *Real-world implementations:*
 - `pool_allocator` or `fast_pool_allocator` provided by ***Boost.Pool library*** [\[link\]](#).

Vectors and SBO

- Recall — *short string optimization* (SSO) = up to some limit length (number of characters), strings are stored
 - in *included storage* of the *string class owner* (its member variable buffer)
 - instead of *in dynamically-allocated storage*.
- *String class* = owner of *dynamic array of characters*.
- *Vector class* = owner of *dynamic array of elements of any type* (value type).
- Could vectors *employ the same optimization technique*?
 - Generally, it is called *small buffer (data/object/size) optimization* (SBO).
- SBO for vectors — we would want:
 - ⇒ Vectors to have *included storage* (member variable buffer) with *capacity* of some small number of vector elements.
 - *Store vector into this buffer* until their count (vector size) does not exceed this “*small buffer*” capacity.

Vectors and SBO (cont.)

- Generally, such optimization technique can be implemented.
- However — unfortunately — not for `std::vector`.
- There are requirements in the C++ standard that hinder application of SBO for `std::vector`:
 - Exemplary problem case — swapping content of two vectors.
 - Corresponding requirement — “no `swap()` function invalidates any references, pointers, or iterators referring to the elements of the containers being swapped” [\[link\]](#).
 - Problem — with SBO, for instance, swapping elements of “short” and “long” vectors do invalidate references, pointers, and iterators to elements in the included storage of the short vector.
- In theory, removing the above-quoted requirement from upcoming C++ standards might enable SBO for `std::vector` in the future.
- However, it could make the existing C++ code incorrect, which is not acceptable.
 - New C++ standards break backward compatibility in extremely rare cases.

Vectors and SBO (cont.)

- There are implementations of *vector-like classes* that **do implement SBO**.
- \Rightarrow They are **not 100% compatible with `std::vector`** and their users need to be aware of it.
- *Exemplary implementation:*
 - `small_vector` class template from *Boost.Container library* [\[link\]](#).
 - In comparison with `std::vector`, it has an **additional template non-type parameter** — *capacity of the included storage* (buffer).
 - \Rightarrow Users **may control short vectors limit sizes** according to their needs.
- Another **commonly-used alternative:**
 - LLVM's `SmallVector` [\[link\]](#).
- *Target use cases* for vectors with SBO:
 - Programs that work with **many vectors with unknown sizes** (until runtime), where **many of them has only few elements**.

Empty base (class) optimization

- Assume we wanted to **make our custom Vector allocator-aware**.
- First consideration — allocator objects **may have state**.
 - \Rightarrow The **same allocator** needs to be used through entire vector lifetime.
- Possible implementation:

```
template <typename T, typename A = std::allocator<T> > class Vector {  
    size_t capacity_, size_;  
    T* data_;  
    A alloc_; // allocator member sub-object  
public:  
    Vector( const A& alloc = A{} )  
        : capacity_(0), size_(0), data_(nullptr), alloc_(alloc) { }
```

- *Problem* — even when allocator is **state-less** (it has no sub-objects), it causes **increase of Vector storage size**:

```
std::cout << sizeof( Vector<int> ); // prints out 32 (originally, it was 24)
```

- Two causes:
 - 1) Generally, **different objects cannot share storage** \Rightarrow `alloc_` member variable **occupies at least 1 byte**.
 - 2) Due to **alignment and padding**, at least **8 bytes are effectively needed**.

Empty base (class) optimization

- Ad 1) Consider the following example:

```
X a[10]; // array of 10 elements (objects) of type X
```

- According to the C++ standard, *size of storage/binary representation of an object of type X* — that is `sizeof(X)` — equals (byte) difference between addresses of two adjacent array elements.
- These two elements **cannot share (be stored at) the same address**.
- \Rightarrow They need to be stored at **different addresses**.
- \Rightarrow **`sizeof(X)` must be at least 1**.
- Generally, this **holds for any type**, even when it is “empty” (has no member/base-class sub-objects):

```
struct X { };
```

```
std::cout << sizeof(X); // prints out 1
```

- This is also the case of the **standard library allocator**, which is *stateless* \Rightarrow does not require any sub-objects:

```
std::cout << sizeof( std::allocator<int> ); // prints out 1
```

Empty base (class) optimization

- \Rightarrow When allocator is a **member sub-object/variable** of Vector, even if it is state-less, it **occupies at least 1 byte**.
- Due to alignment and padding, it **effectively occupies 8 bytes**.
- Another option — **inheritance**:

```
template <typename T, typename A = std::allocator<T> > class Vector : public A {  
    size_t capacity_, size_;  
    T* data_;  
  
public:  
    Vector( const A& alloc = A{} )  
        : capacity_(0), size_(0), data_(nullptr), A(alloc) { }
```

```
std::cout << sizeof( Vector<int> );           // now, prints out 24 (GCC, x86_64)
```

- How is that possible?
 - It is caused by an optimization mechanism called “**empty base (class) optimization**” (EBO/EBCO).
 - *Meaning* — **inheritance** from an “empty” class **does not increase storage/binary representation requirements** of the derived class.
 - This optimization is **not mandatory**, but all mainstream implementations **do apply it**.

Public vs private inheritance

```
template <typename T, typename A = std::allocator<T> > class Vector : public A {
```

- *Note* — *public inheritance* represents the *IS-A relationship*.
- In our case, this is wrong — *vector is NOT an allocator*.
- *Unwanted consequence* — public member functions of allocator type are “*injected*” into the *Vector interface*.
 - Such as *allocate member function* might be called on a *Vector object*.
- Better version — *private inheritance*:

```
template <typename T, typename A = std::allocator<T> > class Vector : private A {
```

- *Notes*:
 - Inheriting Vector from allocator type kind of “*smells*” and can make its *code less understandable*.
 - We implemented it *only because of EBO*.
 - For this reason, *since C++20*, we can have the *same outcome* even with *composition* and *no_unique_address* attribute:

```
template <typename T, typename A = std::allocator<T> > class Vector {  
    size_t capacity_, size_;  
    T* data_;  
    [[no_unique_address]] A alloc_; // better than inheritance, but since C++20
```

Empty base (class) optimization

- Another use case — *unique pointers*:
 - Standard library unique pointers support so-called custom *deleters*.
 - *Deleter* is an object that defines function-call operator, which is applied to the managed pointer when unique pointer itself is destroyed.
- Type of *deleter* makes second template argument of `std::unique_ptr`.
- *By default*, it is `std::default_delete`, which just applies *delete expression* to the managed pointer.
- Illustrative custom implementation:

`std::unique_ptr`

Defined in header `<memory>`

```
template<
    class T,
    class Deleter = std::default_delete<T>
> class unique_ptr;
```

```
template <typename T> struct default_delete { void operator()(T* ptr) { delete ptr; } };
template <typename T, typename D = default_delete<T>> class unique_ptr {
    T* ptr_;
    D del_;
public:
    unique_ptr(T* ptr, const D& del = D{}) : ptr_(ptr), del_(del) { }
    ~unique_ptr() { del_(ptr_); }
};
```

Empty base (class) optimization

```
template <typename T> struct default_delete { void operator()(T* ptr) { delete ptr; } };  
template <typename T, typename D = default_delete<T>> class unique_ptr {  
    T* ptr_;  
    D del_;
```

- Again — with this implementation, *storage size* of unique pointer objects with *default deleter* would *grow from 8 to 16 bytes* (at 64-bit architecture)...
- ... even if this *default deleter* is *state-less* (has no sub-objects).
- This is *highly unwanted* — unique pointers should be just *efficient wrappers about normal (raw) pointers*.
- *Solution* (before C++20):

```
template <typename T, typename D = default_delete<T>> class unique_ptr : private D {  
    T* ptr_;  
public:  
    unique_ptr(T* ptr, const D& del = D{}) : ptr_(ptr), D(del) { }  
    ~unique_ptr() { D::operator()(ptr); }  
};
```


Custom deleters

- Default deleter works for object **dynamically-allocated and initialized by the new expression**:

```
{  
    std::unique_ptr<X> upx = new X{};  
} // delete expression applied automatically to the owned pointer to the allocated object  
// ...in the destructor of upx
```

- Custom deleter **allows applying different operation than delete expression**.
- **Example — MPI I/O** (see NI-PDP for more details):
 - Part of MPI library that provide **parallel access to a single file from multiple MPI processes**.
- **MPI — C API**:
 - **MPI file handler** has type **MPI_File...**
 - ...and its **address** is passed to **MPI_File_open** and **MPI_File_close** functions.

```
MPI_File f;  
MPI_File_open(..., &f);  
... // read/write from/to file simultaneously by multiple MPI processes  
MPI_File_close(&f);
```

Custom deleters (cont.)

- *Problem* — in case of **exception** thrown before `MPI_File_close`, the **MPI file handle will not be closed**, which may result in **loss of data**.
- Solution = *RAII idiom*:
 - *RAII* = *resource acquisition is initialization*.
 - *Meaning* — responsibility for resource release is **put into a destructor** of some block-local object.
 - Its **destructor is called automatically** whenever program leaves that block (reaching its end, return statement, exception, goto,...).
- *Unique pointers* — by default — **apply RAII** for objects **dynamically allocated and initialized by new expression**.
- By defining **custom deleter**, we can use them for **guaranteed correct MPI file handle closing**:

```
MPI_File f;  
MPI_File_open(..., &f);  
  
... // What if throws?  
MPI_File_close(&f);
```

```
struct MPI_deleter {  
    void operator()(MPI_File* ptr) {  
        MPI_File_close(ptr);  
    }  
};
```

```
MPI_File f;  
MPI_File_open(..., &f);  
std::unique_ptr<MPI_File, MPI_deleter> temp = &f;  
...  
// anytime temp is destructed, file is closed
```

Custom deleters — shared pointers

- Same as with *unique pointers*, *shared pointers* also support *custom deleters*.
- In contrast to unique pointers, *type of deleter* does not make a *template argument* of the shared pointer:

```
template <typename T, typename D = default_delete<T>>
class unique_ptr { D del_; ... }; // option #1 - composition (member sub-object)

template <typename T, typename D = default_delete<T>>
class unique_ptr : private D { ... }; // option #2 - inheritance (base-class sub-object)

template <typename T> class shared_ptr { ... }; // ???
```

- *Questions/problems:*
 - 1) *How is a deleter set* for shared pointers?
 - 2) *How to store it* if its type is unknown at a class scope?
- Ad 1) Similarly as for unique pointers, a deleter object is *set in constructor*.
 - ⇒ For shared pointers, *constructor must be parametrized by deleter type*:

```
template <typename D = std::default_delete<T>>
shared_ptr(T* ptr, const D& del = D{}) : ptr_(ptr) ... // ???
```

Custom deleters — shared pointers (*cont.*)

```
template <typename T> class shared_ptr {  
    T* ptr_;  
    ... // reference-counting housekeeping (pointer to control block); omitted later  
public:  
    template <typename D = std::default_delete<T>>  
    shared_ptr(T* ptr, const D& del = D{}) : ptr_(ptr) ... // ???
```

- \Rightarrow Deleter object and its type are **known in the constructor only**.
- **Its copy need to be stored somewhere** and managed by the shared pointer.
- It **cannot be stored as a sub-object** (in the included storage).
- But, **it can be dynamically allocated**:

```
template <typename T> class shared_ptr {  
    T* ptr_;  
    ???* ptr_del_;  
public:  
    template <typename D = std::default_delete<T>>  
    shared_ptr(T* ptr, const D& del = D{}) : ptr_(ptr), ptr_del_(new D{del}) { }
```

- *Problem* — **what type should that pointer be of**, when we don't know the deleter type at the class scope where it is declared?

Custom deleters — shared pointers (*cont.*)

- *Seemingly possible solution:*
 - “Type-less” pointer, that is pointer of type `void*`.
 - Any object may be pointed to by a pointer of type `void*`.

```
template <typename T> class shared_ptr {  
    T* ptr_;  
    void* ptr_del_;  
public:  
    template <typename D = std::default_delete<T>>  
    shared_ptr(T* ptr, const D& del = D{}) : ptr_(ptr), ptr_del_(new D{del}) { }  
    ~shared_ptr() { if (use_count() == 1) { ??? } }  
    long use_count() const  
    { ... /* returns number of shared_ptr instances managing current object */ }  
};
```

- *Problem* — in **destructor of `shared_ptr`**, we possibly need to:
 - *apply the deleter* to the `ptr_` pointer,
 - *apply delete expression* to the deleter itself.
- Neither can be done **without knowing the deleter type**.
 - We need to **cast `ptr_del_` to `D*`** in destructor, where **`D` is unknown**.
- \Rightarrow This is *dead end*.



Custom deleters — shared pointers (*cont.*)

- Another option — storing a deleter in a “wrapper” helper class:

```
template <typename T> class shared_ptr {  
    template <typename D> struct Helper  
    { D del_;  
      Helper(const D& del) : del_(del) { } };  
  
    T* ptr_; ???* ptr_del_; // D unknown here  
  
public:  
    template <typename D = std::default_delete<T>>  
    shared_ptr(T* ptr, const D& del = D{}) : ptr_(ptr), ptr_del_(new Helper<D>{del}) { }
```

- *Problem* — we cannot make `ptr_del_` to have `Helper<D>*` type.
- However, there is a case where **pointer of some type can point to an object of another type** — *inheritance*.
 - Namely, *pointer-to-base class can point to a derived class object*.
- In our case, this means to:
 - create a **common base class**,
 - **inherit Helper** from this base class,
 - make `ptr_del_` a **pointer to base**.

```
template <typename T> class shared_ptr {  
    struct Base { };  
  
    template <typename D>  
    struct Helper : Base {  
        D del_;  
        Helper(const D& del) : del_(del) { }  
    };  
  
    T* ptr_;  
    Base* ptr_del_;
```

Custom deleters — shared pointers (*cont.*)

```
template <typename T> class shared_ptr {
    struct Base { };
    template <typename D>
    struct Helper : Base { D del_; Helper(const D& del) : del_(del) {} };
    T* ptr_; Base* ptr_del_;
public:
    template <typename D = std::default_delete<T>>
    shared_ptr(T* ptr, const D& del = D{}) : ptr_(ptr), ptr_del_(new Helper<D>{del}) { }
    ~shared_ptr() { if (use_count() == 1) {
        // (1) apply operator() of deleter to ptr_
        // (2) delete created Helper<D> class object
    } }
```

- What remains to be done:
 - Object of type `Helper<D>` created in constructor **needs to be “deleted”**.
 - \Rightarrow Application of `delete` expression to `ptr_del_` **must invoke destructor of the `Helper<D>` class**.
 - \Rightarrow This destructor **needs to be virtual**.
 - *Note* — it also **deconstructs deleter itself**.

```
template <typename T>
class shared_ptr {
    struct Base {
        virtual ~Base() = default;
    };
    ...
    ~shared_ptr() {
        if (use_count() == 1) {
            // (1)
            delete ptr_del_;
        }
    }
}
```

Custom deleters — shared pointers (*cont.*)

- *Last problem* — **application of deleter** stored in the derived class object to `ptr_` in **destructor of `shared_ptr`**.
 - Here, only **pointer to base** is available.
 - \Rightarrow We again **need to involve a virtual function mechanism**.

```
template <typename T> class shared_ptr {
    struct Base {
        virtual ~Base() = default;
        virtual void operator()(T* ptr) = 0;
    };

    template <typename D>
    struct Helper : Base {
        D del_;
        Helper(const D& del) : del_(del) {}
        virtual void operator()(T* ptr) override { del_(ptr); }
    };

    T* ptr_; Base* ptr_del_;

public:
    template <typename D = std::default_delete<T>>
    shared_ptr(T* ptr, const D& del = D{}) : ptr_(ptr), ptr_del_(new Helper<D>{del}) { }

    ~shared_ptr() { if (use_count() == 1) {
        ptr_del_->operator()(ptr_); // same as (*ptr_del_)(ptr_);
        delete ptr_del_;
    } }
```

- *Note* — Base is **not meant to be instantiated** \Rightarrow it is made **abstract**.

Custom deleters — shared pointers (*cont.*)

- Exemplary instantiation with **value type X**, **default deleter**, and its custom `default_delete` implementation:

```
shared_ptr<X> ptr = new X;
```

```
struct default_delete<X> { void operator()(X* ptr) { delete ptr; } };
class shared_ptr<X> {
    struct Base {
        virtual ~Base() = default;
        virtual void operator()(X* ptr) = 0;
    };
    struct Helper<default_delete<X>> : Base {
        default_delete<X> del_;
        Helper(const default_delete<X>& del) : del_(del) {}
        virtual void operator()(X* ptr) override { del_(ptr); }
    };
    X* ptr_;
    Base* ptr_del_;
public:
    shared_ptr<default_delete<X>>(X* ptr, const default_delete<X>& del = default_delete<X>{})
        : ptr_(ptr), ptr_del_(new Helper<default_delete<X>>{del}) { }
    ~shared_ptr() { if (use_count() == 1) {
        ptr_del_>operator()(ptr_);
        delete ptr_del_;
    } }
};
```

Custom deleters — shared pointers (*cont.*)

- In case of **additional** instantiation with the **same value type X**, but a **different custom deleter**, the **resulting instance** will look like:

```
shared_ptr<X> ptr = new X;    shared_ptr<X> ptr2 { new X, custom_delete<X>{} };

struct default_delete<X> { void operator()(X* ptr) { delete ptr; } };
struct custom_delete<X> { void operator()(X* ptr) { ... } };

class shared_ptr<X> {
    struct Base {
        virtual ~Base() = default;
        virtual void operator()(X* ptr) = 0;
    };
    struct Helper<default_delete<X>> : Base {
        default_delete<X> del_;
        Helper(const default_delete<X>& del) : del_(del) {}
        virtual void operator()(X* ptr) override { del_(ptr); }
    };
    struct Helper<custom_delete<X>> : Base {
        custom_delete<X> del_;
        Helper(const custom_delete<X>& del) : del_(del) {}
        virtual void operator()(X* ptr) override { del_(ptr); }
    };
    X* ptr_;
    Base* ptr_del_;

public:
    shared_ptr<default_delete<X>>(X* ptr, const default_delete<X>& del = default_delete<X>{})
        : ptr_(ptr), ptr_del_(new Helper<default_delete<X>>{del}) {}
    shared_ptr<custom_delete<X>>(X* ptr, const custom_delete<X>& del = custom_delete<X>{})
        : ptr_(ptr), ptr_del_(new Helper<custom_delete<X>>{del}) {}
    ~shared_ptr() { if (use_count() == 1) { ptr_del_->operator()(ptr_); delete ptr_del_; } }
```


Deleters — unique vs shared ptrs.

- *Observation:*
 - Type of all `shared_ptr` objects with the same value type but different deleter types **is the same**.
 - It is the **same class** = instance of `shared_ptr` class template.
 - Different type of deleter just causes **different instance of shared pointer constructor** to be generated within this class (and, possibly, some implementation inner helper class).
 - On the contrary, **for each type of deleter** `unique_ptr` class template is **instantiated into different class**.
 - Types of unique pointers with different deleters are different and unrelated classes.
- In shared pointers, type of deleter is so-called “*type-erased*”.
 - The technique for hiding its type is called “*type erasure*”.
- Type-erased deleter is **more “powerful”** but has **runtime overhead**.
 - ⇒ Reason why **unique pointers do not employ this technique**.
 - Recall — unique pointer = **lightweight RAII pointer wrapper** with **minimal overhead**.