

# Effective C++ Programming

NIE-EPC (v. 2021):  
EXCEPTION SAFETY, ALLOCATORS, SMALL BUFFER  
OPTIMIZATION, EMPTY BASE OPTIMIZATION, TYPE  
ERASURE  
© 2021 DANIEL LANGR, ČVUT (CTU) FIT

1

## Vector — reallocation — *safety concerns*

- *Not-exactly-final* solution for reserve member function:

```
void reserve(size_t capacity) {
    if (capacity <= capacity_) return;
    T* data = (T*)::operator new(capacity * sizeof(T));
    for (size_t i = 0; i < size_; i++)
        new (data + i) T( std::move( *(data_ + i) ) ); // what if throws ???
    clear();
    ::operator delete(data_);
    data_ = data; capacity_ = capacity;
}
```

- In case of *non-trivially-copyable types*, initialization involves calls of copy or move constructor, which, generally, may throw exceptions.
- Such situations need to be handled:
  - some elements may have been already constructed before,
  - new storage has been already allocated.
- *Note* — for sake of simplicity, we ignored this issue previously.
- *Note* — initialization of objects of *trivially copyable types* with content-copying here cannot result in exception.

2

## Vector — reallocation — *safety (cont.)*

- When initialization of *ith* element throws an exception, we need to “rollback” all the already performed operations:
  - Rollback already constructed elements = *their destruction*.
  - Rollback allocated memory = *deallocation*.

```
void reserve(size_t capacity) {
    if (capacity <= capacity_) return;
    T* data = (T*)::operator new(capacity * sizeof(T));
    size_t i = 0;
    try {
        for ( ; i < size_; i++) new (data + i) T( std::move( *(data_ + i) ) );
    } catch (...) {
        for ( ; i > 0; i--)
            (data + i - 1)->~T(); // destruct already constructed new elements
        ::operator delete(data); // deallocate storage for new elements
        throw; // rethrow caught exception to function caller
    }
    clear();
    ::operator delete(data_);
    data_ = data; capacity_ = capacity;
}
```

- *Question* — will the state/content of the vector itself change?
- *Answer* — it depends.

3

## Vector — reallocation — *safety (cont.)*

- If new elements were initialized while *copying content* from original elements, the answer is **NO**:
  - The original elements — which make vector content (state) — will remain preserved (copying content does not modify copied-from objects).
- *Problem* — if new elements were initialized while *moving content* from original elements, the answer is **YES**:
  - The original elements generally no longer have their original content — it was *moved-from* them.
  - Instead, they are in a “*moved-from*”/“*empty*” state.
- Is there any solution? No.
  - We may try to move content back from new elements (before their destruction) to the original ones.
  - However, since *moving-content* operation may throw, it may throw again during this “rollback”.
  - ⇒ There is no guarantee that such a rollback will succeed.

4

## Vector — reallocation — *safety (cont.)*

- *Conclusion*:
  - *Move constructors* and *move assignment operators* should better generally not throw exceptions.
  - If they do, users of such classes need to be aware of consequences.
  - *Example* — when move constructor throws during Vector reallocation, (some of) its elements may end up in a modified (*empty/moved-from*) state.
- *Alternative approach* — to create vectors more “*safe*” with respect to exceptions, we want, during reallocation:
  - 1) *move constructor* to be used for initialization of new elements if it is available and may not throw (is so-called “*non-throwing*”);
  - 2) otherwise, *copy constructor* to be used if it is available;
  - 3) otherwise — when there is no other possibility — “*throwing*” *move constructor* to be used (as a *last unsafe instance*).

5

## Vector — reallocation — *safety (cont.)*

- Solution in our Vector<T>::reserve code is to change...

```
for ( ; i < size_; i++) new (data + i) T( std::move( *(data_ + i) ) );
```

- ...to:

```
for ( ; i < size_; i++) new (data + i) T( std::move_if_noexcept( *(data_ + i) ) );
```

- *Explanation*:

- `std::move_if_noexcept` “changes a value category” of its argument to *rvalue* (which causes move constructor to be considered first) conditionally — namely, only if move constructor of T is *non-throwing*.
- On the contrary, `std::move` changes a value category or its argument to *rvalue unconditionally* (in all cases).

- How to indicate a non-throwing constructor?

```
struct X {
    X(X&&) noexcept { ... }
};
```

- *Note* — if this “*non-throwing*” specification is violated (exception is thrown) program is exited via `std::terminate` call.

6

## std::vector — reallocation — safety

- This “safer” approach is adopted by std::vector.
- **Illustrative benchmark:**
  - Insertion (“push-back”) of elements into a vector without pre-allocation of storage (reserve).
  - **Value types** — three variants of String class without applied SSO:
    - Variant I. — only copy constructor (CC).
    - Variant II. — copy constructor + throwing move constructor (MC).
    - Variant III. — copy constructor + non-throwing move constructor.
  - **Results** — comparison with Variant I. case:
    - Variant II. — **1.5× faster**.
    - Variant III. — **2.6× faster**.
  - **Link** — <https://quick-bench.com/q/xv8LbXAwIhaFrR6IoQlwwHRuHEQ>.
  - **Explanation:**
    - Variant I. — CC is used for both *insertion* and *reallocation*.
    - Variant II. — MC is used for *insertion*, CC is used for *reallocation*.
    - Variant III. — MC is used for both *insertion* and *reallocation*.

7

## std::vector — realloc. — safety (cont.)

- Formulation from the C++ standard regarding std::vector::reserve (reallocation) — [link]:
  - “If an exception is thrown other than by the move constructor of a non-Cpp17CopyInsertable type, there are no effects.”
- Translation to understandable form:
  - ⇒ “Only if an exception is thrown **other than** by the move constructor of a non-Cpp17CopyInsertable type, there may be some effects.”
  - ⇒ “Only if an exception is thrown by the move constructor of a non-copyable type, there may be some effects.”
  - ⇒ “Only if an exception is thrown by the move constructor of a non-copyable type, the vector state/content may change.”
- **Consequence:**
  - ⇒ The only case when the vector state/content may change during reallocation is when its value type is non-copyable (has no copy constructor) and has throwing move constructor.
  - ⇒ Otherwise — if value types has either copy constructor or non-throwing move constructor, reserve that ends up with exception is guaranteed not to change the vector state.
- The same holds for inserting element at the end (push\_back, emplace\_back), when reallocation may happen — [link].

8

## std::vector — realloc. — safety (cont.)

- **Seemingly surprising consequences:**

```
struct X {
    X(const X&) = delete;    // non-copyable type
};

int main() {
    using T = std::map<int, X>;
    std::vector<T> v;
    v.reserve(100);
}
```

- This code:
  - compiles well with libstdc++ and libc++,
  - results in compilation error with Microsoft STL.
- **Live demo** — <https://godbolt.org/z/j4h14Gfae>.
- **Explanation:**
  - C++ standards do not prescribe move constructor of std::map to be non-throwing (noexcept).
  - However, implementations are allowed to strengthen exception specifications of library (member) functions.
  - libstdc++ and libc++ do make move constructor of std::map non-throwing.
  - Microsoft STL does not.

9

## std::vector — realloc. — safety (cont.)

- **Recall** — initialization of new elements in Vector::reserve:
 

```
for ( ; i < size_; i++) new (data + i) T( std::move_if_noexcept( *(data_ + i) ) );
```
- Implementations of std::vector::reserve basically use effectively the same solution (see later).
  - ⇒ With libstdc++ and libc++, the initialization expression inside reserve is resolved as (non-throwing) **move constructor** call.
  - ⇒ However, with Microsoft STL — due to throwing map (vector value type) move constructor — the initialization expression is resolved as **copy constructor** call.
- Copy constructor of std::map exists and tries to copy its content.
- This is not possible, since content of map include objects of type X, which is non-copyable.
- Namely, copy constructor of map tries to call copy constructor of X, which is deleted.
- **Note** — this is more generic problem:
  - Different compilation behavior with different implementations.
  - Discussion, for example — <https://stackoverflow.com/q/65140603/580083>.

10

## std::vector — realloc. — safety (cont.)

- **Error message of MSVC** — 126 lines, 10909 characters.
- **Important part:**

```
error C2280: 'std::pair<const int,X>::pair(const std::pair<const int,X> &)' : attempting to
reference a deleted function
note: see declaration of 'std::pair<const int,X>::pair'
note: 'std::pair<const int,X>::pair(const std::pair<const int,X> &)' : function was
implicitly deleted because a data member invokes a deleted or inaccessible function
'X::X(const X &)'
note: 'X::X(const X &)' : function was explicitly deleted
```

- It is very hard to find out the cause of the problem here.
  - Which is potentially throwing move constructor of std::map.
- Different behavior with different implementations kind of indicates that this is some bug in Microsoft STL.
  - Actually, it is not.
- Could concepts/constraints help here?
  - In theory yes, but even then it would not be straightforward.
  - At best, copy constructor of map could be “disabled” for non-copyable mapped type (X in our case).

11

## Exception safety guarantees

- Important specification regarding some operation (such as function call) with respect to exceptions is:
  - whether they may throw exceptions,
  - and if they do, what may happen with the program state (typically, state of the involved objects).
- **Possibilities:**
  - **No-throw guarantee** — exceptions may not be thrown.
  - **Strong exception guarantee** — in case of exception, the program state will not change (operation has no effect).
  - **Basic exception guarantee** — in case of exception, the program state may change but remains valid/correct.
  - **No exception guarantee** — in case of exception, anything may happen (no guarantee about the program state ⇒ basically, it cannot continue to run safely).

12

## Exception safety guarantees — examples

- **Example:**
  - With the *new terminology*, `std::vector` member functions `reserve`, `push_back`, and `emplace_back` provide strong exception guarantee if its *value type* has either copy constructor or non-throwing move constructor.
- **Example:**
  - `std::vector::clear` for *non-trivially-copyable* types involve destructor calls.
  - These destructors may, generally, throw exceptions (don't write such ones).
  - However, `clear` has `noexcept` specification  $\Rightarrow$  it provides no-throw guarantee.
  - $\Rightarrow$  If exceptions are thrown by destructors, they must be ignored and not propagated outside of `clear` function call (which is very unsafe).
- **Example:**
  - Destructors of all library types provide no-throw guarantee:
    - "Destructor operations defined in the C++ standard library shall not throw exceptions." [\[link\]](#)
  - **Implication** — for owner types, exceptions from destructors of owning objects are ignored.



33

13

## `std::vector` — reallocation — allocators

- Final version — really?

```
void reserve(size_t capacity) {
    if (capacity <= capacity_) return;
    T* data = (T*)::operator new(capacity * sizeof(T));
    size_t i = 0;
    try {
        for (; i < size_; i++)
            new (data + i) T( std::move_if_noexcept( *(data_ + i) ) ); // can be replaced by...
        // ... std::construct_at<T>(data + i, std::move_if_noexcept(*(data_ + i))); ...
        // ... since C++20
    } catch (...) {
        for (; i > 0; i--)
            (data + i - 1)->~T(); // can be replaced by std::destroy_at(data + i - 1) since C++17
        ::operator delete(data);
        throw;
    }
    clear();
    ::operator delete(data_);
    data_ = data; capacity_ = capacity;
}
```

- Does `std::vector` works this way?
  - **By default** — YES.
  - **Generally** — NO.

34

14

## `std::vector` — realloc. — allocators (cont.)

- `std::vector` is a so-called "*allocator-aware*" container:
  - *Dynamic storage allocation/deallocation* are not performed directly by C++ allocation/deallocation functions (`operator new` and `operator delete`).
  - *Object initialization* is not performed directly by *placement-new* construct.
  - *Object destruction* is not performed directly by (*pseudo*-)destructor call.
- **Instead:**
  - An **allocator type** is provided as a second `std::vector` template argument.
  - This allocator type is then used to resolve how:
    - 1) storage will be allocated/deallocated,
    - 2) elements will be initialized/destructed.
  - **Default template argument** for allocator type template parameter of `std::vector<T>` is `std::allocator<T>`.

```
std::vector
    Defined in header <vector>
    template<
        class T,
        class Allocator = std::allocator<T>
    > class vector;
```

35

15

## `std::vector`—realloc.—`std::allocator`

- With `std::allocator<T>`:
  - *Storage is allocated* by `std::allocator<T>::allocate` member function, which allocates memory by `::operator new`.
  - [Link — https://en.cppreference.com/w/cpp/memory/allocator/allocate](https://en.cppreference.com/w/cpp/memory/allocator/allocate).
  - *Storage is deallocated* with `std::allocator<T>::deallocate` member function, which deallocates storage with `::operator delete`.
  - [Link — https://en.cppreference.com/w/cpp/memory/allocator/deallocate](https://en.cppreference.com/w/cpp/memory/allocator/deallocate).
  - *Elements are initialized* with `std::allocator_traits<std::allocator<T>>::construct`, which initializes object with *placement new* (or, `std::construct_at` since C++20).
  - [Link — https://en.cppreference.com/w/cpp/memory/allocator\\_traits/construct](https://en.cppreference.com/w/cpp/memory/allocator_traits/construct).
  - *Objects are destructed* with `std::allocator_traits<std::allocator<T>>::destroy`, which destructs them with (*pseudo*-)destructor call (or, `std::destroy_at` since C++20).
  - [Link — https://en.cppreference.com/w/cpp/memory/allocator\\_traits/destroy](https://en.cppreference.com/w/cpp/memory/allocator_traits/destroy).
- $\Rightarrow$  **By default** — that is with `std::allocator<T>` — `std::vector` works during reallocation effectively the same as our `Vector`.

36

16

## `std::vector`—realloc.— custom allocators

- However, *customization* of all these "*special*" *allocator operations* regarding storage and object-lifetime management is possible by providing custom allocator types.
- **Illustrative use case I.:**
  - *SIMD processing* — stronger than default alignment is required (such as 64-byte for AVX-512).
  - $\Rightarrow$  Custom allocate function can provide this.
  - Instead of "ordinary" `operator new`, it may use:
    - *alignment-aware heap functions* (such as `posix_memalign`, `aligned_alloc`,...).
    - special versions of `operator new` that support alignment requests (available since C++17).
  - **Relevant discussion:** [\[Making std::vector allocate aligned memory\]](#)

37

17

## `std::vector` — custom allocators (cont.)

- **Illustrative use case II.:**
  - For *trivially copyable* types, standard allocator causes "*zeroing out*" binary representations of vector elements when vector is "*resized*".
  - For large vectors, it may take some time (*runtime overhead*).
  - This is unnecessary if elements will then be rewritten anyway.
    - A use case is, for example, filling vector content from multiple threads (element-insertion member functions cannot be called from multiple threads at once  $\Rightarrow$  vector needs to be "*resized*" in advance).
  - Custom allocator with *do-nothing* `construct` function may resolve this issue.
- **Relevant discussions:**
  - [\[Default-inserting into a vector isn't default initialization?\]](#)
  - [\[Should allocator construct\(\) default initialize instead of value initializing?\]](#)

38

18

## Other containers — *allocators*

- *Illustrative use case III.*:
  - `std::vector` is not the only one allocator-aware data structure from the C++ library.
  - Node-based allocator aware containers (*sets*, *maps*, their *unordered variants*) may benefit from allocators that use **memory pooling**.
  - *Memory pooling* employs the fact that each node requires storage of the same size.
    - Memory pool-based allocators “merge” heap allocations for storage of multiple nodes into single allocation.
    - $\Rightarrow$  Customization of allocate allocator function.
  - Moreover, if a container is *thread-local*, memory-pool allocators may be *thread-local* as well (no need for synchronization between threads).
  - *Benchmark* — **1.3x faster** insertion into *thread-local* hash tables (see NI-MCC for more details.)
  - *Real-world implementations*:
    - `pool_allocator` or `fast_pool_allocator` provided by *Boost.Pool* library [\[link\]](#).

19

## Vectors and SBO

- *Recall* — *short string optimization* (SSO) = up to some limit length (number of characters), strings are stored
  - in *included storage* of the string class owner (its member variable buffer)
  - instead of in *dynamically-allocated storage*.
- *String class* = owner of dynamic array of *characters*.
- *Vector class* = owner of dynamic array of *elements of any type* (value type).
- Could vectors employ the same optimization technique?
  - Generally, it is called *small buffer (data/object/size) optimization* (SBO).
- SBO for vectors — we would want:
  - $\Rightarrow$  Vectors to have included storage (member variable buffer) with *capacity* of some small number of vector elements.
  - Store vector into this buffer until their count (vector size) does not exceed this “small buffer” capacity.

20

## Vectors and SBO (cont.)

- Generally, such optimization technique can be implemented.
- However — unfortunately — not for `std::vector`.
- There are requirements in the C++ standard that hinder application of SBO for `std::vector`:
  - *Exemplary problem case* — swapping content of two vectors.
  - *Corresponding requirement* — “no `swap()` function invalidates any references, pointers, or iterators referring to the elements of the containers being swapped” [\[link\]](#).
  - *Problem* — with SBO, for instance, swapping elements of “short” and “long” vectors do invalidate references, pointers, and iterators to elements in the included storage of the short vector.
- In theory, removing the above-quoted requirement from upcoming C++ standards might enable SBO for `std::vector` in the future.
- However, it could make the existing C++ code incorrect, which is not acceptable.
- New C++ standards break backward compatibility in extremely rare cases.

21

## Vectors and SBO (cont.)

- There are implementations of *vector-like classes* that do implement SBO.
- $\Rightarrow$  They are not 100% compatible with `std::vector` and their users need to be aware of it.
- *Exemplary implementation*:
  - `small_vector` class template from *Boost.Container* library [\[link\]](#).
  - In comparison with `std::vector`, it has an additional template non-type parameter — *capacity of the included storage* (buffer).
  - $\Rightarrow$  Users may control *short* vectors limit sizes according to their needs.
- Another commonly-used alternative:
  - *LLVM’s SmallVector* [\[link\]](#).
- *Target use cases* for vectors with SBO:
  - Programs that work with many vectors with unknown sizes (until runtime), where many of them has **only few elements**.

22

## Empty base (class) optimization

- Assume we wanted to make our custom *Vector* *allocator-aware*.
- First consideration — allocator objects may have state.
  - $\Rightarrow$  The same allocator needs to be used through entire vector lifetime.
- Possible implementation:
 

```
template <typename T, typename A = std::allocator<T> > class Vector {
    size_t capacity_, size_;
    T* data_;
    A alloc_; // allocator member sub-object
public:
    Vector( const A& alloc = A() )
        : capacity_(0), size_(0), data_(nullptr), alloc_(alloc) { }
};
```
- *Problem* — even when allocator is state-less (it has no sub-objects), it causes increase of *Vector* storage size:
 

```
std::cout << sizeof( Vector<int> ); // prints out 32 (originally, it was 24)
```
- Two causes:
  - 1) Generally, different objects cannot share storage  $\Rightarrow$  `alloc_` member variable occupies at least 1 byte.
  - 2) Due to *alignment and padding*, at least 8 bytes are effectively needed.

23

## Empty base (class) optimization

- Ad 1) Consider the following example:
 

```
X a[10]; // array of 10 elements (objects) of type X
```
- According to the C++ standard, *size of storage/binary representation* of an object of type `X` — that is `sizeof(X)` — equals (byte) difference between addresses of two adjacent array elements.
- These two elements cannot share (be stored at) the same address.
- $\Rightarrow$  They need to be stored at different addresses.
- $\Rightarrow$  `sizeof(X)` must be at least 1.
- Generally, this holds for any type, even when it is “empty” (has no member/base-class sub-objects):
 

```
struct X { };
```

```
std::cout << sizeof(X); // prints out 1
```
- This is also the case of the standard library allocator, which is *state-less*  $\Rightarrow$  does not require any sub-objects:
 

```
std::cout << sizeof( std::allocator<int> ); // prints out 1
```

24

## Empty base (class) optimization

- ⇒ When allocator is a member sub-object/variable of Vector, even if it is state-less, it occupies at least 1 byte.
- Due to alignment and padding, it effectively occupies 8 bytes.
- Another option — *inheritance*:

```
template <typename T, typename A = std::allocator<T>> class Vector : public A {
    size_t capacity_, size_;
    T* data_;
public:
    Vector( const A& alloc = A() )
        : capacity_(0), size_(0), data_(nullptr), A(alloc) { }

    std::cout << sizeof( Vector<int> ); // now, prints out 24 (GCC, x86_64)
```

### How is that possible?

- It is caused by an optimization mechanism called "empty base (class) optimization" (EBO/EBCO).
- Meaning* — inheritance from an "empty" class does not increase storage/binary representation requirements of the derived class.
- This optimization is not mandatory, but all mainstream implementations do apply it.

25

## Public vs private inheritance

```
template <typename T, typename A = std::allocator<T>> class Vector : public A {
```

- Note* — *public inheritance* represents the IS-A relationship.
- In our case, this is wrong — vector is NOT an allocator.
- Unwanted consequence* — public member functions of allocator type are "injected" into the Vector interface.
  - Such as allocate member function might be called on a Vector object.
- Better version — *private inheritance*:

```
template <typename T, typename A = std::allocator<T>> class Vector : private A {
```

### Notes:

- Inheriting Vector from allocator type kind of "smells" and can make its code less understandable.
- We implemented it only because of EBO.
- For this reason, since C++20, we can have the same outcome even with *composition* and *no\_unique\_address* attribute:

```
template <typename T, typename A = std::allocator<T>> class Vector {
    size_t capacity_, size_;
    T* data_;
    [[no_unique_address]] A alloc_; // better than inheritance, but since C++20
```

26

## Empty base (class) optimization

- Another use case — *unique pointers*:
  - Standard library unique pointers support so-called custom *deleters*.
  - Deleter* is an object that defines function-call operator, which is applied to the managed pointer when unique pointer itself is destroyed.
- Type of *deleter* makes second template argument of `std::unique_ptr`.
- By default, it is `std::default_delete`, `std::unique_ptr` which just applies *delete expression* to the managed pointer.
- Illustrative custom implementation:

```
template <typename T> struct default_delete { void operator()(T* ptr) { delete ptr; } };
template <typename T, typename D = default_delete<T>> class unique_ptr {
    T* ptr_;
    D del_;
public:
    unique_ptr(T* ptr, const D& del = D()) : ptr_(ptr), del_(del) { }
    ~unique_ptr() { del_(ptr_); }
};
```

27

## Empty base (class) optimization

```
template <typename T> struct default_delete { void operator()(T* ptr) { delete ptr; } };
template <typename T, typename D = default_delete<T>> class unique_ptr {
    T* ptr_;
    D del_;
```

- Again — with this implementation, *storage size* of unique pointer objects with *default deleter* would grow from 8 to 16 bytes (at 64-bit architecture)...
- ... even if this *default deleter* is *state-less* (has no sub-objects).
- This is highly unwanted — unique pointers should be just efficient wrappers about normal (raw) pointers.
- Solution* (before C++20):

```
template <typename T, typename D = default_delete<T>> class unique_ptr : private D {
    T* ptr_;
public:
    unique_ptr(T* ptr, const D& del = D()) : ptr_(ptr), D(del) { }
    ~unique_ptr() { D::operator()(ptr); }
};
```

28

## Custom deleters

- Default deleter works for object dynamically-allocated and initialized by the new *expression*:

```
{
    std::unique_ptr<X> upx = new X();
} // delete expression applied automatically to the owned pointer to the allocated object
// ... in the destructor of upx
```

- Custom deleter allows applying different operation than *delete expression*.
- Example* — MPI I/O (see NI-PDP for more details):
  - Part of MPI library that provide parallel access to a single file from multiple MPI processes.
- MPI — C API:
  - MPI file handler has type `MPI_File`...
  - ...and its address is passed to `MPI_File_open` and `MPI_File_close` functions.

```
MPI_File f;
MPI_File_open(..., &f);
... // read/write from/to file simultaneously by multiple MPI processes
MPI_File_close(&f);
```

29

## Custom deleters (cont.)

- Problem* — in case of exception thrown before `MPI_File_close`, the MPI file handle will not be closed, which may result in loss of data.
- Solution* = *RAII idiom*:
  - RAII* = *resource acquisition is initialization*.
  - Meaning* — responsibility for resource release is put into a destructor of some block-local object.
  - Its destructor is called automatically whenever program leaves that block (reaching its end, return statement, exception, goto,...).
- Unique pointers* — by default — apply RAII for objects dynamically allocated and initialized by new expression.
- By defining custom deleter, we can use them for guaranteed correct MPI file handle closing:

```
struct MPI_deleter {
    void operator()(MPI_File* ptr) {
        MPI_File_close(ptr);
    }
};
```

```
MPI_File f;
MPI_File_open(..., &f);
std::unique_ptr<MPI_File, MPI_deleter> temp = &f;
...
// anytime temp is destructed, file is closed
```

30

## Custom deleters — shared pointers

- Same as with *unique pointers*, *shared pointers* also support *custom deleters*.
- In contrast to unique pointers, type of deleter does not make a template argument of the shared pointer:

```
template <typename T, typename D = default_deleter<T>>
class unique_ptr { D del_; ... }; // option #1 - composition (member sub-object)
template <typename T, typename D = default_deleter<T>>
class unique_ptr; private D { ... }; // option #2 - inheritance (base-class sub-object)
template <typename T> class shared_ptr { ... }; // ???
```

### Questions/problems:

- How is a deleter set for shared pointers?
  - How to store it if its type is unknown at a class scope?
- Ad 1) Similarly as for unique pointers, a deleter object is set in constructor.
- ⇒ For shared pointers, constructor must be parametrized by deleter type:

```
template <typename D = std::default_deleter<T>>
shared_ptr(T* ptr, const D& del = D()) : ptr_(ptr) ... // ???
```

31

## Custom deleters — shared pointers (cont.)

```
template <typename T> class shared_ptr {
    T* ptr_;
    ... // reference-counting housekeeping (pointer to control block); omitted later
public:
    template <typename D = std::default_deleter<T>>
    shared_ptr(T* ptr, const D& del = D()) : ptr_(ptr) ... // ???
```

- ⇒ Deleter object and its type are known in the constructor only.
- Its copy need to be stored somewhere and managed by the shared pointer.
- It cannot be stored as a sub-object (in the included storage).
- But, it can be dynamically allocated:

```
template <typename T> class shared_ptr {
    T* ptr_;
    ???* ptr_del_;
public:
    template <typename D = std::default_deleter<T>>
    shared_ptr(T* ptr, const D& del = D()) : ptr_(ptr), ptr_del_(new D(del)) { }
```

- Problem** — what type should that pointer be of, when we don't know the deleter type at the class scope where it is declared?

32

## Custom deleters — shared pointers (cont.)

- Seemingly possible solution:*
  - "Type-less" pointer, that is pointer of type `void*`.
  - Any object may be pointed to by a pointer of type `void*`.

```
template <typename T> class shared_ptr {
    T* ptr_;
    void* ptr_del_;
public:
    template <typename D = std::default_deleter<T>>
    shared_ptr(T* ptr, const D& del = D()) : ptr_(ptr), ptr_del_(new D(del)) { }
    ~shared_ptr() { if (use_count() == 1) { ??? } }
    long use_count() const
    { ... /* returns number of shared_ptr instances managing current object */ }
```

- Problem** — in destructor of `shared_ptr`, we possibly need to:
  - apply the deleter to the `ptr_` pointer,
  - apply delete expression to the deleter itself.
- Neither can be done without knowing the deleter type.
- We need to cast `ptr_del_` to `D*` in destructor, where `D` is unknown.
- ⇒ This is *dead end*.



33

## Custom deleters — shared pointers (cont.)

- Another option — storing a deleter in a "wrapper" helper class:

```
template <typename T> class shared_ptr {
    template <typename D> struct Helper
    { D del_;
      Helper(const D& del) : del_(del) { } };
    T* ptr_;
    ???* ptr_del_; // D unknown here
public:
    template <typename D = std::default_deleter<T>>
    shared_ptr(T* ptr, const D& del = D()) : ptr_(ptr), ptr_del_(new Helper<D>(del)) { }
```

- Problem** — we cannot make `ptr_del_` to have `Helper<D>*` type.
- However, there is a case where pointer of some type can point to an object of another type — *inheritance*.
  - Namely, *pointer-to-base class* can point to a *derived class object*.
- In our case, this means to:
  - create a common base class,
  - inherit `Helper` from this base class,
  - make `ptr_del_` a pointer to base.

```
template <typename T> class shared_ptr {
    struct Base { };
    template <typename D>
    struct Helper : Base {
        D del_;
        Helper(const D& del) : del_(del) { }
    };
    T* ptr_;
    Base* ptr_del_;
```

34

## Custom deleters — shared pointers (cont.)

```
template <typename T> class shared_ptr {
    struct Base { };
    template <typename D>
    struct Helper : Base { D del_; Helper(const D& del) : del_(del) { } };
    T* ptr_;
    Base* ptr_del_;
public:
    template <typename D = std::default_deleter<T>>
    shared_ptr(T* ptr, const D& del = D()) : ptr_(ptr), ptr_del_(new Helper<D>(del)) { }
    ~shared_ptr() { if (use_count() == 1) {
        // (1) apply operator() of deleter to ptr_
        // (2) delete created Helper<D> class object
    } }
```

- What remains to be done:
  - Object of type `Helper<D>` created in constructor needs to be "deleted".
  - ⇒ Application of `delete` expression to `ptr_del_` must invoke destructor of the `Helper<D>` class.
  - ⇒ This destructor needs to be virtual.
  - Note — it also destructs deleter itself.

```
template <typename T>
class shared_ptr {
    struct Base {
        virtual ~Base() = default;
    };
    ...
    ~shared_ptr() {
        if (use_count() == 1) {
            // (1)
            delete ptr_del_;
        }
    }
```

35

## Custom deleters — shared pointers (cont.)

- Last problem** — application of deleter stored in the derived class object to `ptr_` in destructor of `shared_ptr`.
  - Here, only pointer to base is available.
  - ⇒ We again need to involve a virtual function mechanism.

```
template <typename T> class shared_ptr {
    struct Base {
        virtual ~Base() = default;
        virtual void operator()(T* ptr) = 0;
    };
    template <typename D>
    struct Helper : Base {
        D del_;
        Helper(const D& del) : del_(del) { }
        virtual void operator()(T* ptr) override { del_(ptr); }
    };
    T* ptr_;
    Base* ptr_del_;
public:
    template <typename D = std::default_deleter<T>>
    shared_ptr(T* ptr, const D& del = D()) : ptr_(ptr), ptr_del_(new Helper<D>(del)) { }
    ~shared_ptr() { if (use_count() == 1) {
        ptr_del_>operator()(ptr_); // same as (*ptr_del_)(ptr_);
        delete ptr_del_;
    } }
```

- Note** — Base is not meant to be instantiated ⇒ it is made *abstract*.

36

### Custom deleters — shared pointers (cont.)

- Exemplary instantiation with value type X, default deleter, and its custom default\_delete implementation:

```
shared_ptr<X> ptr = new X;

struct default_delete<X> { void operator()(X* ptr) { delete ptr; } };
class shared_ptr<X> {
    struct Base {
        virtual ~Base() = default;
        virtual void operator()(X* ptr) = 0;
    };
    struct Helper<default_delete<X>> : Base {
        default_delete<X> del;
        Helper(const default_delete<X>& del) : del(del) {}
        virtual void operator()(X* ptr) override { del(ptr); }
    };
    X* ptr_;
    Base* ptr_del_;
public:
    shared_ptr<default_delete<X>>(X* ptr, const default_delete<X>& del = default_delete<X>())
        : ptr_(ptr), ptr_del_(new Helper<default_delete<X>>(del)) {}
    ~shared_ptr() { if (use_count() == 1) {
        ptr_del_>operator()(ptr_);
        delete ptr_del_;
    } }
}
```

37

### Custom deleters — shared pointers (cont.)

- In case of additional instantiation with the same value type X, but a different custom deleter, the resulting instance will look like:

```
shared_ptr<X> ptr = new X; shared_ptr<X> ptr2 { new X, custom_delete<X>() };

struct default_delete<X> { void operator()(X* ptr) { delete ptr; } };
struct custom_delete<X> { void operator()(X* ptr) { ... } };
class shared_ptr<X> {
    struct Base {
        virtual ~Base() = default;
        virtual void operator()(X* ptr) = 0;
    };
    struct Helper<default_delete<X>> : Base {
        default_delete<X> del;
        Helper(const default_delete<X>& del) : del(del) {}
        virtual void operator()(X* ptr) override { del(ptr); }
    };
    struct Helper<custom_delete<X>> : Base {
        custom_delete<X> del;
        Helper(const custom_delete<X>& del) : del(del) {}
        virtual void operator()(X* ptr) override { del(ptr); }
    };
    X* ptr_;
    Base* ptr_del_;
public:
    shared_ptr<default_delete<X>>(X* ptr, const default_delete<X>& del = default_delete<X>())
        : ptr_(ptr), ptr_del_(new Helper<default_delete<X>>(del)) {}
    shared_ptr<custom_delete<X>>(X* ptr, const custom_delete<X>& del = custom_delete<X>())
        : ptr_(ptr), ptr_del_(new Helper<custom_delete<X>>(del)) {}
    ~shared_ptr() { if (use_count() == 1) { ptr_del_>operator()(ptr_); delete ptr_del_; } }
}
```

38

### Deleters — unique vs shared ptrs.

- Observation:**
  - Type of all shared\_ptr objects with the same value type but different deleter types is **the same**.
    - It is the same class = instance of shared\_ptr class template.
  - Different type of deleter just causes different instance of shared pointer constructor to be generated within this class (and, possibly, some implementation inner helper class).
  - On the contrary, for each type of deleter unique\_ptr class template is **instantiated into different class**.
    - Types of unique pointers with different deleters are different and unrelated classes.
- In shared pointers, type of deleter is so-called "type-erased".
  - The technique for hiding its type is called "type erasure".
- Type-erased deleter is more "powerful" but has runtime overhead.
  - ⇒ Reason why unique pointers do not employ this technique.
    - Recall — unique pointer = lightweight RAII pointer wrapper with minimal overhead.

39