

# Effective C++ Programming

NIE-EPC (v. 2021):

TYPE ERASURE, `STD::FUNCTION`, `ANY`, `VARIANT`,  
`CRTP`, `SWAP`, `ADL`,...

© 2021 DANIEL LANGR, ČVUT (CTU) FIT

# Type erasure

- *Recall* — **type erasure** allowed a shared pointer to **own/manage and use a deleter of any type**.
- This technique **may be generalized**:

```
class Any {
    struct Base {
        virtual ~Base() = default;
    };

    template <typename T> // T - type of type-erased object owned/managed by X
    struct Helper : Base {
        T t_;
        Helper(const T& t) : t_(t) { }
    };

    Base* ptr_;

public:
    template <typename T>
    Any(const T& t) : ptr_( new Helper<T>(t) ) { }

    ~Any() { delete ptr_; }
};
```

- Or, better with **RAII**:

```
...
std::unique_ptr<Base> ptr_;
...
~Any() { delete ptr_; }
```

# Type erasure (cont.)

```
class Any {  
    struct Base {  
        virtual ~Base() = default;  
    };  
    template <typename T>  
    struct Helper : Base {  
        T t_;  
        Helper(const T& t) : t_(t) { }  
    };  
    std::unique_ptr<Base> ptr_;  
public:  
    template <typename T>  
    Any(const T& t) : ptr_( new Helper<T>(t) ) { }  
};
```

- Now, we can “store” an object of any type into an Any owner:

```
Any a1( 1 );  
Any a2( true );  
Any a3( std::string("some string") );
```

- *Note* — thanks to type erasure, type of a1, a2, and a3 is the same.
- *Question* — what can be done with such type-erased objects?
  - With the above-provided definition, nothing at all.

# Type erasure (cont.)

- *Question* — **what can be done** with such type-erased objects?
- **Two different cases:**
  - 1) These objects **have some common interface** — they provide some common operation.
  - 2) They **do not have anything in common**.
- *Ad 1)* This was the case of **shared pointer deleters**:
  - All of them are supposed to **provide a function call operator with a value type-pointer parameter**.
  - This operator is then **propagated out of the internal type-erasure classes as a virtual function**.
- *Ad 2)* As with Any class — will be discussed later.

```
template <typename T> class shared_ptr {
    struct Base {
        virtual ~Base() = default;
        virtual void operator()(T* ptr) = 0;
    };
    template <typename D>
    struct Helper : Base {
        D del_;
        Helper(const D& del) : del_(del) {}
        virtual void operator()(T* ptr) override
        { del_(ptr); }
    };
    T* ptr_; Base* ptr_del_;
public:
    template <typename D = std::default_delete<T>>
    shared_ptr(T* ptr, const D& del = D{})
        : ptr_(ptr), ptr_del_(new Helper<D>{del})
        { }
    ~shared_ptr() { if (use_count() == 1) {
        ptr_del_->operator()(ptr_);
        delete ptr_;
    } }
```

# Type erasure (cont.)

- Generally, if type-erased objects **have some common interface**, it may be “propagated” **even out of the owner class**.
  - That is, it may be **made a part of its interface**.
- *Example* — **task owner class** which type-erases **entities callable without arguments**.
- *Motivation* — **thread pool**:
  - Fixed number of running **threads**.
  - **Queue of tasks** (with a thread-safe access, such as protected by a mutex).
  - Once a **task is enqueued** (from outside) and some thread is available (idle), this thread **dequeues the task and starts its execution**.
  - **Task** = any **callable entity** that can be **called without arguments**.
- *Problem* — how to add **callable entities of different types** into the **same queue**?



# Type erasure (cont.)

- Initial attempt:

```
class ThreadPool {  
    std::vector<std::thread> threads_; // running threads  
    std::mutex m_; // protection of q_ shared by threads  
    std::queue< ??? > q_; // queue of tasks  
    ...  
};
```

- Problem* — how to define a queue member variable such that we could enqueue different callable entities into it?
- Example:*

```
void task1() { std::cout << "task1" << std::endl; } // (1) ordinary function  
struct Task2 { void operator()() { std::cout << "task2" << std::endl; } };  
int main() {  
    ThreadPool tp;  
    Task2 task2; // (2) function object (functor)  
    auto task3 = []{ std::cout << "task3" << std::endl; }; // (3) lambda function  
    // How to store task1, task2, and task3 into tp.q_?
```

- We have 3 entities callable without arguments.
- Each one is of a different type.
- How to put them into the queue `q_` of the thread pool `tp`?

# Type erasure (cont.)

- *Analysis:*
  - Each callable entity task1, task2, and task3 has a different type.
  - All elements of standard library containers (such as `std::vector` or `std::list`) and container adaptors (such as `std::queue` or `std::stack`) must have the same (value) type.
  - $\Rightarrow$  We cannot store tasks into the thread pool queue directly.
- Possible solution — *type erasure*:
  - Let us start with Any class:

```
class ThreadPool {  
    std::vector<std::thread> threads_; // running threads  
    std::mutex m_; // protection of q_ shared by threads  
    std::queue< Any > q_; // queue of tasks
```

- Now, we can store all the tasks into the queue.
- *Problem* — we need to eventually dequeue tasks and execute them.
  - $\Rightarrow$  We need to invoke their function-call operator.
  - Such a functionality is not provided by Any class.

# Type erasure (cont.)

- *Alternative:*
  - We define tasks as callable entities without arguments.
  - $\Rightarrow$  All of them **share this common interface feature**.
  - $\Rightarrow$  We can **make this operation available** from the type-erasure owner class:

```
class Task {
    struct Base {
        virtual void operator>() = 0;
        virtual ~Base() = default;
    };
    template <typename T>
    struct Helper : Base {
        T t_;
        Helper(const T& t) : t_(t) { }
        virtual void operator>() override { t_(); }
    };
    std::unique_ptr<Base> ptr_;
public:
    template <typename T>
    Task(const T& t) : ptr_( new Helper<T>(t) ) { }
    void operator>() { ptr_->operator>(); }
};
```



# Type erasure (cont.)

- Now, we can put any callable entity (without arguments) into Task owner and invoke its function call operator:

```
void task1() { std::cout << "task1" << std::endl; }           // (1) ordinary function
struct Task2 { void operator()() { std::cout << "task2" << std::endl; } };
int main() {
    Task2 task2;                                               // (2) function object (functor)
    auto task3 = []{ std::cout << "task3" << std::endl; };    // (3) lambda function

    std::queue<Task> q;
    q.emplace(&task1);    // function is not an object => need to store its pointer instead
    q.emplace(task2);
    q.emplace(task3);

    while (!q.empty()) {
        Task t = std::move(q.front());    // remove task...
        q.pop();                          // ...from the queue...
        t();                              // ...and execute it
    }
}
```

- Live demo: <https://godbolt.org/z/W3zPT5x7j>.
- The **queue** in **ThreadPool** can be now defined and used the **same way**.
- Note* — if task returns something, the **return value is discarded**.

# std::function

- C++ standard library provides a **generic type-erased callable entity owner** — class template **std::function**.
- Our **Task**:
  - Task can own and call an **entity callable without arguments**  $\Rightarrow$  basically callable entities **without parameters**.
  - If the type-erased entity returns, the **return value is discarded** by Task.
- **std::function**:
  - An instance of **std::function** template can **own and call callable entities with particular types or parameters and type of return value**.
  - The types of parameters and type of return value are **determined by std::function template argument**.
- *Examples*:
  - **std::function<void()>** — basically, equivalent of our Task  $\Rightarrow$  stored callable entities have **no parameters** and **do not return values**.
  - **std::function<bool(int,int)>** — callable entities have **two parameters of type int** and **return bool**.

# std::packaged\_task

- `std::packaged_task` — basically `std::function` with some additional functionality.
  - Namely, packaged task allows to obtain a relevant *future* object (of corresponding `std::future` type).
- This future object allows to:
  - **Wait** for the task to be completed (its execution).
  - Obtain its **return value**.
  - **Catch exception** thrown within task execution.
- Moreover, future object may be **accessed by other threads** than the thread executing the task.
- $\Rightarrow$  Packaged tasks are **supposed to be run asynchronously**.
- *Note* — suitable use case = *thread pools*.

# std::packaged\_task (cont.)

- Simple thread pool (*incomplete*) with packaged tasks:

```
class ThreadPool {
    using Task = std::packaged_task<void()>;
    std::mutex m_;
    std::queue<Task> q_;
    std::condition_variable cv_;
    std::vector<std::thread> threads_;
    void worker() {
        while (true) {
            Task task;
            {
                std::unique_lock<std::mutex> lock(m_);
                cv_.wait(lock, [this]{ return !this->q_.empty(); });
                task = std::move(q_.front());
                q_.pop();
            }
            task(); // task executed here
        }
    }
public:
    ThreadPool(int num_threads = std::thread::hardware_concurrency()) {
        for (int i = 0; i < num_threads; i++)
            threads_.emplace_back( [this]{ this->worker(); } );
    }
    std::future<void> enqueue(Task task) {
        std::future<void> future;
        {
            std::lock_guard<std::mutex> lock(m_);
            future = task.get_future();
            q_.emplace(std::move(task));
        }
        cv_.notify_one(); // wake up some thread
        return future;
    }
};
```

# Type erasure — std::any

- *Recall* — **Any** class can store/own type-erased object of **any type**.
- Any type  $\Rightarrow$  there is **no common interface**.
- However, we can **provide access to the stored type-erased object**.
- *First option* — **type-less pointer**:
  - In Base and Any classes, there is **no knowledge of T**.
  - $\Rightarrow$  Only possibility is to return a **pointer to the stored object of type void\***.

```
class Any {
    struct Base {
        virtual ~Base() = default;
        virtual void* get() = 0;
    };
    template <typename T>
    struct Helper : Base {
        T t_;
        Helper(const T& t) : t_(t) { }
        virtual void* get() override { return &t_; }
    };
    std::unique_ptr<Base> ptr_;
public:
    template <typename T>
    Any(const T& t) : ptr_( new Helper<T>(t) ) { }
    void* get() { return ptr_->get(); }
};
```

# Type erasure — std::any (cont.)

- *Consequence* — to use the object, returned pointer **needs to be casted into the pointer-to-object-type**.
- $\Rightarrow$  This type **must be provided from outside!**
  - There is no other way; Any class **cannot “remember” a type**.

```
std::vector<Any> v;  
v.emplace_back(1); // type-erased int  
v.emplace_back(std::string("some string")); // type-erased std::string  
  
int* ptr = static_cast<int*>( v[0].get() ); // pointer to the stored int  
std::string& ref = *static_cast<std::string*>( v[1].get() ); // reference to the stored...  
// ...std::string
```

- *Better option* — wrapping “ugly” casting with a **custom function**:

```
class Any {  
    ...  
public:  
    ...  
    template <typename T> friend T Any_cast(const Any& a); // friend to have access to ptr_  
};  
  
template <typename T> T Any_cast(const Any& a) // free (non-member) function template  
{ return *static_cast< std::remove_reference_t<T>* >( a.ptr_->get() ); }
```

- *Note* — Any\_cast can either return a **reference to the stored object**, or **its copy**.



# Type erasure — `std::any` (cont.)

- Such a **generic type-erased object owner** is provided by the C++ standard library as **`std::any` class**.
  - It is available **since C++17**.

```
std::vector<std::any> v;  
v.emplace_back(1);           // type-erased int  
v.emplace_back(std::string("some string")); // type-erased std::string  
  
int      val = std::any_cast< int      >( v[0] ); // copy of the type-erased object  
std::string& ref = std::any_cast< std::string& >( v[1] ); // reference to the type-erased object
```

- *Notes:*
  - In contrast to *cast operators* (`static_cast`, `reinterpret_cast`, ...), **`std::any_cast` is not an operator**.
  - Instead, it is an “ordinary” **function template**.
- *Use cases:*
  - `std::any` is suitable for **very specific use cases only**.
  - *Relevant discussion:* <https://devblogs.microsoft.com/cppblog/stdany-how-when-and-why/>.

# std::variant

- `std::any` — a single-object owner.
  - The owned object is stored in dynamically allocated storage.
  - Its access involves virtual-function dispatch.
  - $\Rightarrow$  Relatively large runtime and memory overhead.
  - Can store an object of any type, but one needs to “recall” this type during object access.
- Another C++ standard library option — `std::variant`:
  - A single-object owner as well.
  - Is a *variadic template* — template arguments define a list of types.
  - Only objects of these types may be stored in variant.

```
std::any a1 = 1;      // int stored
std::any a2 = true;   // bool stored
std::any a3 = 1.0;    // double stored
a1 = 1.0;             // double stored
```

```
std::variant<int,bool> v1 = 1;      // int stored
std::variant<int,bool> v2 = true;   // bool stored
std::variant<int,bool> v3 = 1.0;    // ERROR
v1 = 1.0;                          // ERROR
```

- Note — all *any* objects have the same type, *variants* with different type list have different types.

# std::variant (cont.)

- List of types **defines their order**.
- $\Rightarrow$  Each type has some **index** in this order.
- **Index of the type of actually stored object** is returned by the **index member function**.

```
std::variant<int, bool, double> v = 1.0; // int has index 0, bool 1, double 2
std::cout << v.index();                 // double is stored => prints out "2"
```

- Stored object is **accessed** with **std::get** free function template.
  - Template argument — **index of the stored object type**, or the **type itself**.
  - **Returns a reference** to the owned object.

```
std::variant<int, bool, double> v = 1.0;
std::cout << std::get< 2 >(); // prints out "1"
std::cout << std::get< double >(); // effectively the same
```

- Note — C++ is a **statically typed language**.
  - $\Rightarrow$  All **types must be resolved at runtime**.
  - $\Rightarrow$  This holds also for the **return type of std::get**.
  - $\Rightarrow$  **index function call cannot be used** as a template argument of **std::get**.

# std::variant (cont.)

- Owned object **needs to be stored in the included storage**.
  - ⇒ Variant is explicitly **disallowed to dynamically allocate memory** [\[link\]](#).

```
std::variant<int, bool> v = 1;  
std::cout << (uintptr_t)&std::get<0>(v) - (uintptr_t)&v << std::endl; // print "0" (GCC)
```

- Implementation — **included buffer aligned and sized suitably** according the type list.
  - Buffer alignment* = **maximum of alignment requirements for of types**.
  - Buffer size* = **maximum of size requirements of types**.
  - Note* — it may be defined with **std::aligned\_union** library helper type.
- ⇒ **No runtime/memory overhead**.
- ⇒ **std::variant** is **much preferred** than **std::any** once there is a fixed list of type alternatives.
- Benchmark* — std::variant vs std::any:
  - insertion** (object storage) — **6.6× faster** with std::variant [\[link\]](#);
  - stored object **access** — **6.1× faster** with std::variant [\[link\]](#).

# std::variant (cont.)

- `std::variant` vs `std::tuple`:
  - `std::tuple` — multiple object owner  $\Rightarrow$  owns/stores a single object for each type from the type list (template arguments).
  - `std::variant` — single object owner  $\Rightarrow$  owns/stores at most one object at a given moment (of a type from the type list).

```
std::variant<int, bool, double> v = 1;           // owns int
std::cout << &std::get<int>(v) << std::endl;     // printed "0x7fffa6a7b030" in experiment
v = true;                                       // owns bool
std::cout << &std::get<bool>(v) << std::endl;    // printed "0x7fffa6a7b030"
v = 1.0;                                       // owns double
std::cout << &std::get<double>(v) << std::endl; // printed "0x7fffa6a7b030"
```

- `std::variant` vs `union`:
  - `std::variant` — basically a type-safe union.
  - `union` is low-level “C” type, which is very unsafe and problematic with classes.
  - *Price* — `std::variant` needs to remember the actual “alternative” (index).

```
union { int i; bool b; double d; } u;
std::cout << sizeof(u);           // printed "8" in experiment

std::variant<int, bool, double> v;
std::cout << sizeof(v);           // printed "16" in experiment
```

# C RTP

- *Example* — vector 2D graphic editor  $\Rightarrow$  **hierarchy of polymorphic classes** that represent various **graphic objects**:

```
class Object2D { // abstract base class
public:
    virtual ~Object2D() = default;
    virtual void draw() const;
    ...
};

class Line : public Object2D { ... };
class Square : public Object2D { ... };
class Circle : public Object2D { ... };
```

- *Scene* is represented as a **collection of graphic objects**:

```
std::vector< std::unique_ptr<Object2D> > scene; // all objects in a scene
```

- *Exemplary functionality* — loading scene from a file:

```
void load_scene(const std::string& filename) {
    std::ifstream f(filename); std::string line;
    while (std::getline(f, line)) {
        if (s == "line") scene.emplace_back( new Line{} );
        else if (s == "circle") scene.emplace_back( new Circle{} );
        else if (s == "square") scene.emplace_back( new Square{} );
        else throw std::invalid_argument("invalid object name");
    }
    for (const auto & obj : scene) obj->draw(); // draw loaded scene
}
```



# CRTP (cont.)

- Another functionality — **selection of objects**:

```
std::vector< Object2D* > selected; // set of selected objects
void select(); // function that put pointers to selected objects into selected vector
```

```
// resolution of user actions:
if (ACTION == SELECT) select();
...
```

- *Note* — selected is a vector of “normal” (raw) pointers.
  - Vectors **cannot hold references**.
  - *Smart pointers do not make sense here*  $\leftarrow$  vector select **does not own selected objects**; only refers to them.
- Yet another functionality — **duplication of selected objects**.

```
// resolution of user actions:
if (ACTION == SELECT) select();
else if (ACTION == DUPLICATE) duplicate();
...
```

```
void duplicate() {
    for (auto p_obj : selected) scene.emplace_back( ??? );
}
```

- *Problem* — **how to create a copy of a pointed-to object?**

# CRTP (cont.)

```
void duplicate() {  
    for (auto p_obj : selected) scene.emplace_back( ??? );  
}
```

- *Analysis:*
  - `p_obj` points to an **object of derived class type** (Line, Square, or Circle).
  - Type of `p_obj` is a **pointer-to-base** (`Object2D*`).
  - $\Rightarrow$  We need to **invoke a copy constructor of a derived class through a pointer to a base class**. *How?*
- Generally, this mechanism is provided by **virtual functions**.
- $\Rightarrow$  We would need a “**virtual copy constructor**”.
- However, C++ **does not support virtual constructors**.
- *Alternative* — **delegation of a copying semantics** into a **separate virtual function**.
  - This function is **typically called** `clone`.
  - The whole solution is sometimes called “**clone pattern/idiom**”.

# CRTP (cont.)

- *Solution* — polymorphic copying/cloning:

```
class Object2D {  
    ...  
    virtual Object2D* clone() const = 0;  
};  
class Line : public Object2D {  
    ...  
    virtual Object2D* clone() const { return new Line(*this); } // copy constructor of...  
}; // ...Line called here  
class Square : public Object2D {  
    ...  
    virtual Object2D* clone() const { return new Square(*this); } // copy constructor of...  
}; // ...Square called here  
class Circle : public Object2D {  
    ...  
    virtual Object2D* clone() const { return new Circle(*this); } // copy constructor of...  
}; // Circle called here
```

- *Application* to scene duplication:

```
void duplicate() {  
    for (auto p_obj : selected) scene.emplace_back( p_obj->clone() );  
}
```

- *Drawback* — clone function needs to be repeatedly / “redundantly” implemented throughout the whole class hierarchy.

# C RTP (cont.)

- *Alternative option* — couldn't we “*inject*” `clone` function into **derived classes** from some base class?
- Obviously, it is **not possible from `Object2D` base class itself**.
  - There is **no knowledge of the cloned type** there.
- The base class where `clone` is defined **must know about the cloned object type**.
- *Solution:*
  - Making **another base class** that is a **template...**
  - ...and “**remembering**” the cloned type in its template parameter.
- *Note:*
  - We **cannot make `Object2D` a template**.
  - *Polymorphism* requires a **single unique base class**.
  - Templated `Object2D` would create a **separate base class for each template argument**.
- *Solution (cont.)* — **templated “middle” class** with `clone` that:
  - inherits from `Object2D` ( $\Rightarrow$  **polymorphism with single base**),
  - serves as a base for graphic object classes ( $\Rightarrow$  **injects `clone` into them**).

# CRTP (cont.)

- *Partial solution:*

```
class Object2D {                                     // single (original) base => polymorphism
    ...
    virtual Object2D* clone() const = 0;
};

template <typename T>                                // template argument = cloned type
class ClonableObject2D : public Object2D {           // middle class => defines clone
    ...
    virtual Object2D* clone() const { ??? }
};
```

- Derived 2d graphic object classes:
  - derive from `ClonableObject2d` class template,
  - where `cloned type` = template argument is `derived class itself`.
- For example, `Line` derives from `ClonableObject2D` and the `cloned type` needs to be `Line`  $\Rightarrow$  template argument needs to be `Line`.

```
class Line : public ClonableObject2d<Line> {
    ...
};
```

- This technique — inheritance from derived template where `base template argument is derived class itself` — is called “*curiously recurring template pattern*” (CRTP).

# CRTP (cont.)

- *Last step* — definition of `clone`:

```
template <typename T> class ClonableObject2D : public Object2D {  
    ...  
    virtual Object2D* clone() const { return new T( ??? ); } // invoke copy constructor of T  
};
```

- `clone` needs to create new object with **copy constructor of T**.
- The **source object** to be copied is **pointed to by this pointer**.
- *Problem* — the **type of this** in `ClonableObject2D<T>::clone()` call is `ClonableObject2D<T>*`.
- But, we know that it **actually points to an object of type T**, which **should be copied**.
- *Example:*

```
Object2D* ptr = new Line;  
ptr->clone(); // inside, type of 'this' is ClonableObject2D<Line>*, but it points to Line
```

- *Consequence* — this needs to be **casted into T\***:

```
virtual Object2D* clone() const {  
    return new T( *static_cast<const T*>(this) ); // or, 'static_cast<const T*>(*this)'  
}
```



# CRTP (cont.)

- *Summary:*

```
class Object2D {                                     // single (original) base => polymorphism
    ...
    virtual Object2D* clone() const = 0;
};

template <typename T>                                // template argument = cloned type
class ClonableObject2D : public Object2D {           // middle class => defines clone
    ...
    virtual Object2D* clone() const {
        return new T( *static_cast<const T*>(this) );
    }
};

// derive 2D graphic object classes:
class Line : public ClonableObject2d<Line> { ... };
class Square : public ClonableObject2d<Square> { ... };
class Circle : public ClonableObject2d<Circle> { ... };
```

- *Application* to scene duplication is the same:

```
void duplicate() {
    for (auto p_obj : selected) scene.emplace_back( p_obj->clone() );
}
```

- *Advantage* — clone is defined only once.
  - $\Rightarrow$  Less code redundancy  $\Rightarrow$  better code maintenance.
- *Note* — generally, CRTP has broader use than just the clone idiom.

# CRTP (cont.)

- *Instantiation:*

```
class Object2D {
    ...
    virtual Object2D* clone() const = 0;
};

class ClonableObject2D<Line> : public Object2D {
    ...
    virtual Object2D* clone() const { return new Line( *static_cast<const Line*>(this) ); }
};

class ClonableObject2D<Square> : public Object2D {
    ...
    virtual Object2D* clone() const { return new Square( *static_cast<const Square*>(this) ); }
};

class ClonableObject2D<Circle> : public Object2D {
    ...
    virtual Object2D* clone() const { return new Circle( *static_cast<const Circle*>(this) ); }
};

class Line : public ClonableObject2D<Line> { ... };
class Square : public ClonableObject2D<Square> { ... };
class Circle : public ClonableObject2D<Circle> { ... };
```

- *Result* — one *base class*, three *middle classes*, three *derived classes*.
- *Exemplary resolution* — if *ptr* points to *Line*, *ptr->clone()* actually calls *ClonableObject2D<Line>::clone()*, which copies *Line* by its copy constructor.

# Function swap

- *Special member functions* related to object content:
  - *default constructor* (sets object to “empty” state),
  - *copy constructor* (copies content into initialized object),
  - *move constructor* (moves content into initialized object),
  - *copy assignment operator* (copies content into existing object),
  - *move assignment operator* (moves content into existing object),
  - *destructor* (destroys content).
- The member functions are *special* in that:
  - they are *automatically generated* under some circumstances,
  - they may be *called without explicit function call syntax*.
- *Example:*

```
struct X {  
    std::string s;  
}; // all 6 special member functions are...  
    // ...automatically generated for TwoStrings class  
  
void f(X);  
  
int main() {  
    X x; // default constructor called  
    f(x); // copy constructor called  
} // destructor called
```

# Function swap (cont.)

- There is one additional “*special*” member function — `swap` for swapping content of two existing objects.
- It is **not truly special**  $\Rightarrow$  **not** automatically generated, **not** called without explicit function call syntax.
- However, it is **widely used inside existing** (not only C++ standard) **library algorithms** (sorting, for instance).
- Suppose such an algorithm **needs to swap content of two objects of type T**.
- Then, there are **two options**:
  - 1) either there **is a custom swap free function** specifically **designed for arguments of type T** (typically defined in the **same namespace** as T),
  - 2) or, there **is no such a custom function**.
- The algorithms then **typically works as follows**:
  - **Ad 1)** They **use that custom function** for required content swapping.
  - **Ad 2)** They **use `std::swap`** instead.

# std::swap

- `std::swap` is a **library function** that is provided in **two forms**:
  - First, there is a **primary template** that works generally for **objects of any type**.
  - Then, there are **"specializations"** for **various library types** [\[link\]](#).

```
struct X { std::string s; };  
  
int main() {  
    int i = 1, j = 2;  
    std::swap(i, j);           // uses generic primary template  
    std::cout << i << j;      // prints out "21" => content-numbers were swapped  
  
    std::string s1("world"), s2("hello ");  
    std::swap(s1, s2);         // uses special overload of std::swap for std::string  
    std::cout << s1 << s2;     // prints out "hello world" => content-strings were swapped  
  
    X x1("world"), x2("hello ");  
    std::swap(x1, x2);         // uses generic primary template  
    std::cout << x1.s << x2.s; // prints out "hello world" => content-strings were swapped  
}
```

- **Questions:**
  - 1) Why are these **"specializations"** defined?
  - 2) Isn't the **primary template** enough?
  - 3) Should we define **swap function** for **custom types**?

# std::swap (cont.)

- How is **generic primary template** of `std::swap` **implemented**?
  - It is **not specified by the C++ standard**, but implementations of C++ standard libraries **typically define primary template *naturally* as follows**:

```
template <typename T>           // primary std::swap template
void swap(T& a, T& b) {
    T temp { std::move(a) }; // moves/copies content from first parameter into a temporary
    a = std::move(b);        // moves/copies content from second parameter into the first
    b = std::move(temp);     // moves/copies content from the temporary into the second param.
}                             // => this effectively swaps content of both parameters
```

- $\Rightarrow$  For **class types**, primary template of `std::swap` is implemented **in terms of its *special member functions***:

```
template <typename T>           // if T is a class type:
void swap(T& a, T& b) {
    T temp { std::move(a) }; // move/copy constructor called (move preferred)
    a = std::move(b);        // move/copy assignment operator called (ditto)
    b = std::move(temp);     // move/copy assignment operator called (ditto)
}                             // destructor called (for temp)
```

- $\Rightarrow$  **Move-content operations** are (naturally) **preferred if they exist**; **otherwise, copy-content operations** are used instead.
- **Note** — for **trivially-copyable types**, there is **no moving of content**.
  - $\Rightarrow$  All operations involve **copying of binary representations**.



# Custom swap vs std::swap (cont.)

- Is it worth defining **special swap overload** for custom types?
  - Obviously, it is **worth for classes without move semantics** and “**expensive**” **copy semantics** (such as *legacy/pre-C++11* classes).
- *Benchmark* — class with a **std::string** subobject.
- *Two versions* — **first** with move semantics:

```
struct X_move {  
    std::string s;  
}; // X_move HAS automatically generated move constructor and move assignment operator
```

- **Second**, with copy but **without move semantics**:

```
struct X_copy { // for instance, pre-C++11 legacy class  
    std::string s;  
    X_copy(const char* s) : s(s) { }  
    X_copy(const X_copy& other) : s(other.s) { } // suppresses auto-generation...  
                                                // ...of move-content operations  
}; // X_copy HAS NOT automatically generated move constructor and move assignment operator
```

- **Comparison of swapping** of two objects of these types with `std::swap` primary template:
  - Swapping of content was **4.3× faster** for `X_move` [\[link\]](#).

# Custom swap vs std::swap (cont.)

- Let us add a **custom swap function** for a move-less version:

```
void swap(X_copy& a, X_copy& b)
{
    std::swap(a.s, b.s); // swapping content = swapping content of subobject s
}
```

- Results** — swapping of content was **12× faster** with custom swap than with `std::swap` [\[link\]](#).
  - ⇒ It is usually **crucial to provide swap function** for **custom classes** if these **do not provide move-content operations** and are **expensive to copy** (typically “non-movable” classes that **own some resources**).
- What about classes that **provide move-content operations**?

```
void swap(X_move& a, X_move& b)
{
    std::swap(a.s, b.s);
}
```

- Results** — swapping of content was **2.6× faster** with custom swap than with `std::swap` [\[link\]](#).

# Custom swap vs std::swap (cont.)

- What is the **difference**:
  - In **all examined cases**, swapping of content of X... class object = **swapping of content of its subobject** (namely, its **s member variable**).
  - In case of **X\_move** and **custom swap**, this **swapping of s** is accomplished with "**specialization**" of **std::swap** for **std::string** arguments:

```
void swap(X_move& a, X_move& b) {  
    std::swap(a.s, b.s); // specialization of std::swap for std::string called here  
}
```

- In case of **std::swap**, **primary template of std::swap** is used and instantiated as follows:

```
void swap<X_move>(X_move& a, X_move& b) {  
    X_move temp { std::move(a) }; // move constructor of X_move called  
    a = std::move(b);             // move assignment operator of X_move called  
    b = std::move(temp);          // move assignment operator of X_move called  
}                                 // destructor of X_move called
```

- Auto-generated special member function of X\_move is **effectively equivalent** with the **same special member function of std::string**.
- $\Rightarrow$  **Swapping of s** is accomplished with the following calls of **std::string special member functions**: 1× move constructor, 2× move assignment operator, 1× destructor.

# Custom swap vs std::swap (cont.)

- *Conclusion:*
  - With a custom swap function, content of s member variable is swapped with the specialization of std::swap for std::string arguments.
  - With std::swap function, content of s member variable is swapped with:
    - single call of std::string move constructor,
    - double call of std::string move assignment operator,
    - and single call of std::string destructor.
- In both cases, the final outcome (observable behavior) of swapping is the very same!
- However, within our experiment and used C++ implementation, the compiler was not able to optimize both cases into the equally efficient machine code.
- *Why not?*

# Custom swap vs std::swap (cont.)

- *Exemplary problem:*
  - After each *move-content* operation, s member variable of source object needs to be put into *moved-from* = “empty string” state, which requires some work.
  - However, within std::swap instance, this is unnecessary since in all 3 cases, the s member variable is then either *assigned new content* or *destroyed*:

```
void swap<X_move>(X_move& a, X_move& b) {  
    X_move temp { std::move(a) }; // a.s put into empty string state...  
    a = std::move(b);             // ...while it is then assigned new content
```

```
    a = std::move(b);             // b.s put into empty string state...  
    b = std::move(temp); // ...while it is then assigned new content
```

```
    b = std::move(temp); // temp.s put into empty string state...  
}                        // ...while it is then destroyed
```

- Compiler optimizers are simply *not perfect/ideal* from our point of view.
- Here, a compiler *may not recognize* that “resetting” strings is unnecessary, and *generate machine code* for it.
- *Note* — due to SSO, *move-content* operations for std::string are *non-trivial but relatively complex*.



# Custom swap vs std::swap (cont.)

- On the contrary, “specialization” of `std::swap` for `std::string` does not require any resetting to empty state.
- If strings are *long*, *swapping of content* effectively just *turns into swapping content of few member variables* (pointer + integers).
- $\Rightarrow$  “Specialization” can be implemented in such a way that it *does not contain any unnecessary operations*.
- With *generic* `std::swap` based on *move-content* operations, *additional operations are involved* and it may be *hard or impossible* for a compiler to *optimize them away* and generate *equally efficient machine code*.
- *Experiment* — swapping content of two `std::string` objects with primary template and specialization of `std::swap`: [\[link\]](#).
- *Consequence* — generally, even for *classes that provide move-content operations*, it may be *worth providing custom swap function*.



# Using swap function

- How to use **swap function**, for instance, in a quicksort algorithm?

```
template <typename T> long partition(T* a, long lo, long hi) {
    long i = lo;
    for (long j = lo; j < hi; j++)
        if (a[j] < a[hi]) ??? ; // need to swap a[i++] and a[j], HOW?
    ??? ; // need to swap a[i] and a[hi], HOW?
    return i;
}

template <typename T> void qs(T* a, long lo, long hi) {
    if (lo < hi) {
        long p = partition(a, lo, hi);
        qs(a, lo, p - 1);
        qs(a, p + 1, hi);
    }
}
```

- Within partitioning, we need to **swap sorted elements**; how?
- We want:
  - 1) to use **custom swap** function **if it exists**,
  - 2) to use **std::swap** otherwise.
- If we called
  - **swap**, then **compilation error** would be generated in the **second case**.
  - **std::swap**, then **custom swap** would not be used in the **first case**.

# Using swap function

- Solution is simple:

```
template <typename T> long partition(T* a, long lo, long hi) {  
    using std::swap;  
    long i = lo;  
    for (long j = lo; j < hi; j++)  
        if (a[j] < a[hi]) swap(a[i++], a[j]);  
    swap(a[i], a[hi]);  
    return i;  
}
```

- Directive `using` “injects” the name `swap` from `std` namespace into *global namespace*.
- Consequently, `swap` call will:
  - use custom `swap` (its “specialization”) if it exists (better overloading candidate),
  - fall back to primary `std::swap` template otherwise.
- Apparent problem:
  - What if custom `swap` is in some custom (for instance, library) *namespace*?
  - Will this solution work?

# Argument dependent lookup

- *Example:*

```
namespace our_library
{
    struct X {
        std::string s;
        bool operator<(const X& rhs) const { return s < rhs.s; } // Lexicographic comparison
    };
    void swap(X& a, X& b) { std::swap(a.s, b.s); }
}
```

```
int main() {
    our_library::X a[2] = { "B", "A" };
    qs(a, 0, 1);
}
```

- In partition, there **is no "knowledge"** of our\_library namespace:

```
template <typename T> long partition(T* a, long lo, long hi) {
    using std::swap;
    using our_library::swap; // we don't know about our_library namespace here

    long i = lo;
    for (long j = lo; j < hi; j++)
        if (a[j] < a[hi]) swap(a[i++], a[j]);
    swap(a[i], a[hi]);
    return i;
}
```

# Argument dependent lookup (cont.)

- Which overload of `swap` will be called when `T` is `our_library::X`?

```
long partition<our_library::X>(our_library::X* a, long lo, long hi) {  
    using std::swap;  
  
    long i = lo;  
    for (long j = lo; j < hi; j++)  
        if (a[j] < a[hi]) swap(a[i++], a[j]);  
    swap(a[i], a[hi]);  
    return i;  
}
```

- Good news — `custom our_library::swap` will be used in this case.
- Reason = “**argument dependent lookup**” (ADL):
  - If types of arguments are from some namespace, candidates for overloading are looked-up in this namespace as well.
- Without ADL, it would be almost impossible to write generic code.
  - Its developer can't know about namespaces from which involved types of objects will come from.