

Effective C++ Programming

NIE-EPC (v. 2021):
TYPE ERASURE, STD::FUNCTION, ANY, VARIANT,
CRTP, SWAP, ADL,...
© 2021 DANIEL LANGR, ČVUT (CTU) FIT

1

Type erasure

- *Recall* — type erasure allowed a shared pointer to own/manage and use a deleter of any type.
- This technique may be generalized:

```
class Any {
    struct Base {
        virtual ~Base() = default;
    };
    template <typename T> // T - type of type-erased object owned/managed by X
    struct Helper : Base {
        T t_;
        Helper(const T& t) : t_(t) {}
    };
    Base* ptr_;
public:
    template <typename T>
    Any(const T& t) : ptr_( new Helper<T>(t) ) {}
    ~Any() { delete ptr_; }
};
```

- Or, better with RAII:

```
...
std::unique_ptr<Base> ptr_;
...
~Any() { delete ptr_; }
```

2

Type erasure (cont.)

```
class Any {
    struct Base {
        virtual ~Base() = default;
    };
    template <typename T>
    struct Helper : Base {
        T t_;
        Helper(const T& t) : t_(t) {}
    };
    std::unique_ptr<Base> ptr_;
public:
    template <typename T>
    Any(const T& t) : ptr_( new Helper<T>(t) ) {}
};
```

- Now, we can "store" an object of any type into an Any owner:


```
Any a1( 1 );
Any a2( true );
Any a3( std::string("some string") );
```
- *Note* — thanks to type erasure, type of a1, a2, and a3 is the same.
- *Question* — what can be done with such type-erased objects?
 - With the above-provided definition, nothing at all.

3

Type erasure (cont.)

- *Question* — what can be done with such type-erased objects?
- Two different cases:

- 1) These objects have some common interface — they provide some common operation.
- 2) They do not have anything in common.

- *Ad 1)* This was the case of shared pointer deleters:
 - All of them are supposed to provide a function call operator with a value type-pointer parameter.
 - This operator is then propagated out of the internal type-erasure classes as a virtual function.

- *Ad 2)* As with Any class — will be discussed later.

```
template <typename T> class shared_ptr {
    struct Base {
        virtual ~Base() = default;
        virtual void operator()(T* ptr) = 0;
    };
    template <typename D>
    struct Helper : Base {
        D del_;
        Helper(const D& del) : del_(del) {}
        virtual void operator()(T* ptr) override {
            del_(ptr);
        }
    };
    T* ptr_; Base* ptr_del_;
public:
    template <typename D = std::default_delete<T>>
    shared_ptr(T* ptr, const D& del = D()) :
        ptr_(ptr), ptr_del_(new Helper<D>(del)) {}
    ~shared_ptr() { if (use_count() == 1) {
        ptr_del_>operator()(ptr_);
        delete ptr_;
    } }
```

4

Type erasure (cont.)

- Generally, if type-erased objects have some common interface, it may be "propagated" even out of the owner class.
 - That is, it may be made a part of its interface.
- *Example* — task owner class which type-erases entities callable without arguments.
- *Motivation* — thread pool:
 - Fixed number of running threads.
 - Queue of tasks (with a thread-safe access, such as protected by a mutex).
 - Once a task is enqueued (from outside) and some thread is available (idle), this thread dequeues the task and starts its execution.
 - Task = any callable entity that can be called without arguments.
- *Problem* — how to add callable entities of different types into the same queue?

5

Type erasure (cont.)

- Initial attempt:

```
class ThreadPool {
    std::vector<std::thread> threads; // running threads
    std::mutex m_; // protection of q_ shared by threads
    std::queue< ??? > q_; // queue of tasks
    ...
};
```

- *Problem* — how to define a queue member variable such that we could enqueue different callable entities into it?

- *Example:*

```
void task1() { std::cout << "task1" << std::endl; } // (1) ordinary function
struct Task2 { void operator()() { std::cout << "task2" << std::endl; } };
int main() {
    ThreadPool tp;
    Task2 task2;
    auto task3 = []() { std::cout << "task3" << std::endl; }; // (2) function object (functor)
    // How to store task1, task2, and task3 into tp.q_?
    ...
}
```

- We have 3 entities callable without arguments.
- Each one is of a different type.
- How to put them into the queue q_ of the thread pool tp?

6

Type erasure (cont.)

- **Analysis:**
 - Each callable entity task1, task2, and task3 has a different type.
 - All elements of standard library containers (such as `std::vector` or `std::list`) and container adaptors (such as `std::queue` or `std::stack`) must have the same (value) type.
 - \Rightarrow We cannot store tasks into the thread pool queue directly.
- Possible solution — **type erasure**:
 - Let us start with Any class:

```
class ThreadPool {
    std::vector<std::thread> threads_; // running threads
    std::mutex m_; // protection of q_ shared by threads
    std::queue<Any> q_; // queue of tasks
}
```

- Now, we can store all the tasks into the queue.
- **Problem** — we need to eventually dequeue tasks and execute them.
 - \Rightarrow We need to invoke their function-call operator.
 - Such a functionality is not provided by Any class.

7

Type erasure (cont.)

- **Alternative:**
 - We define tasks as callable entities without arguments.
 - \Rightarrow All of them share this common interface feature.
 - \Rightarrow We can make this operation available from the type-erasure owner class:

```
class Task {
    struct Base {
        virtual void operator()() = 0;
        virtual ~Base() = default;
    };
    template<typename T>
    struct Helper : Base {
        T t_;
        Helper(const T& t) : t_(t) {}
        virtual void operator()() override { t_(); }
    };
    std::unique_ptr<Base> ptr_;
public:
    template<typename T>
    Task(const T& t) : ptr_(new Helper<T>(t)) {}
    void operator()() { ptr_>operator()(); }
};
```

8

Type erasure (cont.)

- Now, we can put any callable entity (without arguments) into Task owner and invoke its function call operator:

```
void task1() { std::cout << "task1" << std::endl; } // (1) ordinary function
struct Task2 { void operator()() { std::cout << "task2" << std::endl; } };
int main() {
    Task2 task2; // (2) function object (functor)
    auto task3 = []{ std::cout << "task3" << std::endl; }; // (3) Lambda function
    std::queue<Task> q;
    q.emplace(task1); // function is not an object => need to store its pointer instead
    q.emplace(task2);
    q.emplace(task3);
    while (!q.empty()) {
        Task t = std::move(q.front()); // remove task...
        q.pop(); // ...from the queue...
        t(); // ...and execute it
    }
}
```

- Live demo: <https://godbolt.org/z/W3zPT5x7j>.
- The queue in ThreadPool can be now defined and used the same way.
- **Note** — if task returns something, the return value is discarded.

9

std::function

- C++ standard library provides a generic type-erased callable entity owner — class template `std::function`.
- **Our Task:**
 - Task can own and call an entity callable without arguments \Rightarrow basically callable entities without parameters.
 - If the type-erased entity returns, the return value is discarded by Task.
- **std::function:**
 - An instance of `std::function` template can own and call callable entities with particular types or parameters and type of return value.
 - The types of parameters and type of return value are determined by `std::function` template argument.
- **Examples:**
 - `std::function<void()>` — basically, equivalent of our Task \Rightarrow stored callable entities have no parameters and do not return values.
 - `std::function<bool(int,int)>` — callable entities have two parameters of type `int` and return `bool`.

10

std::packaged_task

- `std::packaged_task` — basically `std::function` with some additional functionality.
- Namely, packaged task allows to obtain a relevant *future* object (of corresponding `std::future` type).
- This future object allows to:
 - Wait for the task to be completed (its execution).
 - Obtain its return value.
 - Catch exception thrown within task execution.
- Moreover, future object may be accessed by other threads than the thread executing the task.
- \Rightarrow Packaged tasks are supposed to be run *asynchronously*.
- **Note** — suitable use case = *thread pools*.

11

std::packaged_task (cont.)

- Simple thread pool (*incomplete*) with packaged tasks:

```
class ThreadPool {
    using Task = std::packaged_task<void()>;
    std::mutex m_;
    std::queue<Task> q_;
    std::condition_variable cv_;
    std::vector<std::thread> threads_;
    void worker() {
        while (true) {
            Task task;
            {
                std::unique_lock<std::mutex> lock(m_);
                cv_.wait(lock, [this]{ return !this->q_.empty(); });
                task = std::move(q_.front());
                q_.pop();
            }
            task(); // task executed here
        }
    }
public:
    ThreadPool(int num_threads = std::thread::hardware_concurrency()) {
        for (int i = 0; i < num_threads; i++)
            threads_.emplace_back([this]{ this->worker(); });
    }
    std::future<void> enqueue(Task task) {
        std::future<void> future;
        {
            std::lock_guard<std::mutex> lock(m_);
            future = task.get_future();
            q_.emplace(std::move(task));
        }
        cv_.notify_one(); // wake up some thread
        return future;
    }
};
```

12

Type erasure — std::any

- *Recall* — Any class can store/own type-erased object of any type.
- Any type \Rightarrow there is no common interface.
- However, we can provide access to the stored type-erased object.
- *First option* — type-less pointer:
 - In Base and Any classes, there is no knowledge of T.
 - \Rightarrow Only possibility is to return a pointer to the stored object of type void*.

```
class Any {
    struct Base {
        virtual ~Base() = default;
        virtual void* get() = 0;
    };
    template <typename T>
    struct Helper : Base {
        T t_;
        Helper(const T& t) : t_(t) {}
        virtual void* get() override { return &t_; }
    };
    std::unique_ptr<Base> ptr_;
public:
    template <typename T>
    Any(const T& t) : ptr_(new Helper<T>(t)) {}
    void* get() { return ptr_>get(); }
};
```

13

Type erasure — std::any (cont.)

- *Consequence* — to use the object, returned pointer needs to be casted into the pointer-to-object-type.
- \Rightarrow This type must be provided from outside!
 - There is no other way; Any class cannot "remember" a type.

```
std::vector<Any> v;
v.emplace_back(1); // type-erased int
v.emplace_back(std::string("some string")); // type-erased std::string
int* ptr = static_cast<int*>(v[0].get()); // pointer to the stored int
std::string& ref = static_cast<std::string*>(v[1].get()); // reference to the stored...
// ...std::string
```

- *Better option* — wrapping "ugly" casting with a custom function:

```
class Any {
    ...
public:
    ...
    template <typename T> friend T Any_cast(const Any& a); // friend to have access to ptr_
};

template <typename T> T Any_cast(const Any& a) // free (non-member) function template
{ return *static_cast<std::remove_reference_t<T*>>(a.ptr_>get()); }
```

- *Note* — Any_cast can either return a reference to the stored object, or its copy.

14

Type erasure — std::any (cont.)

- Such a generic type-erased object owner is provided by the C++ standard library as std::any class.
- It is available since C++17.

```
std::vector<std::any> v;
v.emplace_back(1); // type-erased int
v.emplace_back(std::string("some string")); // type-erased std::string
int val = std::any_cast<int>(v[0]); // copy of the type-erased object
std::string& ref = std::any_cast<std::string*>(v[1]); // reference to the type-erased object
```

- *Notes:*
 - In contrast to cast operators (static_cast, reinterpret_cast,...), std::any_cast is not an operator.
 - Instead, it is an "ordinary" function template.
- *Use cases:*
 - std::any is suitable for very specific use cases only.
 - *Relevant discussion:* <https://devblogs.microsoft.com/cppblog/stdany-how-when-and-why/>.

15

std::variant

- std::any — a single-object owner.
 - The owned object is stored in dynamically allocated storage.
 - Its access involves virtual-function dispatch.
 - \Rightarrow Relatively large runtime and memory overhead.
 - Can store an object of any type, but one needs to "recall" this type during object access.
- Another C++ standard library option — std::variant:
 - A single-object owner as well.
 - Is a *variadic template* — template arguments define a list of types.
 - Only objects of these types may be stored in variant.

```
std::any a1 = 1; // int stored
std::any a2 = true; // bool stored
std::any a3 = 1.0; // double stored
a1 = 1.0; // double stored

std::variant<int, bool> v1 = 1; // int stored
std::variant<int, bool> v2 = true; // bool stored
std::variant<int, bool> v3 = 1.0; // ERROR
v1 = 1.0; // ERROR
```

- *Note* — all any objects have the same type, variants with different type list have different types.

16

std::variant (cont.)

- List of types defines their *order*.
- \Rightarrow Each type has some *index* in this order.
- Index of the type of actually stored object is returned by the *index* member function.

```
std::variant<int, bool, double> v = 1.0; // int has index 0, bool 1, double 2
std::cout << v.index(); // double is stored => prints out "2"
```

- Stored object is accessed with std::get free function template.
 - Template argument — index of the stored object type, or the type itself.
 - Returns a reference to the owned object.

```
std::variant<int, bool, double> v = 1.0;
std::cout << std::get<2>(v); // prints out "2"
std::cout << std::get<double>(v); // effectively the same
```

- *Note* — C++ is a statically typed language.
 - \Rightarrow All types must be resolved at runtime.
 - \Rightarrow This holds also for the return type of std::get.
 - \Rightarrow index function call cannot be used as a template argument of std::get.

17

std::variant (cont.)

- Owned object needs to be stored in the included storage.
 - \Rightarrow Variant is explicitly disallowed to dynamically allocate memory [link].

```
std::variant<int, bool> v = 1;
std::cout << (uintptr_t)&std::get<0>(v) - (uintptr_t)&v << std::endl; // print "0" (GCC)
```

- *Implementation* — included buffer aligned and sized suitably according to the type list.
 - *Buffer alignment* = maximum of alignment requirements for of types.
 - *Buffer size* = maximum of size requirements of types.
 - *Note* — it may be defined with std::aligned_union library helper type.
- \Rightarrow No runtime/memory overhead.
- \Rightarrow std::variant is much preferred than std::any once there is a fixed list of type alternatives.
- *Benchmark* — std::variant vs std::any:
 - *insertion* (object storage) — 6.6x faster with std::variant [link];
 - *stored object access* — 6.1x faster with std::variant [link].

18

std::variant (cont.)

- **std::variant** vs **std::tuple**:
 - **std::tuple** — multiple object owner \Rightarrow owns/stores a single object for each type from the type list (template arguments).
 - **std::variant** — single object owner \Rightarrow owns/stores at most one object at a given moment (of a type from the type list).

```
std::variant<int, bool, double> v = 1; // owns int
std::cout << &std::get<int>(v) << std::endl; // printed "0x7ffff6a7b030" in experiment
v = true; // owns bool
std::cout << &std::get<bool>(v) << std::endl; // printed "0x7ffff6a7b030"
v = 1.0; // owns double
std::cout << &std::get<double>(v) << std::endl; // printed "0x7ffff6a7b030"
```

- **std::variant** vs **union**:
 - **std::variant** — basically a type-safe union.
 - **union** is low-level "C" type, which is very unsafe and problematic with classes.
 - **Price** — **std::variant** needs to remember the actual "alternative" (index).

```
union { int i; bool b; double d; } u;
std::cout << sizeof(u); // printed "8" in experiment
std::variant<int, bool, double> v;
std::cout << sizeof(v); // printed "16" in experiment
```

19

CRTP

- **Example** — vector 2D graphic editor \Rightarrow hierarchy of polymorphic classes that represent various graphic objects:

```
class Object2D { // abstract base class
public:
    virtual ~Object2D() = default;
    virtual void draw() const;
    ...
};
class Line : public Object2D { ... };
class Square : public Object2D { ... };
class Circle : public Object2D { ... };
```

- **Scene** is represented as a collection of graphic objects:

```
std::vector< std::unique_ptr<Object2D> > scene; // all objects in a scene
```

- **Exemplary functionality** — loading scene from a file:

```
void load_scene(const std::string& filename) {
    std::ifstream f(filename); std::string line;
    while (std::getline(f, line)) {
        if (s == "line") scene.emplace_back( new Line() );
        else if (s == "circle") scene.emplace_back( new Circle() );
        else if (s == "square") scene.emplace_back( new Square() );
        else throw std::invalid_argument("Invalid object name");
    }
    for (const auto& obj : scene) obj->draw(); // draw loaded scene
}
```

20

CRTP (cont.)

- Another functionality — selection of objects:

```
std::vector< Object2D* > selected; // set of selected objects
void select(); // function that put pointers to selected objects into selected vector
```

```
// resolution of user actions:
if (ACTION == SELECT) select();
...
```

- **Note** — **selected** is a vector of "normal" (raw) pointers.
 - Vectors cannot hold references.
 - *Smart pointers* do not make sense here \Leftarrow vector **selected** does not own selected objects; only refers to them.

- Yet another functionality — duplication of selected objects.

```
// resolution of user actions:
if (ACTION == SELECT) select();
else if (ACTION == DUPLICATE) duplicate();
...
```

```
void duplicate() {
    for (auto p_obj : selected) scene.emplace_back( ??? );
}
```

- **Problem** — how to create a copy of a pointed-to object?

21

CRTP (cont.)

```
void duplicate() {
    for (auto p_obj : selected) scene.emplace_back( ??? );
}
```

- **Analysis:**

- **p_obj** points to an object of derived class type (Line, Square, or Circle).
- Type of **p_obj** is a pointer-to-base (Object2D*).
- \Rightarrow We need to invoke a copy constructor of a derived class through a pointer to a base class. *How?*
- Generally, this mechanism is provided by virtual functions.
- \Rightarrow We would need a "virtual copy constructor".
- However, C++ does not support virtual constructors.
- **Alternative** — delegation of a copying semantics into a separate virtual function.
 - This function is typically called **clone**.
 - The whole solution is sometimes called "clone pattern/idiom".

22

CRTP (cont.)

- **Solution** — polymorphic copying/cloning:

```
class Object2D {
    ...
    virtual Object2D* clone() const = 0;
};
class Line : public Object2D {
    ...
    virtual Object2D* clone() const { return new Line(*this); } // copy constructor of...
    // ...Line called here
};
class Square : public Object2D {
    ...
    virtual Object2D* clone() const { return new Square(*this); } // copy constructor of...
    // ...Square called here
};
class Circle : public Object2D {
    ...
    virtual Object2D* clone() const { return new Circle(*this); } // copy constructor of...
    // ...Circle called here
};
```

- **Application** to scene duplication:

```
void duplicate() {
    for (auto p_obj : selected) scene.emplace_back( p_obj->clone() );
}
```

- **Drawback** — clone function needs to be repeatedly / "redundantly" implemented throughout the whole class hierarchy.

23

CRTP (cont.)

- **Alternative option** — couldn't we "inject" **clone** function into derived classes from some base class?
- Obviously, it is not possible from **Object2D** base class itself.
 - There is no knowledge of the cloned type there.
- The base class where **clone** is defined must know about the cloned object type.
- **Solution:**
 - Making another base class that is a template...
 - ...and "remembering" the cloned type in its template parameter.
- **Note:**
 - We cannot make **Object2D** a template.
 - **Polymorphism** requires a single unique base class.
 - Templated **Object2D** would create a separate base class for each template argument.
- **Solution (cont.)** — templated "middle" class with **clone** that:
 - inherits from **Object2D** (\Rightarrow polymorphism with single base),
 - serves as a base for graphic object classes (\Rightarrow injects **clone** into them).

24

CRTP (cont.)

• *Partial solution:*

```
class Object2D {                // single (original) base => polymorphism
...
    virtual Object2D* clone() const = 0;
};

template <typename T>
class CloneableObject2D : public Object2D {    // template argument = cloned type
...                                           // middle class => defines clone
    virtual Object2D* clone() const { ??? }
};
```

- Derived 2d graphic object classes:
 - derive from CloneableObject2D class template,
 - where cloned type = template argument is derived class itself.
- For example, Line derives from CloneableObject2D and the cloned type needs to be Line \Rightarrow template argument needs to be Line.

```
class Line : public CloneableObject2D<Line> {
...
};
```

- This technique — inheritance from derived template where base template argument is derived class itself — is called "*curiously recurring template pattern*" (CRTP).

25

CRTP (cont.)

• *Last step — definition of clone:*

```
template <typename T> class CloneableObject2D : public Object2D {
...
    virtual Object2D* clone() const { return new T( ??? ); } // invoke copy constructor of T
};
```

- clone needs to create new object with copy constructor of T.
- The source object to be copied is pointed to by this pointer.
- Problem** — the type of this in CloneableObject2D<T>::clone() call is CloneableObject2D<T>*.
- But, we know that it actually points to an object of type T, which should be copied.
- Example:**

```
Object2D* ptr = new Line;
ptr->clone(); // inside, type of 'this' is CloneableObject2D<Line>*, but it points to Line
```

- Consequence** — this needs to be casted into T*:

```
virtual Object2D* clone() const {
    return new T( *static_cast<const T*>(this) ); // or, 'static_cast<const T*>(this)'
}
```

26

CRTP (cont.)

• *Summary:*

```
class Object2D {                // single (original) base => polymorphism
...
    virtual Object2D* clone() const = 0;
};

template <typename T>
class CloneableObject2D : public Object2D {    // template argument = cloned type
...                                           // middle class => defines clone
    virtual Object2D* clone() const {
        return new T( *static_cast<const T*>(this) );
    }
};

// derive 2D graphic object classes:
class Line : public CloneableObject2D<Line> > { ... };
class Square : public CloneableObject2D<Square> { ... };
class Circle : public CloneableObject2D<Circle> { ... };
```

- Application** to scene duplication is the same:

```
void duplicate() {
    for (auto p_obj : selected) scene.emplace_back( p_obj->clone() );
}
```

- Advantage** — clone is defined only once.
 - \Rightarrow Less code redundancy \Rightarrow better code maintenance.
- Note** — generally, CRTP has broader use than just the *clone idiom*.

27

CRTP (cont.)

• *Instantiation:*

```
class Object2D {
...
    virtual Object2D* clone() const = 0;
};

class CloneableObject2D<Line> : public Object2D {
...
    virtual Object2D* clone() const { return new Line( *static_cast<const Line*>(this) ); }
};

class CloneableObject2D<Square> : public Object2D {
...
    virtual Object2D* clone() const { return new Square( *static_cast<const Square*>(this) ); }
};

class CloneableObject2D<Circle> : public Object2D {
...
    virtual Object2D* clone() const { return new Circle( *static_cast<const Circle*>(this) ); }
};

class Line : public CloneableObject2D<Line> > { ... };
class Square : public CloneableObject2D<Square> { ... };
class Circle : public CloneableObject2D<Circle> { ... };
```

- Result** — one base class, three middle classes, three derived classes.
- Exemplary resolution** — if ptr points to Line, ptr->clone() actually calls CloneableObject2D<Line>::clone(), which copies Line by its copy constructor.

28

Function swap

• *Special member functions* related to object content:

- default constructor (sets object to "empty" state),
- copy constructor (copies content into initialized object),
- move constructor (moves content into initialized object),
- copy assignment operator (copies content into existing object),
- move assignment operator (moves content into existing object),
- destructor (destroys content).

- The member functions are *special* in that:

- they are automatically generated under some circumstances,
- they may be called without explicit function call syntax.

• *Example:*

```
struct X {
    std::string s;
}; // all 6 special member functions are...
// ...automatically generated for TwoStrings class

void f(X);

int main() {
    X x; // default constructor called
    f(x); // copy constructor called
} // destructor called
```

29

Function swap (cont.)

- There is one additional "*special*" member function — swap for swapping content of two existing objects.
- It is not truly special \Rightarrow not automatically generated, not called without explicit function call syntax.
- However, it is widely used inside existing (not only C++ standard) library algorithms (sorting, for instance).
- Suppose such an algorithm needs to swap content of two objects of type T.
- Then, there are two options:
 - either there is a custom swap free function specifically designed for arguments of type T (typically defined in the same namespace as T),
 - or, there is no such a custom function.
- The algorithms then typically works as follows:
 - Ad 1) They use that custom function for required content swapping.
 - Ad 2) They use std::swap instead.

30

std::swap

- `std::swap` is a library function that is provided in two forms:
 - First, there is a *primary template* that works generally for objects of any type.
 - Then, there are "specializations" for various library types [\[link\]](#).

```
struct X { std::string s; };
int main() {
    int i = 1, j = 2;
    std::swap(i, j); // uses generic primary template
    std::cout << i << j; // prints out "21" => content-numbers were swapped
    std::string s1("world"), s2("hello");
    std::swap(s1, s2); // uses special overload of std::swap for std::string
    std::cout << s1 << s2; // prints out "hello world" => content-strings were swapped
    X x1("world"), x2("hello");
    std::swap(x1, x2); // uses generic primary template
    std::cout << x1.s << x2.s; // prints out "hello world" => content-strings were swapped
}
```

Questions:

- 1) Why are these "specializations" defined?
- 2) Isn't the primary template enough?
- 3) Should we define swap function for custom types?

32

31

std::swap (cont.)

- How is generic primary template of `std::swap` implemented?
 - It is not specified by the C++ standard, but implementations of C++ standard libraries typically define primary template *naturally* as follows:

```
template <typename T> // primary std::swap template
void swap(T& a, T& b) {
    T temp { std::move(a) }; // moves/copies content from first parameter into a temporary
    a = std::move(b); // moves/copies content from second parameter into the first
    b = std::move(temp); // moves/copies content from the temporary into the second param.
} // => this effectively swaps content of both parameters
```

- \Rightarrow For class types, primary template of `std::swap` is implemented in terms of its *special member functions*:

```
template <typename T> // if T is a class type:
void swap(T& a, T& b) {
    T temp { std::move(a) }; // move/copy constructor called (move preferred)
    a = std::move(b); // move/copy assignment operator called (ditto)
    b = std::move(temp); // move/copy assignment operator called (ditto)
} // destructor called (for temp)
```

- \Rightarrow *Move-content operations* are (naturally) preferred if they exist; otherwise, *copy-content operations* are used instead.
- **Note** — for *trivially-copyable types*, there is no moving of content.
 - \Rightarrow All operations involve copying of binary representations.

32

32

Custom swap vs std::swap (cont.)

- Is it worth defining special swap overload for custom types?
 - Obviously, it is worth for classes without move semantics and "expensive" copy semantics (such as *legacy/pre-C++11* classes).
- **Benchmark** — class with a `std::string` subobject.
- **Two versions** — first with move semantics:

```
struct X_move {
    std::string s;
}; // X_move HAS automatically generated move constructor and move assignment operator
```

- **Second**, with copy but without move semantics:

```
struct X_copy { // for instance, pre-C++11 legacy class
    std::string s;
    X_copy(const char* s) : s(s) {}
    X_copy(const X_copy& other) : s(other.s) {} // suppresses auto-generation...
}; // X_copy HAS NOT automatically generated move constructor and move assignment operator
```

- Comparison of swapping of two objects of these types with `std::swap` primary template:
 - Swapping of content was **4.3x faster** for `X_move` [\[link\]](#).

33

33

Custom swap vs std::swap (cont.)

- Let us add a custom swap function for a move-less version:

```
void swap(X_copy& a, X_copy& b)
{
    std::swap(a.s, b.s); // swapping content = swapping content of subobject s
}
```

- **Results** — swapping of content was **12x faster** with custom swap than with `std::swap` [\[link\]](#).
 - \Rightarrow It is usually crucial to provide swap function for custom classes if these do not provide move-content operations and are expensive to copy (typically "non-movable" classes that own some resources).
- What about classes that provide move-content operations?

```
void swap(X_move& a, X_move& b)
{
    std::swap(a.s, b.s);
}
```

- **Results** — swapping of content was **2.6x faster** with custom swap than with `std::swap` [\[link\]](#).

34

34

Custom swap vs std::swap (cont.)

- What is the difference:
 - In all examined cases, swapping of content of `X...` class object = swapping of content of its subobject (namely, its `s` member variable).
 - In case of `X_move` and custom swap, this *swapping of s* is accomplished with "specialization" of `std::swap` for `std::string` arguments:

```
void swap(X_move& a, X_move& b) {
    std::swap(a.s, b.s); // specialization of std::swap for std::string called here
}
```

- In case of `std::swap`, primary template of `std::swap` is used and instantiated as follows:

```
void swap(X_move& a, X_move& b) {
    X_move temp { std::move(a) }; // move constructor of X_move called
    a = std::move(b); // move assignment operator of X_move called
    b = std::move(temp); // move assignment operator of X_move called
} // destructor of X_move called
```

- Auto-generated special member function of `X_move` is effectively equivalent with the same special member function of `std::string`.
- \Rightarrow *Swapping of s* is accomplished with the following calls of `std::string` special member functions: 1x move constructor, 2x move assignment operator, 1x destructor.

35

35

Custom swap vs std::swap (cont.)

Conclusion:

- With a custom swap function, content of `s` member variable is swapped with the specialization of `std::swap` for `std::string` arguments.
- With `std::swap` function, content of `s` member variable is swapped with:
 - single call of `std::string` move constructor,
 - double call of `std::string` move assignment operator,
 - and single call of `std::string` destructor.
- In both cases, the final outcome (observable behavior) of swapping is the very same!
- However, within our experiment and used C++ implementation, the compiler was not able to optimize both cases into the equally efficient machine code.
- **Why not?**

36

36

Custom swap vs std::swap (cont.)

• Exemplary problem:

- After each *move-content* operation, *s* member variable of source object needs to be put into *moved-from* = "empty string" state, which requires some work.
- However, within `std::swap` instance, this is unnecessary since in all 3 cases, the *s* member variable is then either *assigned new content* or *deconstructed*:

```
void swap(X_move& X_move& a, X_move& b) {
    X_move temp { std::move(a) }; // a.s put into empty string state...
    a = std::move(b);           // ...while it is then assigned new content
}
```

```
a = std::move(b); // b.s put into empty string state...
b = std::move(temp); // ...while it is then assigned new content
```

```
b = std::move(temp); // temp.s put into empty string state...
// ...while it is then destructed
```

- Compiler optimizers are simply not *perfect/ideal* from our point of view.
- Here, a compiler may not recognize that "resetting" strings is unnecessary, and generate machine code for it.
- Note** — due to SSO, *move-content* operations for `std::string` are non-trivial but relatively complex.

37

Custom swap vs std::swap (cont.)

- On the contrary, "specialization" of `std::swap` for `std::string` does not require any resetting to empty state.
- If strings are *long*, swapping of content effectively just turns into swapping content of few member variables (pointer + integers).
- ⇒ "Specialization" can be implemented in such a way that it does not contain any unnecessary operations.
- With generic `std::swap` based on *move-content* operations, additional operations are involved and it may be hard or impossible for a compiler to optimize them away and generate equally efficient machine code.
- Experiment** — swapping content of two `std::string` objects with primary template and specialization of `std::swap`: [link](#).
- Consequence** — generally, even for classes that provide *move-content* operations, it may be worth providing custom swap function.

38

Using swap function

• How to use swap function, for instance, in a quicksort algorithm?

```
template <typename T> long partition(T* a, long lo, long hi) {
    long i = lo;
    for (long j = lo; j < hi; j++)
        if (a[j] < a[hi]) ???; // need to swap a[i++] and a[j], HOW?
    ???; // need to swap a[i] and a[hi], HOW?
    return i;
}

template <typename T> void qs(T* a, long lo, long hi) {
    if (lo < hi) {
        long p = partition(a, lo, hi);
        qs(a, lo, p - 1);
        qs(a, p + 1, hi);
    }
}
```

- Within partitioning, we need to swap sorted elements; *how?*
- We want:
 - to use custom swap function if it exists,
 - to use `std::swap` otherwise.
- If we called
 - `swap`, then compilation error would be generated in the *second case*.
 - `std::swap`, then custom swap would not be used in the *first case*.

39

Using swap function

• Solution is simple:

```
template <typename T> long partition(T* a, long lo, long hi) {
    using std::swap;
    long i = lo;
    for (long j = lo; j < hi; j++)
        if (a[j] < a[hi]) swap(a[i++], a[j]);
    swap(a[i], a[hi]);
    return i;
}
```

- Directive using "injects" the name `swap` from `std` namespace into *global namespace*.
- Consequently**, `swap` call will:
 - use custom swap (its "specialization") if it exists (better overloading candidate),
 - fall back to primary `std::swap` template otherwise.
- Apparent problem:**
 - What if custom swap is in some custom (for instance, library) *namespace*?
 - Will this solution work?

40

Argument dependent lookup

• Example:

```
namespace our_library {
    struct X {
        std::string s;
        bool operator<(const X& rhs) const { return s < rhs.s; } // Lexicographic comparison
    };
    void swap(X& a, X& b) { std::swap(a.s, b.s); }
}
```

```
int main() {
    our_library::X a[2] = { "B", "A" };
    qs(a, 0, 1);
}
```

- In partition, there is no "knowledge" of `our_library` namespace:

```
template <typename T> long partition(T* a, long lo, long hi) {
    using std::swap;
    using our_library::swap; // we don't know about our_library namespace here
    long i = lo;
    for (long j = lo; j < hi; j++)
        if (a[j] < a[hi]) swap(a[i++], a[j]);
    swap(a[i], a[hi]);
    return i;
}
```

41

Argument dependent lookup (cont.)

- Which overload of `swap` will be called when `T` is `our_library::X`?

```
long partition(our_library::X* a, long lo, long hi) {
    using std::swap;
    long i = lo;
    for (long j = lo; j < hi; j++)
        if (a[j] < a[hi]) swap(a[i++], a[j]);
    swap(a[i], a[hi]);
    return i;
}
```

- Good news** — custom `our_library::swap` will be used in this case.
- Reason = "**argument dependent lookup**" (ADL):
 - If types of arguments are from some namespace, candidates for overloading are looked-up in this namespace as well.
- Without ADL, it would be almost impossible to write generic code.
 - Its developer can't know about namespaces from which involved types of objects will come from.

42